



departamento de informática
FACULDADE DE CIÊNCIAS E TECNOLOGIA
UNIVERSIDADE NOVA DE LISBOA

Parallel Algorithms

Concurrency and Parallelism — 2018-19


Master in Computer Science

(Mestrado Integrado em Eng. Informática)

Joao Lourenço <joao.lourenco@fct.unl.pt>


Outline

- Parallel computations as DAGs
 - Parallel computing by divide-and-conquer
 - Maps and reductions on tree-like DAGs
 - The Prefix-Sum (Scan) problem and its parallel solution
 - An implementation for the Pack parallel pattern
- Bibliography:
 - **Chapter 3, 4 and 5** of book
McCool M., Arch M., Reinders J.;
Structured Parallel Programming: Patterns for
Efficient Computation;
Morgan Kaufmann (2012);
ISBN: 978-0-12-415993-8




The DAG

- A program execution using fork and join can be seen as a DAG
 - Nodes: Pieces of work
 - Edges: Source must finish before destination starts
- A fork “ends a node” and makes two outgoing edges
 - New thread
 - Continuation of current thread
- A join “ends a node” and makes a node with two incoming edges
 - Node just ended
 - Last node of thread joined on




A simple example

- **fork** and **join** are very flexible, but divide-and-conquer use them in a very basic way:
 - A tree on top of an upside-down tree



Another example: reduce

- Summing an array went from $O(n)$ sequential to $O(\log n)$ parallel (assuming a lot of processors and very large n)



- Anything that can use results from two halves and merge them in $O(1)$ time has the same properties and exponential speed-up (in theory)

Applications of “reduce”

- Maximum or minimum element
- Is there an element satisfying some property?
 - e.g., is there a 17?
- Left-most element satisfying some property?
 - e.g., index of first 17
- Corners of a rectangle containing all points (a “bounding box”)
- Counts
 - e.g., # of strings that start with a vowel
 - This is just summing with a different base case

More Interesting DAGs?

- Of course, the DAGs are not always so simple (and neither are the related parallel problems)
- Example:
- Suppose combining two results might be expensive enough that we want to parallelize each one
- Then each node in the inverted tree on the previous slide would itself expand into another set of nodes for that parallel computation

Reductions


- Such computations of this simple form are common enough to have a name: reductions (or reduces?)
- Produce single answer from collection via an associative operator
 - Examples: max, count, leftmost, rightmost, sum, ...
 - Non-example: median
- Recursive results don't have to be single numbers or strings and can be arrays or objects with fields
 - Example: Histogram of test results
- But some things are inherently sequential
 - How we process `arr[i]` may depend entirely on the result of processing `arr[i-1]`

Maps and Reductions on Trees

- Work just fine on balanced trees
 - Divide-and-conquer each child
 - Example:
Finding the minimum element in an unsorted but balanced binary tree takes $O(\log n)$ time given enough processors
- How to do you implement the sequential cut-off?
 - Each node stores number-of-descendants (easy to maintain)
 - Or approximate it (e.g., AVL tree height)
- Parallelism also correct for unbalanced trees but obviously one gets worse speed-ups

Linked Lists

- Can you parallelize maps or reduces over linked lists?
 - Example: Increment all elements of a linked list
 - Example: Sum all elements of a linked list



- Nope. Once again, data structures matter!
- For parallelism, balanced trees are generally better than lists so that we can get to all the data exponentially faster $O(\log n)$ vs. $O(n)$
 - Trees have the same flexibility as lists compared to arrays (i.e., no shifting for insert or remove)

Division of Responsibility

- Parallel Framework users (e.g., Cilk+, Java ForkJoin)
 - Pick a good parallel algorithm and implement it
 - Its execution creates a DAG of things to do
 - Make all the nodes small(ish) and approximately equal amount of work
- The framework-writer's job:
 - Assign work to available processors to avoid idling
 - Keep constant factors low
 - Give the expected-time optimal guarantee assuming framework-user did his/her job
- $T_P = O((T_1 / P) + T_\infty)$

Examples: $T_P = O((T_1 / P) + T_\infty)$

- Sum an array
 - $T_1 = O(n)$ and $T_\infty = O(\log n)$ $\Rightarrow T_P = O(n / P + \log n)$
- Suppose
 - $T_1 = O(n^2)$ and $T_\infty = O(n)$ $\Rightarrow T_P = O(n^2 / P + n)$
- Of course, these expectations ignore any overhead or memory issues

The Prefix (Scan) Sum Problem

- Given `int[] input`, produce `int[] output` such that:

`output[i] = input[0] + input[1] + ... + input[i]`

- A sequential solution in a typical exam problem:

```
int[] prefix_sum(int[] input){
    int[] output = new int[input.length];
    output[0] = input[0];
    for(int i=1; i < input.length; i++)
        output[i] = output[i-1] + input[i];
    return output;
}
```

The Prefix (Scan) Sum Problem

```
int[] prefix_sum(int[] input){
    int[] output = new int[input.length];
    output[0] = input[0];
    for(int i=1; i < input.length; i++)
        output[i] = output[i-1]+input[i];
    return output;
}
```

- Above algorithm does not seem to be parallelizable!
 - Work (T_1): $O(n)$ Span (T_∞): $O(n)$
- It isn't. The above algorithm is sequential.
- But a different algorithm gives a span of $O(\log n)$

Parallel Prefix-Sum

- The parallel-prefix algorithm does two passes
 - Each pass has $O(n)$ work and $O(\log n)$ span
 - In total there is $O(n)$ work and $O(\log n)$ span
 - Just like array summing, parallelism is $O(n / \log n)$
 - An exponential speedup
- The first pass builds a tree bottom-up
- The second pass traverses the tree top-down

Historical note:

*Original algorithm due to R. Ladner
and M. Fischer at the UW in 1977*



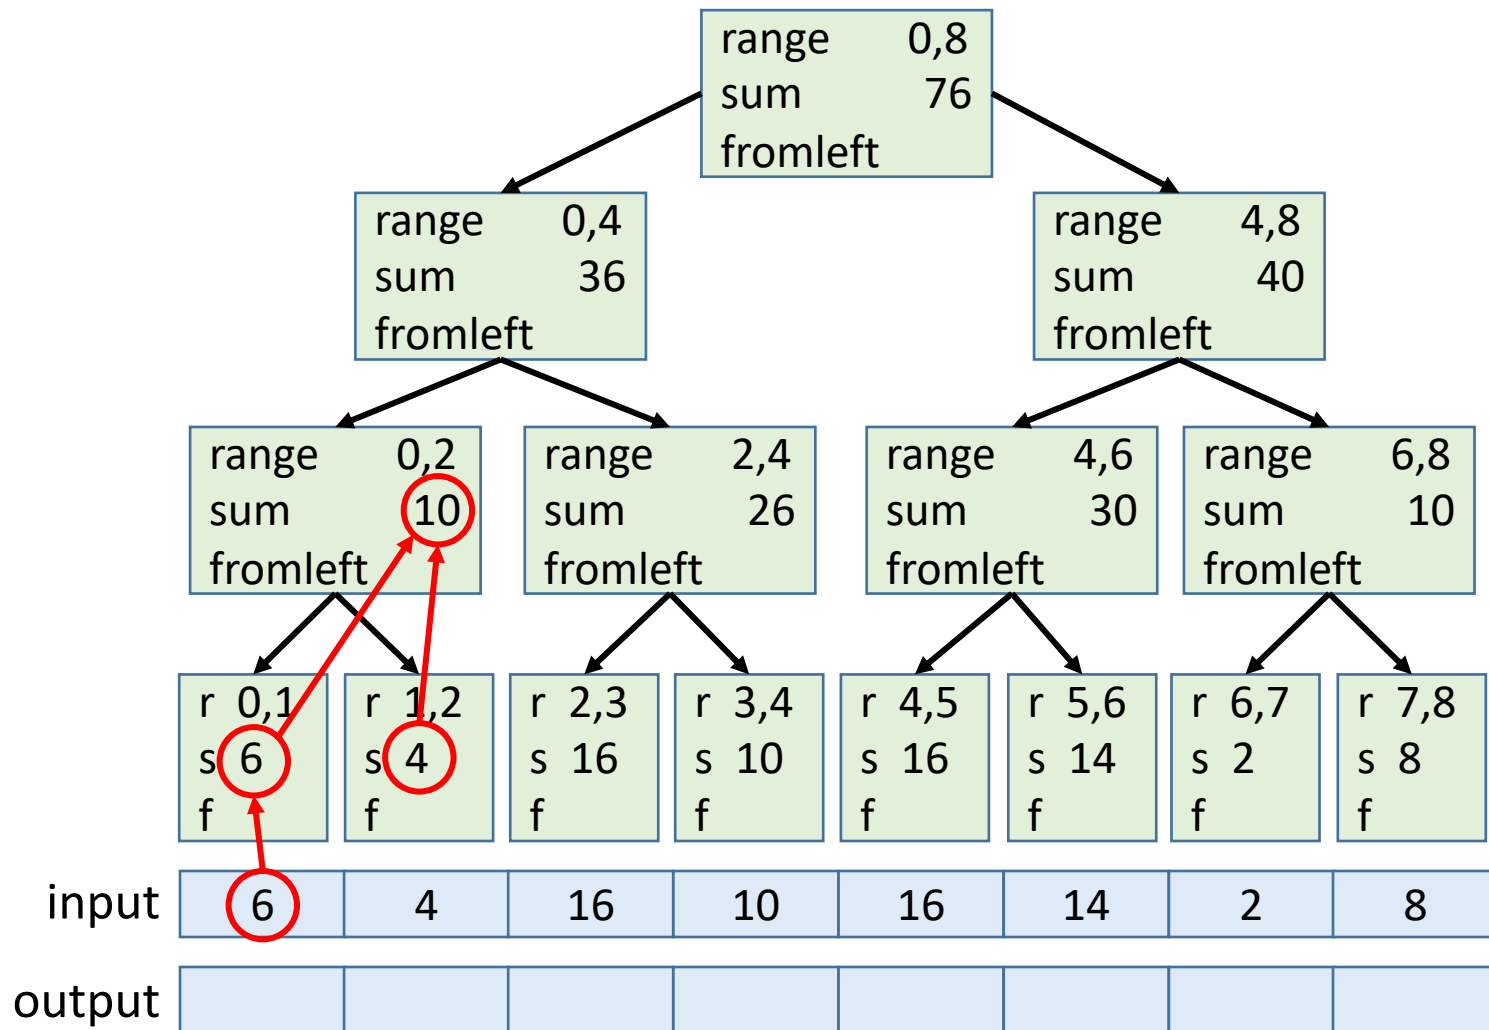
Parallel Prefix: The Up Pass

- We build want to a binary tree where
 - Root has sum of the range $[x,y]$
 - If a node has sum of $[lo,hi]$ and $hi > lo$,
 - Left child has sum of $[lo,middle]$
 - Right child has sum of $[middle,hi]$
 - A leaf has sum of $[i,i+1]$, which is simply $input[i]$
- It is critical that we actually create the tree as we will need it for the down pass
 - We do not need an actual linked structure
 - We could use an array as we did with heaps

Parallel Prefix: The Up Pass

- This is an easy fork-join computation:
- `buildRange(arr,lo,hi)`
 - If `lo+1 == hi`, create new node with sum `arr[lo]`
 - Else, create two new threads:
 - `buildRange(arr,lo,mid)`
 - `buildRange(arr,mid+1,high)`
 - Where `mid = (low+high)/2`
 - When threads complete, make new node with
 - `sum = left.sum + right.sum`
- Performance Analysis:
 - Work: $O(n)$
 - Span: $O(\log n)$

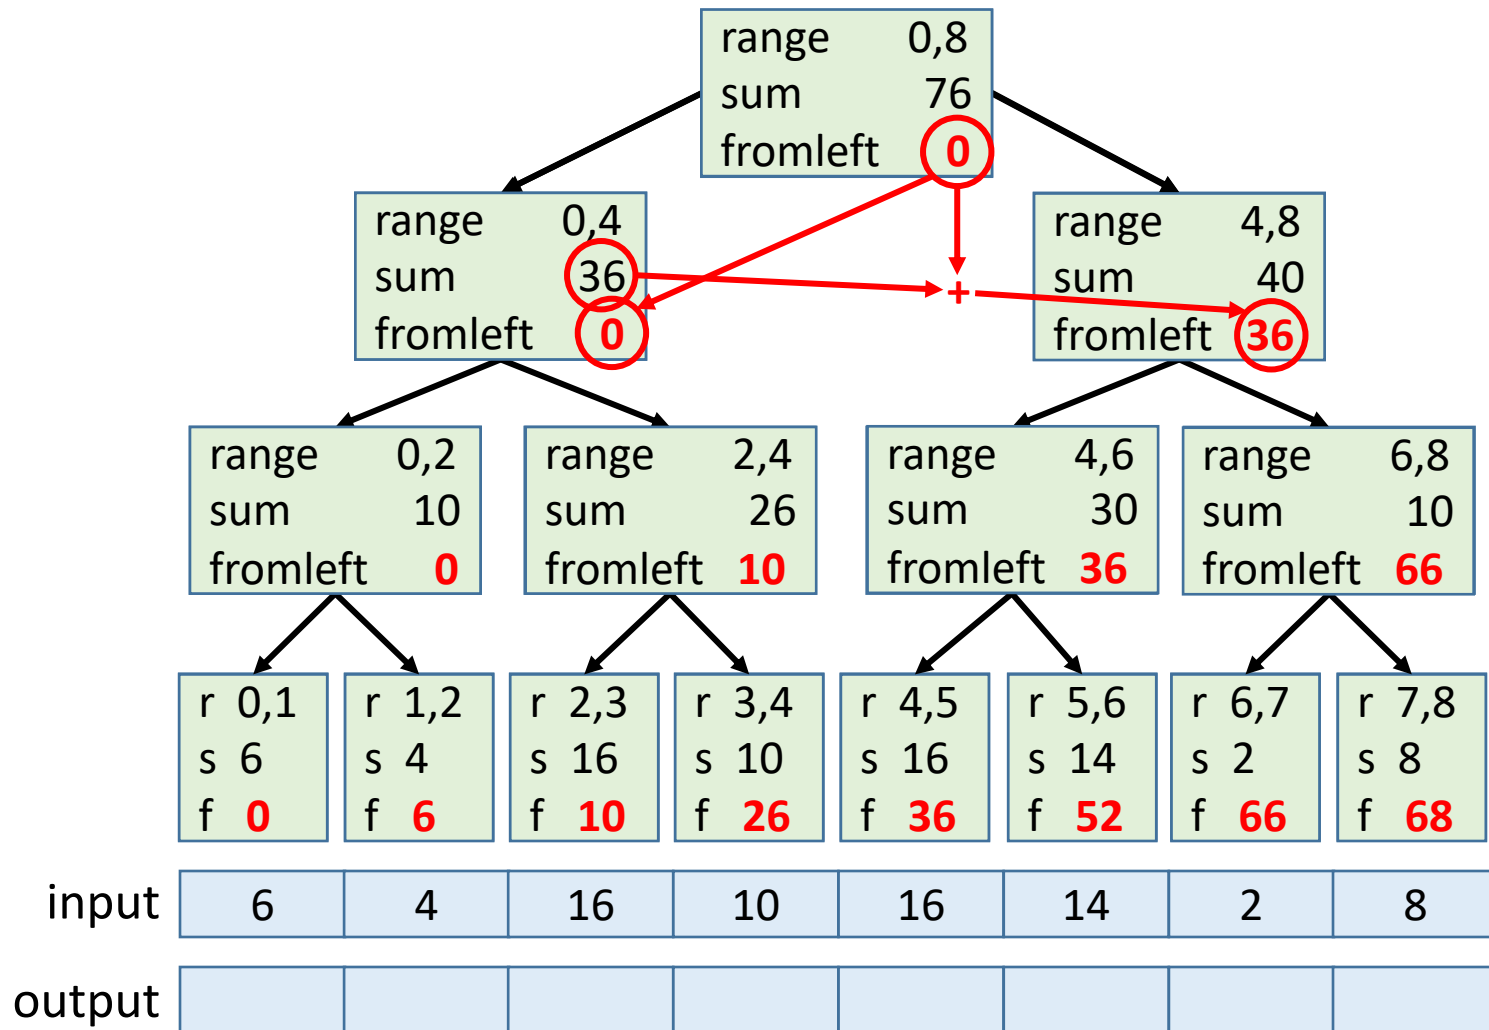
Up Pass Example



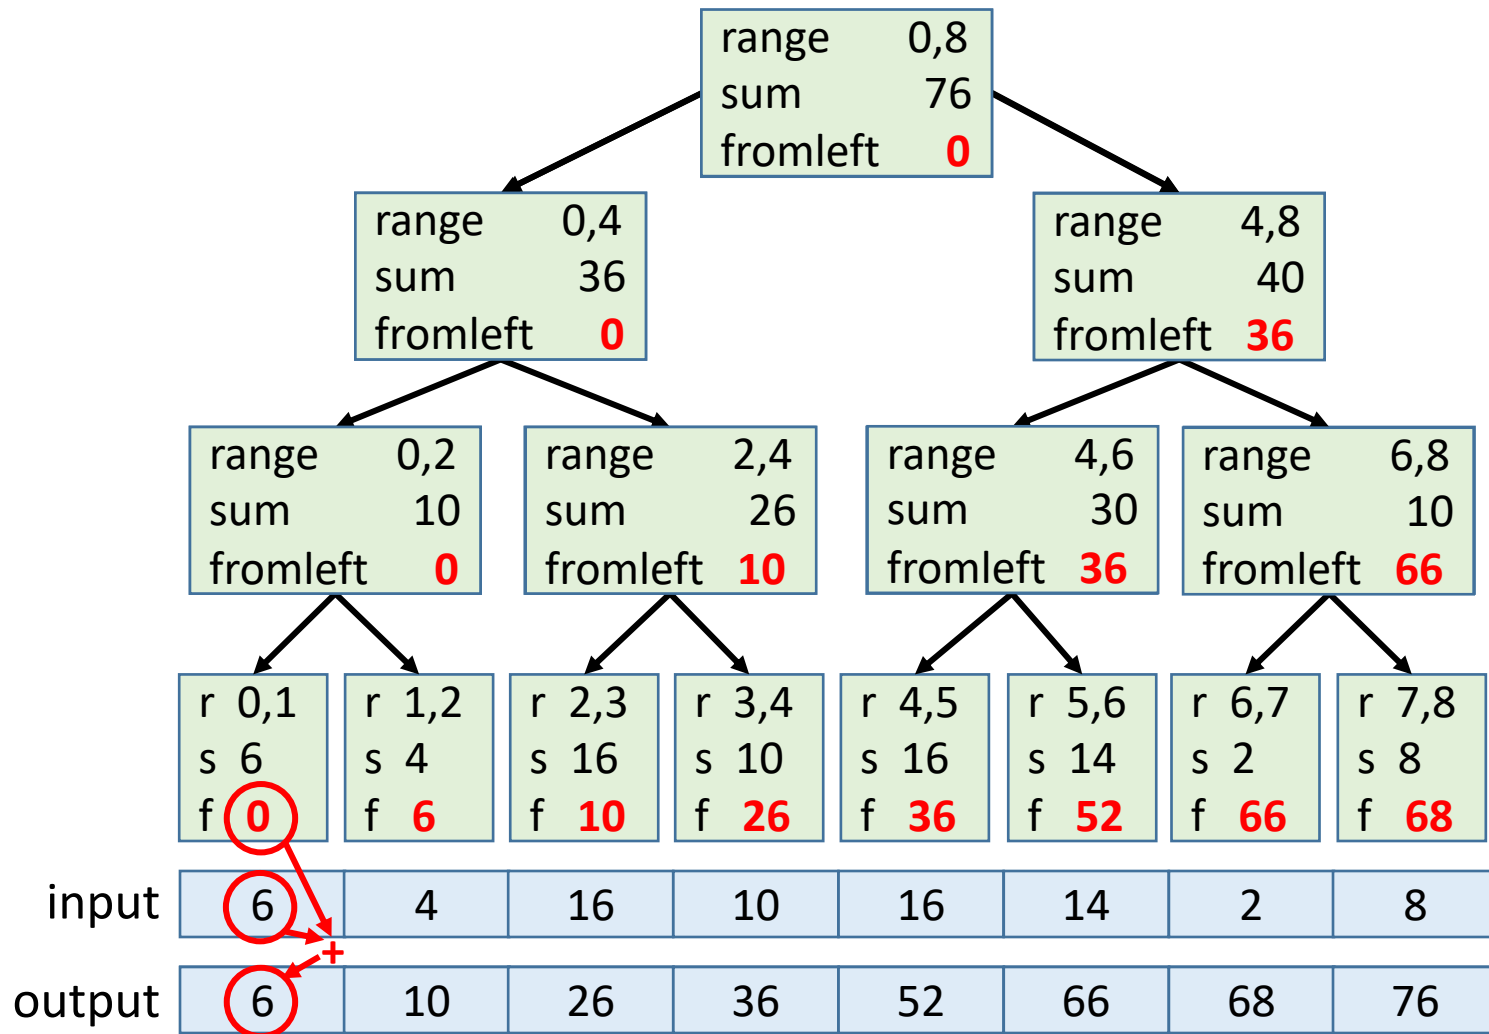
Parallel Prefix: The Down Pass

- We now use the tree to get the prefix sums using an easy fork-join computation
- Starting at the root:
 - Root is given a fromLeft of 0
 - Each node takes its fromLeft value and:
 - Passes to the left child: fromLeft
 - Passes to the right child: fromLeft + left.sum
 - At leaf for position i , $\text{output}[i] = \text{fromLeft} + \text{input}[i]$
- Invariant: **fromLeft is sum of elements left of the node's range**

Down Pass Example



Down Pass Example



Parallel Prefix: The Down Pass

- Note that this parallel algorithm does not return any values
 - Leaves assign to output array
 - **This is a map, not a reduction**
- Performance Analysis:
 - Work: $O(n)$
 - Span: $O(\log n)$

Generalizing Parallel Prefix

- Prefix-sum illustrates a pattern that can be used in many problems
 - Minimum, maximum of all elements to the left of i
 - Is there an element to the left of i satisfying some property?
 - Count of elements to the left of i satisfying some property!
- That last one is perfect for an efficient parallel pack that builds on top of the “parallel prefix trick”

Pack (Think Filtering)

- Given an array `input` and boolean function $f(e)$ produce an array `output` containing only elements e such that $f(e)$ is `true`
- Example:
input [17, 4, 6, 8, 11, 5, 13, 19, 0, 24]
 $f(e)$: is $e > 10$?
output [17, 11, 13, 19, 24]
- Is this parallelizable? Of course!
 - Finding elements for the output is easy
 - But getting them in the right place seems hard

Pack: Parallel Map + Parallel Prefix + Parallel Map

1. Use a parallel map to compute a bit-vector for true elements

```
input    [17, 4, 6, 8, 11, 5, 13, 19, 0, 24]
bits     [ 1, 0, 0, 0,  1, 0,  1,  1, 0,  1]
```

2. Parallel-prefix sum on the bit-vector

```
bitsum   [ 1, 1, 1, 1,  2, 2,  3,  4, 4,  5]
```

3. Parallel map to produce the output

```
bitsum   [ 17, 11, 13, 19, 24]
```

```
output = new array of size bitsum[n-1]
FORALL(i=0; i < input.length; i++){
    if(bits[i]==1)
        output[bitsum[i]-1] = input[i];
}
```

Pack Comments

- First two steps can be combined into one pass
 - Will require changing base case for the prefix sum
 - No effect on asymptotic complexity
- Can also combine third step into the down pass of the prefix sum
 - Again no effect on asymptotic complexity
- Analysis: Work: $O(n)$; Span: $O(\log n)$
 - Multiple passes, but this is a constant

The END

- Sources:
 - Parallel Computing, CIS 410/510, Department of Computer and Information Science
 - <https://courses.cs.washington.edu/courses/cse332/12su/slides/lecture12-parallelism-work-span.pdf>