



departamento de informática
FACULDADE DE CIÊNCIAS E TECNOLOGIA
UNIVERSIDADE NOVA DE LISBOA

Monitoring Concurrency Errors: Detection of Deadlocks, Atomicity Violations, and Data Races (2)

Concurrency and Parallelism — 2018-19
Master in Computer Science
(Mestrado Integrado em Eng. Informática)

Joao Lourenço <joao.lourenco@fct.unl.pt>

Agenda

- Concurrency Anomalies
- Assigning Semantics to Concurrent Programs
- Concurrency Errors
 - Detection of data races
 - Detection of high-level data races and stale value errors
 - Detection of deadlocks

Concurrency Errors

Data Race Detection

Overview

- Static program analysis
- Dynamic program analysis
 - Lock-set algorithm
 - Happens-Before
 - Noise-Injection

Static Data Race Detection

- Advantages:
 - Reason about all inputs/interleavings
 - No run-time overhead
 - Adapt well-understood static-analysis techniques
 - Possibly with annotations to document concurrency invariants
- Example Tools:
 - RCC/Java type-based
 - ESC/Java "functional verification"
 (theorem proving-based)

Static Data Race Detection

- Advantages:
 - Reason about all inputs/interleavings
 - No run-time overhead
 - Adapt well-understood static-analysis techniques
 - Possibly with annotations to document concurrency invariants
- Disadvantages of static approach:
 - Tools produce “false positives” and/or “false negatives”
 - May be slow, require programmer annotations
 - May be hard to interpret results
 - May not scale to large or complex programs

Dynamic Data Race Detection

- Advantages

- Soundness
 - Every actual data race is reported
- Completeness
 - All reported warnings are actually races (avoid “false positives”)

- Disadvantages

- Run-time overhead (5-20x for best tools)
- Memory overhead for analysis state
- Reasons only about observed executions
 - sensitive to test coverage
 - (some generalization possible...)

Approaches

- Happens-Before
- Lock-set algorithm
 - Learns which shared memory locations are protected by which locks
 - Issues warning if finds no lock protects a shared memory location
- (...)

Concurrency Errors

Dynamic Data Race Detection Using
Happens-before [Lamport '78]

Lock Definition

- **Lock**: a synchronization object that is either available, or owned (by a thread)
 - Operations: **lock(mu)** and **unlock(mu)**
 - *(We are assuming no explicit initialize operation)*
 - A lock can only be unlocked by its current owner
 - The **lock()** operation is blocking if the lock is owned by another thread

The Happens-before Relation

- *happens-before* defines a partial order for events in a set of concurrent threads
 - In a single thread, *happens-before* reflects the temporal order of event occurrence
 - Between threads, **A** happens before **B** if A is an unlock access in one thread, and **B** is a lock access in a **different** thread (*assuming the threads obey the semantics of the lock , i.e., can't have two successive locks, or two successive unlocks, or a lock in one thread and an unlock in a different thread*)


The Happens-before Relation

- Let **event a** be in thread 1 and **event b** be in thread 2

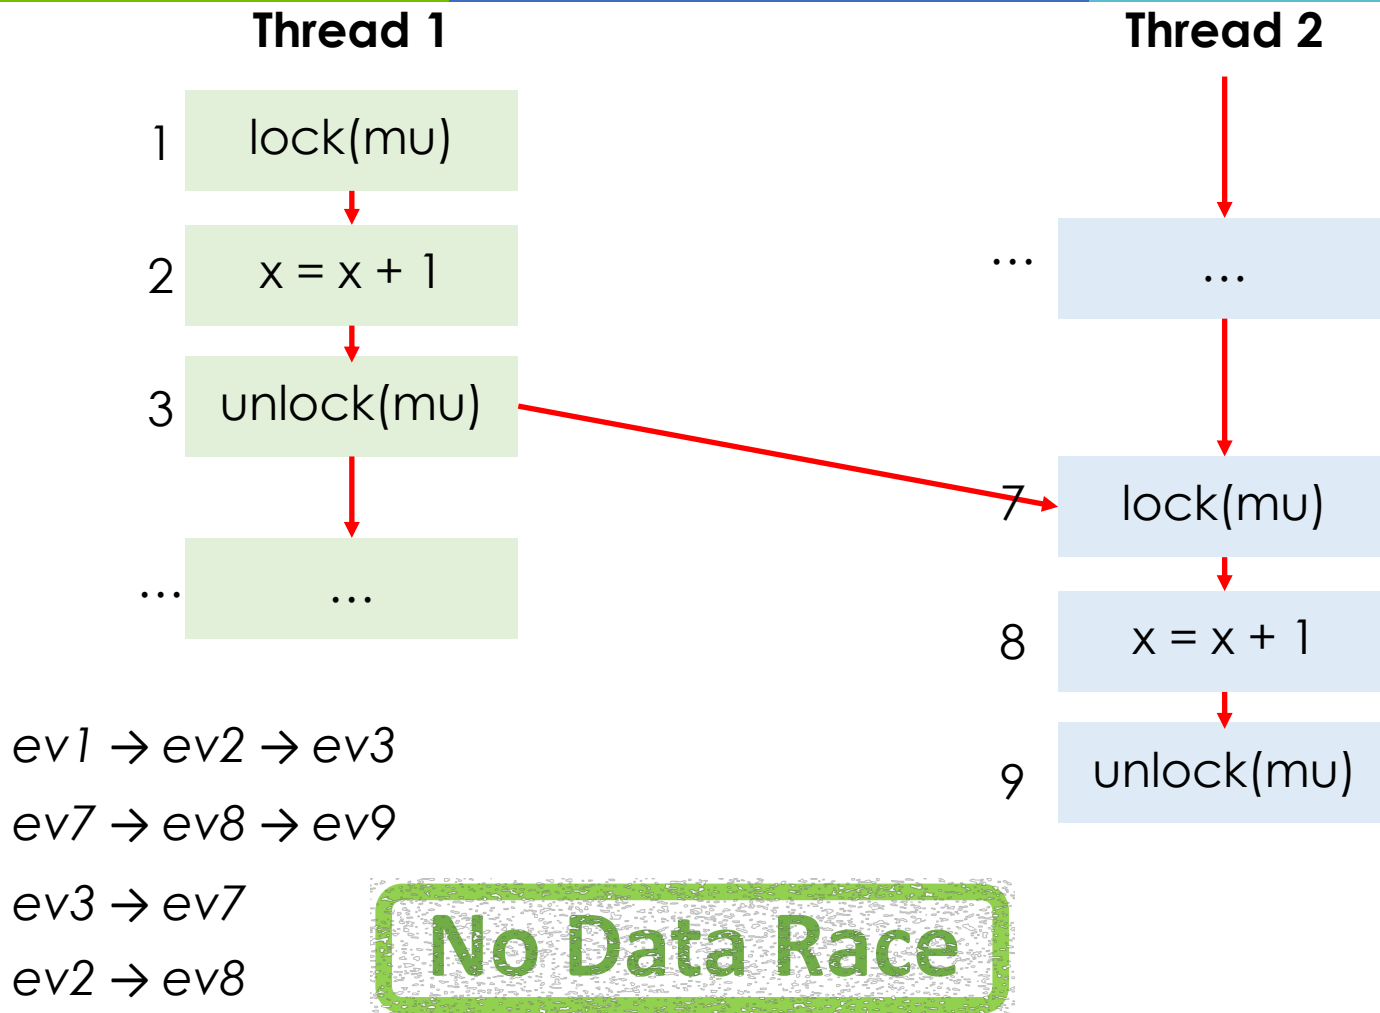
*If $a = \text{unlock}(\mu)$ and $b = \text{lock}(\mu)$ then
 $a \rightarrow b$ (a happens-before b)*

- Data races between threads are *possible* if accesses to shared variables are not ordered by *happens-before*

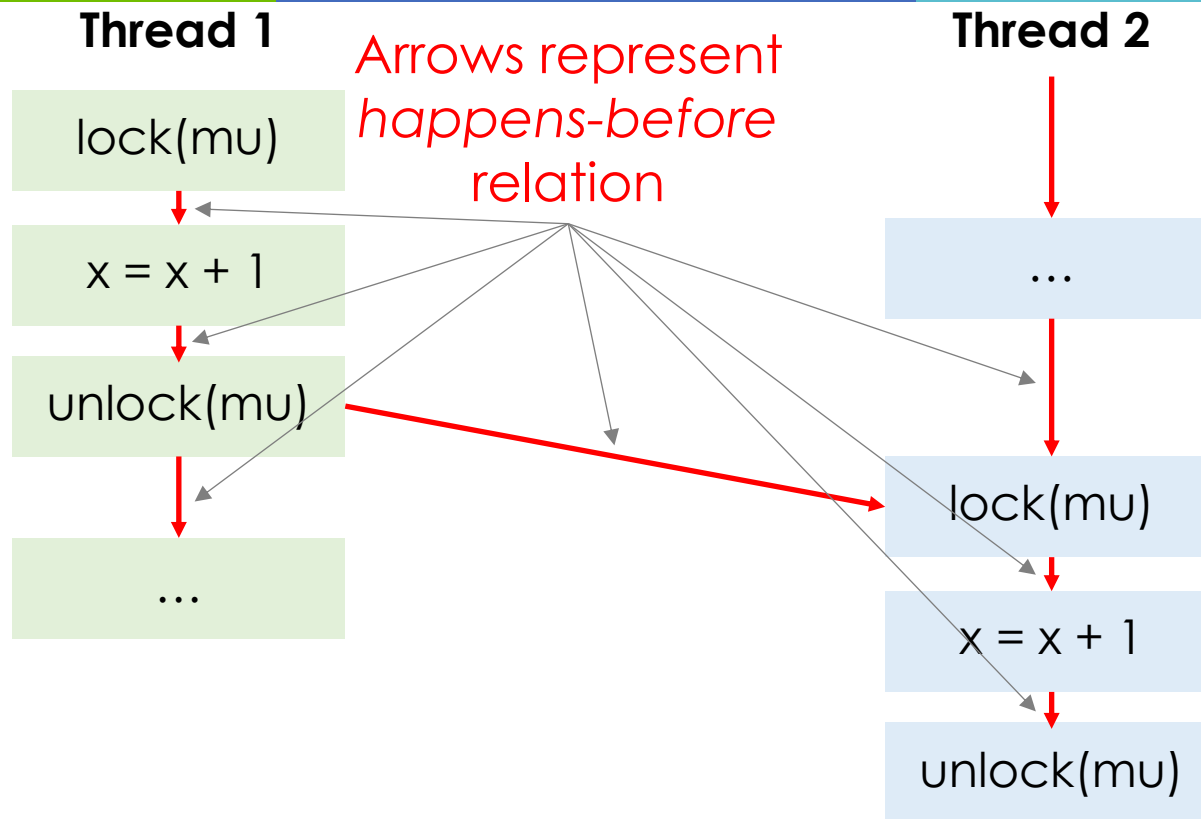
Example 1




Example 1



Example 1



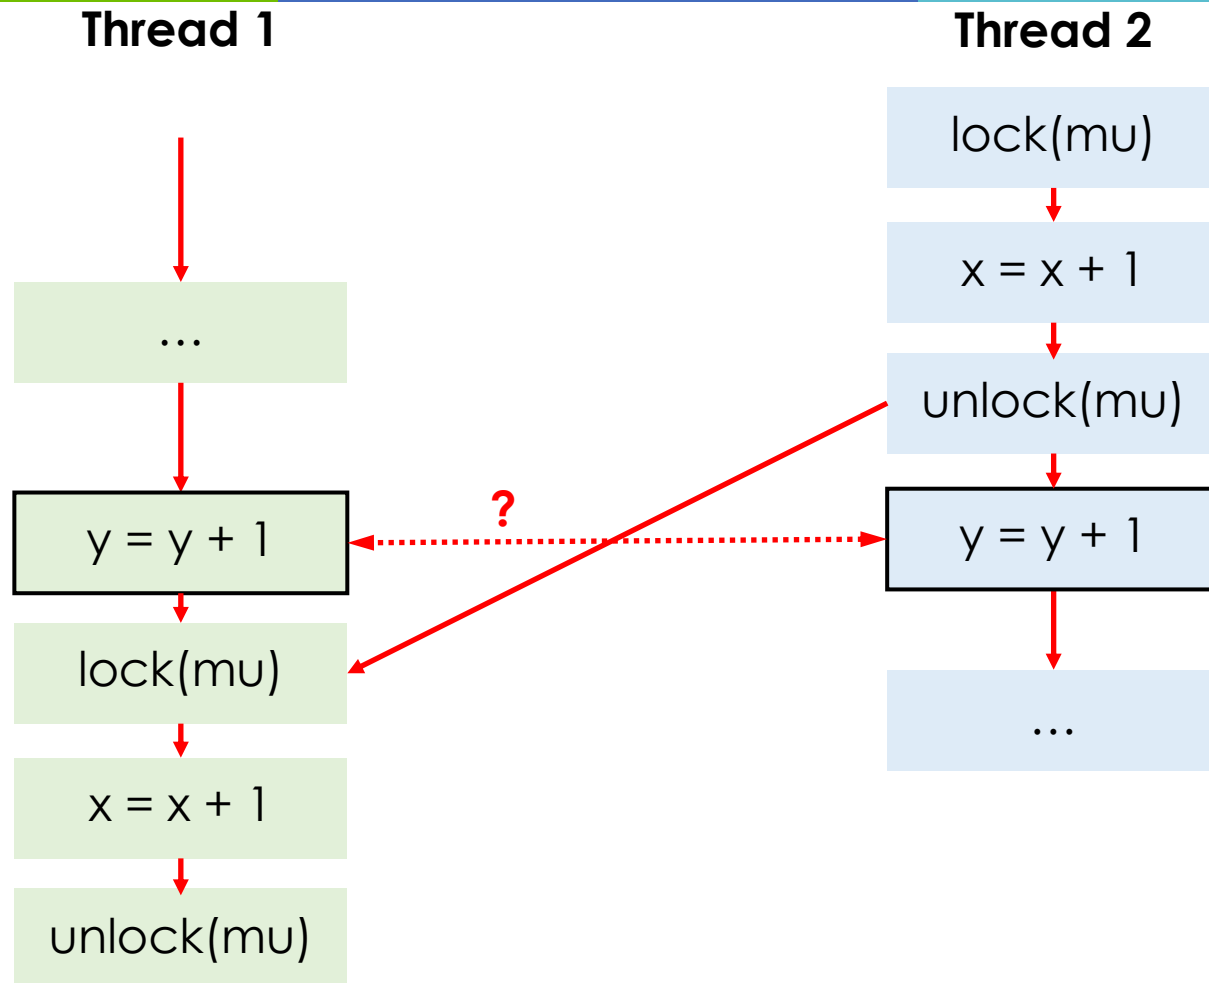
Example 2




Accesses to both 'x' and 'y' are ordered by happens-before, so no data race occurred.

But ... a different execution ordering could get different results?! Happens-before only detects data races if the incorrect order shows up in the execution trace.

Example 3



Example 3



Concurrency Errors

The Lock-Set Algorithm — Eraser [Savage et.al. '97]

Approaches

- Checks a sufficient condition for data-race freedom
- Consistent locking discipline
 - Every data structure is protected by a single lock
 - All accesses to the data structure are made while holding the lock

Thread 1

```
void Bank::Deposit(int a) {  
  
    int t = bal;  
    bal = t + a;  
  
}
```

Thread 2

```
void Bank::Withdraw(int a) {  
  
    int t = bal;  
    bal = t - a;  
  
}
```

Approaches

- Checks a sufficient condition for data-race freedom
- Consistent locking discipline
 - Every data structure is protected by a single lock
 - All accesses to the data structure are made while holding the lock

Thread 1

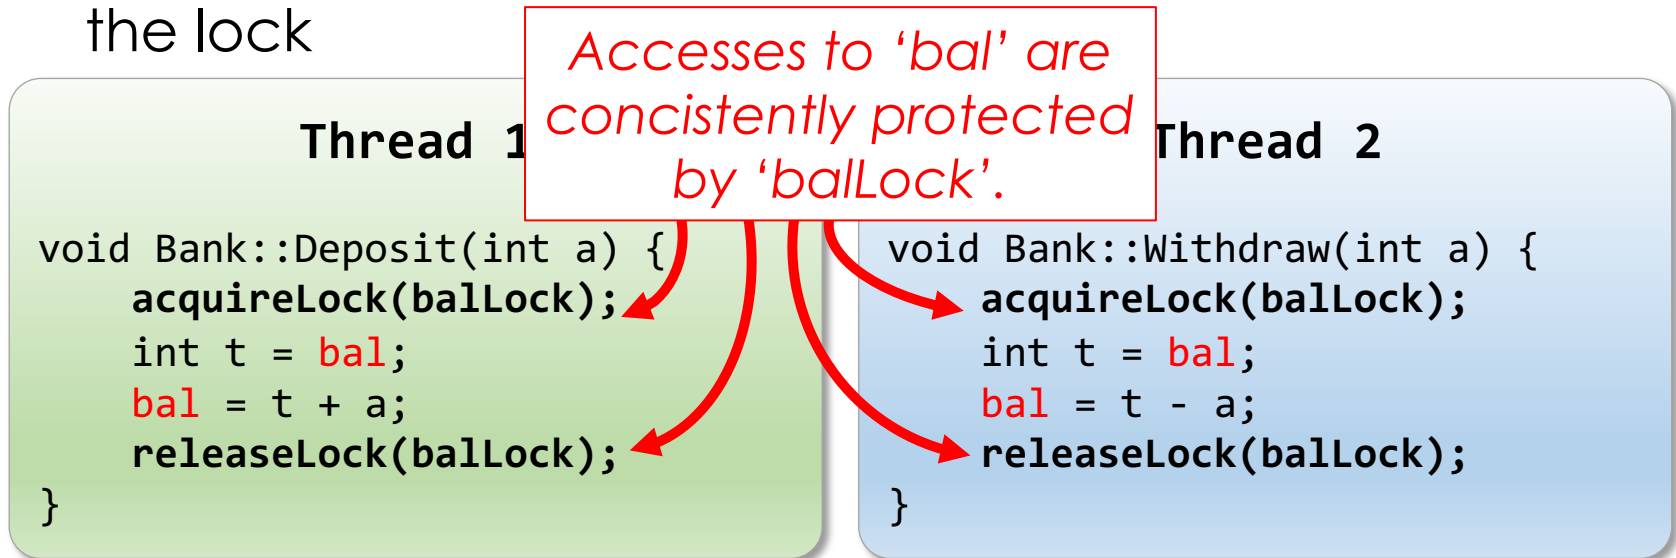
```
void Bank::Deposit(int a) {  
    acquireLock(balLock);  
    int t = bal;  
    bal = t + a;  
    releaseLock(balLock);  
}
```

Thread 2

```
void Bank::Withdraw(int a) {  
    acquireLock(balLock);  
    int t = bal;  
    bal = t - a;  
    releaseLock(balLock);  
}
```

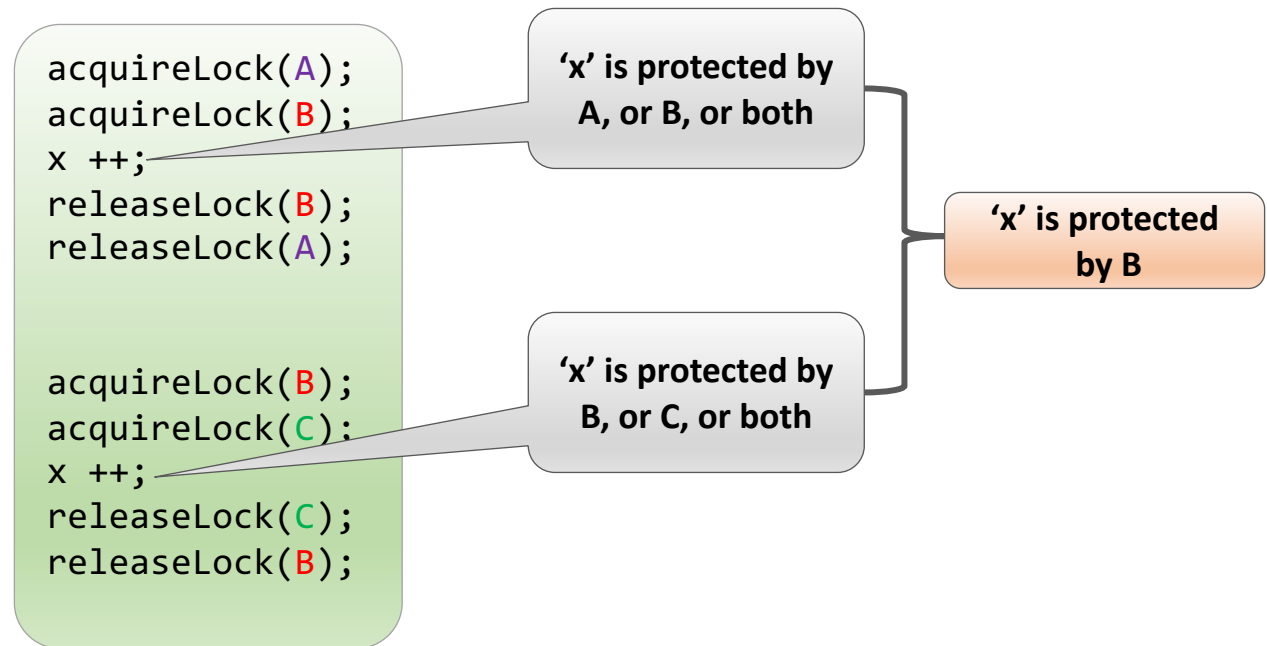
Approach

- Checks a sufficient condition for data-race freedom
- Consistent locking discipline
 - Every data structure is protected by a single lock
 - All accesses to the data structure are made while holding the lock



Approach

- How to know which locks protect each memory location?
 - Ask the programmer? Cumbersome!
 - Infer from the program code? Is it effective?



The Lock-Set Algorithm

- Two data structures:
 - $\text{LocksHeld}(t)$ = set of locks held currently by thread t
 - Initially set to Empty
 - $\text{LockSet}(x)$ = set of locks that could potentially be protecting x
 - Initially set to the universal set
- When thread ' t ' acquires lock ' l '
 - $\text{LocksHeld}(t) = \text{LocksHeld}(t) \cup \{l\}$
- When thread ' t ' releases lock ' l '
 - $\text{LocksHeld}(t) = \text{LocksHeld}(t) \setminus \{l\}$
- When thread ' t ' accesses location ' x '
 - $\text{LockSet}(x) = \text{LockSet}(x) \cap \text{LocksHeld}(t)$
- “Data race” warning if $\text{LockSet}(x)$ becomes empty


Another Example

Program Code	LocksHeld	LockSet
	{ }	{m1, m2}
lock (m1)		
lock(m2)		
v = v + 1		
unlock(m2)		
v = v + 2		
unlock(m1)		
lock(m2)		
v = v + 1		
unlock(m2)		

Another Example

Program Code	LocksHeld	LockSet
	{ }	{m1, m2}
lock (m1) → U		
lock(m2)		
v = v + 1		
unlock(m2)		
v = v + 2		
unlock(m1)		
lock(m2)		
v = v + 1		
unlock(m2)		

Another Example

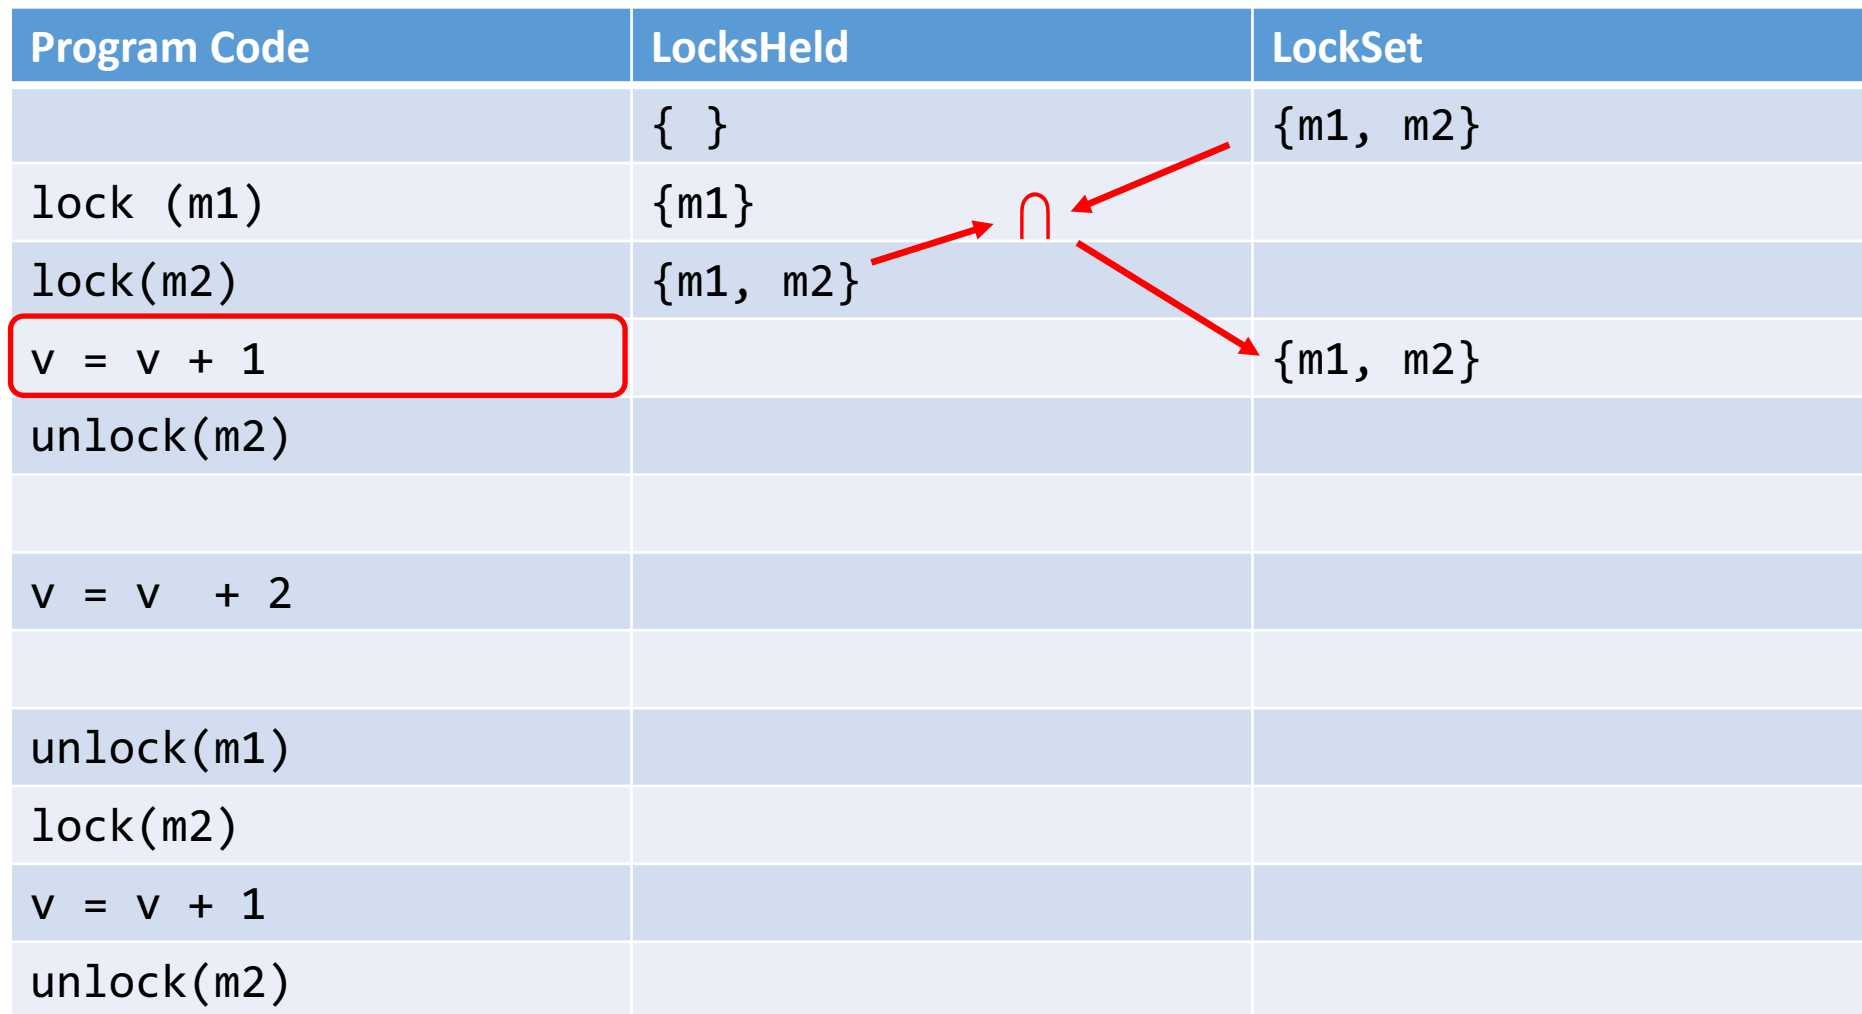
Program Code	LocksHeld	LockSet
	{ }	{m1, m2}
 lock (m1)	{m1}	
lock(m2)		
v = v + 1		
unlock(m2)		
v = v + 2		
unlock(m1)		
lock(m2)		
v = v + 1		
unlock(m2)		

Another Example

Program Code	LocksHeld	LockSet
	{ }	{m1, m2}
lock (m1)	{m1}	
lock(m2)	{m1, m2}	
v = v + 1		
unlock(m2)		
v = v + 2		
unlock(m1)		
lock(m2)		
v = v + 1		
unlock(m2)		

Another Example

Program Code	LocksHeld	LockSet
	{ }	{m1, m2}
lock (m1)	{m1}	
lock(m2)	{m1, m2}	
v = v + 1		{m1, m2}
unlock(m2)		
v = v + 2		
unlock(m1)		
lock(m2)		
v = v + 1		
unlock(m2)		



Another Example

Program Code	LocksHeld	LockSet
	{ }	{m1, m2}
lock (m1)	{m1}	
lock(m2)	{m1, m2}	
v = v + 1		{m1, m2}
unlock(m2) → \	{m1}	
v = v + 2		
unlock(m1)		
lock(m2)		
v = v + 1		
unlock(m2)		

Another Example

Program Code	LocksHeld	LockSet
	{ }	{m1, m2}
lock (m1)	{m1}	
lock(m2)	{m1, m2}	
v = v + 1		{m1, m2}
unlock(m2)	{m1}	
v = v + 2		
unlock(m1)		
lock(m2)		
v = v + 1		
unlock(m2)		

Another Example

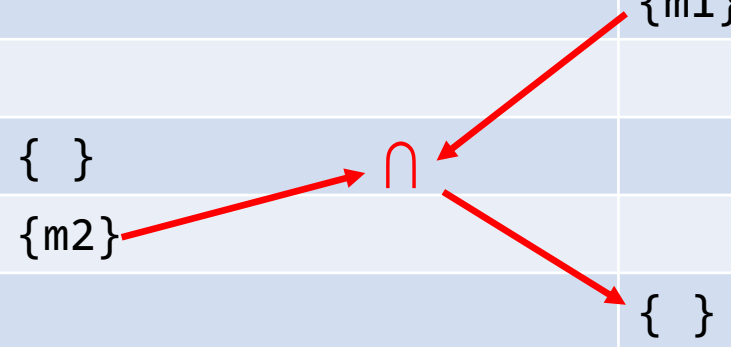
Program Code	LocksHeld	LockSet
	{ }	{m1, m2}
lock (m1)	{m1}	
lock(m2)	{m1, m2}	
v = v + 1		{m1, m2}
unlock(m2)	{m1}	
v = v + 2		{m1}
unlock(m1)	{ }	
lock(m2)		
v = v + 1		
unlock(m2)		

Another Example

Program Code	LocksHeld	LockSet
	{ }	{m1, m2}
lock (m1)	{m1}	
lock(m2)	{m1, m2}	
v = v + 1		{m1, m2}
unlock(m2)	{m1}	
v = v + 2		{m1}
unlock(m1)	{ }	
lock(m2) → U →	{m2}	
v = v + 1		
unlock(m2)		

Another Example

Program Code	LocksHeld	LockSet
	{ }	{m1, m2}
lock (m1)	{m1}	
lock(m2)	{m1, m2}	
v = v + 1		{m1, m2}
unlock(m2)	{m1}	
v = v + 2		{m1}
unlock(m1)	{ }	
lock(m2)	{m2}	
v = v + 1		{ }
unlock(m2)		



Another Example

Program Code	LocksHeld	LockSet
	{ }	{m1, m2}
lock (m1)	{m1}	
lock(m2)	{m1, m2}	
v = v + 1		{m1, m2}
unlock(m2)	{m1}	
v = v + 2		{m1}
unlock(m1)	{ }	
lock(m2)	{m2}	
v = v + 1		{ } — ALARM
unlock(m2)		

Another Example

Program Code	LocksHeld	LockSet
	{ }	{m1, m2}
lock (m1)	{m1}	
lock(m2)	{m1, m2}	
v = v + 1		{m1, m2}
unlock(m2)	{m1}	
v = v + 2		{m1}
unlock(m1)	{ }	
lock(m2)	{m2}	
v = v + 1		{ } - ALARM
unlock(m2) → \	{ }	

Algorithm Guarantees

- No warnings \Rightarrow no data races on the current execution
 - The program followed consistent locking discipline in this execution
- Warnings does not imply a data race
 - Thread-local initialization or Bad locking discipline

Algorithm Guarantees

- No warnings \Rightarrow no data races on the current execution
 - The program followed consistent locking discipline in this execution
- Warnings does not imply a data race
 - Thread-local initialization or **Bad locking discipline**

Thread 1

```
acquireLock(m1);  
acquireLock(m2);  
x = x + 1;  
releaseLock(m2);  
releaseLock(m1);
```

Thread 2

```
acquireLock(m2);  
acquireLock(m3);  
x = x + 1;  
releaseLock(m3);  
releaseLock(m2);
```

Thread 3

```
acquireLock(m1);  
acquireLock(m3);  
x = x + 1;  
releaseLock(m3);  
releaseLock(m1);
```



Acknowledgments

- Some parts of this presentation was based in publicly available slides and PDFs
 - www.cs.cornell.edu/courses/cs4410/2011su/slides/lecture10.pdf
 - www.microsoft.com/en-us/research/people/madanm/
 - williamstallings.com/OperatingSystems/
 - codex.cs.yale.edu/avi/os-book/OS9/slide-dir/

The END
