

# Parallel Algorithms Patterns: Map

---

COMPUTAÇÃO DE ALTO DESEMPENHO 2018/2019

HERVÉ PAULINO

SLIDES ADAPTED FROM

“STRUCTURED PARALLEL PROGRAMMING - CIS 410/510, UNIVERSITY OF OREGON”

REZAUL A. CHOWDHURY - CSE 613: PARALLEL PROGRAMMING



# Bibliography

---

**Sections 3.5 to 3.7 and Chapter 4 of Structured Parallel Programming, Michael McCool, Arch D. Robison and James Reinders. Morgan Kaufmann (2012)**

# Parallel Control Patterns

---

Serial control patterns

- Loop, recursion, and more

Parallel control patterns extend serial control patterns

Each parallel control pattern is related to at least one serial control pattern, but relaxes assumptions of serial control patterns

Parallel control patterns: fork-join, map, stencil, reduction, scan, recurrence

# Parallel Control Patterns: Fork-Join

---

**Fork-join:** allows control flow to fork into multiple parallel flows, then rejoin later

CUDA implements this when executing a kernel

A “join” is different than a “barrier

- After join only one thread continues
- Barrier – all threads continue


# Parallel Control Patterns: Map

---

**Map:** performs a function over every element of a collection

Map replicates a serial iteration pattern where each iteration is independent of the others, the number of iterations is known in advance, and computation only depends on the iteration count and data from the input collection

The replicated function is referred to as an “elemental function”



# Mapping

---

“Do the same thing many times”

```
foreach i in foo:  
    do something
```


Well-known higher order function in languages like ML, Haskell, Scala

map:  $\forall ab.(a \rightarrow b)List\langle a \rangle \rightarrow List\langle b \rangle$


applies a function each element in a list and returns a list of results

# Example Maps

Add 1 to every item in an array



Double every item in an array



**Key Point:** An operation is a map if it can be applied to each element without knowledge of neighbors.

# Key Idea

---

Map is a “foreach loop” where each iteration is **independent**

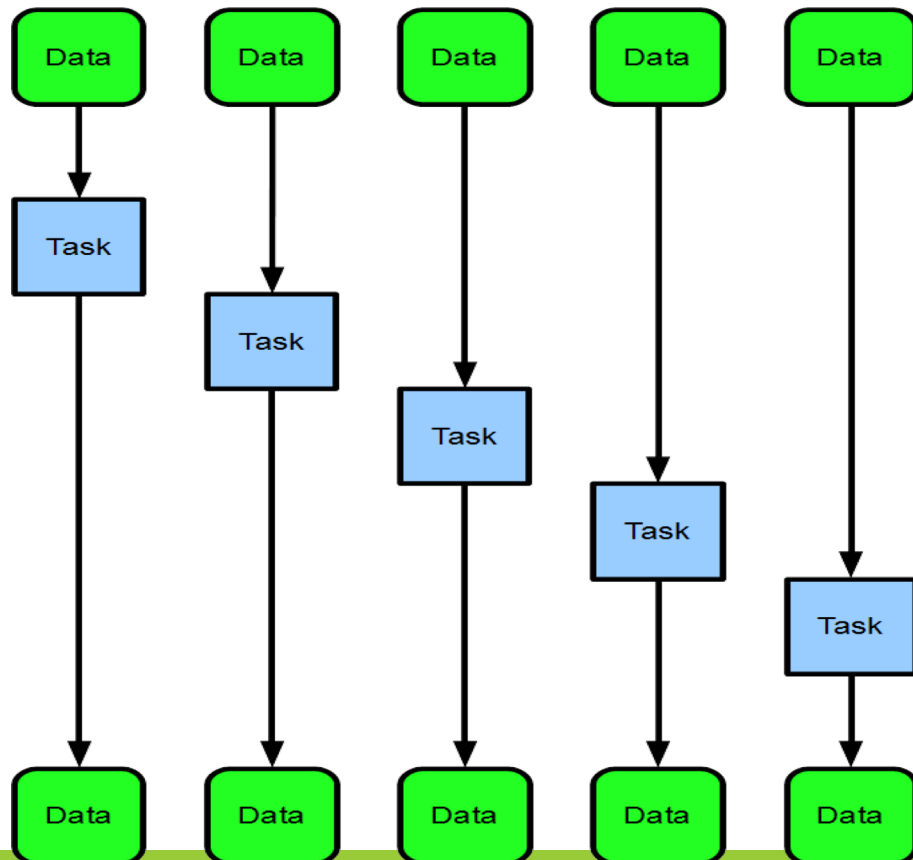
## Embarrassingly Parallel

Independence is a big win. We can run map completely in parallel.  
Significant speedups! More precisely:  $T(\infty)$  is  $O(1)$  plus implementation overhead




# Comparing Maps

SERIAL MAP



PARALLEL MAP



## Speedup

The space here is speedup. With the parallel map, our program finished execution early, while the serial map is still running.

# Independence

---

The key to (embarrassing) parallelism is independence

**Warning: No shared state!**

Map function should be “pure” (or “pure-ish”) and should not modify shared states

Modifying shared state breaks perfect independence

Results of accidentally violating independence:

- non-determinism
- data-races
- undefined behavior
- segfaults

# Unary Map: Map with 1 Input, 1 Output

	0	1	2	3	4	5	6	7	8	9	10	11
x	3	7	0	1	4	0	0	4	5	3	1	0
	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓
result	6	14	0	2	8	0	0	8	10	6	2	0

**CUDA  
implementation**

```
__global__ void mult2(int* c, int n) {  
    int index = threadIdx.x + blockIdx.x * blockDim.x;  
    if (index < n)  
        c[index] = c[index] * 2;  
}
```

# N-Array Map: Map with N Inputs, 1 Output

	0	1	2	3	4	5	6	7	8	9	10	11
x	3	7	0	1	4	0	0	4	5	3	1	0
y	2	4	2	1	8	3	9	5	5	1	2	1
	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓
result	5	11	2	2	12	3	9	9	10	4	3	1

**CUDA  
implementation**


```
__global__ void add(int* a, int* b, int* c, int n) {  
    int index = threadIdx.x + blockIdx.x * blockDim.x;  
    if (index < n)  
        c[index] = a[index] + b[index];  
}
```

# Optimization Sequences of Maps

Often several map operations occur in sequence

- Vector math consists of many small operations such as additions and multiplications applied as maps

A naïve implementation may write each intermediate result to memory, wasting memory BW and likely overwhelming the cache




# Optimization Code Fusion

Can sometimes “fuse”  
together the operations to  
perform them at once

Adds arithmetic intensity,  
reduces memory/cache  
usage

Ideally, operations can be  
performed using registers  
alone



# Optimization

## Code Fusion

---

$C = A + B$

$C = C - 1$

```
__global__ void add(int* a, int* b,
                    int* c, int n) {
    int index = threadIdx.x + blockIdx.x *
                blockDim.x;
    if (index < n)
        c[index] = a[index] + b[index];
}
```

```
__global__ void sub1(int* c, int n) {
    int index = threadIdx.x + blockIdx.x *
                blockDim.x;
    if (index < n)
        c[index] = c[index] - 1;
}
```

$C = A + B - 1$

```
__global__ void addsub1(int* a, int* b,
                        int* c, int n) {
    int index = threadIdx.x + blockIdx.x *
                blockDim.x;
    if (index < n) {
        c[index] = a[index] + b[index];
        c[index] = c[index] - 1;
    }
}
```

# Optimization

## In GPU computing is not always possible

---

$C = A + B$


FORALL  $I > 1$   $C[I] = C[I-1]$

```
__global__ void add(int* a, int* b,
                    int* c, int n) {
    int index = threadIdx.x + blockIdx.x *
                blockDim.x;
    if (index < n)
        c[index] = a[index] + b[index];
}

__global__ void sub1(int* c, int n) {
    int index = threadIdx.x + blockIdx.x *
                blockDim.x;
    if (index > 1 && index < n)
        c[index] = c[index-1];
}
```

$C = A + B - 1$

```
__global__ void addsub1(int* a, int* b,
                        int* c, int n) {
    int index = threadIdx.x + blockIdx.x *
                blockDim.x;
    if (index < n) {
        c[index] = a[index] + b[index];
        if (index > 1)
            c[index] = c[index-1];
    }
}
```



In the general case, it is not possible to know if  $c[index-1]$  has already been computed



# Related Patterns

---

Three patterns related to map are discussed here:

- Stencil
- Workpile
- Divide-and-Conquer

More detail presented in a later lecture

SAXPY:  $y \leftarrow ax + y$

---

	0	1	2	3	4	5	6	7	8	9	10	11
a * x	4	4	4	4	4	4	4	4	4	4	4	4
	2	4	2	1	8	3	9	5	5	1	2	1
+												
y	3	7	0	1	4	0	0	4	5	3	1	0
	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓
y	11	23	8	5	36	12	36	49	50	7	9	4

SAXPY:  $y \leftarrow ax + y$

---

	0	1	2	3	4	5	6	7	8	9	10	11
a	4	4	4	4	4	4	4	4	4	4	4	4
*	2	4	2	1	8	3	9	5	5	1	2	1
x												
+												
y	3	7	0	1	4	0	0	4	5	3	1	0
	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓
y	11	23	8	5	36	12	36	49	50	7	9	4

SAXPY:  $y \leftarrow ax + y$

	0	1	2	3	4	5	6	7	8	9	10	11
a	4	4	4	4	4	4	4	4	4	4	4	4
*												
x	2	4	2	1	8	3	9	5	5	1	2	1
+												
y	3	7	0	1	4	0	0	4	5	3	1	0
	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓	↓
y	11	23	8	5	36	12	36	49	50	7	9	4