

# COMPUTAÇÃO DE ALTO DESEMPENHO

2018/2019

## Test 1

1. [2 values] Is a GPU a shared memory MIMD (Multiple Instruction Multiple Data) architecture? Justify your answer.
2. Consider the following sequential C function:

```
1. int* some_func(int* a, int* b, int x, int y, int size) {
2.     int* result = malloc (size * sizeof(int));
3.     int z = sqrt (x * y);
4.     for (int i = 0; i < size; i++) result[i] = a[i] + b[i]*z;
5. }
```

- a. [1 value] Is there a potential hotspot in this algorithm? Where and why?
  - b. [2 values] Parallelize the identified hotspot by providing a CUDA implementation (please provide only the CUDA kernel not the host code).
  - c. [2 values] Provide a rewriting of function some\_func that makes use of the kernel implemented in the previous question. You must include all memory allocation and transfers operations.
3. Consider a sequential implementation of an algorithm A with execution time of 100 time-units. Consider also a parallel implementation (PA) of that same algorithm with the following execution times:

Number of threads	Execution time (in time-units)
1	120
2	60
4	40
8	24
16	20
32	12

- a. [0.5 values] What is the speedup yield by PA when executed by 8 threads?
  - b. [1.5 values] Roughly plot PA's efficiency curve. What is your conclusion regarding the PA's efficiency? Justify.
  - c. [1 value] Is PA cost-optimal? Why?
4. Consider that you have the task of parallelizing, with CUDA, the search of an element in a huge array.

```
int contains(int* array, size_t size, int elementToSearch);
```

- a. [1.5 values] Which would be your implementation strategy? In other words, how many kernels would you use, what would be their behavior (map, reduce, scan), and which operation/function would they apply?
  - b. [1.5 values] Does any of your kernels benefit from the use of shared memory? Justify.
  - c. [1.5 values] Consider now that you have a framework that allows you to execute your kernel in multiple GPUs. Which modifications do you need to apply to your strategy? Justify.
  - d. [1.5 values] What is the lower and upper bound of the expected speedup of a N-GPU execution of your algorithm when compared against a 1-GPU execution? Consider that all N GPUs are of the same model.
5. You now have the task of parallelizing the following algorithm:

```
let a, b and c be arrays in
array e = a + b
e = (a*e - c*e)4;
for i in [1, (e.size()-2)/2] do f[i] = (e[i-1] + e[i] + e[i+1])/3
for i in [(e.size()-2)/2, e.size()-2] do f[i] = (2*e[i-1] + e[i] + 2*e[i+1])/5
```

- a. [1.5 values] In the context of the BSP abstract computer model, how many super steps are needed to parallelize this algorithm? What would each of these super steps do?

- b. [1 value] Can the solution you gave to the previous question be implemented on CUDA (or other GPGPU framework)? Justify your answer.
- c. [1.5 values] Does your CUDA solution uses kernels that have thread divergence (for instance **ifs**)? If not, why? If so, how harmful will this divergence be in the performance of your kernel(s)?

## Appendix

### **shared a**

allocates variable a in shared (thread block local) memory

### **syncthreads()**

synchronizes all threads in the current thread block

### **kernel\_function<<<NUMBER\_BLOCKS, BLOCK\_SIZE, LOCAL\_MEMORY, STREAM>>>(arguments)**

signature of kernel launching in CUDA

**cudaError\_t cudaMalloc (void\*\* dev\_ptr, size\_t count)**

**cudaError\_t cudaFree (void\* dev\_ptr)**

**cudaError\_t cudaMemcpy (void\* dst, const void\* src, size\_t count, cudaMemcpyKind kind)**

**cudaError\_t cudaMemcpyAsync (void\* dst, const void\* src,  
size\_t count, cudaMemcpyKind kind, cudaStream\_t stream = 0 )**

### **cudaMemcpyKind:**

<i>cudaMemcpyHostToHost</i>	Host -> Host.
<i>cudaMemcpyHostToDevice</i>	Host -> Device.
<i>cudaMemcpyDeviceToHost</i>	Device -> Host.
<i>cudaMemcpyDeviceToDevice</i>	Device -> Device