

Confiabilidade de Sistemas Distribuídos Dependable Distributed Systems

DI-FCT-UNL, Nuno Preguiça


Lecture: Introduction
Dependable Distributed Systems
Principles, Concepts, Properties and Technology

2018/2019, 2nd SEM

MIEI
Mestrado Integrado em Engenharia Informática

Distributed Systems


Distributed Systems and Components



- A **system** is an entity that interacts with other entities, i.e., other systems, including hardware, software, humans, and the physical world with its natural phenomena
- These other systems are the **environment** of the given system
- The **system boundary (perimeter)** is the common frontier between the system and its environment
- A system may consists of one or more **components**, such as nodes or processes

Distributed Components:

System of distributed components Environment,
System Boundary (Perimeter)




DS Communication
Environment

- A **system** is an entity that interacts with other entities, i.e., other systems, including hardware, software, humans, and the physical world with its natural phenomena
- These other systems are the **environment** of the given system
- The **system boundary (perimeter)** is the common frontier between the system and its environment
- A system may consists of one or more **components**, such as nodes or processes


Distributed System and Granularity

Vision as a System of Systems:
Scale and Complexity




Distributed System and Granularity

Vision as a System of Systems:
Scale and Complexity



Distributed systems and the vision of systems of systems



- DS: a **set of distributed nodes, executing autonomous computations (processes) that need to interact to coordinate their actions**
- **Baseline for discussion:**
 - Synchronous or asynchronous by nature ?
 - Can we expect independent failures ?
 - How to identify (define and circumvent) the trust computing base and its foundations ?
 - The same for “dependability base”

Failure Model Definition

- A definition expressing the typology of faults in the system model design assumptions
- Need very-well defined typology of faults considered in the system design model
 - Usually, looked as accidental failures

Threat or Adversary Model Definition

- A definition expressing the typology of threats or attacks (attack types as the concretization of defined threats) in the system model design assumptions
- Need very-well defined typology of attacks (vectors, adversary conditions and hypothesis) as the concretization of the threats potential considered in the system design model
 - Initial discussion: is it very different from the failure model?
 - See this as the injection of malicious failures by attackers

Dependability

Dependable Distributed Systems

What is “Dependability” ?

- Context (**Dependable Distributed System**):
 - We have components providing services to clients (or other components).
 - To provide services, components may require services from other components
 - We say that the component **may depend** on some other component

We say that a **component C depends on C*** if the **correctness of C's behaviour depends on the correctness of C*'s behaviour**.

Then C* is a Dependable System

C* is used by C as a DEPENDABLE COMPUTING BASE


DEPENDABLE vs. TRUSTWORTHY

What are “these” components about ? : Different levels in a Distributed System Stack

A definition for Dependability

- Ability of a Distributed System to provide correct services (computations) to its users (or to other services depending from it), **despite** various **threats**, including undetected software defects or failures, or hardware failures, whether they are accidental or caused by malicious attacks
 - Note: **this includes intrusions causing malicious failures, induction of incorrect behaviour (SW or HW)**

DS Stack: where is the TCB ?



Computation, Processing
(Support Levels)

Application
Static Libraries

Runtime Libraries
Middleware Services

OS Services and Resources

Data Repository Services
(Local NVS)

Virtualized Environment
(Virtualization Layer)


Communication Channels

Dependability Property

- Conjugation of properties, mainly aggregated in two dimensions:
 - Fault Tolerance Properties
 - Security Properties
 - Communication attacks + intrusions
- Under a Well-Defined System Model and Software Architecture
 - A Well-Defined Fault Model
 - A Well-Defined Adversary Model

Dependability Dimensions

- As two “Faces” of the same coin
- Challenge/Trend: addressable “all-in-one” solution ?



Fault-Tolerance Dimension

Dependability Properties: Fault-Tolerance Dimension

- **Base concepts**
 - **Availability**
 - Readiness for usage
 - **Reliability**
 - Continuity of service delivery
 - **Safety**
 - Very low probability of catastrophes/disasters
 - Accidental Failures (in typical approaches)
 - **Maintainability**
 - How easily can a failed system be repaired or recovered

Reliability vs. Availability (2)

- **Availability: $A(t)$:**
 - **Average fraction of time** that a component has been up and running in the interval $[0, t]$
- Long-Term Availability (or Always Available):
 - $A(\infty)$

Relating:

- **$A = \text{MTTF} / \text{MTBF}$**
 $\Rightarrow A = \text{MTTF} / (\text{MTTF} + \text{MTTR})$

Reliability vs. Availability (1)

- **Reliability $R(t)$:**
 - **probability** that a component has been **up and running** (**correctly and continuously**) in the time interval $[0, t]$

Conventional Metrics:

- **MTTF:** Mean Time To Failure:
 - Average time until a component fails
- **MTTR:** Average time it takes to repair (recover) a failed component.
- **MTBF:** Mean Time Between Failures
 - $MTTF + MTTR$

Reliability vs. Availability (3)

- Important Observation:
 - Reliability and availability make sense:
 - **If we have an accurate notion of what a failure actually is**
 - **This Requires a “very well-defined” Failure Model, related to the System Model and Design Assumptions**
 - => Must address Reliability vs. Availability vs. Efficiency or performance Tradeoffs BY DESIGN !

Safety

- **Safety.** Level of tolerance against catastrophic failures or accidents
 - Faults mainly as “accidental” failures (in the more typical approach of FAULT TOLERANT DISTRIBUTED SYSTEMS)
 - $RISK = FAULTS \times \text{Operation Exposition}$

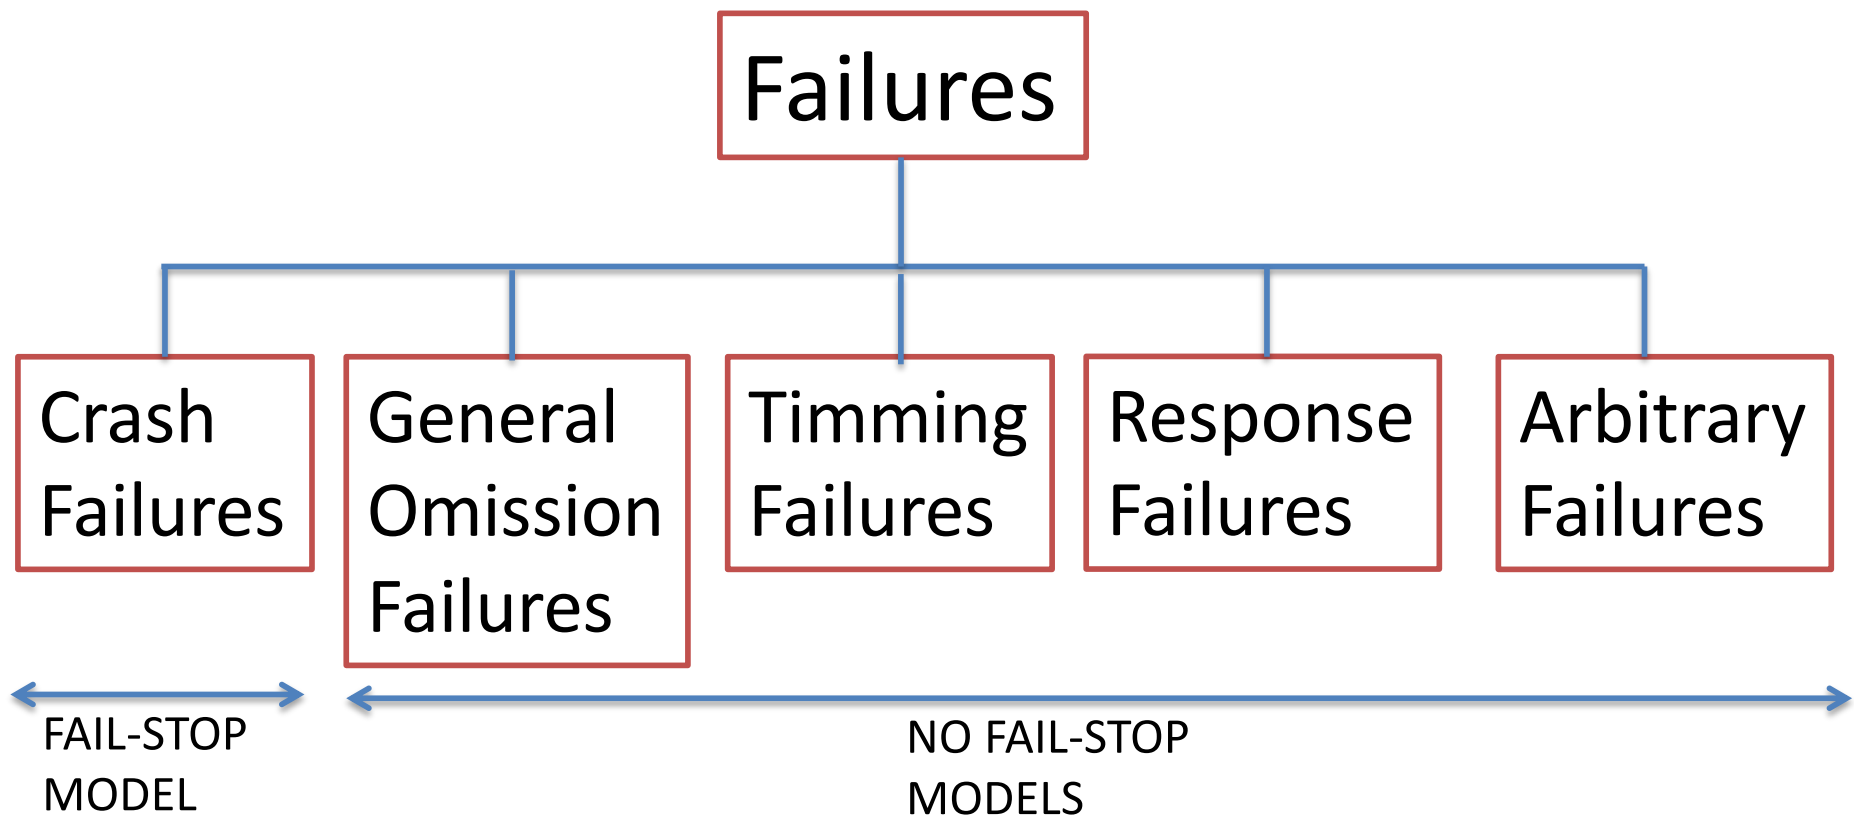
Need to establish concepts and metrics for FAULTS and Operation Exposition / Related to previous metrics

Maintainability

- Maintainability
 - How easy to recover a failed system (manual intervention or automatic mechanisms)
 - Some times only the manual intervention is considered
 - But fault-tolerance (for availability or operation continuity) requires automatic recovery mechanisms
 - Ex., Replication (w/ consistency)
 - Ex., Diversity
 - Ex., Virtualization
 - Possible “Onion” Supported Approaches

Failure Models and Types of Failures

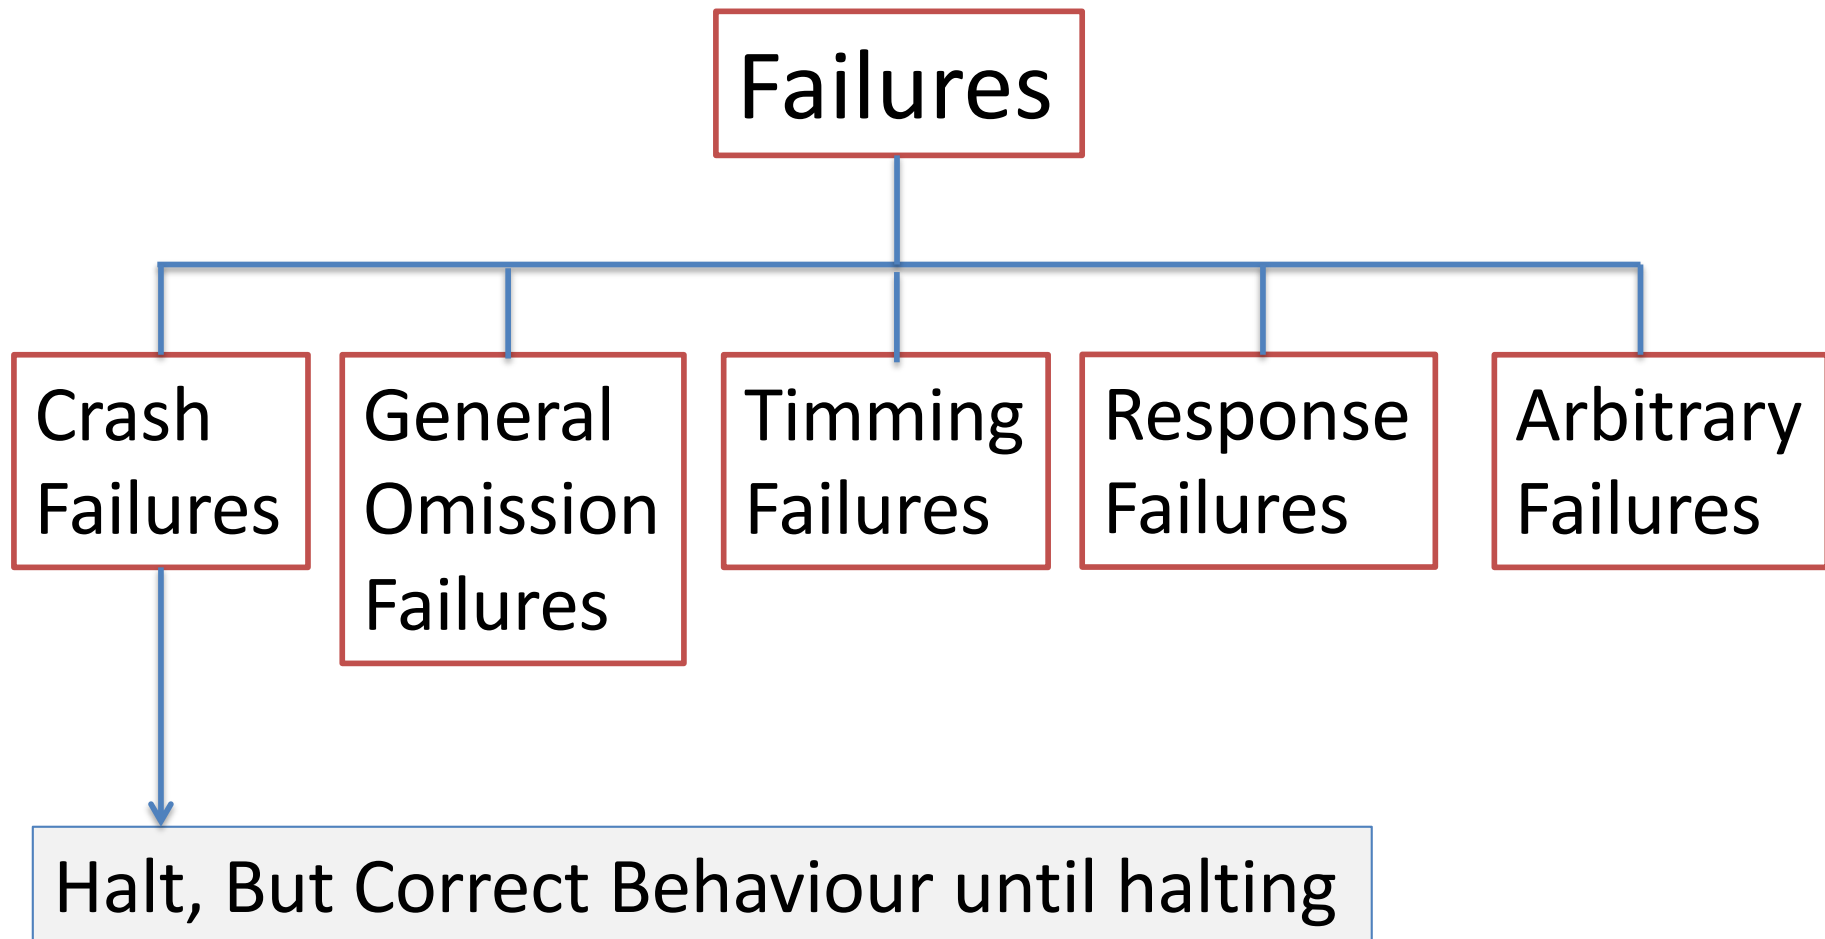
Generic Characterization



Typology (as ref. in Andrew Tanenbaum, Maarten Van Steen,
Distributed Systems - Principles and Paradigms,
Chap. 7 – Fault Tolerance (2nd Edition,

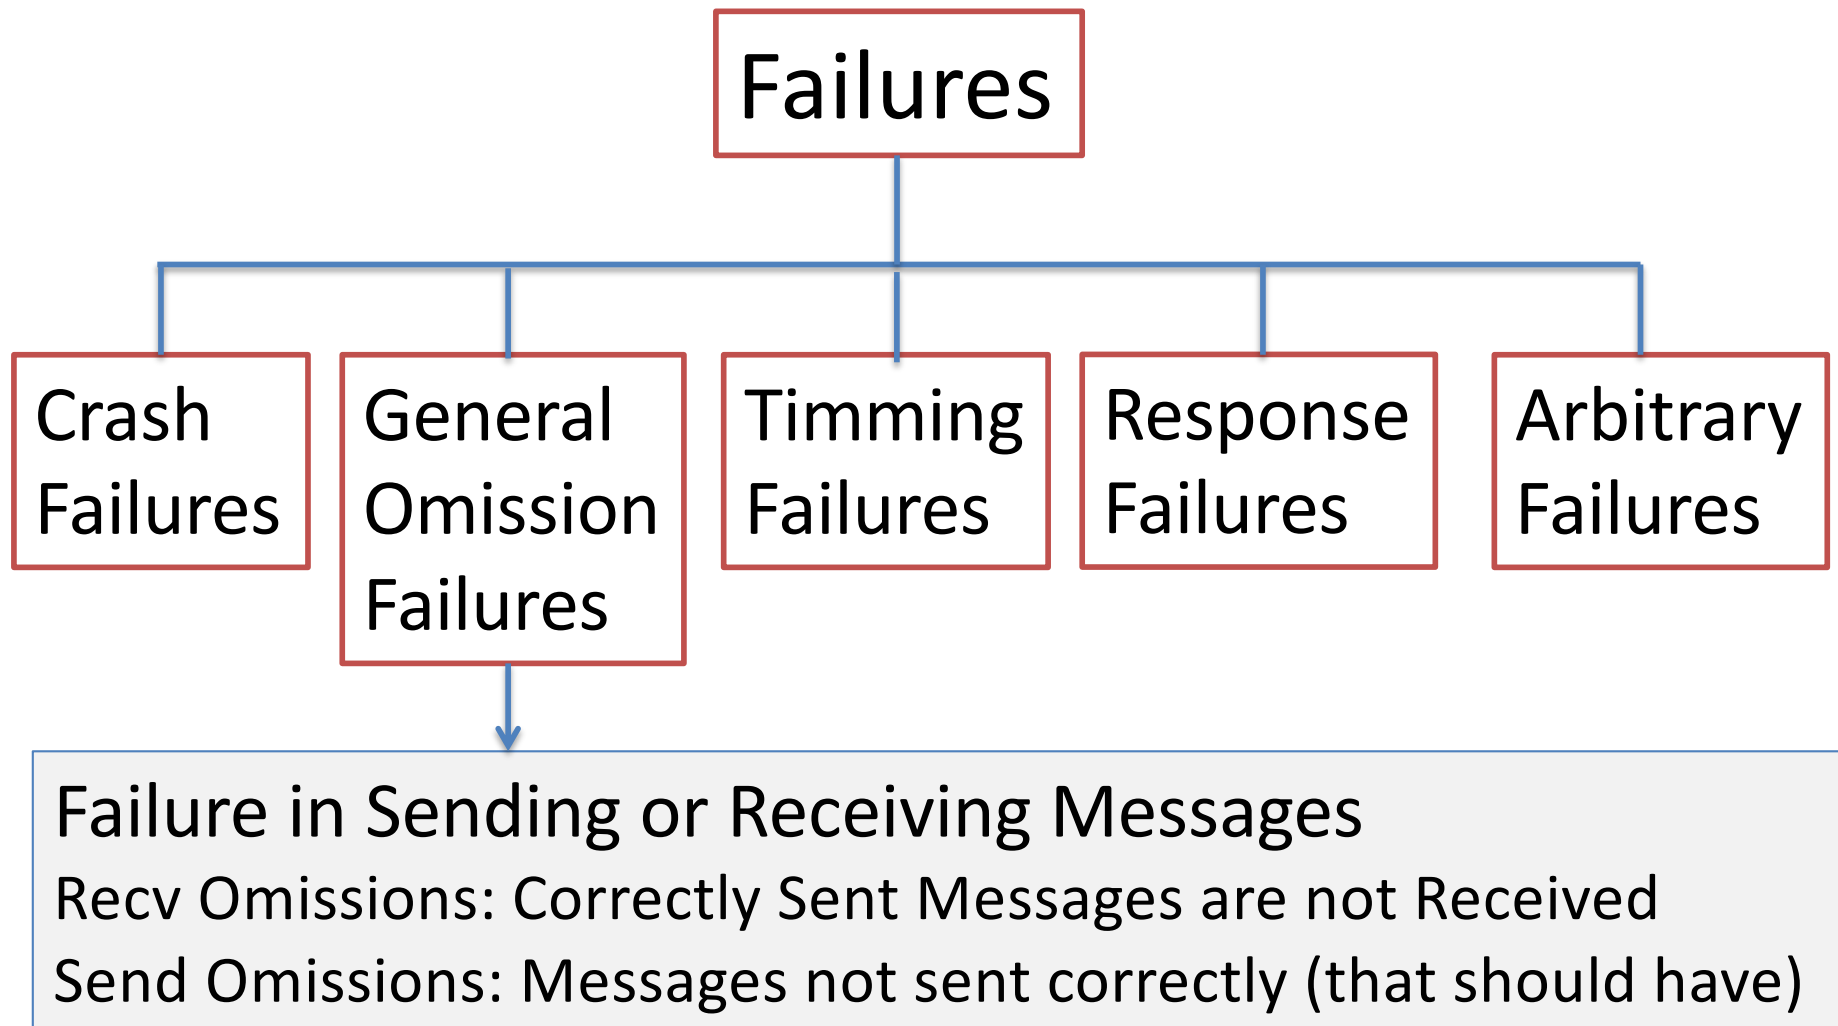
Failure Models

> Crash Failures



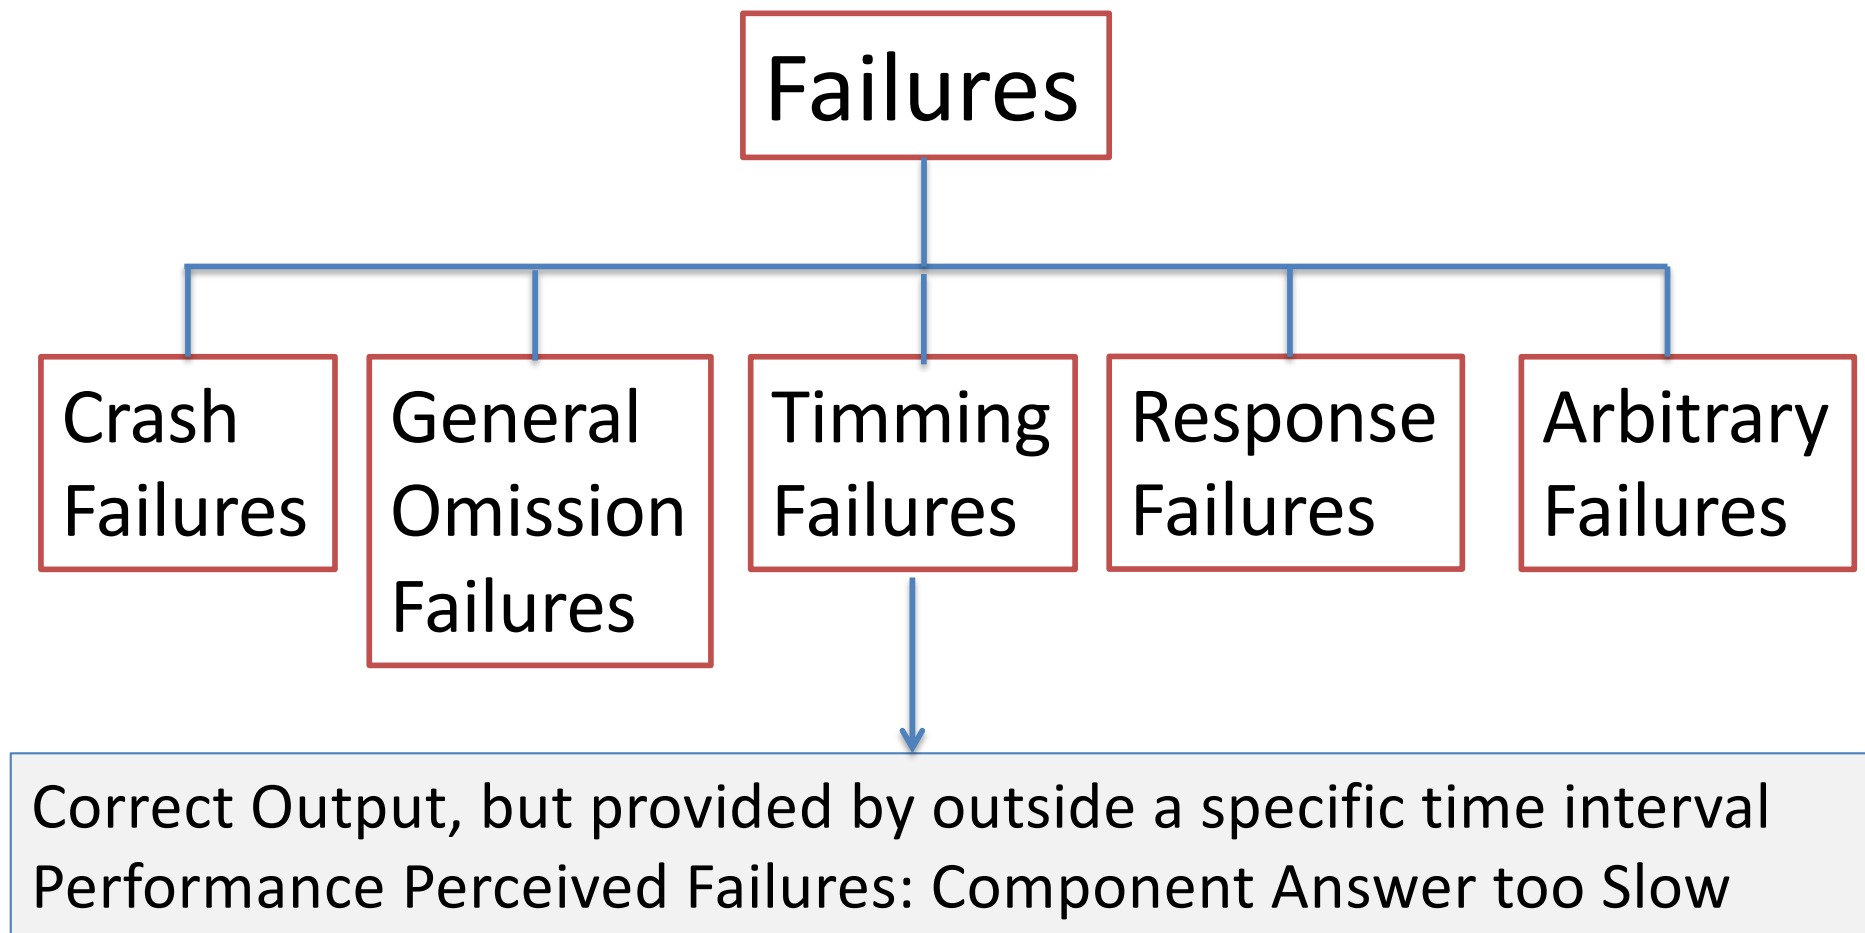
Failure Models

> Omission Failures



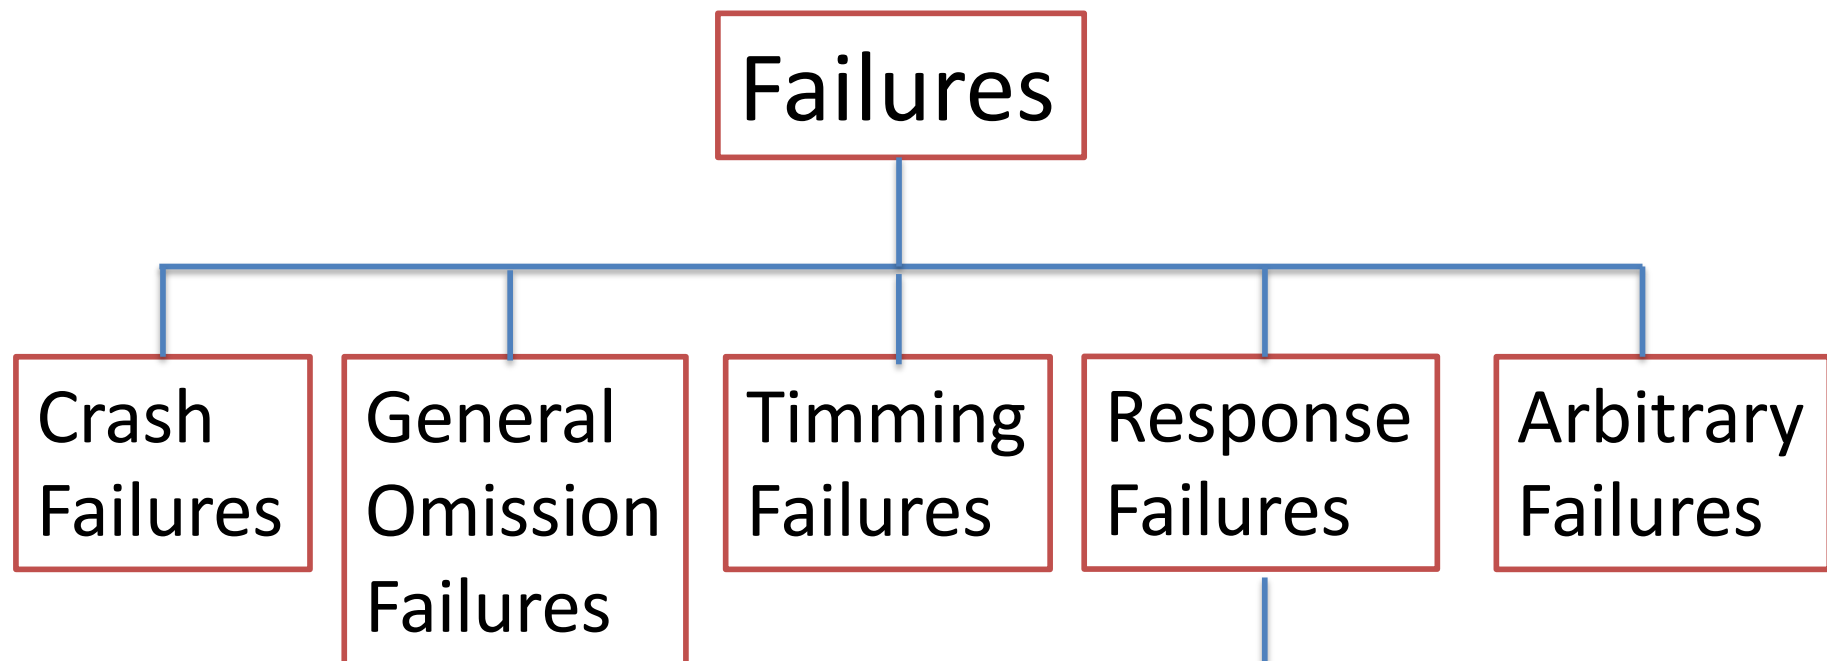
Failure Models

> Timing Failures



Failure Models

> Response Failures

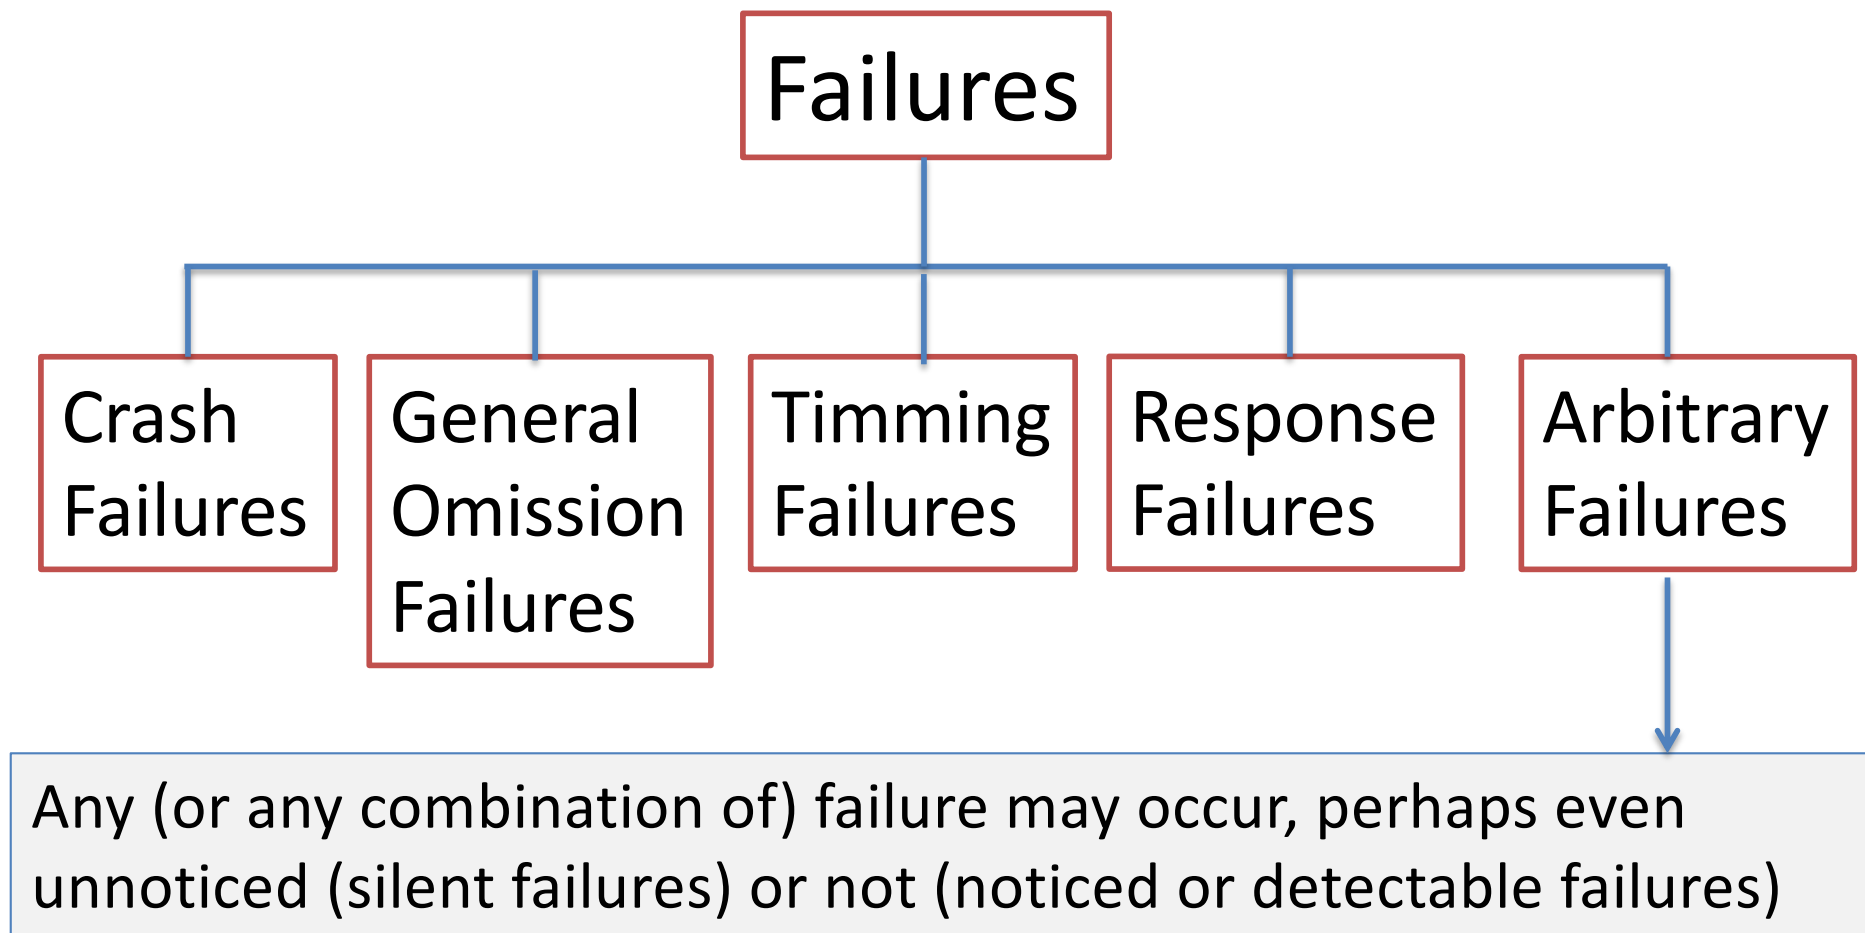


Incorrect output, but cannot be accounted to another component

- > **Value Failures**: wrong output values
- > **State-Transition Failures**: deviation from correct flow of control
(Note: this failure may initially not even be observable)

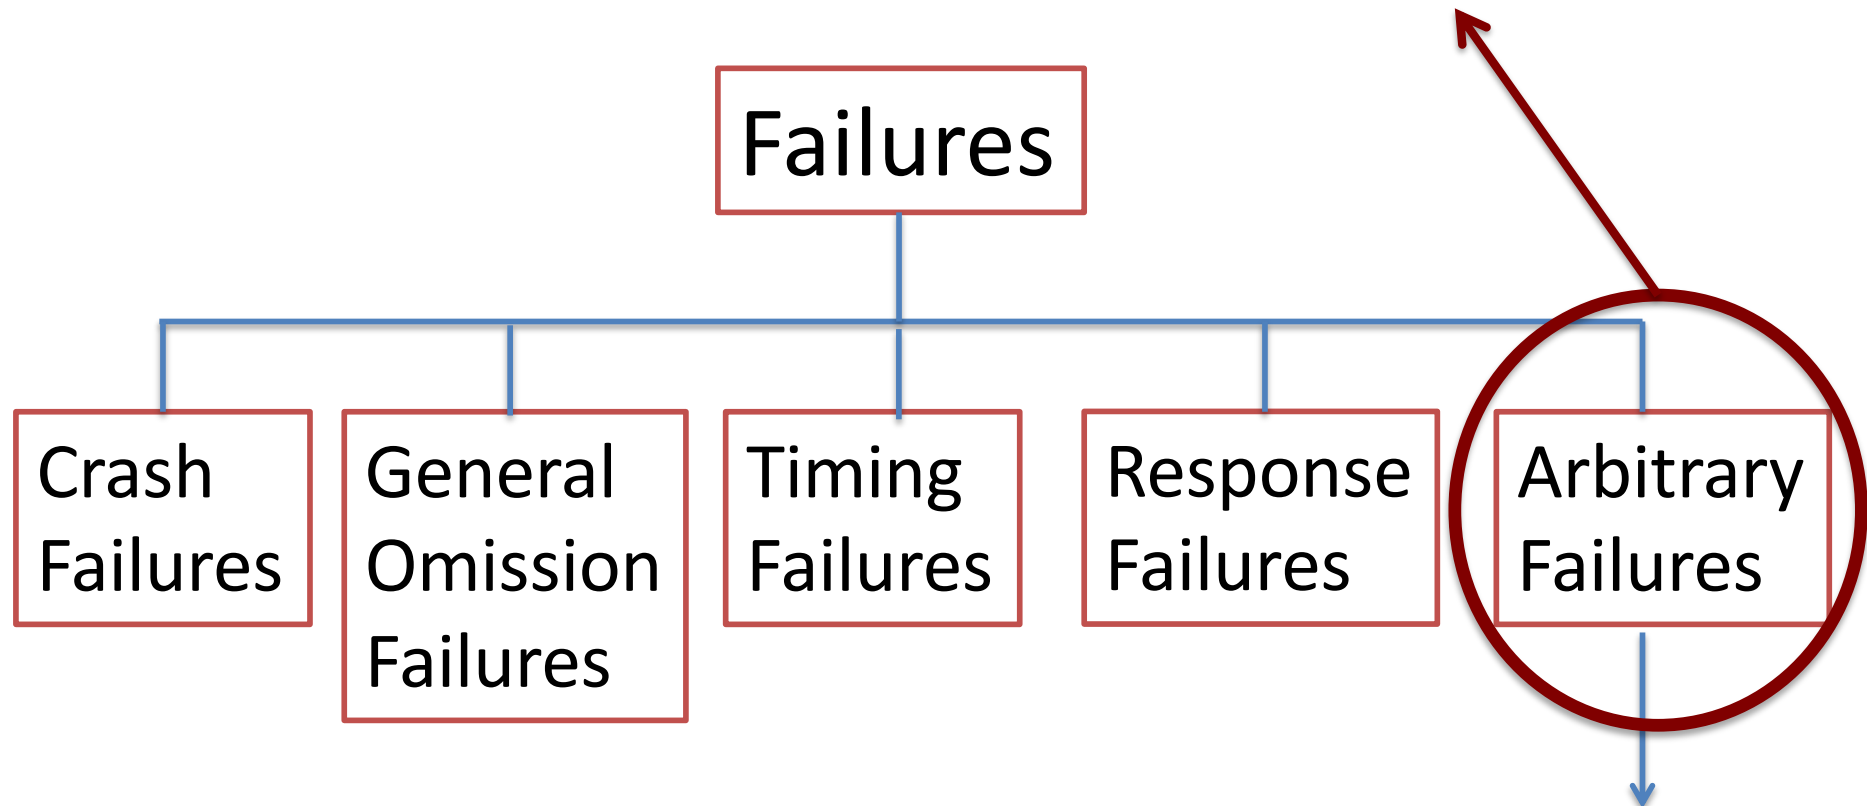
Failure Models

> Arbitrary Failures



Failure Models

> Arbitrary Failures => Byzantine Faults



Any (or any combination of) failure may occur, perhaps even unnoticed (silent failures) or not (noticed or detectable failures)

BYZANTINE FAULT MODEL ASSUMPTIONS: Interesting for Dependability Assumptions. Why ?

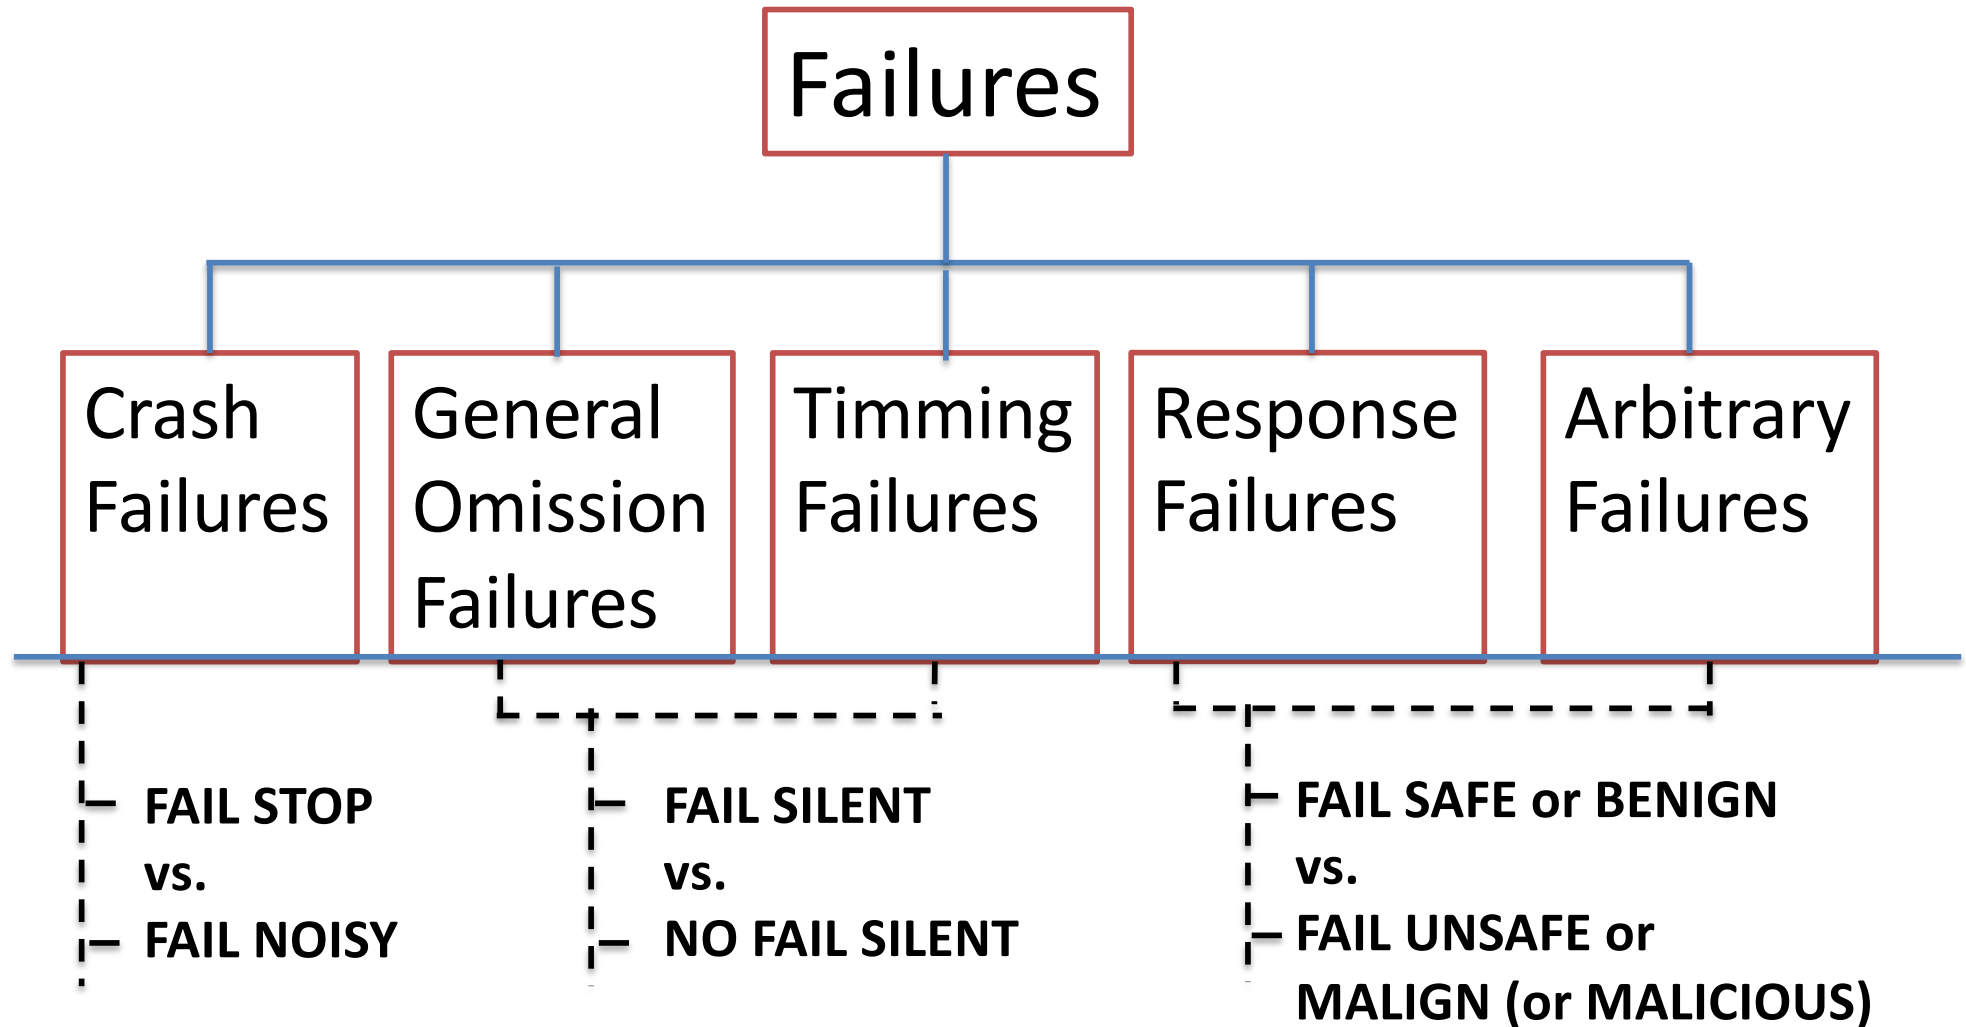
Can we detect halting failures?

- **Synchronous system model**: process execution speeds and message delivery times are bounded
 - → More easy and sometimes we can reliably detect omission and timing failures.
- **Asynchronous system**: no time-bound no assumptions about process execution speeds or message delivery times
 - → Cannot reliably detect crash failures, neither general omission failures or timing failures

Fault Tolerance: Halting Failures

- Sometimes we can consider **quasi or partially synchronous systems**:
 - Most of the time (quasi) or in some moments or parts (partially), we can assume the system to be synchronous
 - As example is also when we consider “synchronous behaviour in some component or in a component in our Trust Computing Base)
 - Yet there is no bound on the time that a system is asynchronous (rest of the time for the rest of the system)

Failure Models and Refinements



Failures: other classification criteria

- Nature of output : malicious vs. non-malicious failures
- Permanent vs. transient (or intermittent) failures
- Independent vs Correlated (or Colluding Failures)

Nature of output

Assumptions we can make (some authors):

- **Malicious fault:** The fault that causes a unit to behave arbitrarily or malicious. Also referred or associated by some authors to a **Byzantine fault**
 - A sensor sending conflicting outputs to different processors
 - Compromised software system that attempts to cause service failure
- **Non-malicious faults:** the opposite of malicious faults
 - Faults that are not caused with malicious intention
 - Faults that exhibit themselves consistently to all observers, e.g., fail-stop
 - A fail-stop system simply stops executing once it fails
- Of course: Malicious faults are much harder to detect than non-malicious faults

Permanent vs. Transient Faults

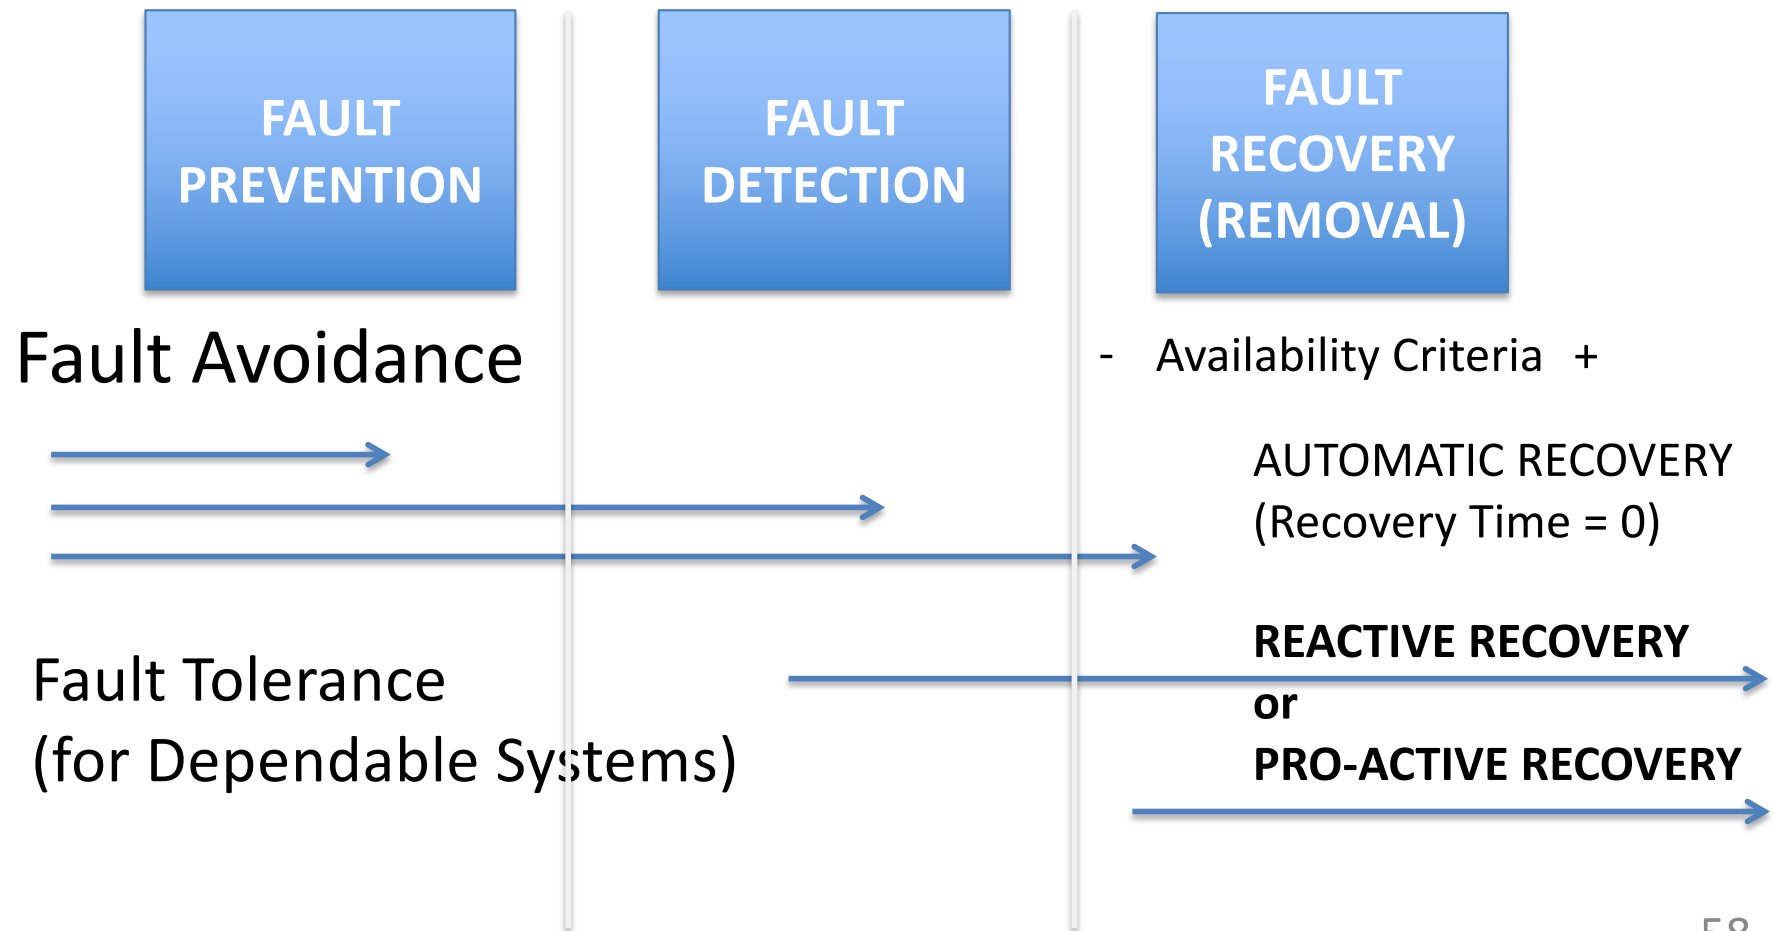
- **Permanent faults** are caused by irreversible device/software failures within a component due to damage, fatigue, or improper manufacturing, or bad design and implementation
 - Permanent software faults are also called **Bohrbugs**
 - Easier to detect
- **Transient/intermittent faults** are triggered by environmental disturbances (external excitation) or incorrect design
 - Transient software faults are also referred to as **Heisenbugs (ex., temporary memory corruption)**
 - Many studies show that Heisenbugs are the majority software faults
 - Harder to detect

Types by correlation criteria: independent vs. correlated failures

- Components fault may be independent of one another or correlated
 - A fault is said to be **independent** if it does not directly or indirectly cause another fault
 - Faults are said to be **correlated** if they are related. Faults could be correlated due to physical or electrical coupling of components
- Correlated faults are more difficult to detect than independent faults

Means for Fault-Tolerance

Means to achieve dependability (fault-model dimension)



Failure Masking and Redundancy

- To be Fault-Tolerant we must **hide the possible occurrence of failures** in distributed processes (**notion of failure masking**)
- **Key-Technique: Redundancy**
 - **Data** (or information) redundancy
 - Integrity codes (EC/FEC, Hamming codes, Integrity Checks, Data Replication and Erasure Coding...)
 - **Time** Redundancy (transient/intermittent failures)
 - Retry strategies (ex., distributed transactions)
 - **SW** redundancy
 - State and Processing (Replication services, Replicas, Process Groups)
 - **Physical redundancy w/ HW**

Failure Masking and Resilience: k-Fault Tolerant Groups

- **k-Fault-tolerant group:**
 - When a **group** can mask any **k** concurrent **member failures**
 - **k** is called **degree of fault tolerance**
 - Resilience Group Metrics for the Group Cardinality and the inherent Replication Protocol, consistency model and synchronicity model
 - Ex: **$N = 3f + 1$** for **f** byzantine failures, etc