

03

Continuous Query Language

Notice

■ Author

- ◆ José Júlio Alferes (jja@fct.unl.pt)
- ◆ João Moura Pires (jmp@fct.unl.pt)

- This material can be freely used for personal or academic purposes without any previous authorization from the author, provided that this notice is maintained/kept.
- For commercial purposes the use of any part of this material requires the previous authorization from the author(s).

Bibliography

- **Many examples are extracted and adapted from:**
 - ◆ Arvind Arasu, Shivnath Babu, Jennifer Widom: The CQL continuous query language: semantic foundations and query execution. VLDB J. 15(2): 121-142 (2006)

Table of Contents

- CQL Basic Concepts
- CQL language
- On CQL implementation
- CQL in Oracle CEP

CQL Basic Concepts

Continuous Query Language

■ CQL is a declarative query language for DSMS

- ◆ It was developed in 2006 for STREAM, the **Stanford Stream Data Management System**
- ◆ It is an extension of SQL, tailored to deal with streams, besides dealing with relations
- ◆ It is part of Oracle's system since version 11g (in Oracle CEP)
 - You may try CQL in Oracle,
 - or by installing the (deprecated) STREAM prototype

Continuous Query Language

■ CQL is a declarative query language for DSMS

- ◆ It was developed in 2006 for STREAM, the **Stanford Stream Data Management System**
- ◆ It is an extension of SQL, tailored to deal with streams, besides dealing with relations
- ◆ It is part of Oracle's system since version 11g (in Oracle CEP)
 - You may try CQL in Oracle,
 - or by installing the (deprecated) STREAM prototype

■ We will cover

- ◆ CQL motivation and design principles
- ◆ CQL formal abstract semantics
- ◆ CQL language, and usage examples
- ◆ CQL query plans and execution (and other implementation issues)

Continuous Query Language

■ CQL is a declarative query language for DSMS

- ◆ It was developed in 2006 for STREAM, the **Stanford Stream Data Management System**
- ◆ It is an extension of SQL, tailored to deal with streams, besides dealing with relations
- ◆ It is part of Oracle's system since version 11g (in Oracle CEP)
 - You may try CQL in Oracle,
 - or by installing the (deprecated) STREAM prototype

■ We will cover

- ◆ CQL motivation and design principles
 - ◆ CQL formal abstract semantics
 - ◆ CQL language, and usage examples
 - ◆ CQL query plans and execution (and other implementation issues)
- Context

Continuous Query Language

■ CQL is a declarative query language for DSMS

- ◆ It was developed in 2006 for STREAM, the **Stanford Stream Data Management System**
- ◆ It is an extension of SQL, tailored to deal with streams, besides dealing with relations
- ◆ It is part of Oracle's system since version 11g (in Oracle CEP)
 - You may try CQL in Oracle,
 - or by installing the (deprecated) STREAM prototype

■ We will cover

- ◆ CQL motivation and design principles Context
- ◆ CQL formal abstract semantics Relational algebra
- ◆ CQL language, and usage examples
- ◆ CQL query plans and execution (and other implementation issues)

Continuous Query Language

■ CQL is a declarative query language for DSMS

- ◆ It was developed in 2006 for STREAM, the **Stanford Stream Data Management System**
- ◆ It is an extension of SQL, tailored to deal with streams, besides dealing with relations
- ◆ It is part of Oracle's system since version 11g (in Oracle CEP)
 - You may try CQL in Oracle,
 - or by installing the (deprecated) STREAM prototype

■ We will cover

- ◆ CQL motivation and design principles Context
- ◆ CQL formal abstract semantics Relational algebra
- ◆ CQL language, and usage examples SQL and its usage
- ◆ CQL query plans and execution (and other implementation issues)

CQL Basic Concepts

CQL Basic Concepts

- Two main data types:

- ◆ Relations
- ◆ Streams

CQL Basic Concepts

- Two main data types:
 - ◆ Relations
 - ◆ Streams
- Queries can return:
 - ◆ a **relation** based on other relations (as in SQL)
 - ◆ a **stream** based on other streams and relations (extends SQL)

CQL Basic Concepts

- Two main data types:
 - ◆ Relations
 - ◆ Streams
- Queries can return:
 - ◆ a **relation** based on other relations (as in SQL)
 - ◆ a **stream** based on other streams and relations (extends SQL)
- This is done with an algebra that has 3 kinds of operators
 - ◆ *stream-to-relation, relation-to-relation, relation-to-stream*

CQL Basic Concepts

- Two main data types:
 - ◆ Relations
 - ◆ Streams
- Queries can return:
 - ◆ a **relation** based on other relations (as in SQL)
 - ◆ a **stream** based on other streams and relations (extends SQL)
- This is done with an algebra that has 3 kinds of operators
 - ◆ *stream-to-relation, relation-to-relation, relation-to-stream*
- There are no **stream-to-stream** operators, and this makes things closer to SQL
 - ◆ Can be obtained by combining *stream-to-relation - relation-to-relation - relation-to-stream*

A motivating example

- Simplified version of the Linear Road Benchmark

A motivating example

- Simplified version of the Linear Road Benchmark



A motivating example

- **Simplified version of the Linear Road Benchmark**
- **A system that computes variable tolling in a highway**
 - ◆ The tolls depend on the distance travelled in the highway (as usual), but also on traffic conditions
 - The more congested the highway is, the more toll is charged (in order to discourage cars to go to congested highways)



A motivating example

- **Simplified version of the Linear Road Benchmark**
- **A system that computes variable tolling in a highway**
 - ◆ The tolls depend on the distance travelled in the highway (as usual), but also on traffic conditions
 - The more congested the highway is, the more toll is charged (in order to discourage cars to go to congested highways)
- **Each car has a sensor that relays its position and speed every 30s, while in the highway**
 - ◆ This information is continuously passed to the system, as a stream of data



A motivating example

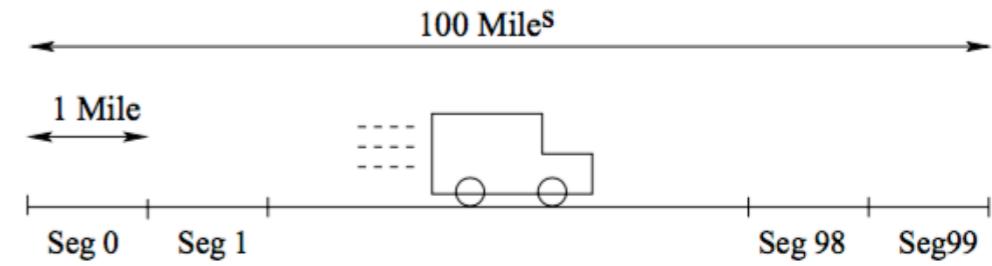
- **Simplified version of the Linear Road Benchmark**
- **A system that computes variable tolling in a highway**
 - ◆ The tolls depend on the distance travelled in the highway (as usual), but also on traffic conditions
 - The more congested the highway is, the more toll is charged (in order to discourage cars to go to congested highways)
- **Each car has a sensor that relays its position and speed every 30s, while in the highway**
 - ◆ This information is continuously passed to the system, as a stream of data
- **The system computes tolls in real-time, and continuously outputs a stream with the prices to pay by each car, upon their exit**



Simplified Linear Road (details)

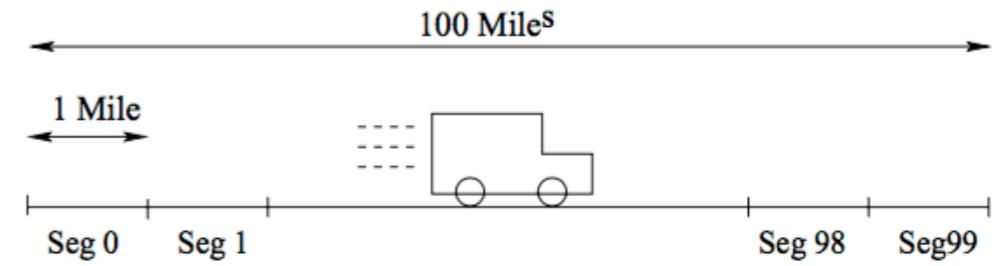
- The road is **divided into 100 segments** (of 1 Mile each), possibly with an entrance at the beginning of each segment, and/or an exit at the end.

Simplified Linear Road (details)



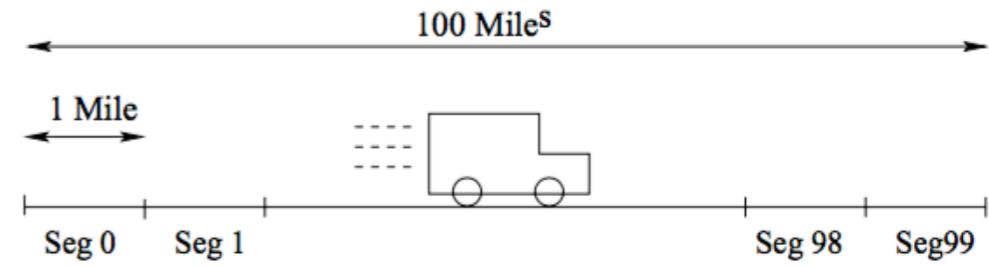
- The road is **divided into 100 segments** (of 1 Mile each), possibly with an entrance at the beginning of each segment, and/or an exit at the end.

Simplified Linear Road (details)



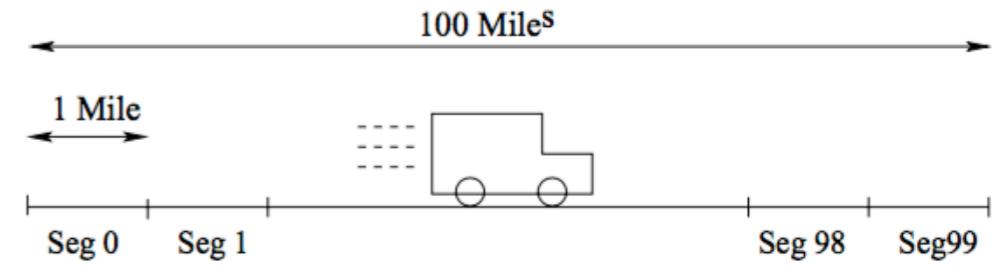
- The road is **divided into 100 segments** (of 1 Mile each), possibly with an entrance at the beginning of each segment, and/or an exit at the end.
- **The input stream has tuples with carID, position, and speed**
 - ◆ The position is given as the distance in number of feet from the beginning of the road; the speed in miles per hour

Simplified Linear Road (details)



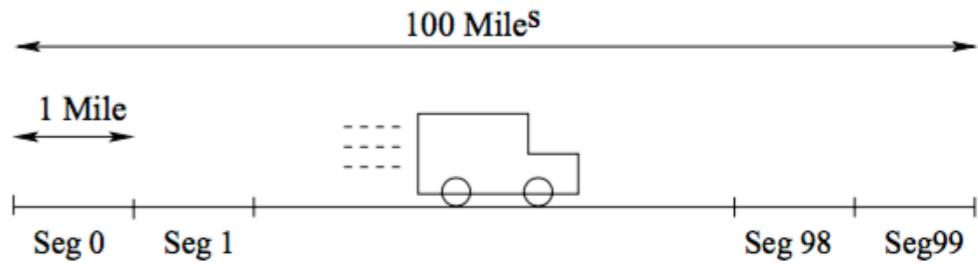
- The road is **divided into 100 segments** (of 1 Mile each), possibly with an entrance at the beginning of each segment, and/or an exit at the end.
- **The input stream has tuples with carID, position, and speed**
 - ◆ The position is given as the distance in number of feet from the beginning of the road; the speed in miles per hour
- **The output stream has tuples with carID and price to pay for a segment**, any time a car enters that segment

Simplified Linear Road (details)



- The road is **divided into 100 segments** (of 1 Mile each), possibly with an entrance at the beginning of each segment, and/or an exit at the end.
- **The input stream has tuples with carID, position, and speed**
 - ◆ The position is given as the distance in number of feet from the beginning of the road; the speed in miles per hour
- **The output stream has tuples with carID and price to pay for a segment**, any time a car enters that segment
- **The total price to pay for a segment is as follows:**
 - ◆ a non-congested segment pays nothing
 - ◆ a congested segment pays $2 \times (\text{numCars}-50)^2$, where *numCars* is the number of cars in the segment at the moment. A segment is **congested** if the **average speed of all cars in the last 5 minutes** < than 40 Miles/h.

Simplified Linear Road (details)



- The road is **divided into 100 segments** (of 1 Mile each), possibly with an entrance at the beginning of each segment, and/or an exit at the end.
- **The input stream has tuples with carID, position, and speed**
 - ◆ The position is given as the distance in number of feet from the beginning of the road; the speed in miles per hour
- **The output stream has tuples with carID and price to pay for a segment**, any time a car enters that segment
- **The total price to pay for a segment is as follows:**
 - ◆ a non-congested segment pays nothing
 - ◆ a congested segment pays $2 \times (\text{numCars}-50)^2$, where *numCars* is the number of cars in the segment at the moment. A segment is **congested** if the **average speed of all cars in the last 5 minutes** < than 40 Miles/h.
- Note that all of this changes dynamically, and the **continuous query must compute the tolls at all times**

Streams and Relations

- There is a notion of **schemas**, which is exactly as in relational DBs (**a set of named attributes**)
- There is also a notion of **timestamps** (or time instants) which is simply a **discrete totally ordered set T** (e.g. nonnegative integers, or datetime)

Streams and Relations

- There is a notion of **schemas**, which is exactly as in relational DBs (**a set of named attributes**)
- There is also a notion of **timestamps** (or time instants) which is simply a **discrete totally ordered set T** (e.g. nonnegative integers, or datetime)
- A **stream S** is a possible infinite **bag of elements $\langle s, t \rangle$** where **s** is a tuple in the schema of **S** , and $t \in T$ is the timestamp of the element.

Streams and Relations

- There is a notion of **schemas**, which is exactly as in relational DBs (**a set of named attributes**)
- There is also a notion of **timestamps** (or time instants) which is simply a **discrete totally ordered set T** (e.g. nonnegative integers, or datetime)
- A **stream S** is a possible infinite **bag of elements $\langle s, t \rangle$** where **s** is a tuple in the schema of **S** , and $t \in T$ is the timestamp of the element.
 - ◆ I.e. **stream are timestamped sequences of tuples**

Streams and Relations

- There is a notion of **schemas**, which is exactly as in relational DBs (**a set of named attributes**)
- There is also a notion of **timestamps** (or time instants) which is simply a **discrete totally ordered set T** (e.g. nonnegative integers, or datetime)
- A **stream S** is a possible infinite **bag of elements $\langle s, t \rangle$** where s is a tuple in the schema of S , and $t \in T$ is the timestamp of the element.
 - ◆ I.e. **stream are timestamped sequences of tuples**
- A **relation R** is a mapping from each time instant in T to a finite bag of tuples in the schema of R

Instantaneous Relations and Relations

- Instantaneous relation $R(t)$ is the bag of tuples of R at time t
 - ◆ This exactly corresponds to a relation in DBs

Instantaneous Relations and Relations

- Instantaneous relation $R(t)$ is the bag of tuples of R at time t
 - ◆ This exactly corresponds to a relation in DBs
- Note that these relations are an extension of relations in DBs
 - ◆ In DBs a relation is a (permanent) bag of tuples - it does not take into account the changes in relations
 - ◆ DSMS must take into account that relations change over time

Instantaneous Relations and Relations

- Instantaneous relation $R(t)$ is the bag of tuples of R at time t
 - ◆ This exactly corresponds to a relation in DBs
- Note that these relations are an extension of relations in DBs
 - ◆ In DBs a relation is a (permanent) bag of tuples - it does not take into account the changes in relations
 - ◆ DSMS must take into account that relations change over time
- A Relation is a set of Instantaneous relation $R(t)$

Instantaneous Relations and Relations

- **S up to t** denotes the bag of elements in stream S with timestamps $\leq t$
 - ◆ $S \text{ up to } t = \{\langle s, t' \rangle \in S : t' \leq t\}$

Instantaneous Relations and Relations

- **S up to t** denotes the bag of elements in stream S with timestamps $\leq t$
 - ◆ $S \text{ up to } t = \{\langle s, t' \rangle \in S : t' \leq t\}$
- **S at t** denotes the bag of elements of S with timestamp t
 - ◆ $S \text{ at } t = \{\langle s, t' \rangle \in S : t' = t\}$

Instantaneous Relations and Relations

- **S up to t** denotes the bag of elements in stream S with timestamps $\leq t$
 - ◆ $S \text{ up to } t = \{\langle s, t' \rangle \in S : t' \leq t\}$
- **S at t** denotes the bag of elements of S with timestamp t
 - ◆ $S \text{ at } t = \{\langle s, t' \rangle \in S : t' = t\}$
- **R up to t** denotes the collection of instantaneous relations $R(0), \dots, R(t)$

Instantaneous Relations and Relations

- **S up to t** denotes the bag of elements in stream S with timestamps $\leq t$
 - ◆ $S \text{ up to } t = \{\langle s, t' \rangle \in S : t' \leq t\}$
- **S at t** denotes the bag of elements of S with timestamp t
 - ◆ $S \text{ at } t = \{\langle s, t' \rangle \in S : t' = t\}$
- **R up to t** denotes the collection of instantaneous relations $R(0), \dots, R(t)$
- **R at t** denotes the instantaneous relation $R(t)$.

Linear Road Example

- The Linear Road application there is just **one base stream** containing vehicle speed and position measurements, with schema:

PosSpeedStr (vehicleId, speed, xPos)

Linear Road Example

- The Linear Road application there is just **one base stream** containing vehicle speed and position measurements, with schema:

PosSpeedStr (vehicleId, speed, xPos)

- The Linear Road application contains **no base relations**, but several **derived relations** are useful in toll computation. For example, the toll for a congested segment depends on the current number of vehicles in the segment

SegVolRel (segNo, numVehicles)

Linear Road Example

- The Linear Road application there is just **one base stream** containing vehicle speed and position measurements, with schema:

PosSpeedStr (vehicleId, speed, xPos)

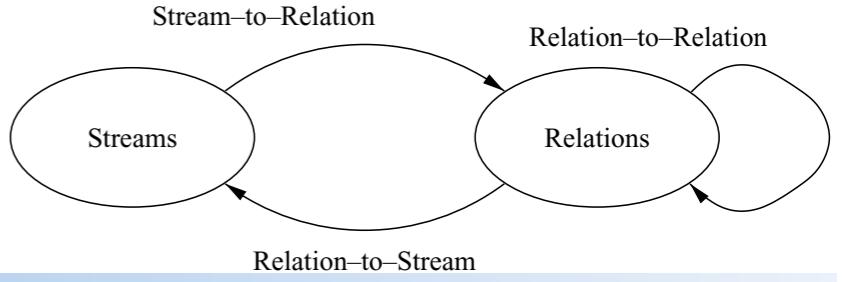
- The Linear Road application contains **no base relations**, but several **derived relations** are useful in toll computation. For example, the toll for a congested segment depends on the current number of vehicles in the segment

SegVolRel (segNo, numVehicles)

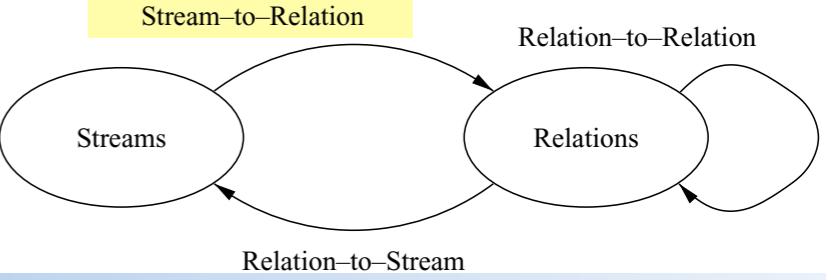
- At time t , SegVolRel(t) contains the count of vehicles (numVehicles) in each highway segment (segNo) as of time t .

It seems more natural to model “the current number of vehicles in a segment” as a **time-varying relation**, rather than as a stream of the latest values.

Abstract operators



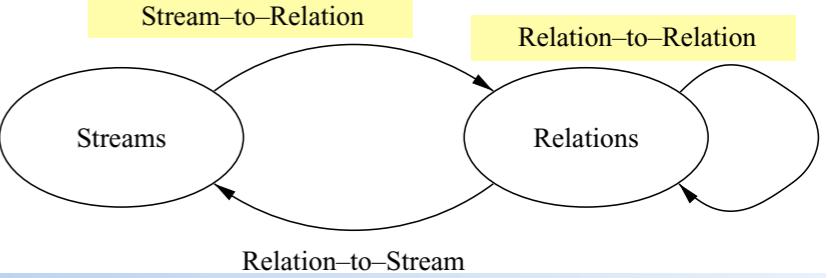
Abstract operators



■ Stream-to-relation (s2r) operators

- ◆ Receive a stream S and produce a relation R with the same schema of S
- ◆ At any instant t , $R(t)$ must be computable from S up to t
 - (i.e. from the tuples of S with timestamp $\leq t$)

Abstract operators



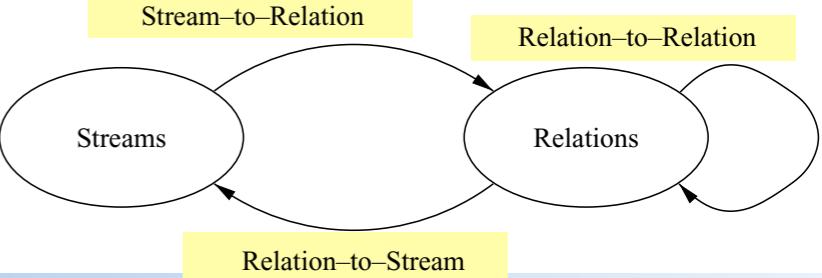
■ Stream-to-relation (s2r) operators

- ◆ Receive a stream S and produce a relation R with the same schema of S
- ◆ At any instant t , $R(t)$ must be computable from S up to t
 - (i.e. from the tuples of S with timestamp $\leq t$)

■ Relation-to-relation (r2r) operators

- ◆ Receive one or more relations $R_1 \dots R_n$ and produce a relation R
- ◆ At any instant t , $R(t)$ must be computable from $R_1(t) \dots R_n(t)$

Abstract operators



■ Stream-to-relation (s2r) operators

- ◆ Receive a stream S and produce a relation R with the same schema of S
- ◆ At any instant t , $R(t)$ must be computable from S up to t
 - (i.e. from the tuples of S with timestamp $\leq t$)

■ Relation-to-relation (r2r) operators

- ◆ Receive one or more relations $R_1 \dots R_n$ and produce a relation R
- ◆ At any instant t , $R(t)$ must be computable from $R_1(t) \dots R_n(t)$

■ Relation-to-stream (r2s) operators

- ◆ Receive one relation R and produce a stream S
- ◆ At any instant t , S at t must be computable from $R(t)$
 - (where S at t is the bag of all tuples of S with timestamp t)

Continuous Semantics

- A **query Q** is a well typed algebraic expression (or **tree**) formed with the previously mentioned operators.

Continuous Semantics

- A **query Q** is a well typed algebraic expression (or **tree**) formed with the previously mentioned operators.
- Suppose the set of **all inputs** to the innermost (leaf) operators of Q are **streams S_1, \dots, S_n** ($n \geq 0$) and **relations R_1, \dots, R_m** ($m \geq 0$)

Continuous Semantics

- A **query Q** is a well typed algebraic expression (or **tree**) formed with the previously mentioned operators.
- Suppose the set of **all inputs** to the innermost (leaf) operators of Q are **streams S_1, \dots, S_n** ($n \geq 0$) and **relations R_1, \dots, R_m** ($m \geq 0$)
- The **result of continuous query Q at a time t** , which denotes the result of Q once all inputs up to t are “available”. There are two cases:
 - ◆ If the outermost operator of Q is a **r2s** operator producing **S**
 - ◆ If the outermost operator is a **s2r** or **r2r** operator producing **R**

Continuous Semantics

- The **result of continuous query Q at a time t**, which denotes the result of Q once all inputs up to t are “available”. There are two cases:
- **If the outermost operator of Q is a r2s operator producing S**
 - ◆ the result of Q at t is **S up to t**
 - ◆ it is produced by recursively applying the operators in Q to the leaf streams and relations *up to t*

Continuous Semantics

- The **result of continuous query Q at a time t** , which denotes the result of Q once all inputs up to t are “available”. There are two cases:
- **If the outermost operator of Q is a $r2s$ operator producing S**
 - ◆ the result of Q at t is **S up to t**
 - ◆ it is produced by recursively applying the operators in Q to the leaf streams and relations *up to t*
- **If the outermost operator is a $s2r$ or $r2r$ operator producing R**
 - ◆ the result of Q at t is **$R(t)$** ,
 - ◆ it is produced by recursively applying the operators in Q to the leaf streams and relations *up to t*

Continuous Queries

- How to operationally understand the continuous semantics?

Continuous Queries

■ How to operationally understand the continuous semantics?

- ◆ Let time advance in T

Continuous Queries

■ How to operationally understand the continuous semantics?

- ◆ Let time advance in T
- ◆ At time $t \in T$, all inputs **up to t** are processed
 - This requires that, after processing the outputs at time t , no further input with timestamp t can arrive
 - *heartbeats* are timestamps t that prevent the system from receiving further input with timestamps $\leq t$

Continuous Queries

■ How to operationally understand the continuous semantics?

- ◆ Let time advance in T
- ◆ At time $t \in T$, **all inputs up to t are processed**
 - This requires that, after processing the outputs at time t , no further input with timestamp t can arrive
 - *heartbeats* are timestamps t that prevent the system from receiving further input with timestamps $\leq t$
- ◆ If the query produces a **stream**, the continuous query **emits new results with timestamp t**

Continuous Queries

■ How to operationally understand the continuous semantics?

- ◆ Let time advance in T
- ◆ At time $t \in T$, all inputs **up to t** are processed
 - This requires that, after processing the outputs at time t , no further input with timestamp t can arrive
 - *heartbeats* are timestamps t that prevent the system from receiving further input with timestamps $\leq t$
- ◆ If the query produces a **stream**, the continuous query **emits new results with timestamp t**
- ◆ If the query produces a **relation R** , it **updates the relation from $R(t-1)$ to $R(t)$**

Continuous Semantics Example

- Query: `Istream(Filter(LastRow(S)))`
- S has a_i every instant i

Time	S
0	$\langle(a_0), 0 \rangle$
1	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$
2	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$ $\langle(a_2), 2 \rangle$
3	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$ $\langle(a_2), 2 \rangle$ $\langle(a_3), 3 \rangle$
4	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$ $\langle(a_2), 2 \rangle$ $\langle(a_3), 3 \rangle$ $\langle(a_4), 4 \rangle$
:	:

Continuous Semantics Example

- Query: `Istream(Filter>LastRow(S))`
- S has a_i every instant i

CQL

```
Select Istream(*)
From S [Rows 1]
Where <filter>.
```

Time	S
0	$\langle(a_0), 0 \rangle$
1	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$
2	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$ $\langle(a_2), 2 \rangle$
3	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$ $\langle(a_2), 2 \rangle$ $\langle(a_3), 3 \rangle$
4	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$ $\langle(a_2), 2 \rangle$ $\langle(a_3), 3 \rangle$ $\langle(a_4), 4 \rangle$
:	:

Continuous Semantics Example

- Query: `Istream(Filter(LastRow(S)))`
- S has a_i every instant i
- `LastRow` is `s2r` that at time t produces a relation with the last tuples that arrived in S

CQL

```
Select Istream(*)
From S [Rows 1]
Where <filter>.
```

Time	S
0	$\langle(a_0), 0 \rangle$
1	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$
2	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$ $\langle(a_2), 2 \rangle$
3	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$ $\langle(a_2), 2 \rangle$ $\langle(a_3), 3 \rangle$
4	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$ $\langle(a_2), 2 \rangle$ $\langle(a_3), 3 \rangle$ $\langle(a_4), 4 \rangle$
:	:

Continuous Semantics Example

- Query: `Istream(Filter(LastRow(S)))`
- S has a_i every instant i
- `LastRow` is `s2r` that at time t produces a relation with the last tuples that arrived in S

CQL

```
Select Istream(*)
From S [Rows 1]
Where <filter>.
```

Time	S	LastRow
0	$\langle(a_0), 0 \rangle$	(a_0)
1	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$	(a_1)
2	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$ $\langle(a_2), 2 \rangle$	(a_2)
3	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$ $\langle(a_2), 2 \rangle$ $\langle(a_3), 3 \rangle$	(a_3)
4	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$ $\langle(a_2), 2 \rangle$ $\langle(a_3), 3 \rangle$ $\langle(a_4), 4 \rangle$	(a_4)
:	:	:

Continuous Semantics Example

- Query: `Istream(Filter(LastRow(S)))`
- S has a_i every instant i
- `LastRow` is s2r that at time t produces a relation with the last tuples that arrived in S
- `Filter` is r2r and filters way all a_i where i is odd

CQL `Select Istream(*)
From S [Rows 1]
Where <filter>.`

Time	S	LastRow
0	$\langle(a_0), 0 \rangle$	(a_0)
1	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$	(a_1)
2	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$ $\langle(a_2), 2 \rangle$	(a_2)
3	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$ $\langle(a_2), 2 \rangle$ $\langle(a_3), 3 \rangle$	(a_3)
4	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$ $\langle(a_2), 2 \rangle$ $\langle(a_3), 3 \rangle$ $\langle(a_4), 4 \rangle$	(a_4)
:	:	:

Continuous Semantics Example

- Query: `Istream(Filter(LastRow(S)))`
- S has a_i every instant i
- `LastRow` is s2r that at time t produces a relation with the last tuples that arrived in S
- `Filter` is r2r and filters way all a_i where i is odd

CQL `Select Istream(*)
From S [Rows 1]
Where <filter>.`

Time	S	LastRow	Filter
0	$\langle(a_0), 0 \rangle$	(a_0)	(a_0)
1	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$	(a_1)	ϕ
2	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$ $\langle(a_2), 2 \rangle$	(a_2)	(a_2)
3	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$ $\langle(a_2), 2 \rangle$ $\langle(a_3), 3 \rangle$	(a_3)	ϕ
4	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$ $\langle(a_2), 2 \rangle$ $\langle(a_3), 3 \rangle$ $\langle(a_4), 4 \rangle$	(a_4)	(a_4)
:	:	:	:

Continuous Semantics Example

- Query: `Istream(Filter(LastRow(S)))`
- S has a_i every instant i
- `LastRow` is `s2r` that at time t produces a relation with the last tuples that arrived in S
- `Filter` is `r2r` and filters way all a_i where i is odd
- `Istream` is `r2s` that “streams” every **new** tuple in R

CQL

```
Select Istream(*)
From S [Rows 1]
Where <filter>.
```

Time	S	LastRow	Filter
0	$\langle(a_0), 0 \rangle$	(a_0)	(a_0)
1	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$	(a_1)	ϕ
2	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$ $\langle(a_2), 2 \rangle$	(a_2)	(a_2)
3	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$ $\langle(a_2), 2 \rangle$ $\langle(a_3), 3 \rangle$	(a_3)	ϕ
4	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$ $\langle(a_2), 2 \rangle$ $\langle(a_3), 3 \rangle$ $\langle(a_4), 4 \rangle$	(a_4)	(a_4)
:	:	:	:

Continuous Semantics Example

- Query: `Istream(Filter(LastRow(S)))`
- S has a_i every instant i
- `LastRow` is `s2r` that at time t produces a relation with the last tuples that arrived in S
- `Filter` is `r2r` and filters way all a_i where i is odd
- `Istream` is `r2s` that “streams” every **new** tuple in R

CQL `Select Istream(*)
From S [Rows 1]
Where <filter>.`

Time	S	LastRow	Filter	Istream
0	$\langle(a_0), 0 \rangle$	(a_0)	(a_0)	$\langle(a_0), 0 \rangle$
1	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$	(a_1)	ϕ	$\langle(a_0), 0 \rangle$
2	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$ $\langle(a_2), 2 \rangle$	(a_2)	(a_2)	$\langle(a_0), 0 \rangle$ $\langle(a_2), 2 \rangle$
3	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$ $\langle(a_2), 2 \rangle$ $\langle(a_3), 3 \rangle$	(a_3)	ϕ	$\langle(a_0), 0 \rangle$ $\langle(a_2), 2 \rangle$
4	$\langle(a_0), 0 \rangle$ $\langle(a_1), 1 \rangle$ $\langle(a_2), 2 \rangle$ $\langle(a_3), 3 \rangle$ $\langle(a_4), 4 \rangle$	(a_4)	(a_4)	$\langle(a_0), 0 \rangle$ $\langle(a_2), 2 \rangle$ $\langle(a_4), 4 \rangle$
:	:	:	:	:

Requirements for processing a query

- In the Linear Road application the sequence of operators producing derived relation **SegVolRel** conceptually produces, **at every time instant t** , the instantaneous relation **SegVolRel(t)** containing the current number of vehicles in each segment.

Requirements for processing a query

- In the Linear Road application the sequence of operators producing derived relation **SegVolRel** conceptually produces, **at every time instant t** , the instantaneous relation **SegVolRel(t)** containing the current number of vehicles in each segment.
- In a DSMS implementing this semantics, **SegVolRel(t) cannot be produced until it is known that all elements on input stream**

PosSpeedStr(vehicleId, speed, xPos)

with timestamp $\leq t$ have been received !

Requirements for processing a query

- In the Linear Road application the sequence of operators producing derived relation **SegVolRel** conceptually produces, **at every time instant t** , the instantaneous relation **SegVolRel(t)** containing the current number of vehicles in each segment.
- In a DSMS implementing this semantics, **SegVolRel(t) cannot be produced until it is known that all elements on input stream**

PosSpeedStr(vehicleId, speed, xPos)

with timestamp $\leq t$ have been received !

- Furthermore, once they have been received, **there may be additional lag before the relation is actually updated** due to query processing time.

CQL language

CQL language

- CQL is obtained by a precise instantiation of the abstract operators

CQL language

- CQL is obtained by a precise instantiation of the abstract operators

- CQL is obtained by a precise instantiation of the abstract operators
- General Idea
 - ◆ Support a **large class** of **relation-to-relation** operators that **perform the bulk of data manipulation** in a typical CQL query, along with a **small set** of **stream-to-relation** and **relation-to-stream** operators that **convert streams to relations and back.**

CQL language

- CQL is obtained by a precise instantiation of the abstract operators
- General Idea
 - ◆ Support a **large class** of **relation-to-relation** operators that **perform the bulk of data manipulation** in a typical CQL query, along with a **small set** of **stream-to-relation** and **relation-to-stream** operators that **convert streams to relations and back**.
 - ◆ The primary advantage of this approach is the ability to **reuse** the formal foundations and **huge body of implementation techniques for relation-to-relation languages** such as relational algebra and SQL !

CQL language

- CQL is obtained by a precise instantiation of the abstract operators
- r2r operators are those of relational algebra
 - ◆ Formally, they have to be adapted to time-varying relations, but that is easy:
 - ◆ If \oplus_r is a traditional operator (e.g. \bowtie , Π , σ , ...) over R_1, \dots, R_n then \oplus in CQL produces a time-varying relations R , such that at each t , $R(t) = \oplus(R_1(t), \dots, R_n(t))$

CQL language

- CQL is obtained by a precise instantiation of the abstract operators
- r2r operators are those of relational algebra
 - ◆ Formally, they have to be adapted to time-varying relations, but that is easy:
 - ◆ If \oplus_r is a traditional operator (e.g. \bowtie , Π , σ , ...) over R_1, \dots, R_n then \oplus in CQL produces a time-varying relations R , such that at each t , $R(t) = \oplus(R_1(t), \dots, R_n(t))$
- For writing the queries, CQL uses an extension of SQL syntax
 - ◆ By now, it should be more than clear to you how to switch from SQL syntax to relational algebra, and back!

CQL language

- CQL is obtained by a precise instantiation of the abstract operators
- r2r operators are those of relational algebra
 - ◆ Formally, they have to be adapted to time-varying relations, but that is easy:
 - ◆ If \oplus_r is a traditional operator (e.g. \bowtie , Π , σ , ...) over R_1, \dots, R_n then \oplus in CQL produces a time-varying relations R , such that at each t , $R(t) = \oplus(R_1(t), \dots, R_n(t))$
- For writing the queries, CQL uses an extension of SQL syntax
 - ◆ By now, it should be more than clear to you how to switch from SQL syntax to relational algebra, and back!
- Incorporating user-defined procedures, aggregates, and windows, is straightforward in CQL, at least from the semantics perspective.

CQL's **s2r** operators

CQL's **s2r** operators

- They are all based on the concept of **sliding window over a stream**
 - ◆ at any point t a window is an instantaneous relation $R(t)$ that contains a historical snapshot of a finite portion of the stream
 - Note that in the relation the timestamps disappear

CQL's **s2r** operators

- They are all based on the concept of **sliding window over a stream**
 - ◆ at any point t a window is an instantaneous relation $R(t)$ that contains a historical snapshot of a finite portion of the stream
 - Note that in the relation the timestamps disappear
- 4 types of sliding window s2r operators
 - ◆ **time-based** (e.g. last minute)
 - ◆ **tuple-based** (e.g. last five tuples)
 - ◆ **partitioned** (similar to group by in SQL)
 - ◆ **with slide parameter** (e.g. last minute, picked every minute)

Time-based windows

- **S [Range T]** takes a stream **S** and a time-interval **T**, and produces the relation **R** such that

$$R(t) = \{s | \langle s, t' \rangle \in S \wedge \max(t - T, 0) \leq t' \leq t\}$$

Time-based windows

- **S [Range T] takes a stream S and a time-interval T, and produces the relation R such that**

$$R(t) = \{s | \langle s, t' \rangle \in S \wedge \max(t - T, 0) \leq t' \leq t\}$$

- ◆ Produces a relation with all tuples of stream S with timestamps in the last T time
 - E.g. S [Range 30s] produces a relation with all tuples from S in the last 30 seconds
- Note that the tuples in R do not have timestamps. They are “time-flat”!

Time-based windows

- **S [Range T]** takes a stream S and a time-interval T, and produces the relation R such that

$$R(t) = \{s | \langle s, t' \rangle \in S \wedge \max(t - T, 0) \leq t' \leq t\}$$

- ◆ Produces a relation with all tuples of stream S with timestamps in the last T time
 - E.g. S [Range 30s] produces a relation with all tuples from S in the last 30 seconds
- Note that the tuples in R do not have timestamps. They are “time-flat”!
- Special cases:
 - ◆ $T=0$ (write S [Now])
 - at any time t , $R(t)$ has the tuples that appeared in the stream at t , possibly empty
 - ◆ $T=\infty$ (write S [Range Unbounded])
 - at any time R , has all the tuples that ever appeared in S, so far

Example

- Let **PosSpeedStr** be the **input stream** of the Linear Road

Example

- Let **PosspeedStr** be the **input stream** of the Linear Road
- $\sigma_{\text{speed} \geq 60}(\text{PosspeedStr} \text{ [Range } 30\text{s}])$ is the relation that, at any instant t , has all the tuples $(\text{carID}, \text{position}, \text{speed})$ of cars detected in the 30 seconds just before t , with speed higher than 60

Example

- Let **PosspeedStr** be the **input stream** of the Linear Road
- $\sigma_{\text{speed} \geq 60}(\text{PosspeedStr} \text{ [Range } 30\text{s}])$ is the relation that, at any instant t , has all the tuples $(\text{carID}, \text{position}, \text{speed})$ of cars detected in the 30 seconds just before t , with speed higher than 60
- In **SQL-like syntax**:

```
select *
  from PosspeedStr [ Range 30s ]
 where speed >= 60;
```

Example

- Let **PosspeedStr** be the **input stream** of the Linear Road
- $\sigma_{\text{speed} \geq 60}(\text{PosspeedStr} \text{ [Range } 30\text{s}])$ is the relation that, at any instant t , has all the tuples $(\text{carID}, \text{position}, \text{speed})$ of cars detected in the 30 seconds just before t , with speed higher than 60
- In **SQL-like syntax**:

```
select *
from PosspeedStr [ Range 30s ]
where speed >= 60;
```

- The **ids of cars that were ever detected with a speed higher than 200**

```
select distinct carID
from PosspeedStr [ Range Unbounded ]
where speed >= 200;
```

Tuple-based windows

- **S [Rows N]** takes a stream S and an integer N, and produces the relation R such that

Tuple-based windows

- **S [Rows N]** takes a stream S and an integer N, and produces the relation R such that
 - ◆ $R(t)$ contains exactly the N tuples in stream S with highest timestamps $\leq t$

Tuple-based windows

- **S [Rows N]** takes a stream S and an integer N, and produces the relation R such that
 - ◆ $R(t)$ contains exactly the N tuples in stream S with **highest timestamps $\leq t$**
 - if S has less than N tuples with timestamps $\leq t$, then $R(t)$ has all tuples in S

Tuple-based windows

- **S [Rows N]** takes a stream S and an integer N, and produces the relation R such that
 - ◆ $R(t)$ contains exactly the N tuples in stream S with **highest timestamps $\leq t$**
 - if S has less than N tuples with timestamps $\leq t$, then R(t) has all tuples in S
 - if there are ties, then the N tuples are determined arbitrarily

Tuple-based windows

- **S [Rows N]** takes a stream S and an integer N, and produces the relation R such that
 - ◆ $R(t)$ contains exactly the N tuples in stream S with **highest timestamps $\leq t$**
 - if S has less than N tuples with timestamps $\leq t$, then R(t) has all tuples in S
 - if there are ties, then the N tuples are determined arbitrarily
 - ◆ Produces a relation with the last N tuples of stream S

Tuple-based windows

- **S [Rows N]** takes a stream S and an integer N, and produces the relation R such that
 - ◆ $R(t)$ contains exactly the N tuples in stream S with highest timestamps $\leq t$
 - if S has less than N tuples with timestamps $\leq t$, then R(t) has all tuples in S
 - if there are ties, then the N tuples are determined arbitrarily
 - ◆ Produces a relation with the last N tuples of stream S
- E.g. **PosSpeedStr [Rows 1]** denotes, at any instant, the latest position-speed measurement

Tuple-based windows

- **S [Rows N]** takes a stream S and an integer N, and produces the relation R such that
 - ◆ $R(t)$ contains exactly the N tuples in stream S with highest timestamps $\leq t$
 - if S has less than N tuples with timestamps $\leq t$, then $R(t)$ has all tuples in S
 - if there are ties, then the N tuples are determined arbitrarily
 - ◆ Produces a relation with the last N tuples of stream S
- E.g. **PosSpeedStr [Rows 1]** denotes, at any instant, the latest position-speed measurement
- Special case:
 - ◆ $N=\infty$ (write S [Rows Unbounded])
 - It is equivalent to S [Range Unbounded]

Partitioned windows

■ What if we wanted to have *the last measurement for each car?*

- ◆ similar to *group by* each carID and then getting the last measurement for each group

Partitioned windows

- What if we wanted to have *the last measurement for each car?*
 - ◆ similar to *group by* each carID and then getting the last measurement for each group
- **S [Partition By A₁, ..., A_n Rows N]** takes a stream S, a set of attributes A₁, ..., A_n, and an integer N, and produces the relation R where:
 - ◆ a tuple s with values a₁...a_n in attributes A₁, ..., A_n belongs to R(t) iff there exists <s,t'> ∈ S, with t' ≤ t, such that t' is among the N largest timestamps of elements in S with those values in those attributes

Partitioned windows

- What if we wanted to have *the last measurement for each car?*
 - ◆ similar to *group by* each carID and then getting the last measurement for each group
- **S [Partition By A₁, ..., A_n Rows N]** takes a stream S, a set of attributes A₁, ..., A_n, and an integer N, and produces the relation R where:
 - ◆ a tuple s with values a₁...a_n in attributes A₁, ..., A_n belongs to R(t) iff there exists <s,t'> ∈ S, with t' ≤ t, such that t' is among the N largest timestamps of elements in S with those values in those attributes
- E.g. PosSpeedStr [Partition by carID Rows 1]

Partitioned windows

- What if we wanted to have *the last measurement for each car?*
 - ◆ similar to *group by* each carID and then getting the last measurement for each group
- **S [Partition By A₁, ..., A_n Rows N]** takes a stream S, a set of attributes A₁, ..., A_n, and an integer N, and produces the relation R where:
 - ◆ a tuple s with values a₁...a_n in attributes A₁, ..., A_n belongs to R(t) iff there exists <s,t'> ∈ S, with t' ≤ t, such that t' is among the N largest timestamps of elements in S with those values in those attributes
- E.g. **PossSpeedStr [Partition by carID Rows 1]**
 - ◆ denotes, at any instant, the latest position-speed measurement of each car

Windows with a slide parameter

- **What if we only want to look at the measures once every minute?**
 - ◆ A sliding windows that only changes once every minute, and, in a change, takes everything from the last minute

Windows with a slide parameter

- **What if we only want to look at the measures once every minute?**
 - ◆ A sliding windows that only changes once every minute, and, in a change, takes everything from the last minute
- **Windows can optionally contain a **slide** parameter, indicating the **granularity at which the window slides**.**
 - ◆ The slide parameter is a **time interval** for **time-based windows**
 - ◆ The slide parameter is a **positive integer** for **row-based** and **partitioned windows**

Windows with a slide parameter

- **What if we only want to look at the measures once every minute?**
 - ◆ A sliding windows that only changes once every minute, and, in a change, takes everything from the last minute
- **Windows can optionally contain a **slide** parameter, indicating the **granularity at which the window slides**.**
 - ◆ The slide parameter is a **time interval** for **time-based windows**
 - ◆ The slide parameter is a **positive integer** for **row-based** and **partitioned windows**
- A **time-based window** over stream S with **window size T** and **slide parameter L** is denoted as
 - ◆ S [Range T Slide L]

Windows with a slide parameter

- What if we only want to look at the measures once every minute?
- S [**Range** T **Slide** L] produces relation R , where:

$$R(t) = \begin{cases} \{\} & \text{if } t < L - 1 \\ \{s | \langle s, t' \rangle \in S \wedge \max(t_{start} - T, 0) \leq t' \leq t_{start}\} & \text{otherwise} \end{cases}$$

where t_{start} is the largest multiple of L smaller than t

Windows with a slide parameter

- What if we only want to look at the measures once every minute?
- S [**Range** T **Slide** L] produces relation R , where:

$$R(t) = \begin{cases} \{\} & \text{if } t < L - 1 \\ \{s | \langle s, t' \rangle \in S \wedge \max(t_{start} - T, 0) \leq t' \leq t_{start}\} & \text{otherwise} \end{cases}$$

where t_{start} is the largest multiple of L smaller than t

- So it gets the elements with a time-stamp between t_{start} and $t_{end} = \max(t_{start} - T, 0)$

Windows with a slide parameter

- What if we only want to look at the measures once every minute?
- **S [Range T Slide L]** produces relation **R**, where:

$$R(t) = \begin{cases} \{\} & \text{if } t < L - 1 \\ \{s | \langle s, t' \rangle \in S \wedge \max(t_{start} - T, 0) \leq t' \leq t_{start}\} & \text{otherwise} \end{cases}$$

where t_{start} is the largest multiple of L smaller than t

- So it gets the elements with a time-stamp between t_{start} and $t_{end} = \max(t_{start} - T, 0)$
- E.g. **PosSpeedStr [Range 30s Slide 1min]**
 - ◆ every minute, get the measurements from the last clock half minute
 - up to second 59, it outputs nothing
 - from time 60s to 119s, contains the measurement made in the 2nd half minute (from second 31 to 60)
 - from time 120s to 179s, contains the measurements made in the 4th half minute

Windows with a slide parameter

- What if we only want to look at the measures once every minute?
- **S [Range T Slide L]** produces relation **R**, where:

$$R(t) = \begin{cases} \{\} & \text{if } t < L - 1 \\ \{s | \langle s, t' \rangle \in S \wedge \max(t_{start} - T, 0) \leq t' \leq t_{start}\} & \text{otherwise} \end{cases}$$

where t_{start} is the largest multiple of L smaller than t

- So it gets the elements with a time-stamp between t_{start} and $t_{end} = \max(t_{start} - T, 0)$
- E.g. **PosSpeedStr [Range 30s Slide 1min]**

- ◆ every minute, get the measurements from the last clock half minute
 - up to second 59, it outputs nothing
 - from time 60s to 119s, contains the measurement made in the 2nd half minute (from second 31 to 60)
 - from time 120s to 179s, contains the measurements made in the 4th half minute



Windows with a slide parameter

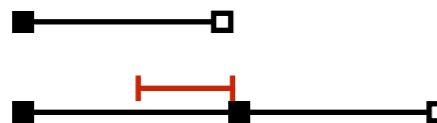
- What if we only want to look at the measures once every minute?
- S [**Range** T **Slide** L] produces relation R , where:

$$R(t) = \begin{cases} \{\} & \text{if } t < L - 1 \\ \{s | \langle s, t' \rangle \in S \wedge \max(t_{start} - T, 0) \leq t' \leq t_{start}\} & \text{otherwise} \end{cases}$$

where t_{start} is the largest multiple of L smaller than t

- So it gets the elements with a time-stamp between t_{start} and $t_{end} = \max(t_{start} - T, 0)$
- E.g. **PosSpeedStr** [**Range** 30s **Slide** 1min]

- ◆ every minute, get the measurements from the last clock half minute
 - up to second 59, it outputs nothing
 - from time 60s to 119s, contains the measurement made in the 2nd half minute (from second 31 to 60)
 - from time 120s to 179s, contains the measurements made in the 4th half minute



Windows with a slide parameter

- What if we only want to look at the measures once every minute?
- S [**Range** T **Slide** L] produces relation R , where:

$$R(t) = \begin{cases} \{\} & \text{if } t < L - 1 \\ \{s | \langle s, t' \rangle \in S \wedge \max(t_{start} - T, 0) \leq t' \leq t_{start}\} & \text{otherwise} \end{cases}$$

where t_{start} is the largest multiple of L smaller than t

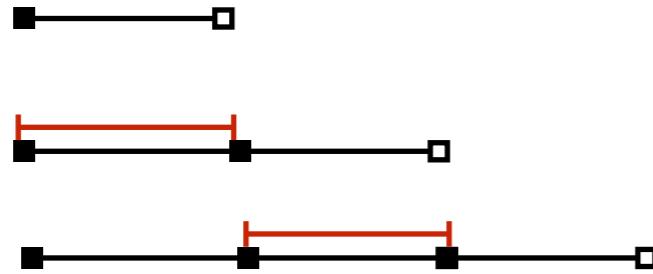
- So it gets the elements with a time-stamp between t_{start} and $t_{end} = \max(t_{start} - T, 0)$
- E.g. **PosSpeedStr** [**Range** 30s **Slide** 1min]

- ◆ every minute, get the measurements from the last clock half minute
 - up to second 59, it outputs nothing
 - from time 60s to 119s, contains the measurement made in the 2nd half minute (from second 31 to 60)
 - from time 120s to 179s, contains the measurements made in the 4th half minute



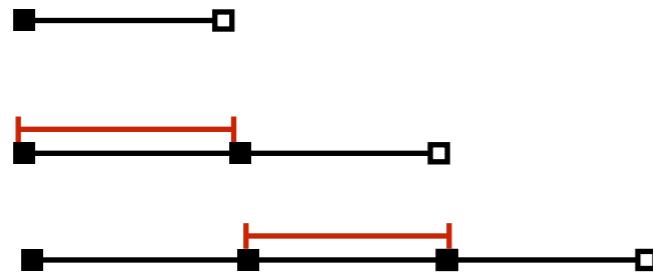
Windows with a slide parameter

- This type of window in which $L = T$ has been referred to as a **tumbling** window



Windows with a slide parameter

- This type of window in which $L = T$ has been referred to as a **tumbling** window



- Tuple-based and **partitioned windows** with a **slide parameter** are defined analogously.

CQL's **r2s** operators

- They take a **varying-time relation R** , and build a **stream S** based on the changes that occur in R over the time
- Three r2s operators

CQL's **r2s** operators

- They take a **varying-time relation R** , and build a **stream S** based on the changes that occur in R over the time
- Three **r2s** operators
 - ◆ **Istream** (*insert stream*): the stream “registers” the **insertions** made in the relation

CQL's **r2s** operators

- They take a **varying-time relation R** , and build a **stream S** based on the changes that occur in R over the time
- Three **r2s** operators
 - ◆ **Istream (insert stream)**: the stream “registers” the **insertions** made in the relation
 - ◆ **Dstream (delete stream)**: “registers” the **deletions**

CQL's **r2s** operators

- They take a **varying-time relation R** , and build a **stream S** based on the changes that occur in R over the time
- Three **r2s** operators
 - ◆ **Istream (insert stream)**: the stream “registers” the **insertions** made in the relation
 - ◆ **Dstream (delete stream)**: “registers” the **deletions**
 - ◆ **Rstream (relation stream)**: at each instant, outputs to the stream **all the tuples** that are in the relation

Insert stream

$$\text{Istream}(R) = \bigcup_{t \geq 0} ((R(t) - R(t-1)) \times \{t\})$$

- ◆ I.e. S has $\langle s, t \rangle$ whenever s is in $R(t)$ and was not in $R(t-1)$

Insert stream

$$\text{Istream}(R) = \bigcup_{t \geq 0} ((R(t) - R(t-1)) \times \{t\})$$

- ◆ I.e. S has $\langle s, t \rangle$ whenever s is in $R(t)$ and was not in $R(t-1)$
- **Example: Output in a stream the id and position of a car whenever that car is detected with a speed greater than 150**

```
select Istream(carID, position)
from PosSpeedStr [Range Unbounded]
where speed >= 150;
```

Insert stream

$$\text{Istream}(R) = \bigcup_{t \geq 0} ((R(t) - R(t-1)) \times \{t\})$$

- ◆ I.e. S has $\langle s, t \rangle$ whenever s is in $R(t)$ and was not in $R(t-1)$
- **Example: Output in a stream the id and position of a car whenever that car is detected with a speed greater than 150**

```
select Istream(carID, position)
from PosSpeedStr [Range Unbounded]
where speed >= 150;
```

- Note that, this query has as **input a stream** and as **output another stream**, and continuously produces the output upon what it receives from the input - i.e., it is a continuous query

Delete stream

$$\text{Dstream}(R) = \bigcup_{t \geq 0} ((R(t-1) - R(t)) \times \{t\})$$

- ◆ I.e. S has $\langle s, t \rangle$ whenever s was in $R(t-1)$ and is no longer in $R(t)$

Delete stream

$$\text{Dstream}(R) = \bigcup_{t \geq 0} ((R(t-1) - R(t)) \times \{t\})$$

- ◆ I.e. S has $\langle s,t \rangle$ whenever s was in $R(t-1)$ and is no longer in $R(t)$
- **Example: Output in a stream the ids of cars that exit the highway**

```
select Dstream(carID)
      from PosSpeedStr [Range 30s];
```

- ◆ An id of a car is output in the stream whenever the car reported its position and speed before 30 seconds ago, and in the last 30 seconds did not send anything

Relation stream

$$\text{Rstream}(R) = \bigcup_{t \geq 0} (R(t) \times \{t\})$$

- ◆ I.e. S has $\langle s, t \rangle$ whenever s is $R(t)$ - at any instant, outputs to the stream the whole content of the relation

Relation stream

$$\text{Rstream}(R) = \bigcup_{t \geq 0} (R(t) \times \{t\})$$

- ◆ I.e. S has $\langle s, t \rangle$ whenever s is $R(t)$ - at any instant, outputs to the stream the whole content of the relation
- **Example: What does this query compute?**

```
select Rstream(carID, position)
from PosSpeedStr [Now]
where speed >= 150;
```

Relation stream

$$\text{Rstream}(R) = \bigcup_{t \geq 0} (R(t) \times \{t\})$$

- ◆ I.e. S has $\langle s, t \rangle$ whenever s is $R(t)$ - at any instant, outputs to the stream the whole content of the relation
- Example: What does this query compute?

```
select Rstream(carID, position)
from PosSpeedStr [Now]
where speed >= 150;
```

- Exactly the same as this other one!

```
select Istream(carID, position)
from PosSpeedStr [Range Unbounded]
where speed >= 150;
```

- We will see that this is a general equivalence result. This is the general form of a stream filter

Shortcuts

- To makes things easier to write, there are some shortcut and default assumptions (though this is not usual in SQL-like languages)

1. When a stream is referred in a place where a relation is expected (e.g. in the from clause) assume an unbounded range
2. If one has a monotonic query based on streams, assume that the result should be a stream, and use the Istream operator

Shortcuts

- To makes things easier to write, there are some shortcut and default assumptions (though this is not usual in SQL-like languages)

1. When a stream is referred in a place where a relation is expected (e.g. in the from clause) assume an unbounded range
2. If one has a monotonic query based on streams, assume that the result should be a stream, and use the Istream operator

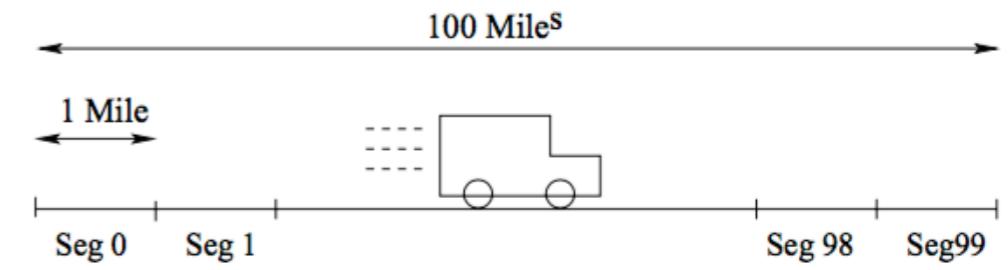
- Example: With this, one can simply write

```
select carID, position  
from PosSpeedStr  
where speed ≥ 150;
```

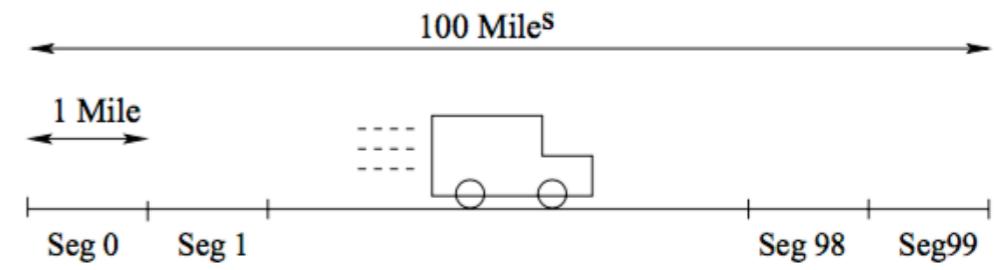
- ◆ to denote

```
select Istream(carID, position)  
from PosSpeedStr [Range Unbounded]  
where speed ≥ 150;
```

Simplified Linear Road (details)

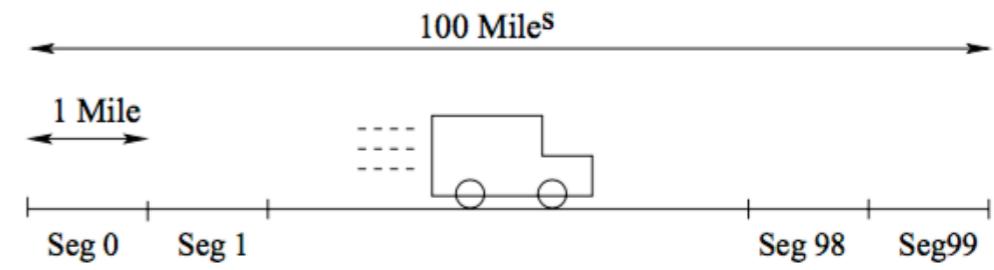


Simplified Linear Road (details)



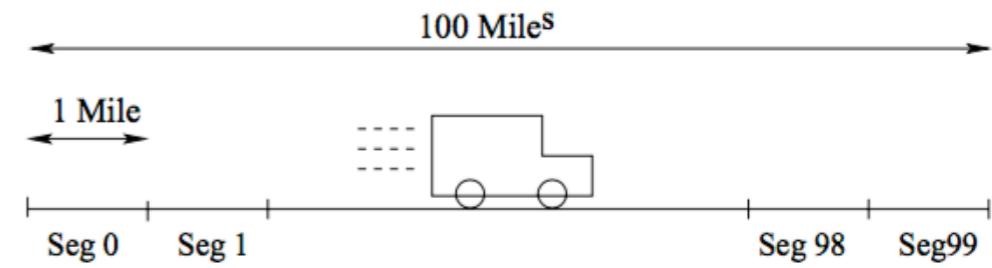
- The road is divided into 100 segments (of Mile each)

Simplified Linear Road (details)



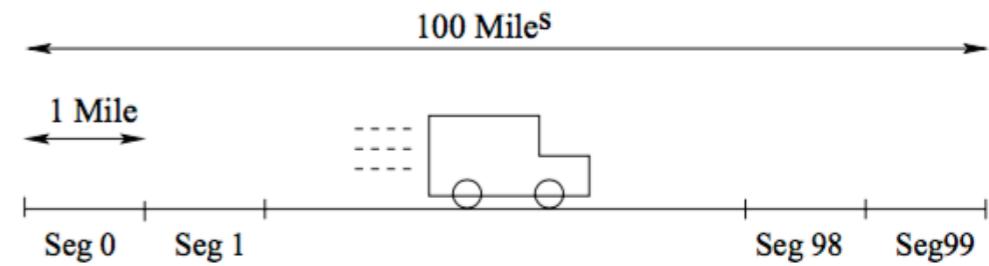
- The road is divided into 100 segments (of Mile each)
- The input stream has tuples with **carID**, **position**, and **speed**, every 30s
 - ◆ The position given as the distance in feet from the beginning of the road

Simplified Linear Road (details)



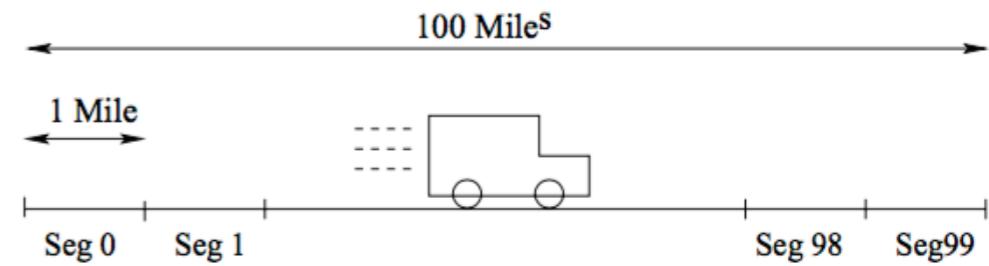
- The road is divided into 100 segments (of Mile each)
- The **input stream** has tuples with **carID**, **position**, and **speed**, every 30s
 - ◆ The position given as the distance in feet from the beginning of the road
- The **output stream** has tuples with **carID** and **price to pay for a segment**, any time a **car** enters that segment

Simplified Linear Road (details)



- The road is divided into 100 segments (of Mile each)
- The **input stream has tuples with carID, position, and speed, every 30s**
 - ◆ The position given as the distance in feet from the beginning of the road
- The **output stream has tuples with carID and price to pay for a segment, any time a car enters that segment**
- The total price to pay for a segment is as follows:
 - ◆ a non-congested segment pays nothing
 - ◆ a congested segment pays $2 \times (\text{numCars}-50)^2$, where *numCars* is the number of cars in the segment at the moment
 - ◆ a seg. is congested if the average speed of all cars in the last 5 minutes is less than 40 Miles/h

Simplified Linear Road (details)



- The road is divided into 100 segments (of Mile each)
- The **input stream has tuples with carID, position, and speed, every 30s**
 - ◆ The position given as the distance in feet from the beginning of the road
- The **output stream has tuples with carID and price to pay for a segment, any time a car enters that segment**
- The total price to pay for a segment is as follows:
 - ◆ a non-congested segment pays nothing
 - ◆ a congested segment pays $2 \times (\text{numCars}-50)^2$, where numCars is the number of cars in the segment at the moment
 - ◆ a seg. is congested if the average speed of all cars in the last 5 minutes is less than 40 Miles/h
- Lets implement this in CQL!

- Start by defining a query **SegSpeedStr** that in the input stream **replaces the position by the number of the segment**

```
select Istream(carID, speed, position/5280 as segNo)  
from PosspeedStr [ Range Unbounded];
```

S ■ Start by defining a query **SegSpeedStr** that in the input stream **replaces the position by the number of the segment**

```
select Istream(carID, speed, position/5280 as segNo)  
from PosspeedStr [ Range Unbounded];
```

- S ■ Start by defining a query **SegSpeedStr** that in the input stream replaces the position by the number of the segment

```
select Istream(carID, speed, position/5280 as segNo)  
from PosSpeedStr [ Range Unbounded];
```

- Define **ActCarSegRel** as the relation with the current segments of cars in the highway

```
select carID, segNo  
from SegSpeedStr [ Range 30s];
```

- S ■ Start by defining a query **SegSpeedStr** that in the input stream replaces the position by the number of the segment

```
select Istream(carID, speed, position/5280 as segNo)  
from PosSpeedStr [ Range Unbounded];
```

- R ■ Define **ActCarSegRel** as the relation with the current segments of cars in the highway

```
select carID, segNo  
from SegSpeedStr [ Range 30s];
```

- S ■ Start by defining a query **SegSpeedStr** that in the input stream **replaces the position by the number of the segment****

```
select Istream(carID, speed, position/5280 as segNo)  
from PosSpeedStr [ Range Unbounded];
```

- R ■ Define **ActCarSegRel** as the relation with the **current segments of cars** in the highway**

```
select carID, segNo  
from SegSpeedStr [ Range 30s];
```

- Define the stream **CarSegEntryStr** signalling cars that enter the highway**

```
select Istream(*)  
from ActCarSegRel;
```

- S ■ Start by defining a query **SegSpeedStr** that in the input stream replaces the position by the number of the segment

```
select Istream(carID, speed, position/5280 as segNo)  
from PosSpeedStr [ Range Unbounded];
```

- R ■ Define **ActCarSegRel** as the relation with the current segments of cars in the highway

```
select carID, segNo  
from SegSpeedStr [ Range 30s];
```

- S ■ Define the stream **CarSegEntryStr** signalling cars that enter the highway

```
select Istream(*)  
from ActCarSegRel;
```

Linear Road in CQL

- Define **CongSegRel** as the relation containing, at any instant, this **list of all congested segments**

Linear Road in CQL

- Define **CongSegRel** as the relation containing, at any instant, this **list of all congested segments**

a seg. is congested if the **average speed** of all cars in the **last 5 minutes** is less than 40 Miles/h

Linear Road in CQL

R ■ Define **CongSegRel** as the relation containing, at any instant, this **list of all congested segments**

```
select segNo  
from SegSpeedStr [Range 5min]           (carID, speed, segNo)  
group by segNo  
having avg(speed) < 40;
```

a seg. is congested if the **average speed** of all cars in the **last 5 minutes** is less than 40 Miles/h

Linear Road in CQL

R ■ Define **CongSegRel** as the relation containing, at any instant, this list of all congested segments

```
select segNo  
from SegSpeedStr [Range 5min]           (carID, speed, segNo)  
group by segNo  
having avg(speed) < 40;
```

a seg. is congested if the **average speed** of all cars in the **last 5 minutes** is less than 40 Miles/h

R ■ Define **SegVolRel** as the relation with the current **number of cars** in each segment

Linear Road in CQL

R ■ Define **CongSegRel** as the relation containing, at any instant, this list of all congested segments

a seg. is congested if the **average speed** of all cars in the **last 5 minutes** is less than 40 Miles/h

```
select segNo  
from SegSpeedStr [Range 5min]           (carID, speed, segNo)  
group by segNo  
having avg(speed) < 40;
```

R ■ Define **SegVolRel** as the relation with the current **number of cars** in each segment

```
select segNo, count(carID) as numCars  
from ActCarSegRel                      (segNo, numCars)  
group by segNo
```

Linear Road in CQL

PosSpeedStr(carID, speed, segNo)

Linear Road in CQL

PosSpeedStr(carID, speed, segNo)

- **SegSpeedStr - (carID, speed, segNo)**

Linear Road in CQL

PosSpeedStr(carID, speed, segNo)

- **SegSpeedStr - (carID, speed, segNo)**
- **ActCarSegRel - current segments of cars - (carID, speed, segNo) - Last 30s**

Linear Road in CQL

PosSpeedStr(carID, speed, segNo)

- **SegSpeedStr - (carID, speed, segNo)**
- **ActCarSegRel - current segments of cars - (carID, speed, segNo) - Last 30s**
- **CarSegEntryStr - signalling cars that enter the highway - (carID, speed, segNo)**

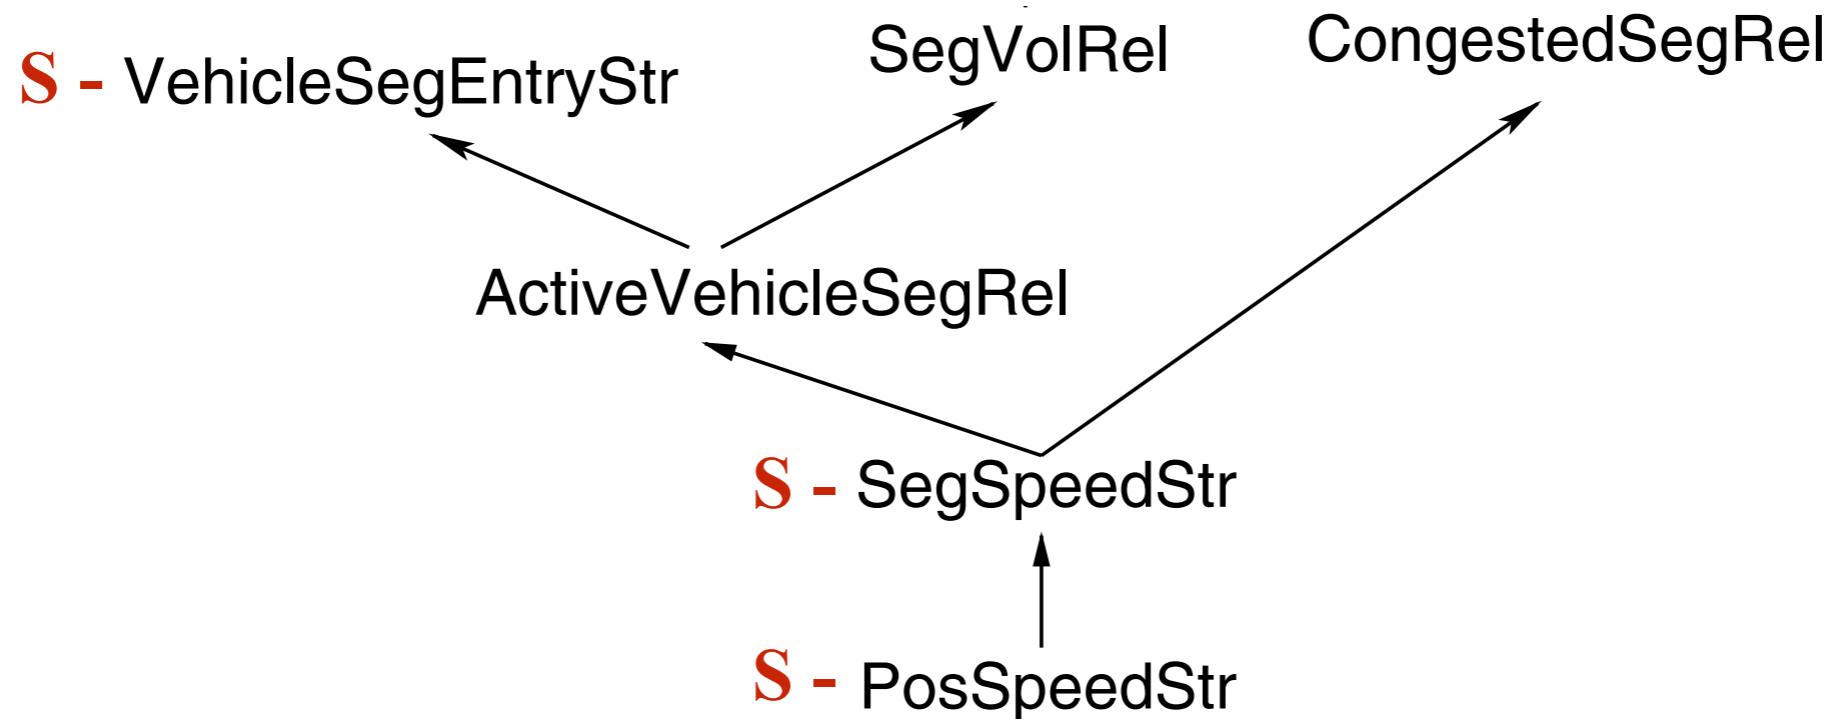
- **SegSpeedStr - (carID, speed, segNo)**
- **ActCarSegRel - current segments of cars - (carID, speed, segNo) - Last 30s**
- **CarSegEntryStr - signalling cars that enter the highway - (carID, speed, segNo)**
- **CongestedSegRel - at any instant, this list of all congested segments - (segNo)**

- **SegSpeedStr - (carID, speed, segNo)**
- **ActCarSegRel - current segments of cars - (carID, speed, segNo) - Last 30s**
- **CarSegEntryStr - signalling cars that enter the highway - (carID, speed, segNo)**
- **CongestedSegRel - at any instant, this list of all congested segments - (segNo)**
- **SegVolRel - the current number of cars in each segment - (segNo, numCars)**

Linear Road in CQL

PosSpeedStr(carID, speed, segNo)

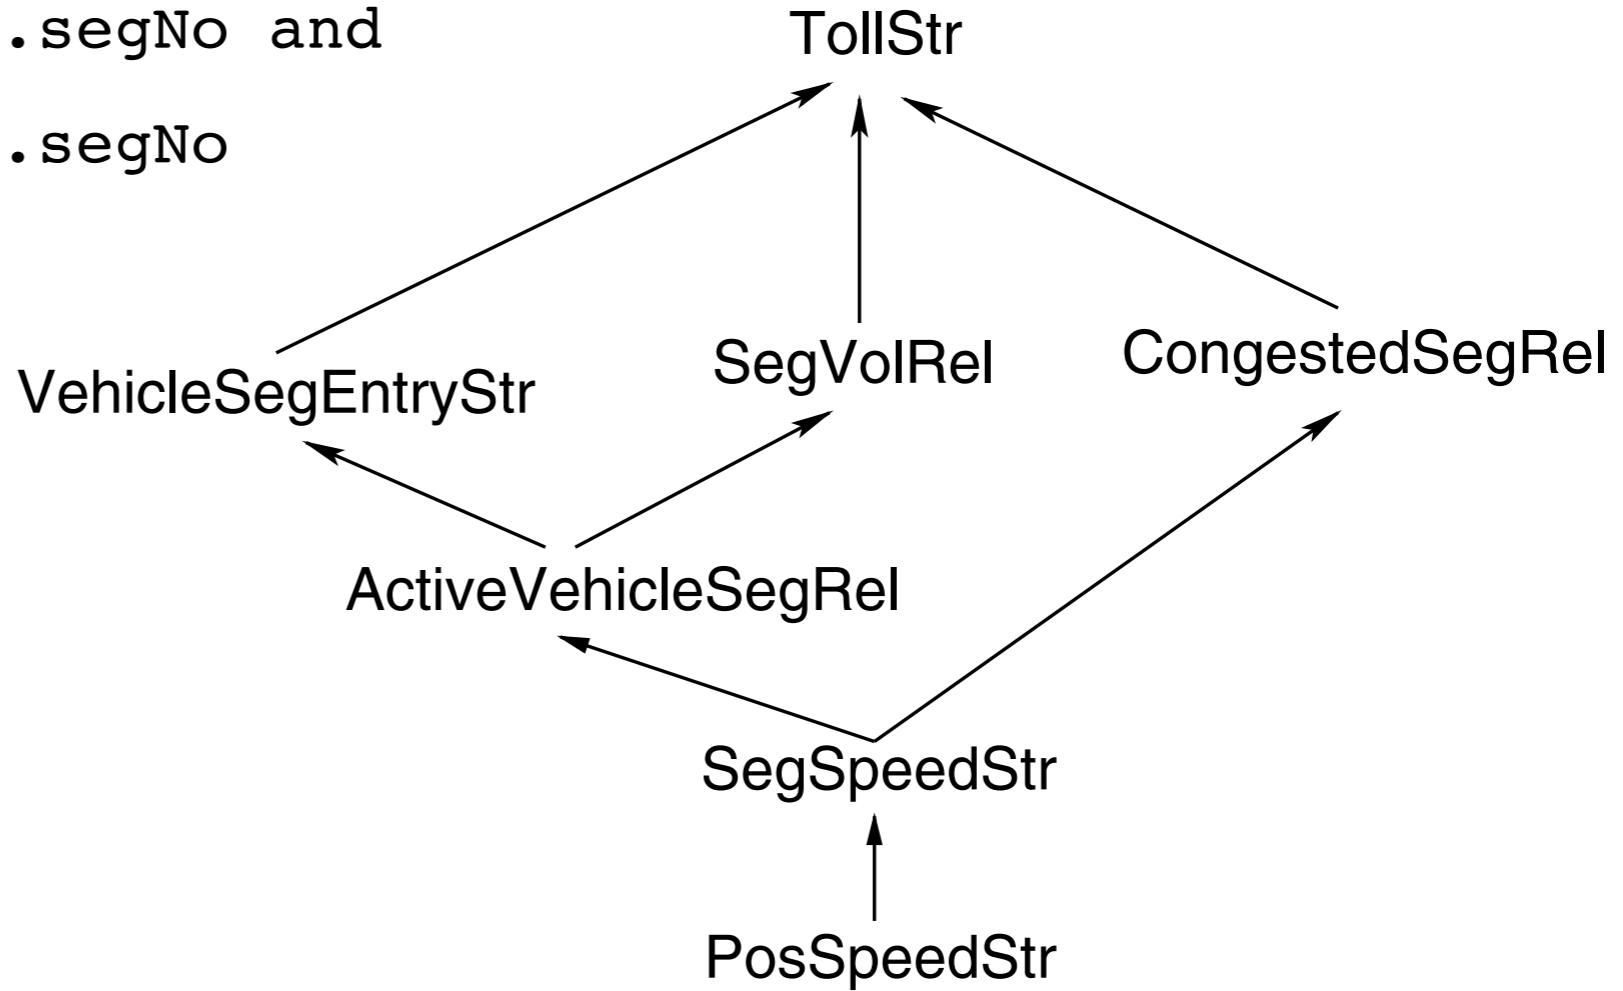
- **SegSpeedStr - (carID, speed, segNo)**
- **ActCarSegRel - current segments of cars - (carID, speed, segNo) - Last 30s**
- **CarSegEntryStr - signalling cars that enter the highway - (carID, speed, segNo)**
- **CongestedSegRel - at any instant, this list of all congested segments - (segNo)**
- **SegVolRel - the current number of cars in each segment - (segNo, numCars)**



Linear Road in CQL

- The query that outputs the desired stream is:

```
select Rstream(E.carID, 2*(V.numCars-50)**2)
from CarSegEntryStr [Now] as E,
     CongestedSegRel as C, SegVolRel as V
where E(segNo) = C(segNo) and
      C(segNo) = V(segNo)
```



The complete query (just for the sake of completeness)

```
select Rstream(E.carID,2*(V.numCars-50)**2)
from
( select Istream(*)
  from
( select carID, segNo
    from
      ( select Istream(carID, speed,position/1000 as segNo) from PosSpeedStr [Range Unbounded]
        ) [Range 30s]
    )
  ) [Now] as E,
( select segNo
  from
    ( select Istream(*)
      from
        ( select Istream(carID, speed,position/1000 as segNo) from PosSpeedStr [Range Unbounded]
          ) [Range 30s]
        ) [Range 5min]
    group by segNo
    having avg(speed) < 40
  ) as C,
( select segNo, count(carID) as numCar
  from (select carID, segNo
        from (select Istream(carID, speed,position/1000 as segNo) from PosSpeedStr [Range Unbounded]
          )[Range 30s] )
    group by segNo
  ) as V
where E(segNo = C(segNo) and C(segNo = V(segNo);
```

Some variants

■ Additional **fixed price per segment** (when segment is congested)

```
select Rstream(E.carID,P.price+2*(V.numCars-50)**2 )  
from CarSegEntryStr [Now] as E,  
      CongSegRel as C, SegVolRel as V, SegPrice as P  
where E.segNo = C.segNo and  
      C.segNo = V.segNo and V.segNo = P.segNo;
```

```
select Rstream(E.carID,2*(V.numCars-50)**2 )  
from CarSegEntryStr [Now] as E,  
      CongestedSegRel as C, SegVolRel as V  
where E.segNo = C.segNo and  
      C.segNo = V.segNo
```

Some variants

■ Charging the customer: having the prices being output upon exit of the cars

```
select Rstream(A.NIF, P.price+2*(V.numCars-50)**2)  
from CarSegEntryStr [Now] as E,  
    CongSegRel as C, SegVolRel as V,  
    SegPrice as P, Cars, Accounts as A  
where E.segNo = C.segNo and  
      C.segNo = V.segNo and V.segNo = P.segNo  
      E.carID = Cars.ID and Cars.owner = A.client;
```

Another variant

- Fixed price per segment (charged always), plus variable toll:

```
select Rstream(E.carID,P.price+2*(V.numCars-50)**2)  
from  
(CarSegEntryStr [Now] E inner join SegPrice S  
on(E(segNo = S(segNo))  
left outer join  
(CongSegRel C inner join SegVolRel V  
on(C(segNo = V(segNo))  
on(E(segNo = C(segNo))
```

Other variants

- If a car is detected in segment N and then in segment N+2, within a time frame of 30 minutes, and in-between it is never detected in segment N+1, then signal a possible failure

Other variants

- If a car is detected in segment N and then in segment N+2, within a time frame of 30 minutes, and in-between it is never detected in segment N+1, then signal a possible failure
 - ◆ This may be problematic...

Other variants

- If a car is detected in segment N and then in segment N+2, within a time frame of 30 minutes, and in-between it is never detected in segment N+1, then signal a possible failure
 - ◆ This may be problematic...
- If a car is detected in a segment N, and the last two times it was detected was in segment N+1, then warn the police (it may be driving in the wrong way)

Other variants

- If a car is detected in segment N and then in segment N+2, within a time frame of 30 minutes, and in-between it is never detected in segment N+1, then signal a possible failure
 - ◆ This may be problematic...
- If a car is detected in a segment N, and the last two times it was detected was in segment N+1, then warn the police (it may be driving in the wrong way)
 - ◆ Again this may be problematic
 - ◆ Note that when one applies a s2r operator, the relation loses the timestamps

Some common queries

- **Stream filters:**

- ◆ Either use
 - `Rstream` and `[Now]`, or
 - `Istream` and `[Range Unbounded]`

Some common queries

- **Stream filters:**
 - ◆ Either use
 - **Rstream** and [Now], or
 - **Istream** and [Range Unbounded]
 - ◆ Note the **Istream-Unbounded** window combination is the **default** for streams whose window specification is omitted

Some common queries

■ Stream-relation **joins**:

- ◆ When a stream is joined with a relation, it is usually **most meaningful to apply a Now window over the stream and an Rstream operator over the join result**
 - Rstream and [Now], or
 - Istream and [Range Unbounded]

Some common queries

■ Stream-relation **joins**:

- ◆ When a stream is joined with a relation, it is usually **most meaningful to apply a Now window over the stream and an Rstream operator over the join result**
 - **Rstream and [Now]**, or
 - **Istream and [Range Unbounded]**
- ◆ But these two ways may have different meanings!

Some common queries

■ Stream-relation joins:

- ◆ When a stream is joined with a relation, it is usually **most meaningful to apply a Now window over the stream and an Rstream operator over the join result**
 - Rstream and [Now], or
 - Istream and [Range Unbounded]
- ◆ But these two ways may have different meanings!

■ Example: Get prices for purchased items:

```
select Rstream(item.id, prices.price)  
from item [Now], prices  
where item.id=prices.id;
```

```
select Istream(item.id, prices.price)  
from item [Range Unbounded], prices  
where item.id=prices.id;
```

■ Differences?

Some common queries

■ Stream-relation joins:

- ◆ When a stream is joined with a relation, it is usually **most meaningful to apply a Now window over the stream and an Rstream operator over the join result**
 - Rstream and [Now], or
 - Istream and [Range Unbounded]
- ◆ But these two ways may have different meanings!

■ Example: Get prices for purchased items:

```
select Rstream(item.id, prices.price)
from item [Now], prices
where item.id=prices.id;
```

```
select Istream(item.id, prices.price)
from item [Range Unbounded], prices
where item.id=prices.id;
```

■ Differences?

- ◆ The 1st outputs an item with price, whenever a new item arrives

Some common queries

■ Stream-relation joins:

- ◆ When a stream is joined with a relation, it is usually **most meaningful to apply a Now window over the stream and an Rstream operator over the join result**
 - Rstream and [Now], or
 - Istream and [Range Unbounded]
- ◆ But these two ways may have different meanings!

■ Example: Get prices for purchased items:

```
select Rstream(item.id, prices.price)
from item [Now], prices
where item.id=prices.id;
```

```
select Istream(item.id, prices.price)
from item [Range Unbounded], prices
where item.id=prices.id;
```

■ Differences?

- ◆ The 1st outputs an item with price, whenever a new item arrives
- ◆ The 2nd does that plus, whenever a price changes, it outputs all the items purchased so far, with the new price

Some common queries

■ Sliding windows joins:

```
select Istream(*)  
from S1 [Rows 5], S2 [Rows 10]  
where S1.A=S2.A;
```

Some common queries

■ Sliding windows joins:

```
select Istream( * )  
from S1 [Rows 5], S2 [Rows 10]  
where S1.A=S2.A;
```

- ◆ A new tuple is produced whenever a tuple in the last 5 elements of S1 matches with one in the last 10 elements of S2

Some common queries

■ Sliding windows joins:

```
select Istream( * )  
from S1 [Rows 5], S2 [Rows 10]  
where S1.A=S2.A;
```

- ◆ A new tuple is produced whenever a tuple in the last 5 elements of S1 matches with one in the last 10 elements of S2
- But if both S1 and S2 can have duplicates in attribute A, it may work differently!

Some common queries

■ Sliding windows joins:

```
select Istream( * )  
from S1 [Rows 5], S2 [Rows 10]  
where S1.A=S2.A;
```

- ◆ A new tuple is produced whenever a tuple in the last 5 elements of S1 matches with one in the last 10 elements of S2
- **But if both S1 and S2 can have duplicates in attribute A, it may work differently!**
 - ◆ Suppose a new tuple of S2 is identical to the tuple 10 positions earlier

Some common queries

■ Sliding windows joins:

```
select Istream( * )  
from S1 [Rows 5], S2 [Rows 10]  
where S1.A=S2.A;
```

- ◆ A new tuple is produced whenever a tuple in the last 5 elements of S1 matches with one in the last 10 elements of S2
- But if both S1 and S2 can have duplicates in attribute A, it may work differently!
 - ◆ Suppose a new tuple of S2 is identical to the tuple 10 positions earlier
 - ◆ Then the new S2 tuple does not produce any join result tuples, even if it joins with one of the last 5 tuples of S1.

Stream aggregations

Aggregation produces relations, not streams, in CQL since aggregations are relation-to-relation operators

- Stream the value of aggregation whenever it changes:

```
select Istream(Count(*) )  
from PosSpeedStr [Range 1min];
```

- ◆ Counts the measurement over the previous minute, and outputs the count to the stream whenever it changes

Stream aggregations

Aggregation produces relations, not streams, in CQL since aggregations are relation-to-relation operators

- Stream the value of aggregation whenever it changes:

```
select Istream(Count(*))  
from PosSpeedStr [Range 1min];
```

- ◆ Counts the measurement over the previous minute, and outputs the count to the stream whenever it changes

- Stream the value of aggregation periodically:

```
select Istream(Count(*))  
from PosSpeedStr [Range 1min Slide 1min];
```

- ◆ Outputs the number of measurements over the last minute, once a minute

Stream aggregations

Aggregation produces relations, not streams, in CQL since aggregations are relation-to-relation operators

- Stream the value of aggregation whenever it changes:

```
select Istream(Count(*))  
from PosSpeedStr [Range 1min];
```

- ◆ Counts the measurement over the previous minute, and outputs the count to the stream whenever it changes

- Stream the value of aggregation periodically:

```
select Istream(Count(*))  
from PosSpeedStr [Range 1min Slide 1min];
```

- ◆ Outputs the number of measurements over the last minute, once a minute
- ◆ If the count does not change, then it output nothing!
 - It is possible, with a more complex query, also to output in that case.

Some equivalences

L is any **select-list**, S is any **stream** (including a subquery producing a stream), and C is any **condition**.

Some equivalences

L is any **select-list**, *S* is any **stream** (including a subquery producing a stream), and *C* is any **condition**.

■ Window reduction - the two are equivalent

```
select Istream(L)
```

```
from S [Range Unbounded]
```

```
where Cond;
```

```
select Rstream(L)
```

```
from S [Now]
```

```
where Cond;
```

- ◆ the execution of the query on the right should be more efficient (it doesn't require buffering of stream results)

Some equivalences

L is any **select-list**, S is any **stream** (including a subquery producing a stream), and C is any **condition**.

■ Window reduction - the two are equivalent

```
select Istream(L)
```

```
from S [Range Unbounded]
```

```
where Cond;
```

```
select Rstream(L)
```

```
from S [Now]
```

```
where Cond;
```

- ◆ the execution of the query on the right should be more efficient (it doesn't require buffering of stream results)

■ Filter-Window commutativity - the two are equivalent

```
(select L from S where C) [Range T]
```

```
select L from S [Range T] where C
```

- ◆ the execution of the 1st query should be more efficient, because it filters before applying the window
 - If the system uses a query strategy that materialises the window, materialising only the filtered tuples requires less steady-state memory

On CQL implementation

On CQL implementation

- **CQL was first implemented in the prototype system STREAM**
 - ◆ Currently there are other implementations (e.g. Oracle)
- **One of the nice things about declarative languages is that they hide quite a lot of the implementation details**
 - ◆ We can use them without having any idea of how things are implemented

On CQL implementation

- **CQL was first implemented in the prototype system STREAM**
 - ◆ Currently there are other implementations (e.g. Oracle)
- **One of the nice things about declarative languages is that they hide quite a lot of the implementation details**
 - ◆ We can use them without having any idea of how things are implemented
- **But, in the end, the declarative language must be implemented!**
 - ◆ If one wants to use the system efficiently, it is good to have an idea of how things are implemented, and run

CQL Query Plans

- When a CQL query is registered, a query plan is generated

CQL Query Plans

- When a CQL query is registered, a query plan is generated
 - ◆ The query plan is a **tree/graph** with **operators**, roughly corresponding to the operators in the query's algebraic expression (pretty much as in DBMSs)

CQL Query Plans

- When a CQL query is registered, a query plan is generated
 - ◆ The query plan is a **tree/graph** with **operators**, roughly corresponding to the operators in the query's algebraic expression (pretty much as in DBMSs)
 - ◆ The first plan is generated taking into account several simple heuristics
 - But while running, with concrete data, it may be adapted

CQL Query Plans

■ When a CQL query is registered, a query plan is generated

- ◆ The query plan is a **tree/graph** with **operators**, roughly corresponding to the operators in the query's algebraic expression (pretty much as in DBMSs)
- ◆ The first plan is generated taking into account several simple heuristics
 - But while running, with concrete data, it may be adapted
- ◆ A new query plan is then merged with existing ones
 - This way, **parts of the plan that are common to more than one query are only (continuously) executed once**
 - At least the parts of incoming streams, are usually common

CQL Query Execution

- **Query execution strategy is based mostly on *incremental processing*:**

CQL Query Execution

- **Query execution strategy is based mostly on *incremental processing*:**
 - ◆ **s2r** operators transform **new items in the input stream** into changes (**insertions and deletions**) in the **output relation**

CQL Query Execution

- **Query execution strategy is based mostly on *incremental processing*:**
 - ◆ **s2r** operators transform **new items in the input stream** into changes (**insertions and deletions**) in the **output relation**
 - ◆ **r2r** operators transform **changes** in the input relations into **changes** in the output relation.

CQL Query Execution

- **Query execution strategy is based mostly on *incremental processing*:**
 - ◆ **s2r** operators transform **new items in the input stream** into changes (**insertions and deletions**) in the **output relation**
 - ◆ **r2r** operators transform **changes** in the input relations into **changes** in the output relation.
 - ◆ **r2s** operators transform changes in the input relations into new items in the output stream

CQL Query Execution

- **Query execution strategy is based mostly on *incremental processing*:**
 - ◆ **s2r** operators transform **new items in the input stream** into changes (**insertions and deletions**) in the **output relation**
 - ◆ **r2r** operators transform **changes** in the input relations into **changes** in the output relation.
 - ◆ **r2s** operators transform changes in the input relations into new items in the output stream
- **We need components associated to operators, to:**
 - ◆ store **intermediate changes** - inserts and deletes (**queues**)
 - ◆ store **intermediate results** - e.g. with materialised windows (**synopses**)

Internal representation

- Both streams and relations are internally represented as sequences of *tagged tuples*

Internal representation

- Both streams and relations are internally represented as sequences of *tagged tuples*
- A tagged tuple is a tuple plus a tag with:
 - ◆ a timestamp
 - ◆ an indication whether it is an insertion or a deletion

Internal representation

- Both streams and relations are internally represented as sequences of *tagged tuples*
- A tagged tuple is a tuple plus a tag with:
 - ◆ a timestamp
 - ◆ an indication whether it is an insertion or a deletion
- Streams can only have **insertion** tags

Internal representation

- Both streams and relations are internally represented as sequences of *tagged tuples*
- A tagged tuple is a tuple plus a tag with:
 - ◆ a timestamp
 - ◆ an indication whether it is an insertion or a deletion
- Streams can only have **insertion** tags
- Relation can have both.
 - ◆ Relations are thus represented as a set of insert/delete operations that give rise to the current state of the relation

Queues and Synopses

- **A queue connects two operators O_i and O_o**

- ◆ At any time, contains a sequence of tagged tuples
 - ◆ The queue buffers the insertions and deletions output by O_i until they are processed by O_o

Queues and Synopses

■ A queue connects two operators O_i and O_o

- ◆ At any time, contains a sequence of tagged tuples
- ◆ The queue buffers the insertions and deletions output by O_i until they are processed by O_o

■ A synopsis is associated to one operator

- ◆ It stores an intermediate state needed by the operator
 - E.g. when joining two windowed streams, the join needs, at any time, to have the contents of the whole two windows
- ◆ Synopses mostly materialise the content of relations resulting from windows, or store intermediate summaries (aggregations) of data

Queues and Synopses

- **A queue connects two operators O_i and O_o**
 - ◆ At any time, contains a sequence of tagged tuples
 - ◆ The queue buffers the insertions and deletions output by O_i until they are processed by O_o
- **A synopsis is associated to one operator**
 - ◆ It stores an intermediate state needed by the operator
 - E.g. when joining two windowed streams, the join needs, at any time, to have the contents of the whole two windows
 - ◆ Synopses mostly materialise the content of relations resulting from windows, or store intermediate summaries (aggregations) of data
- **Two operators are always connected by a queue, but many operators need no synopsis**

Example plan

■ Q1

```
select B, max(A)  
from S1 [Rows 50000]  
group by B;
```

■ Q2

```
select Istream(*)  
from S1 [Rows 40000]  
S2 [Range 600s]  
where S1.A = S2.A;
```

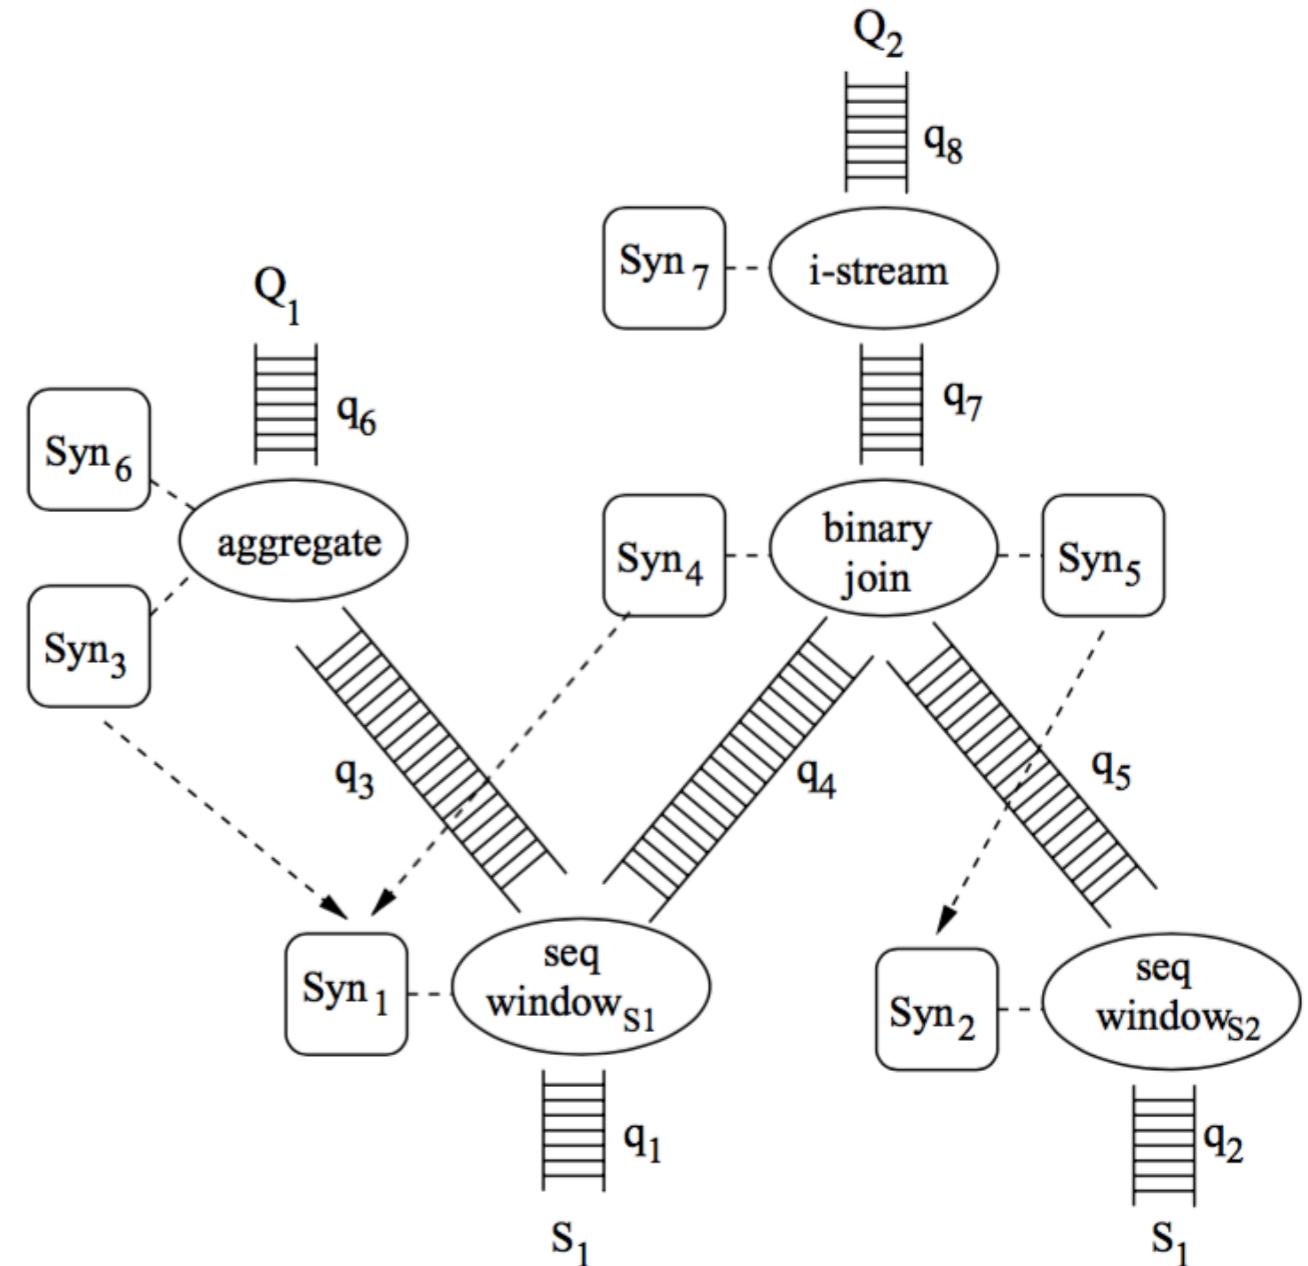
Example plan

■ Q1

```
select B, max(A)  
from S1 [Rows 50000]  
group by B;
```

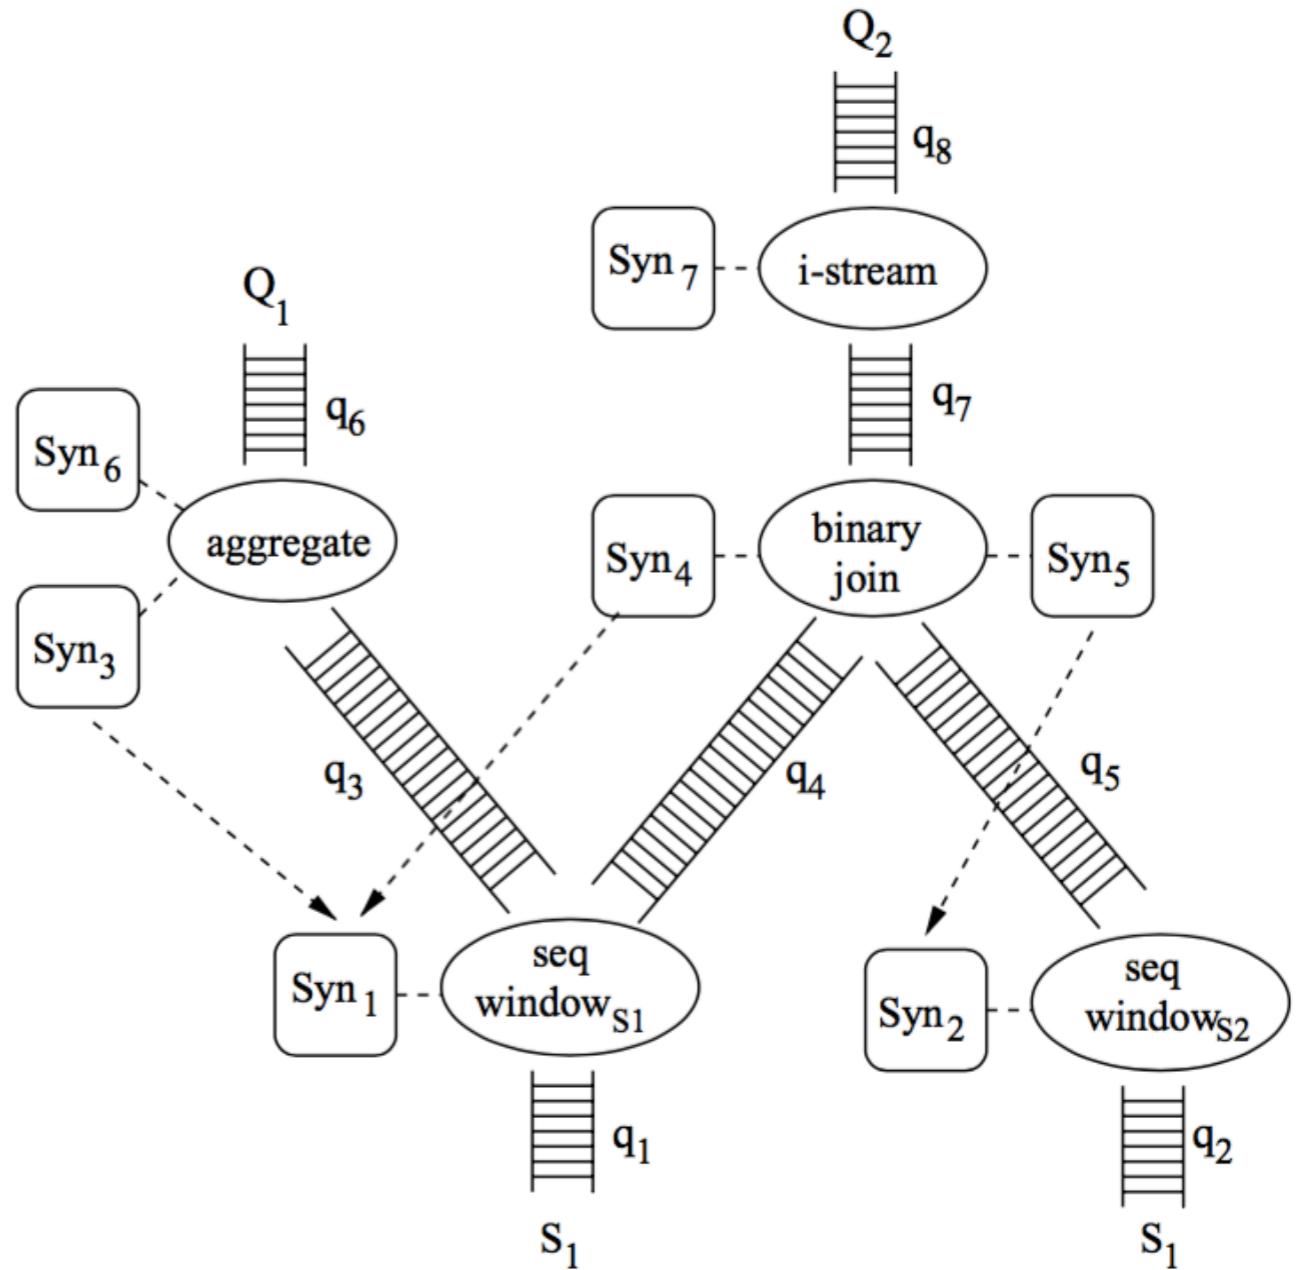
■ Q2

```
select Istream(*)  
from S1 [Rows 40000]  
S2 [Range 600s]  
where S1.A = S2.A;
```



Example plan

This plan contains



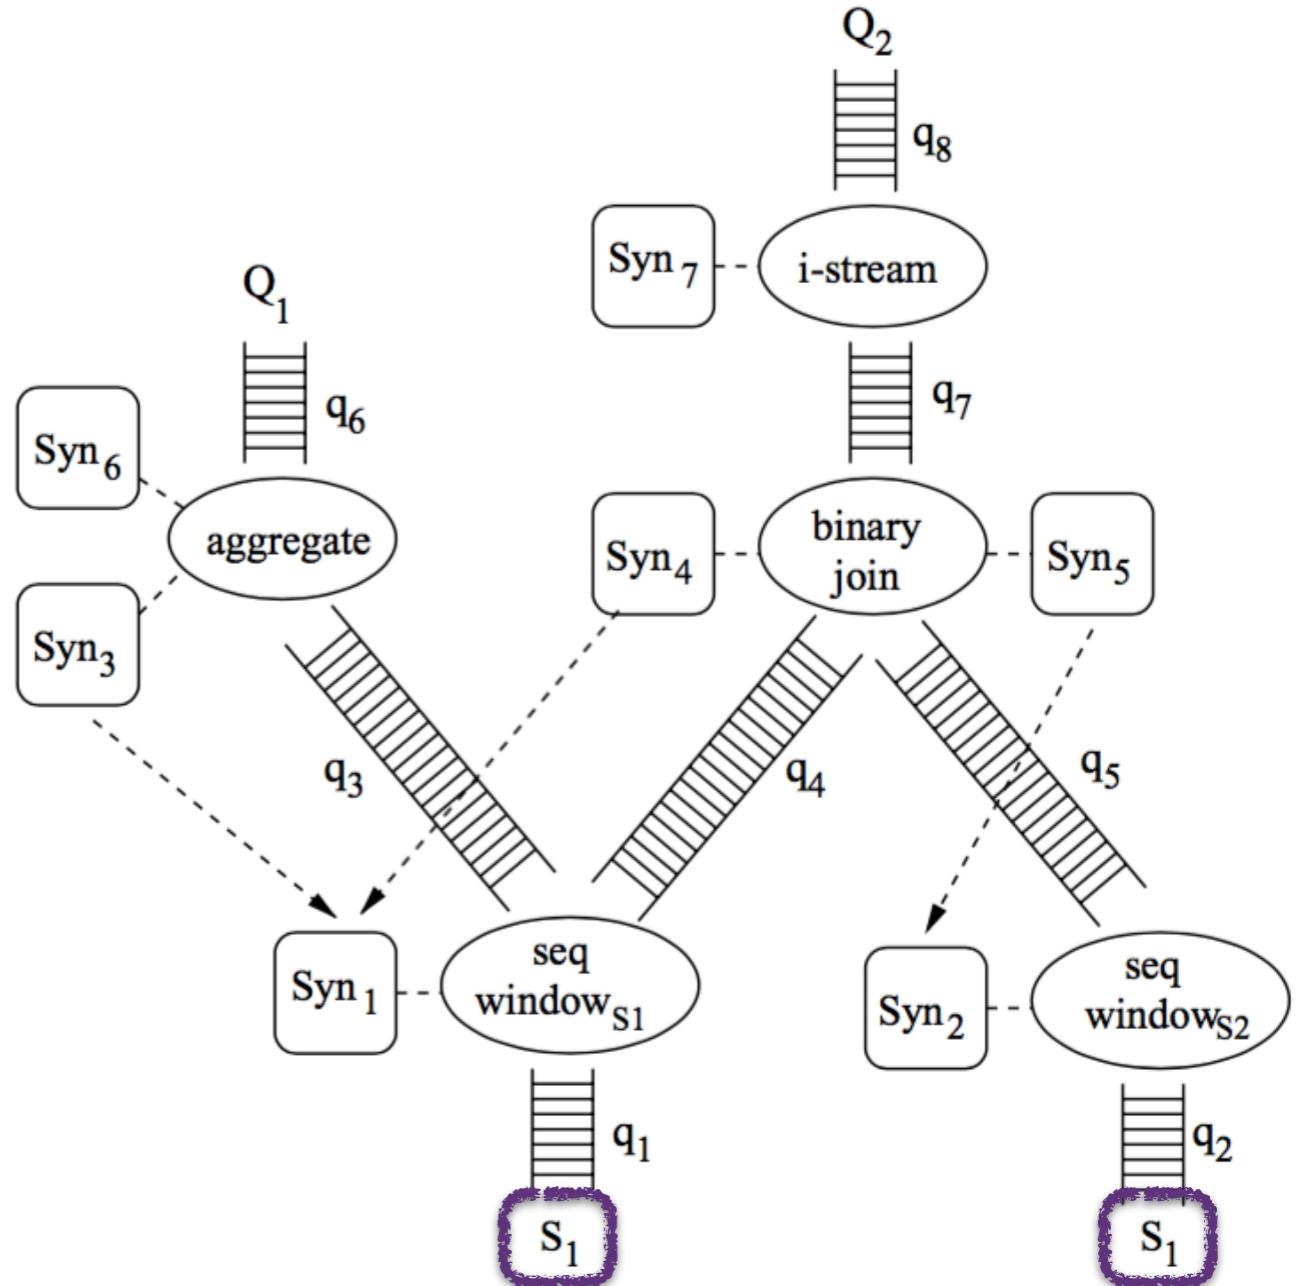
Q1: `select B, max(A)
from S1 [Rows 50000]
group by B;`

Q2: `select Istream(*)
from S1 [Rows 40000]
S2 [Range 600s]
where S1.A = S2.A;`

Example plan

This plan contains

- ◆ 2 input streams



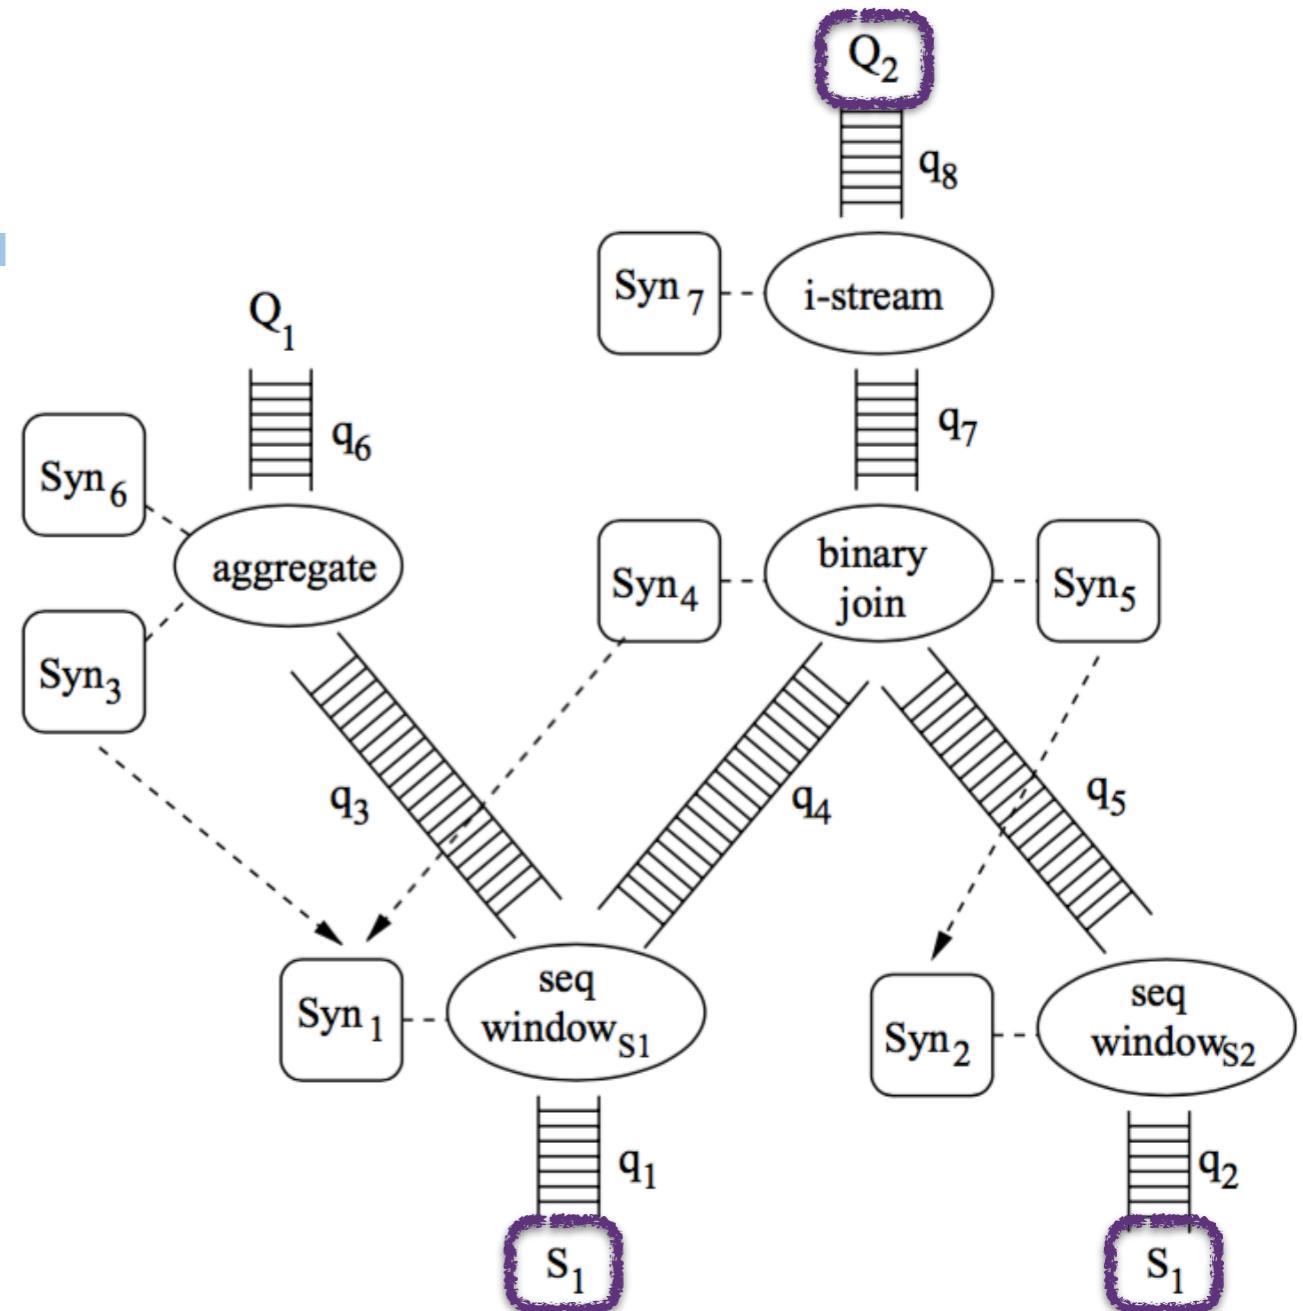
Q1: `select B, max(A)
from S1 [Rows 50000]
group by B;`

Q2: `select Istream(*)
from S1 [Rows 40000]
S2 [Range 600s]
where S1.A = S2.A;`

Example plan

This plan contains

- ◆ 2 input streams
- ◆ 1 output stream



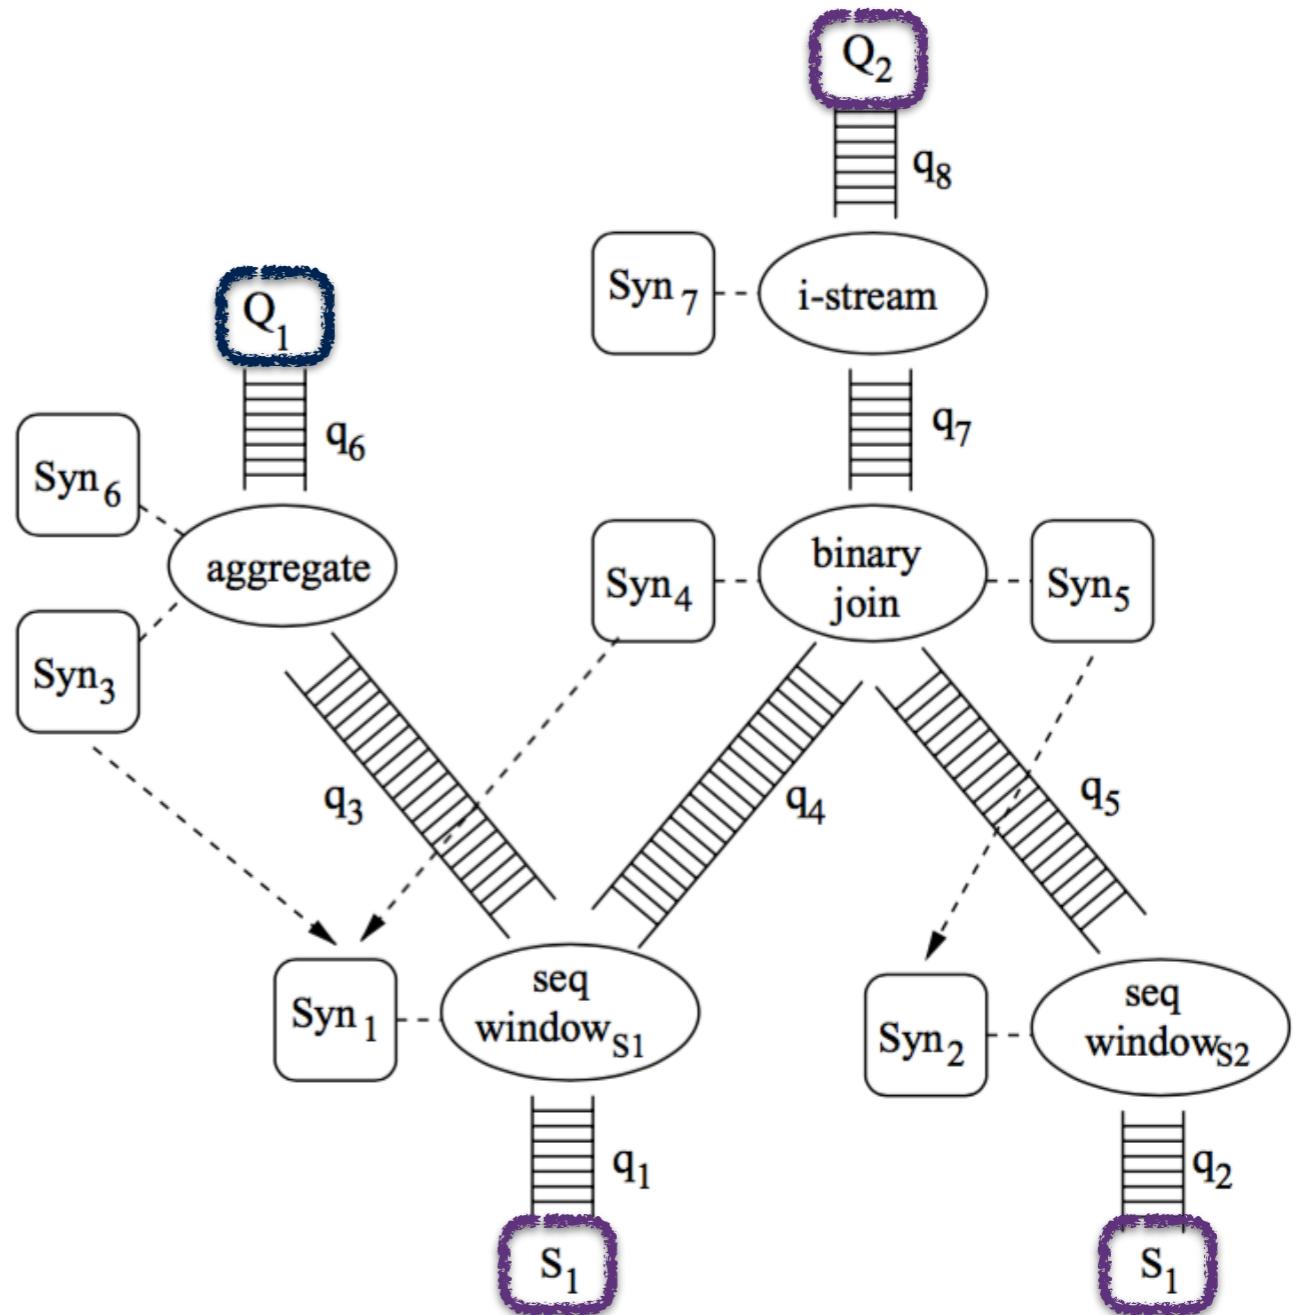
Q1: `select B, max(A)
from S1 [Rows 50000]
group by B;`

Q2: `select Istream(*)
from S1 [Rows 40000]
S2 [Range 600s]
where S1.A = S2.A;`

Example plan

This plan contains

- ◆ 2 input streams
- ◆ 1 output stream
- ◆ 1 output relation



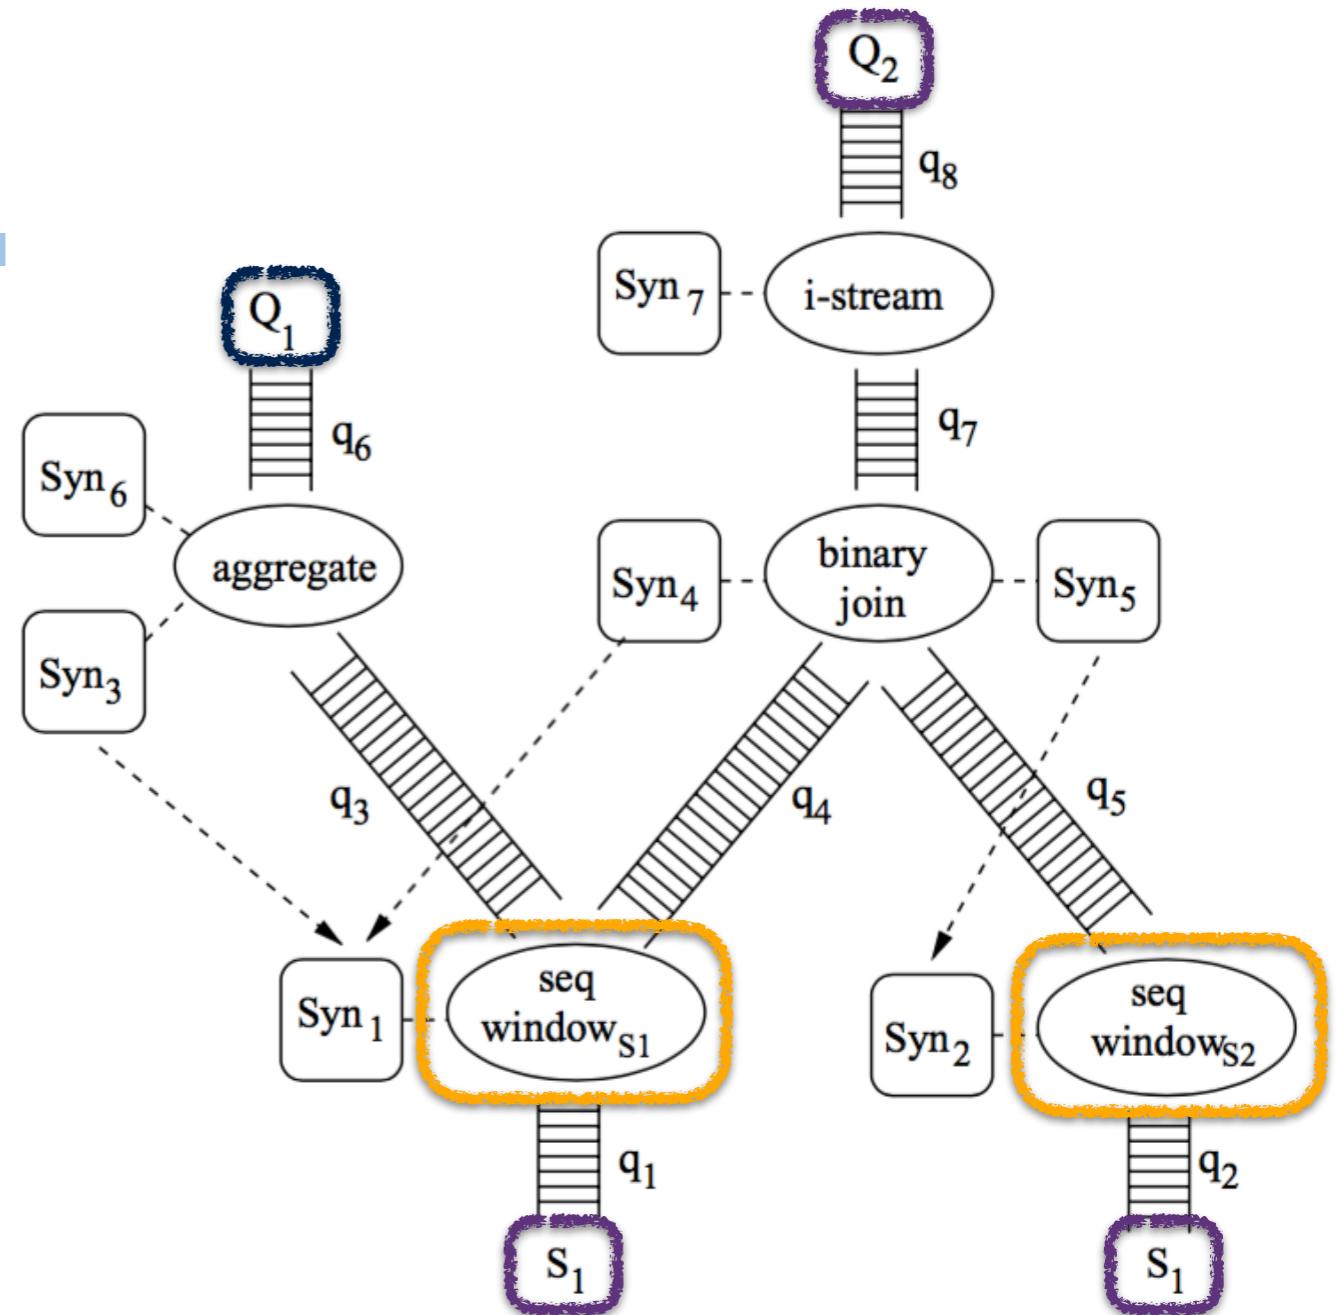
Q1: `select B, max(A)
from S1 [Rows 50000]
group by B;`

Q2: `select Istream(*)
from S1 [Rows 40000]
S2 [Range 600s]
where S1.A = S2.A;`

Example plan

This plan contains

- ◆ 2 input streams
- ◆ 1 output stream
- ◆ 1 output relation
- ◆ 2 s2r operators



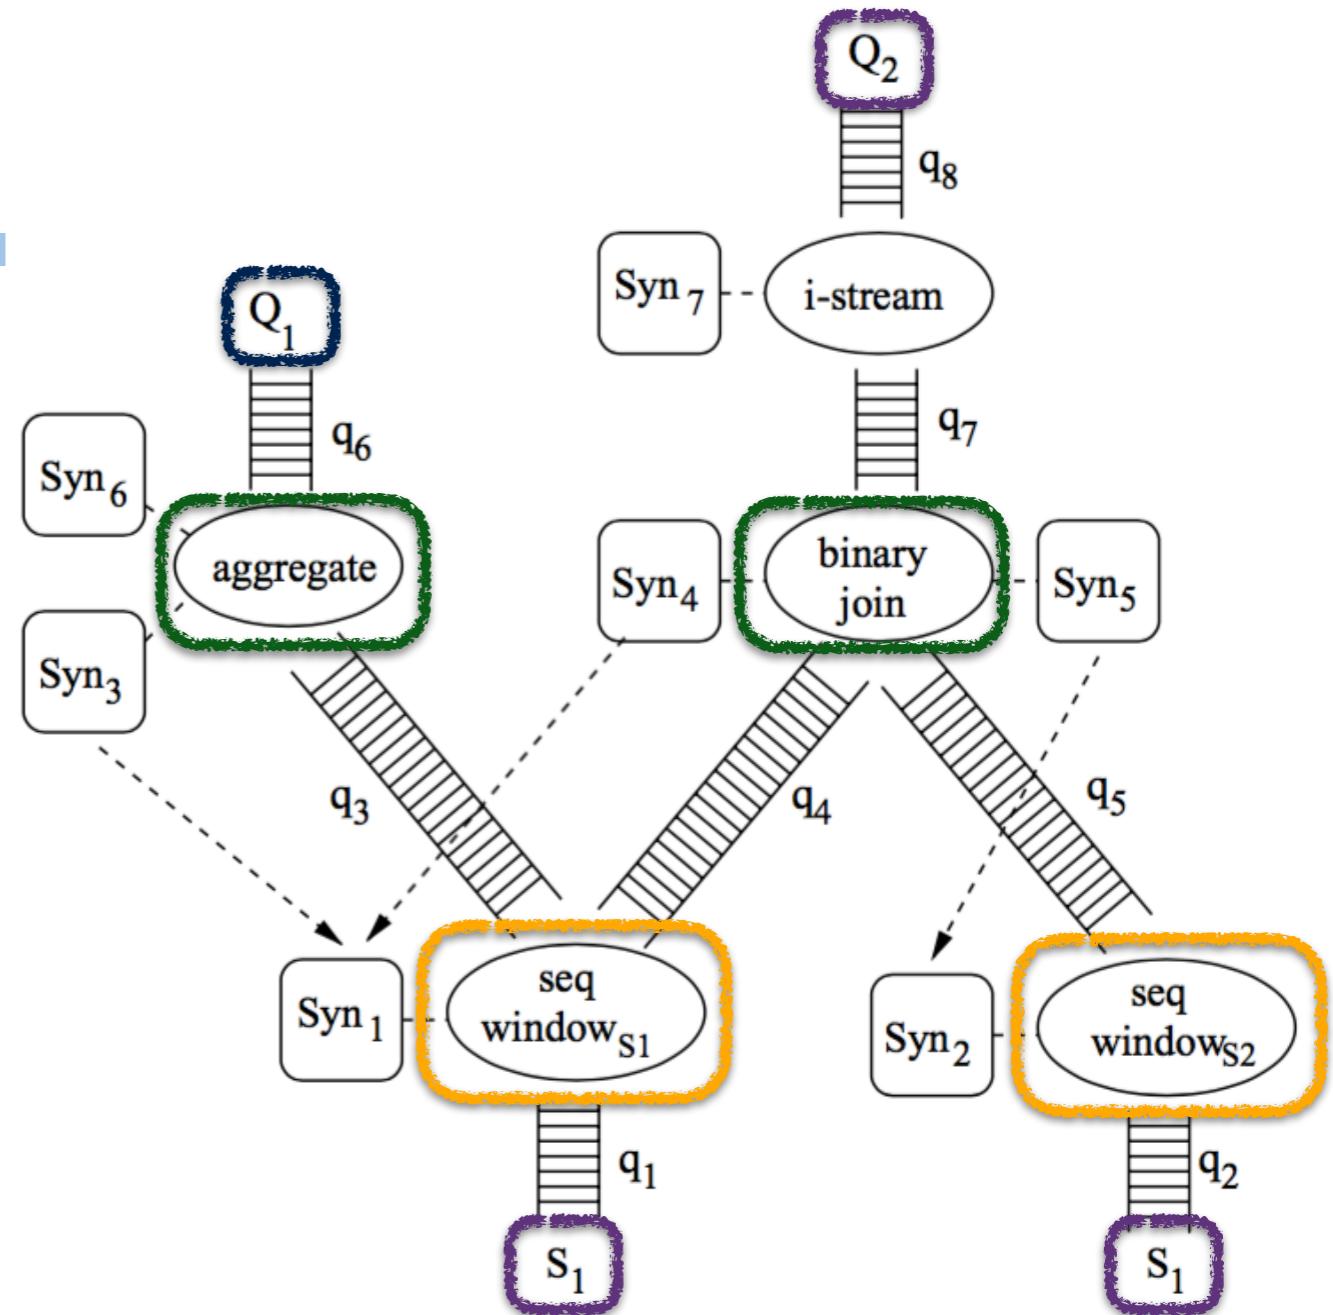
Q1: `select B, max(A)
from S1 [Rows 50000]
group by B;`

Q2: `select Istream(*)
from S1 [Rows 40000]
S2 [Range 600s]
where S1.A = S2.A;`

Example plan

This plan contains

- ◆ 2 input streams
- ◆ 1 output stream
- ◆ 1 output relation
- ◆ 2 s2r operators
- ◆ 2 r2r operators



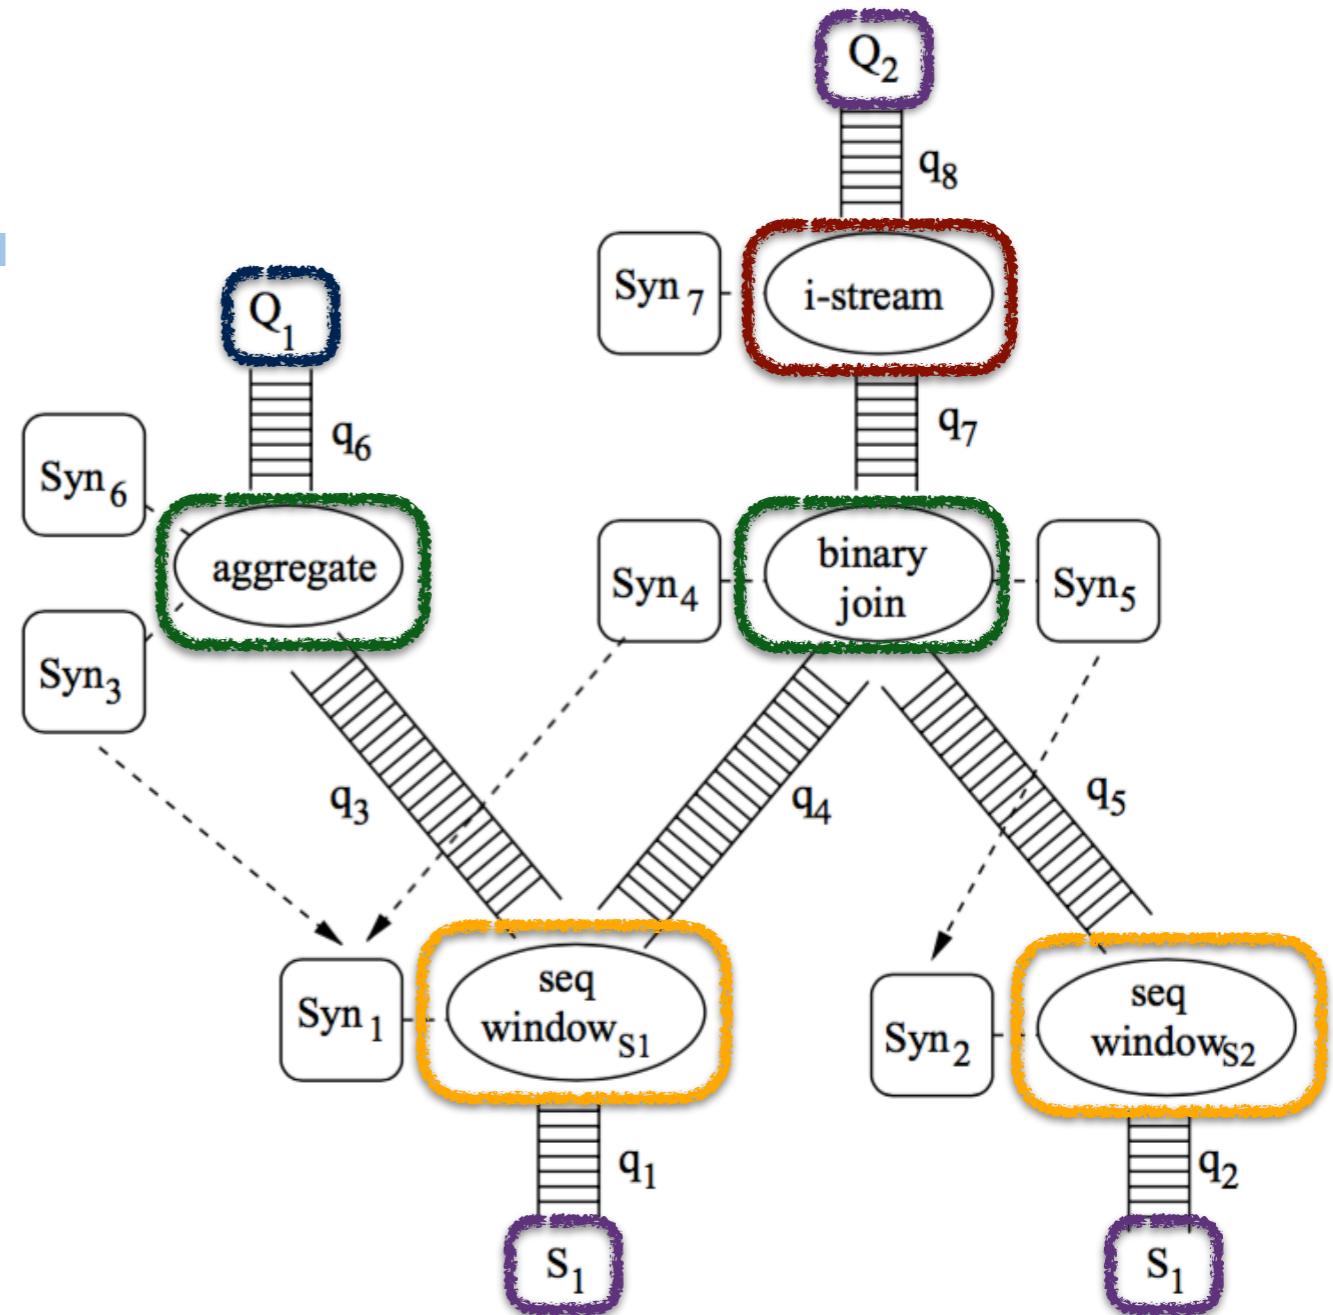
Q1: `select B, max(A)
from S1 [Rows 50000]
group by B;`

Q2: `select Istream(*)
from S1 [Rows 40000]
S2 [Range 600s]
where S1.A = S2.A;`

Example plan

This plan contains

- ◆ 2 input streams
- ◆ 1 output stream
- ◆ 1 output relation
- ◆ 2 s2r operators
- ◆ 2 r2r operators
- ◆ 1 r2s operator



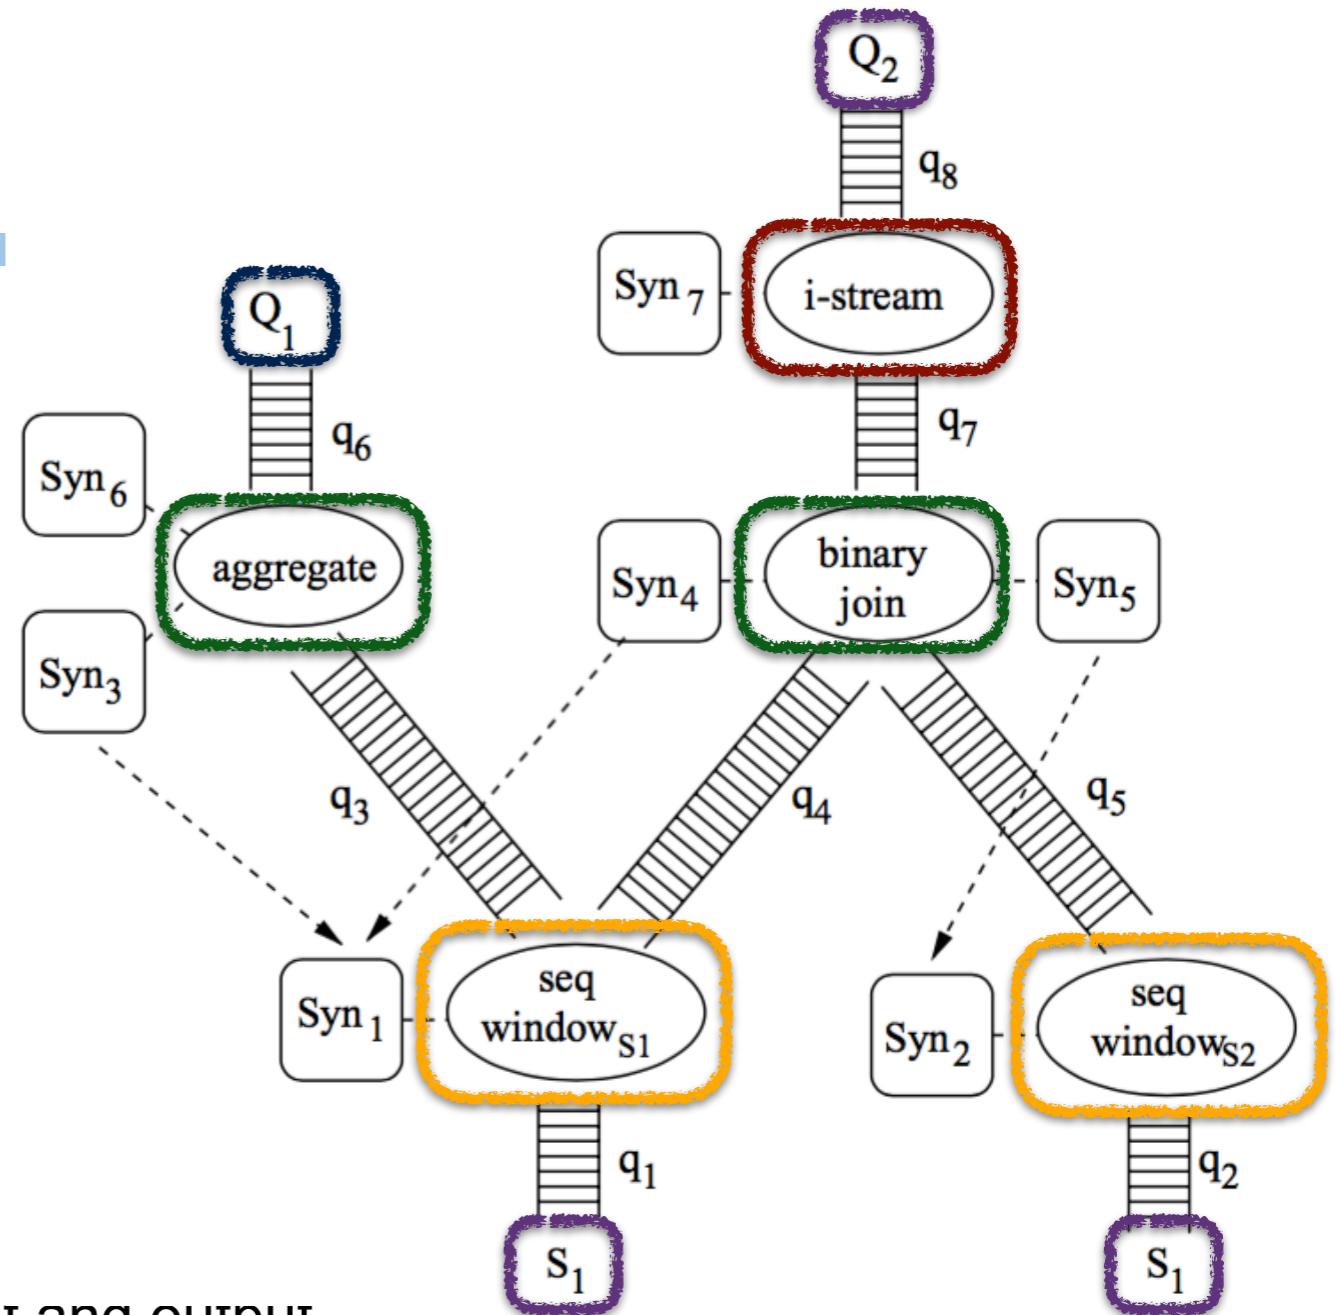
Q1: `select B, max(A)
from S1 [Rows 50000]
group by B;`

Q2: `select Istream(*)
from S1 [Rows 40000]
S2 [Range 600s]
where S1.A = S2.A;`

Example plan

This plan contains

- ◆ 2 input streams
- ◆ 1 output stream
- ◆ 1 output relation
- ◆ 2 s2r operators
- ◆ 2 r2r operators
- ◆ 1 r2s operator
- ◆ 8 queues connecting the operators, input and output



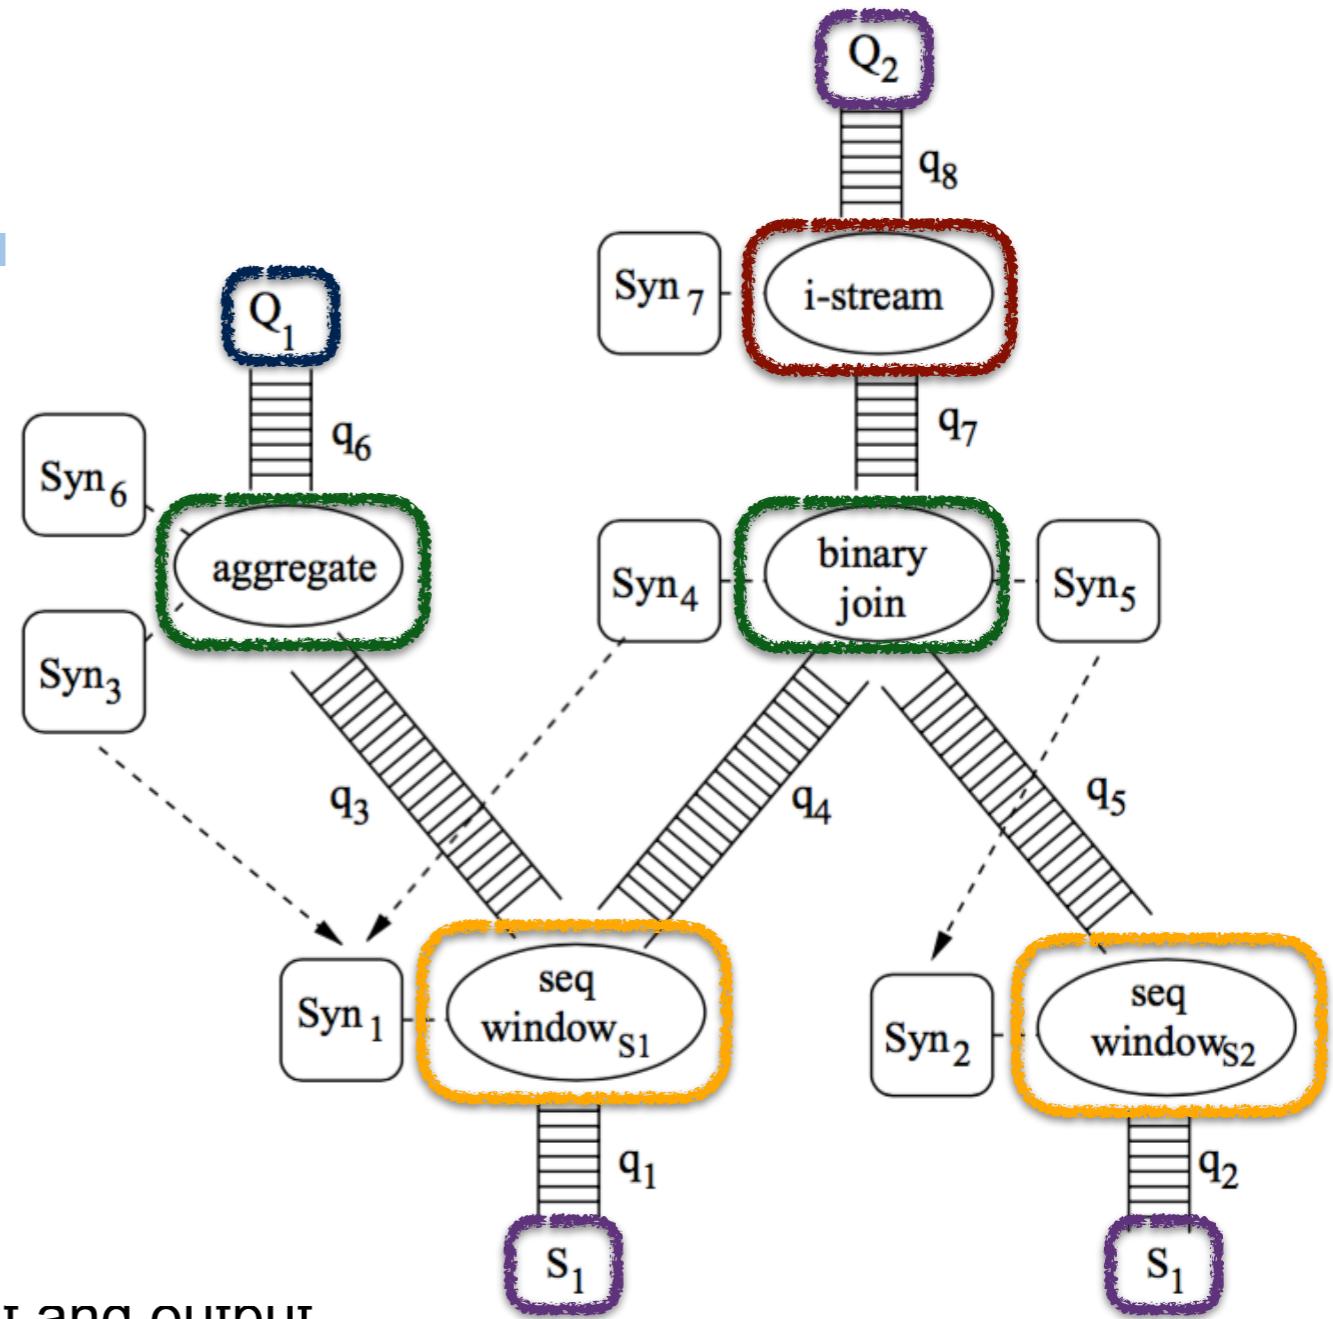
Q1: `select B, max(A)
from S1 [Rows 50000]
group by B;`

Q2: `select Istream(*)
from S1 [Rows 40000]
S2 [Range 600s]
where S1.A = S2.A;`

Example plan

This plan contains

- ◆ 2 input streams
- ◆ 1 output stream
- ◆ 1 output relation
- ◆ 2 s2r operators
- ◆ 2 r2r operators
- ◆ 1 r2s operator
- ◆ 8 queues connecting the operators, input and output
- ◆ 7 synopses associated to operators



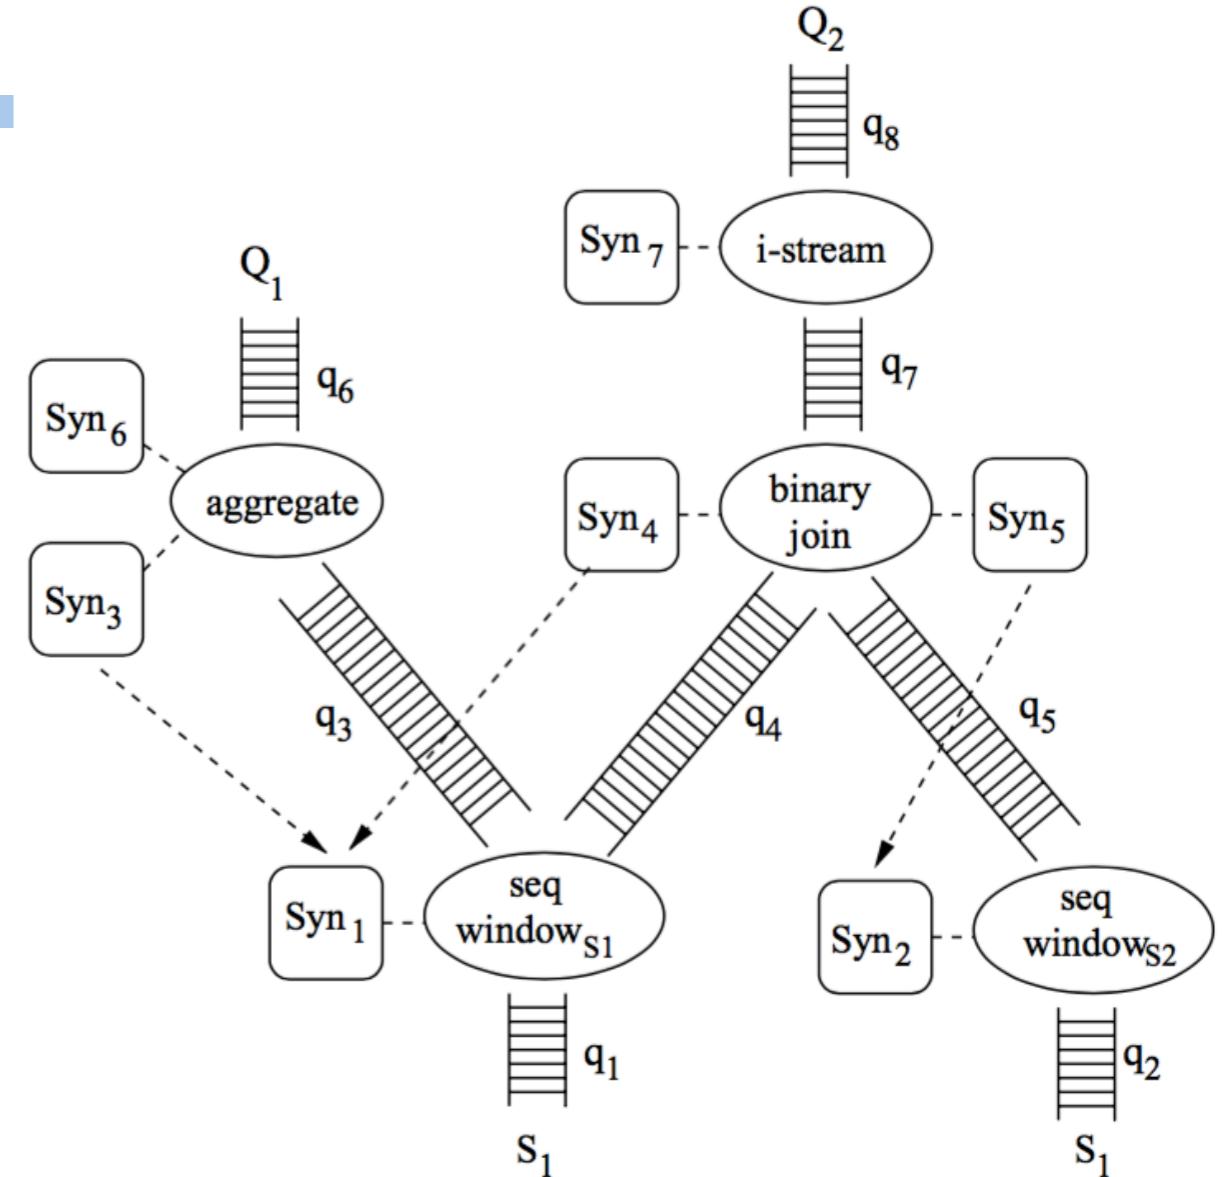
Q1: `select B, max(A)
from S1 [Rows 50000]
group by B;`

Q2: `select Istream(*)
from S1 [Rows 40000]
S2 [Range 600s]
where S1.A = S2.A;`

Example plan

Sliding window operator seq-windows₁

- ◆ Reads stream S₁'s tuples
- ◆ Outputs *insertions and deletions in the window in both q₃ and q₄*



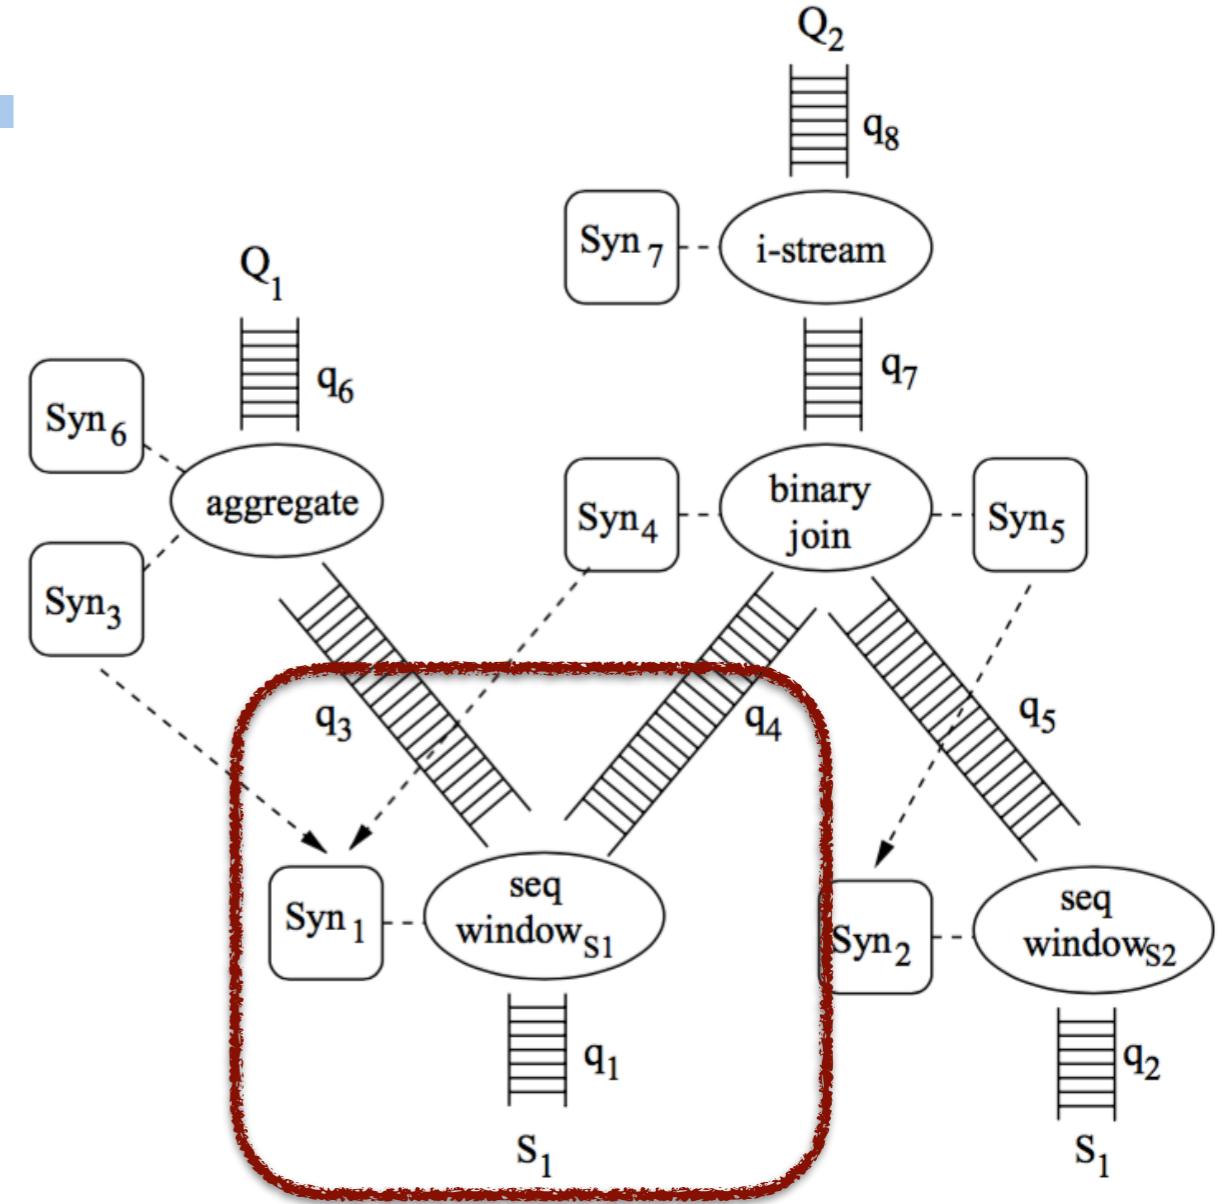
Q1: select B, max(A)
from S1 [Rows 50000]
group by B;

Q2: select Istream(*)
from S1 [Rows 40000]
S2 [Range 600s]
where S1.A = S2.A;

Example plan

Sliding window operator seq-windows₁

- ◆ Reads stream S₁'s tuples
- ◆ Outputs *insertions and deletions in the window in both q₃ and q₄*



Q1: select B, max(A)
from S1 [Rows 50000]
group by B;

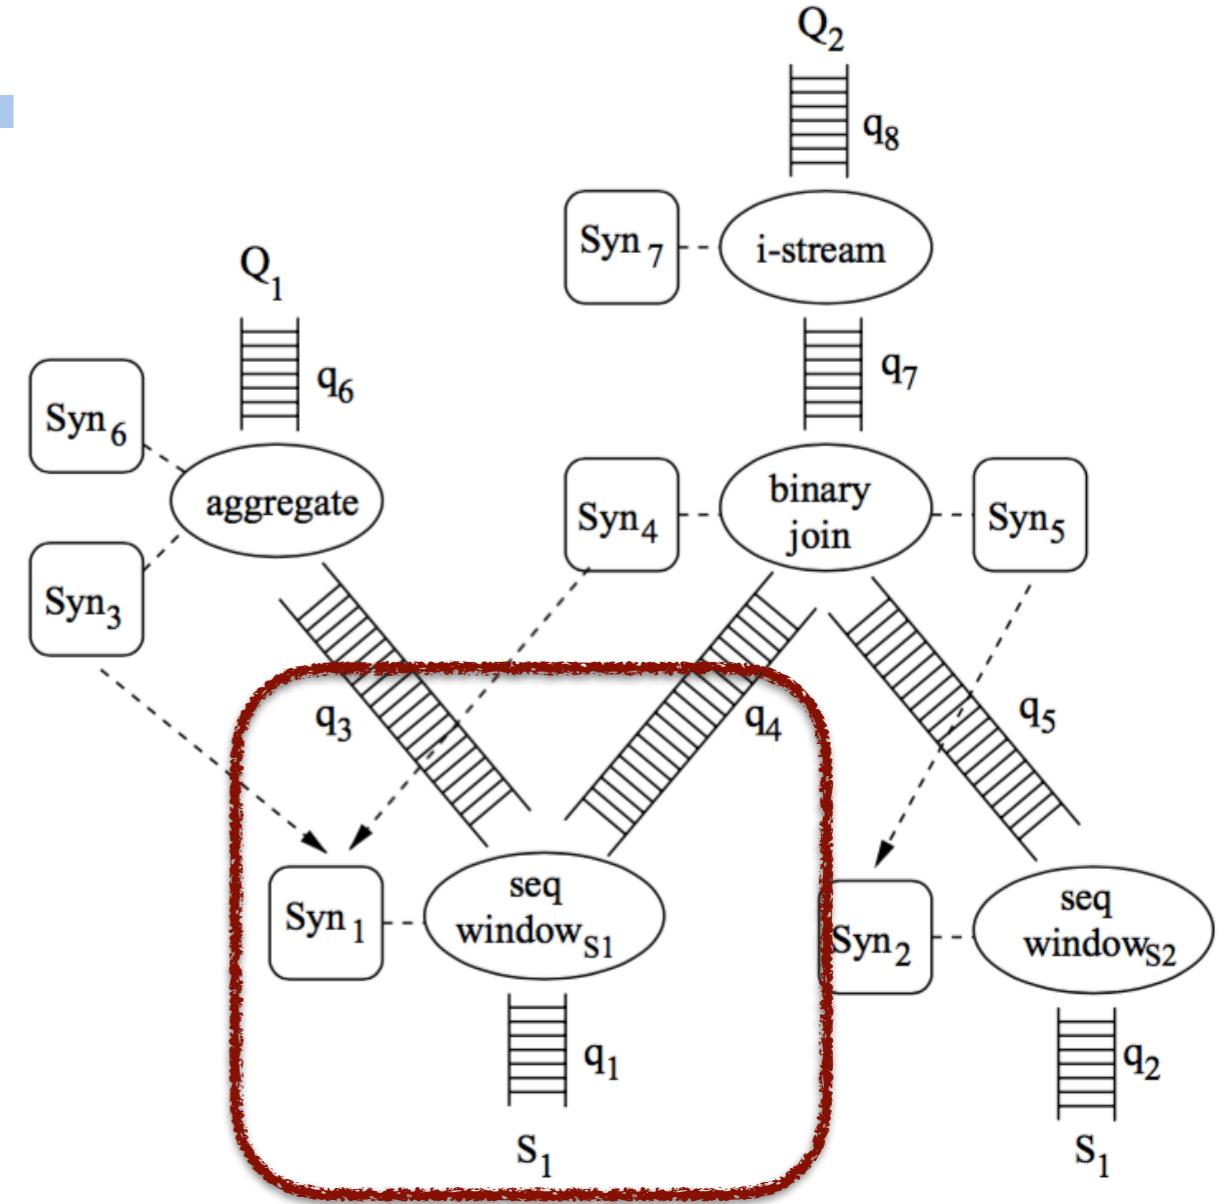
Q2: select Istream(*)
from S1 [Rows 40000]
S2 [Range 600s]
where S1.A = S2.A;

Example plan

■ Sliding window operator seq-windows₁

- ◆ Reads stream S_1 's tuples
- ◆ Outputs *insertions and deletions in the window in both q_3 and q_4*

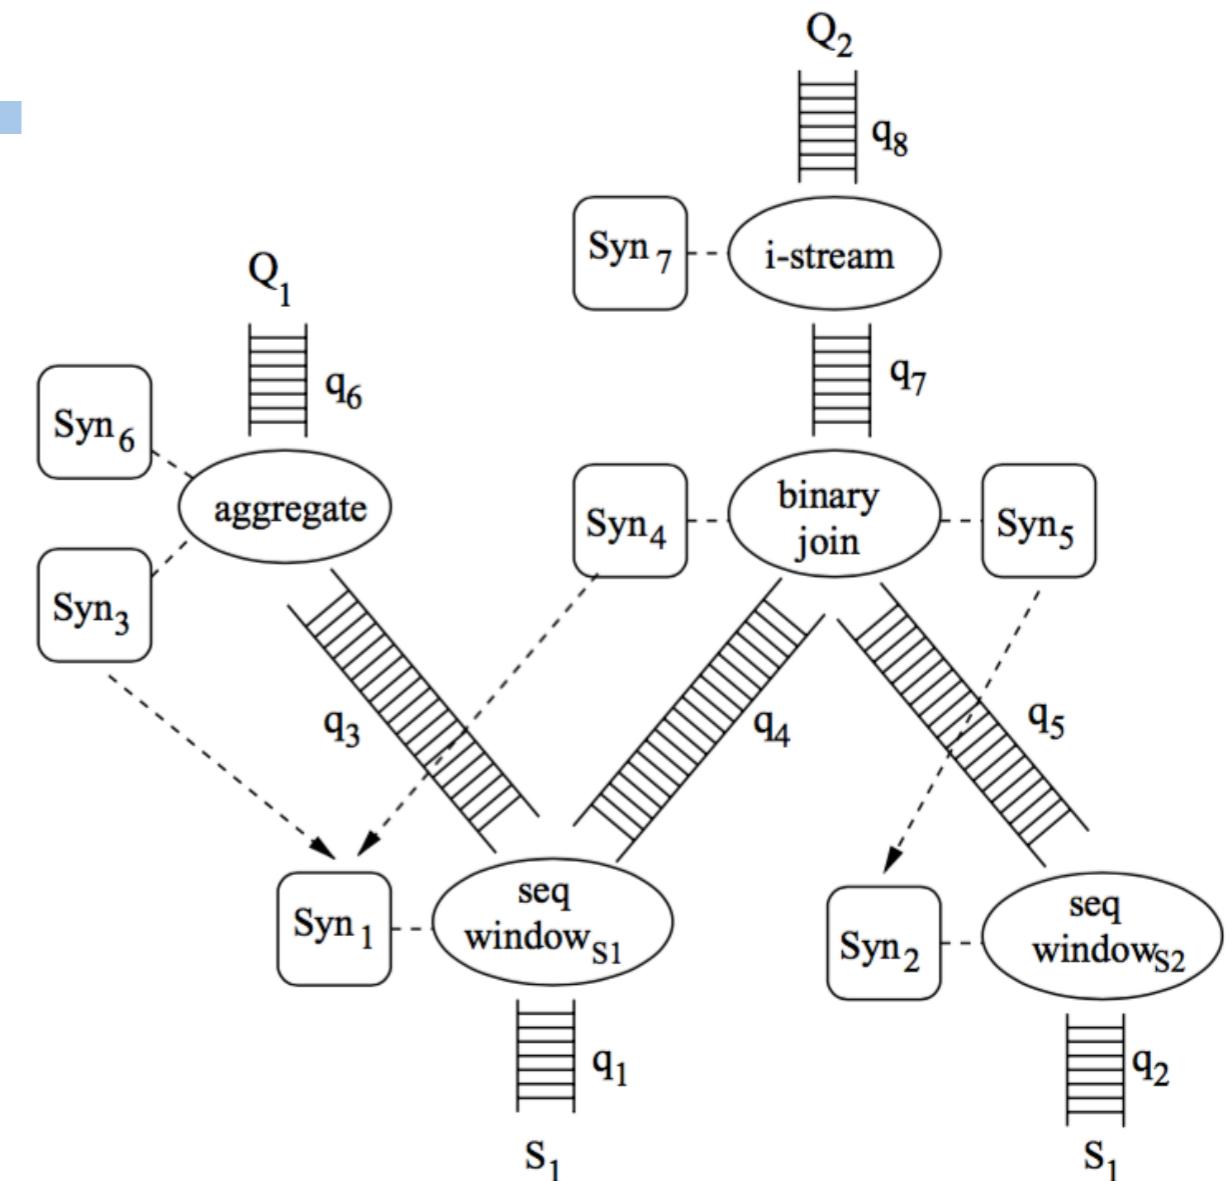
■ Synopsis Syn_1 stores the last 50.000 tuples that arrived in Stream 1



Q1: select B, max(A)
from S1 [Rows 50000]
group by B;

Q2: select Istream(*)
from S1 [Rows 40000]
S2 [Range 600s]
where S1.A = S2.A;

Example plan



Q1: select B, max(A)
from S1 [Rows 50000]
group by B;

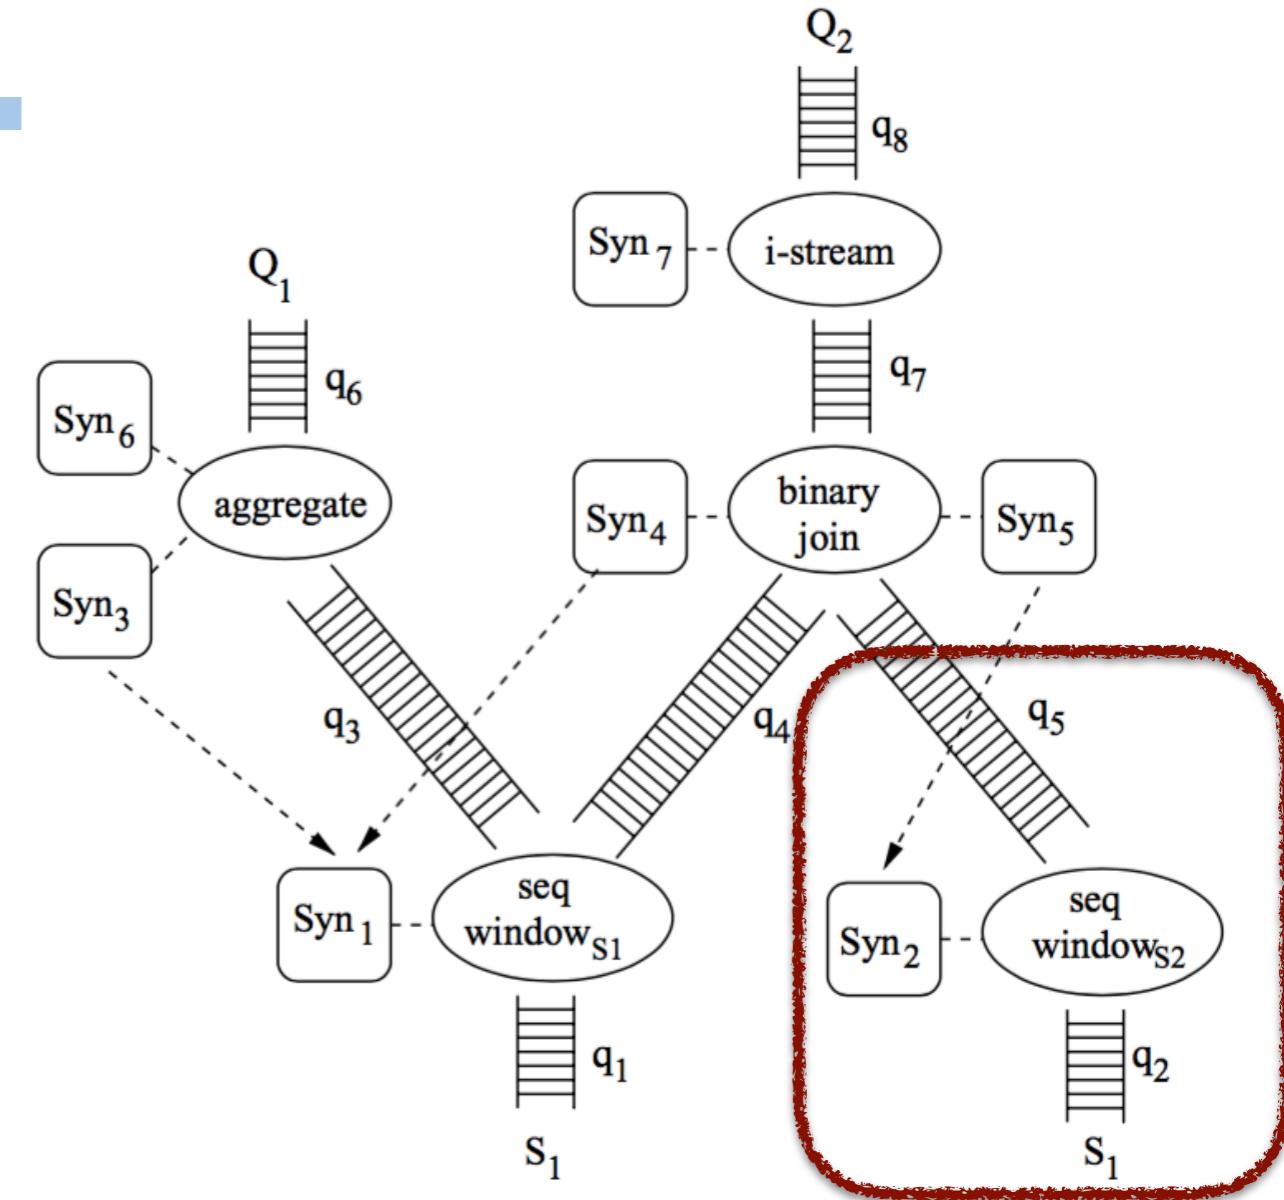
Q2: select Istream(*)
from S1 [Rows 40000]
S2 [Range 600s]
where S1.A = S2.A;

Example plan

■ Sliding window operator seq-windows2

- ◆ Reads stream S_2 's tuples
- ◆ Processes them to know the ones that arrived in the last 10 minutes
- ◆ Outputs in q_5 , the insertions and deletions in the window

■ Synopsis Syn_2 stores the items that arrived at S_2 in the last 10 minutes



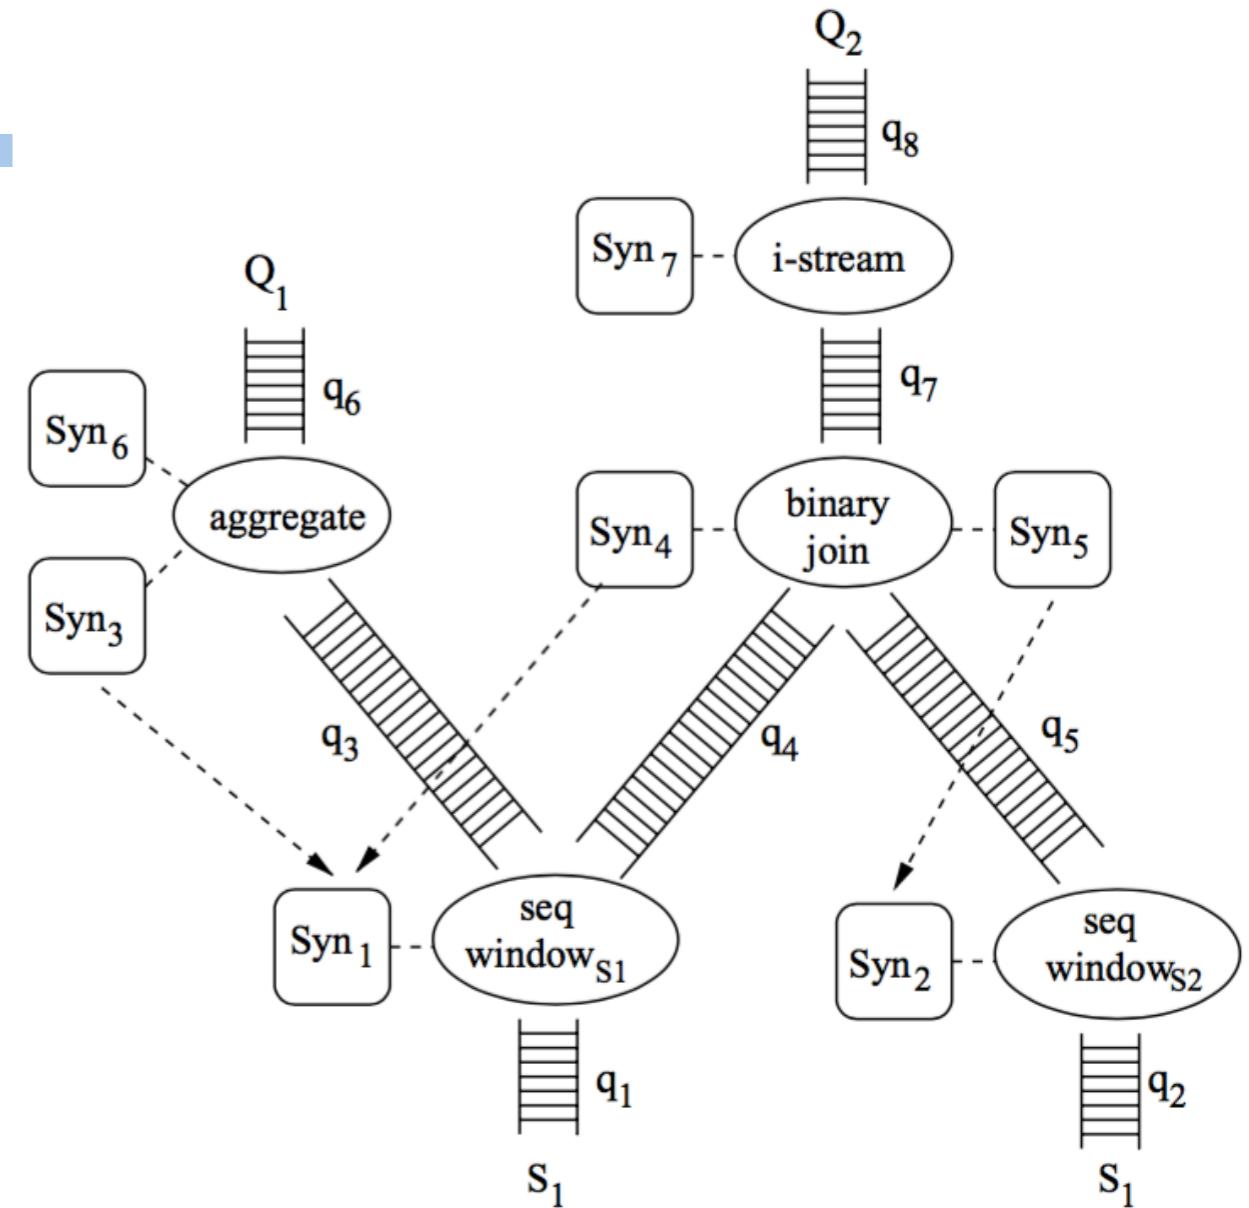
Q1: select B, max(A)
from S1 [Rows 50000]
group by B;

Q2: select Istream(*)
from S1 [Rows 40000]
S2 [Range 600s]
where S1.A = S2.A;

Example plan

operator binary-join

- ◆ Incrementally computes the join of two relations
 - Uses insertions and deletions in q_4 and q_5 , to produce insertions and deletions in q_7



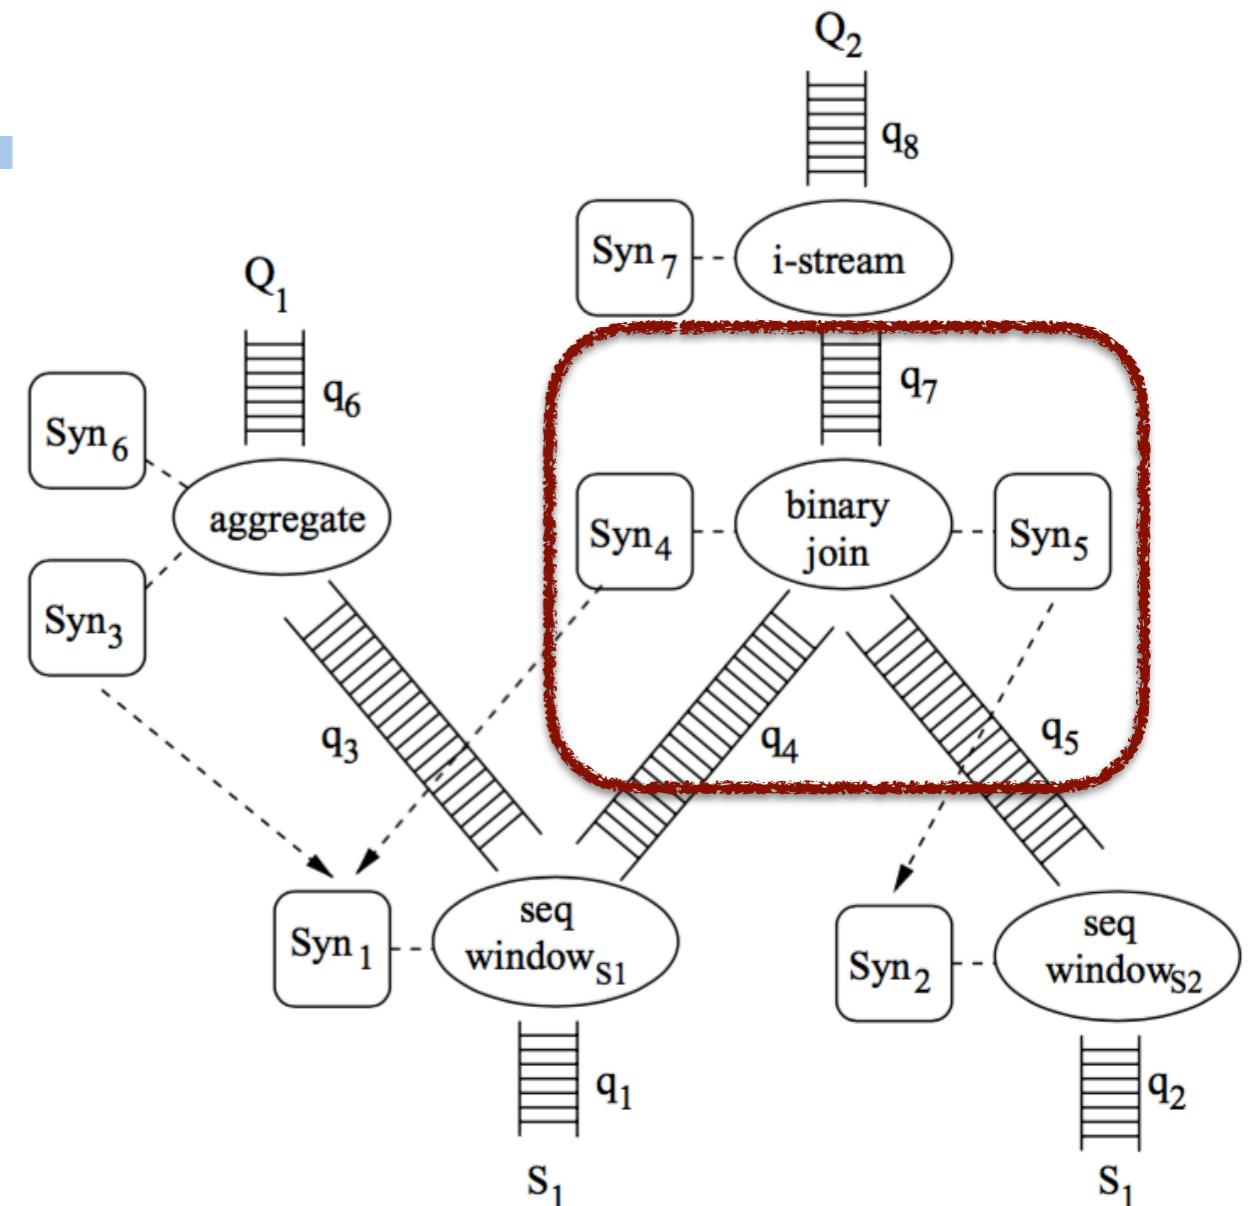
```
Q1: select B, max(A)
      from S1 [Rows 50000]
      group by B;
```

```
Q2: select Istream(*)
      from S1 [Rows 40000]
            S2 [Range 600s]
      where S1.A = S2.A;
```

Example plan

operator binary-join

- ◆ Incrementally computes the join of two relations
 - Uses insertions and deletions in q_4 and q_5 , to produce insertions and deletions in q_7



Q1: `select B, max(A)
from S1 [Rows 50000]
group by B;`

Q2: `select Istream(*)
from S1 [Rows 40000]
S2 [Range 600s]
where S1.A = S2.A;`

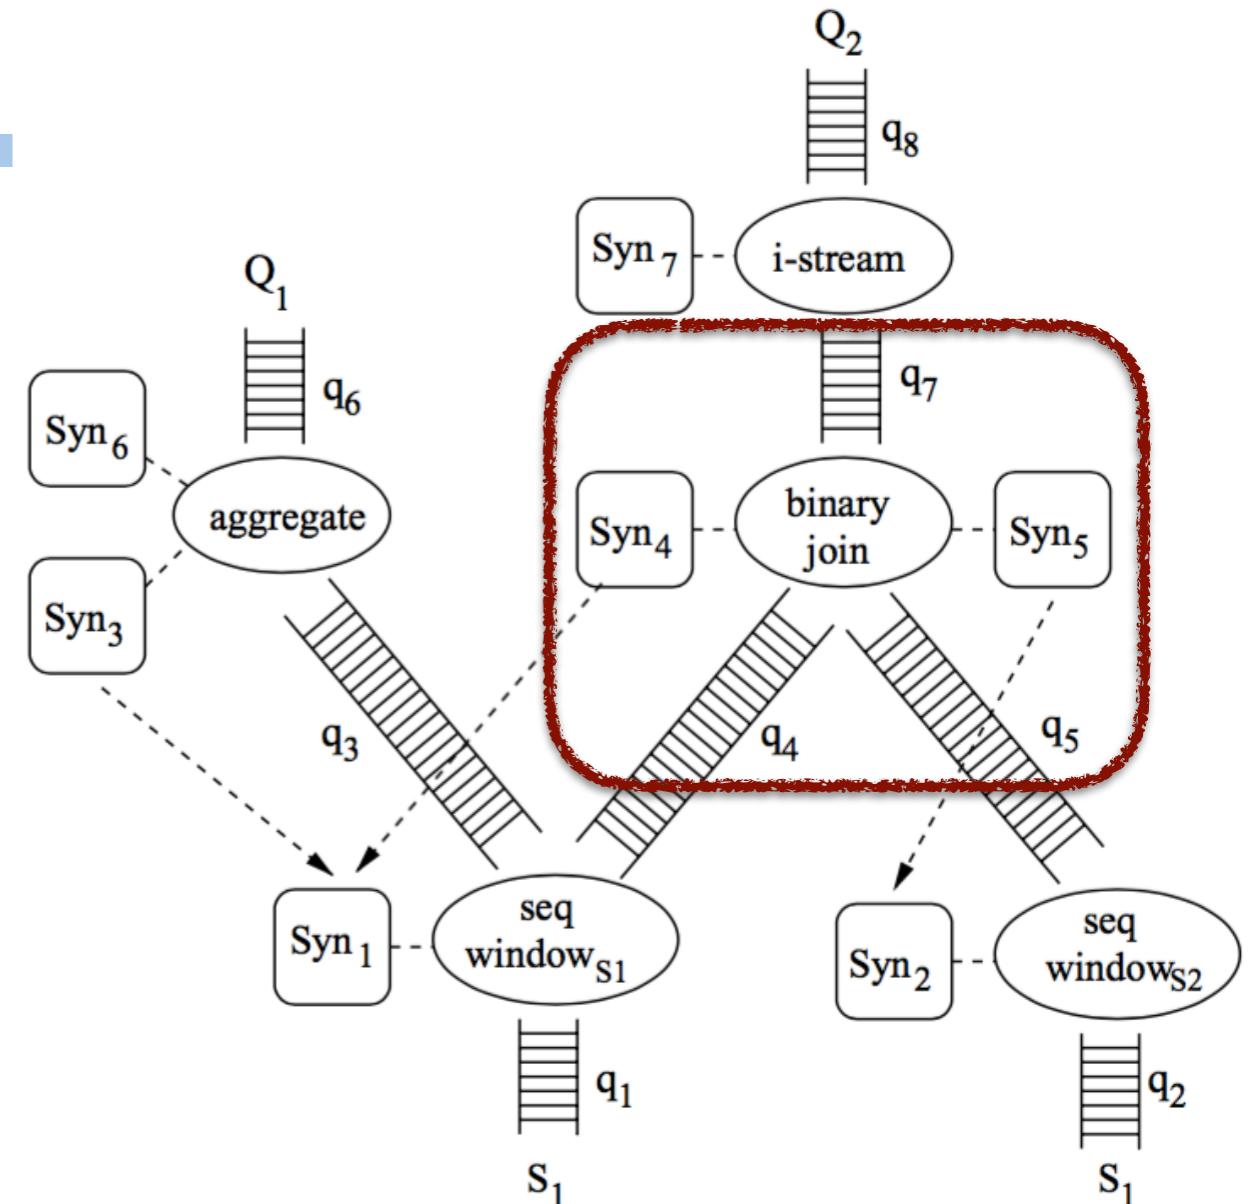
Example plan

operator binary-join

- ◆ Incrementally computes the join of two relations
 - Uses insertions and deletions in q_4 and q_5 , to produce insertions and deletions in q_7

Synopsis Syn_4 (resp. Syn_5) is shared with Syn_1 (resp Syn_2)

- ◆ An inserted tuple in q_4 is joined with the tuples in Syn_5 to compute an insertion in q_7
- ◆ A deleted tuple in q_4 is joined with the tuples in Syn_5 to compute a deletion in q_7
- ◆ The same for q_5 and Syn_4 to produce tuples in q_7



Q1: `select B, max(A)
from S1 [Rows 50000]
group by B;`

Q2: `select Istream(*)
from S1 [Rows 40000]
S2 [Range 600s]
where S1.A = S2.A;`

Example plan

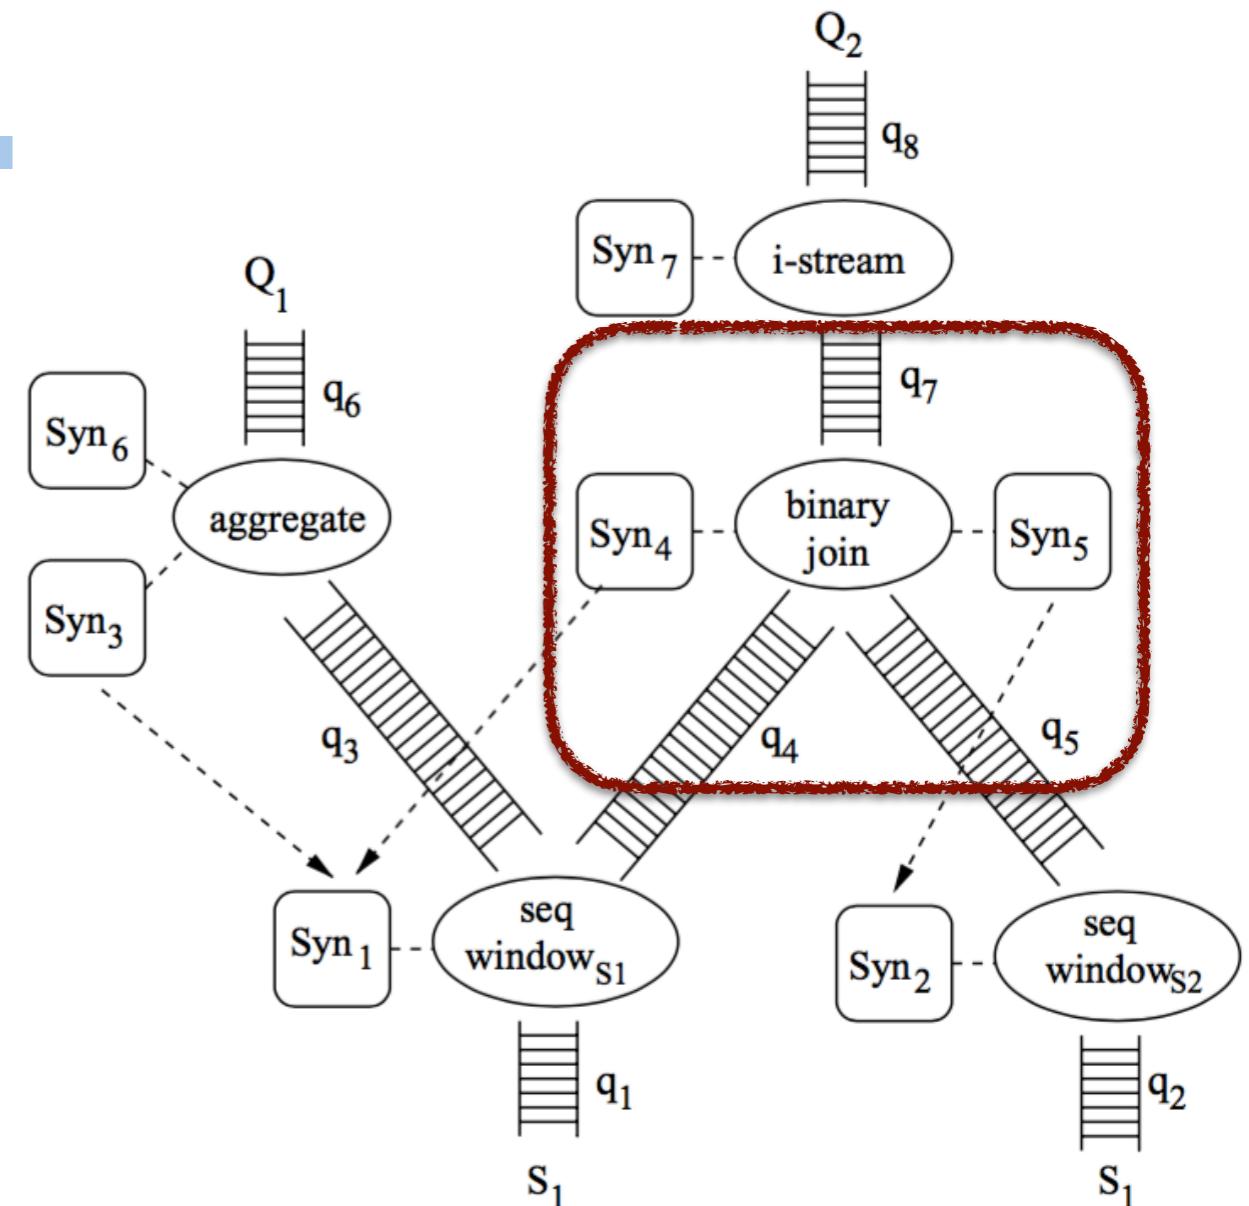
operator binary-join

- ◆ Incrementally computes the join of two relations
 - Uses insertions and deletions in q_4 and q_5 , to produce insertions and deletions in q_7

Synopsis Syn_4 (resp. Syn_5) is shared with Syn_1 (resp Syn_2)

- ◆ An inserted tuple in q_4 is joined with the tuples in Syn_5 to compute an insertion in q_7
- ◆ A deleted tuple in q_4 is joined with the tuples in Syn_5 to compute a deletion in q_7
- ◆ The same for q_5 and Syn_4 to produce tuples in q_7

When joining with a table, the *queue reports changes*, and the synopsis is the table content



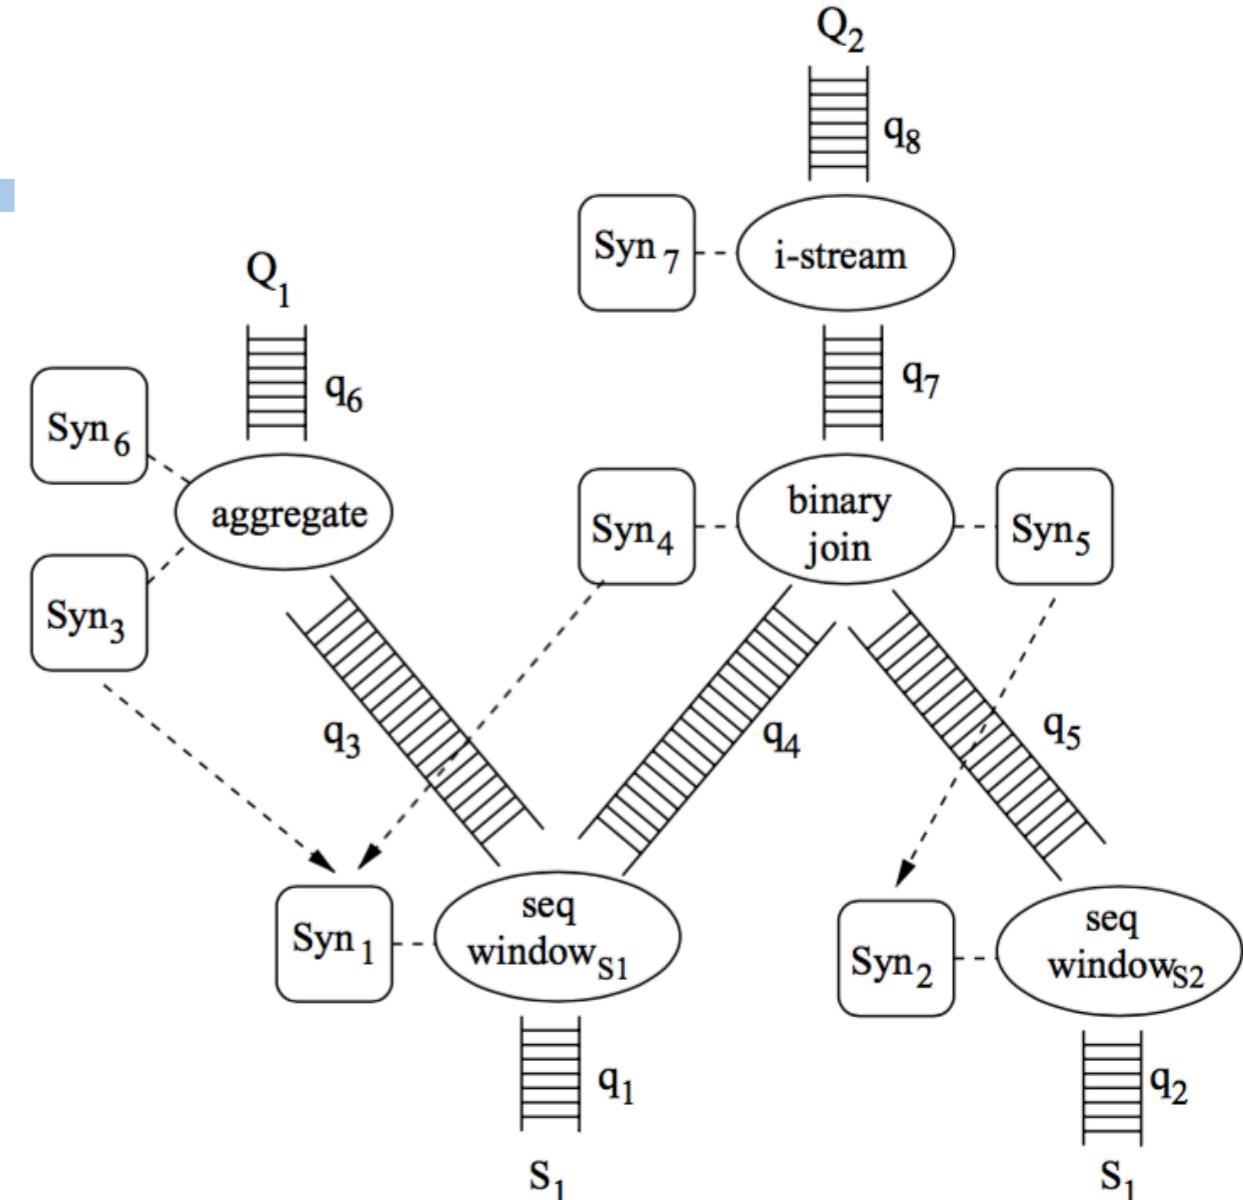
Q1: `select B, max(A)
from S1 [Rows 50000]
group by B;`

Q2: `select Istream(*)
from S1 [Rows 40000]
S2 [Range 600s]
where S1.A = S2.A;`

Example plan

operator aggregate

- ◆ (Incrementally) maintains the max value of A for each distinct B
- ◆ Synopsis Syn_6 stores those max values



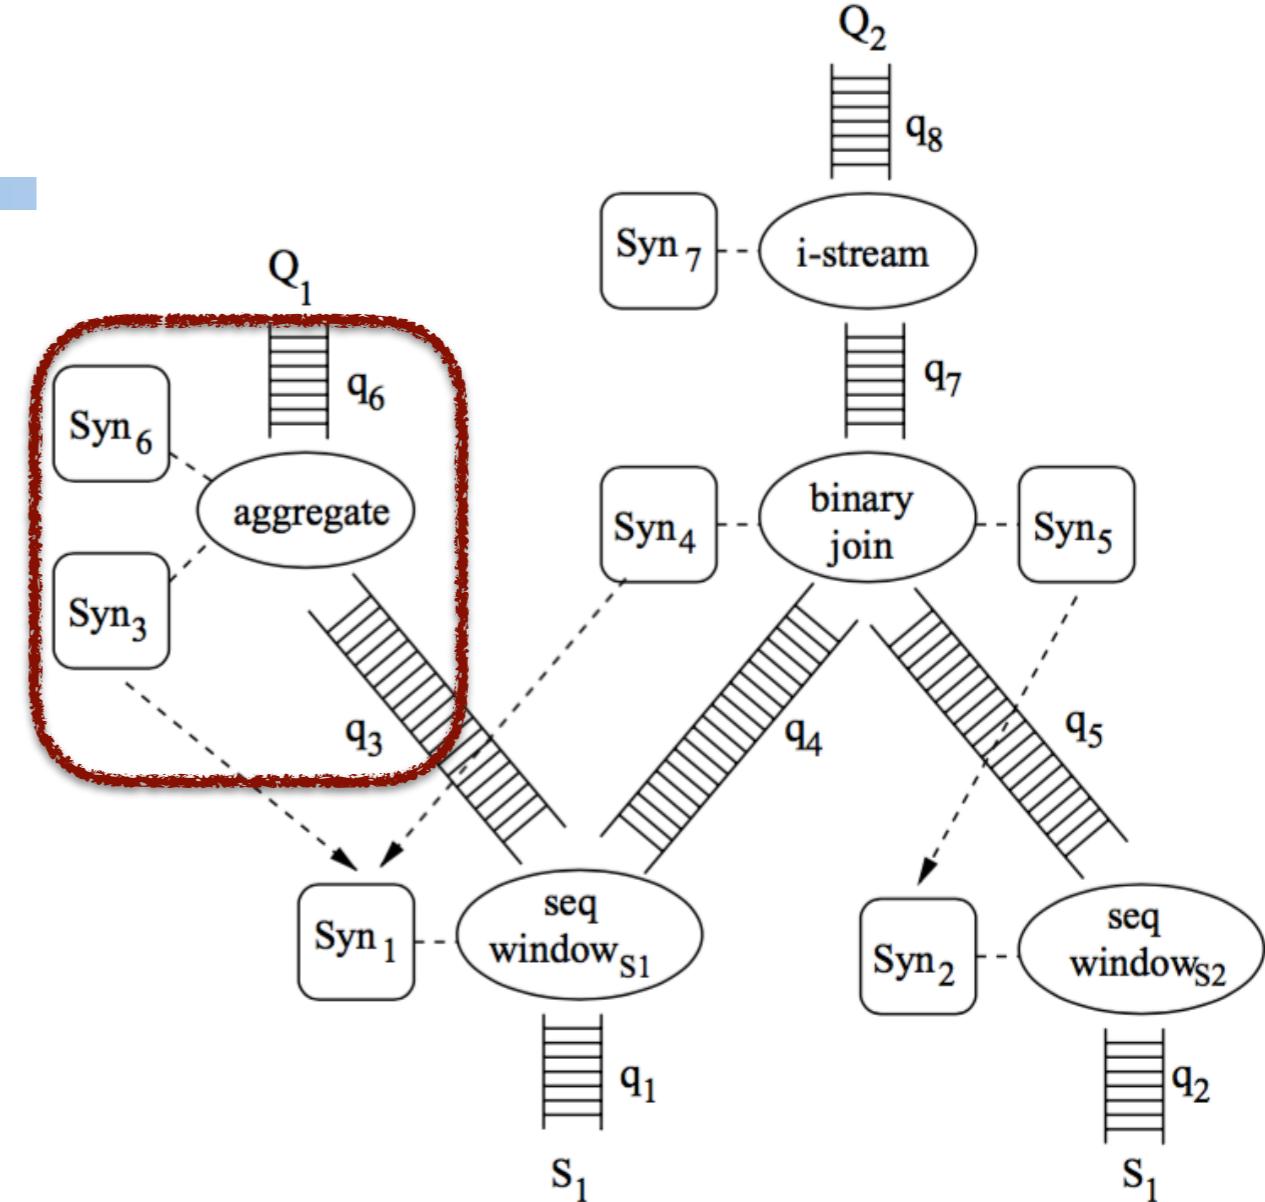
Q1: `select B, max(A)
from S1 [Rows 50000]
group by B;`

Q2: `select Istream(*)
from S1 [Rows 40000]
S2 [Range 600s]
where S1.A = S2.A;`

Example plan

operator aggregate

- ◆ (Incrementally) maintains the max value of A for each distinct B
- ◆ Synopsis Syn_6 stores those max values



```
Q1: select B, max(A)
     from S1 [Rows 50000]
     group by B;
```

```
Q2: select Istream(*)
     from S1 [Rows 40000]
           S2 [Range 600s]
     where S1.A = S2.A;
```

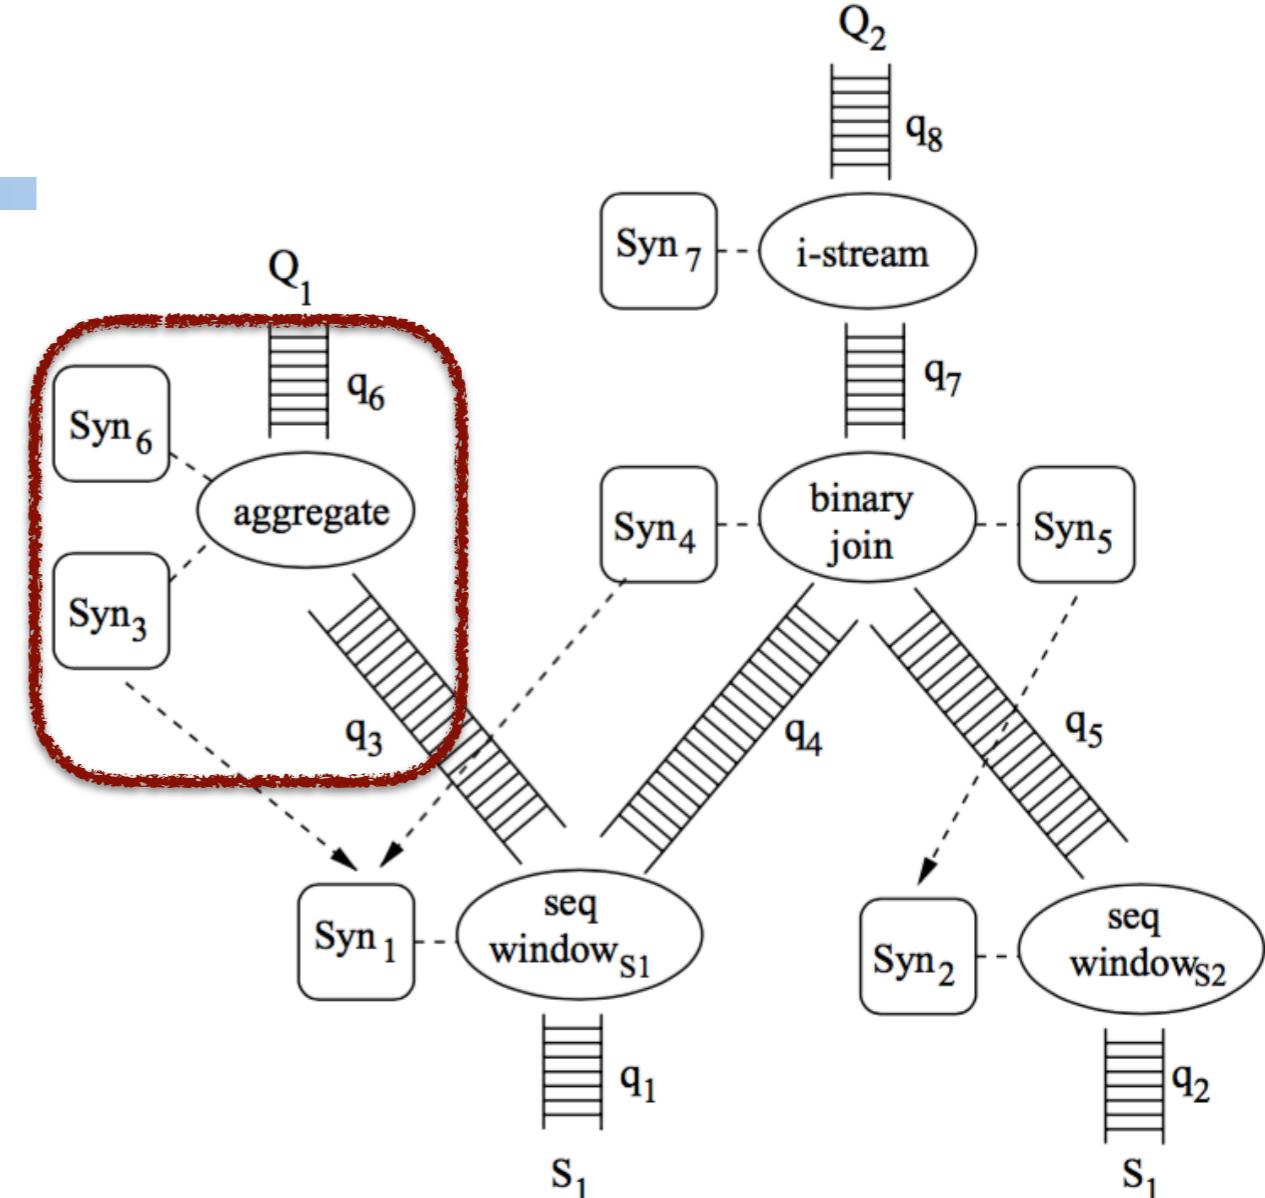
Example plan

operator aggregate

- ◆ (Incrementally) maintains the max value of A for each distinct B
- ◆ Synopsis Syn_6 stores those max values

Any time an insertion (resp. deletion) appears in q_3 , the operator must look at Syn_6 , and see whether that changes a max value

- ◆ When a max is deleted, it must consult the values in the window; that is in Syn_3
 - Some of this info may be shared with Syn_1
- ◆ If a max value changes, it outputs in q_6 both a deletion (of the old max) and an insertion (of the new max)



$Q1: \text{select } B, \max(A)$
from $S1$ [Rows 50000]
group by B ;

$Q2: \text{select Istream(*)}$
from $S1$ [Rows 40000]
 $S2$ [Range 600s]
where $S1.A = S2.A$;

Example plan

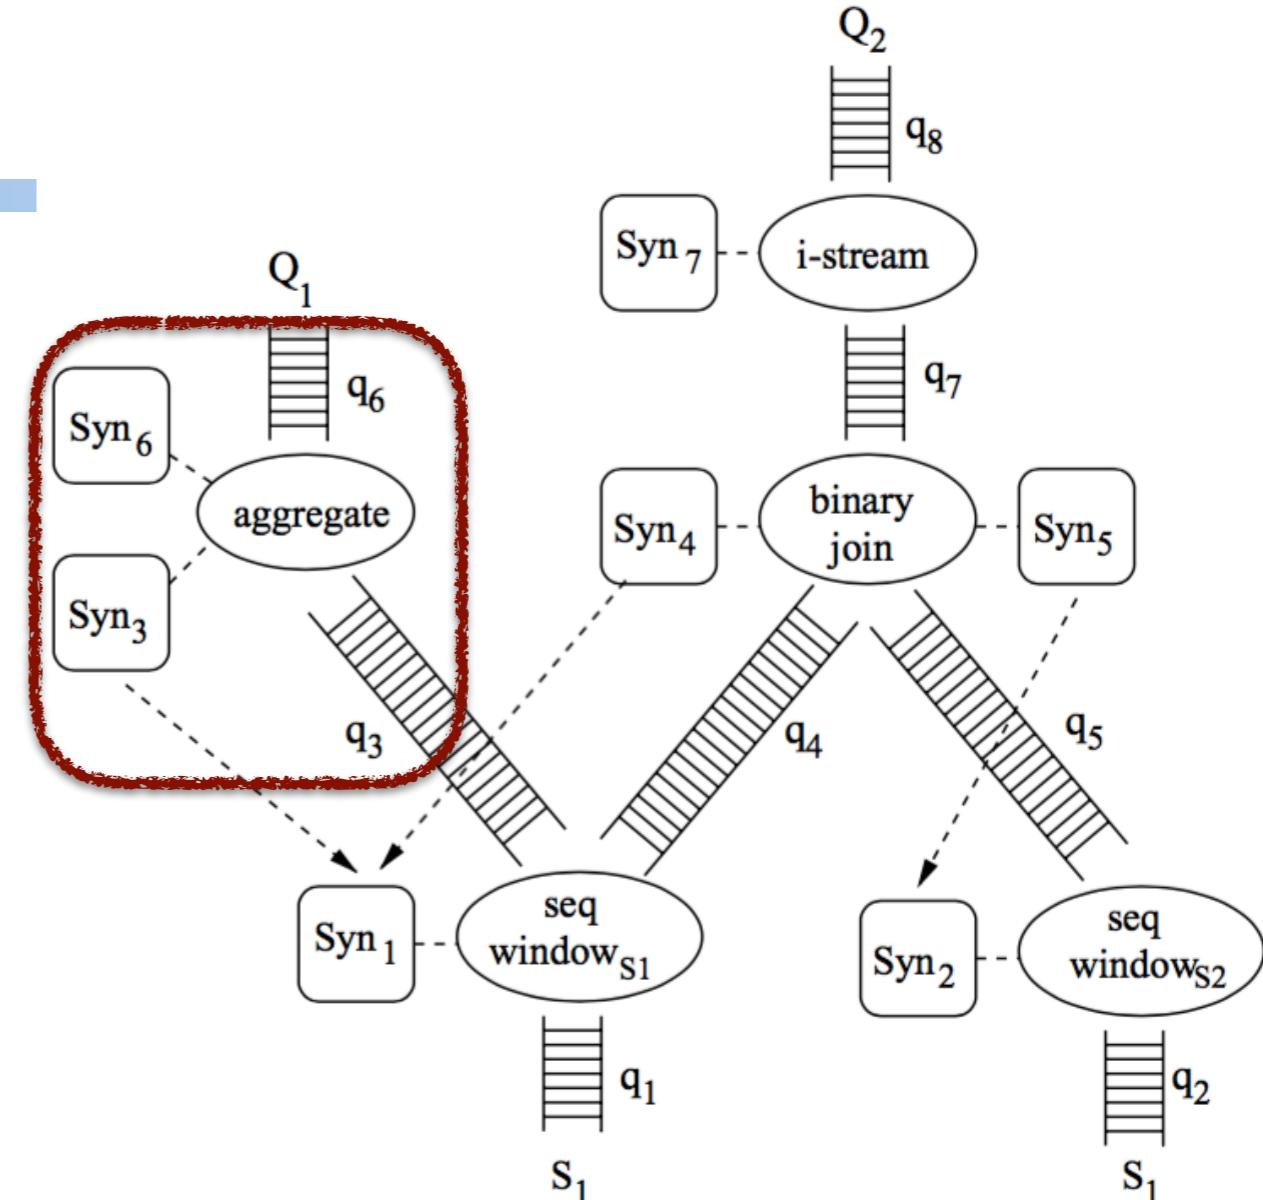
operator aggregate

- ◆ (Incrementally) maintains the max value of A for each distinct B
- ◆ Synopsis Syn_6 stores those max values

Any time an insertion (resp. deletion) appears in q_3 , the operator must look at Syn_6 , and see whether that changes a max value

- ◆ When a max is deleted, it must consult the values in the window; that is in Syn_3
 - Some of this info may be shared with Syn_1
- ◆ If a max value changes, it outputs in q_6 both a deletion (of the old max) and an insertion (of the new max)

Note that this operator is non-monotonic



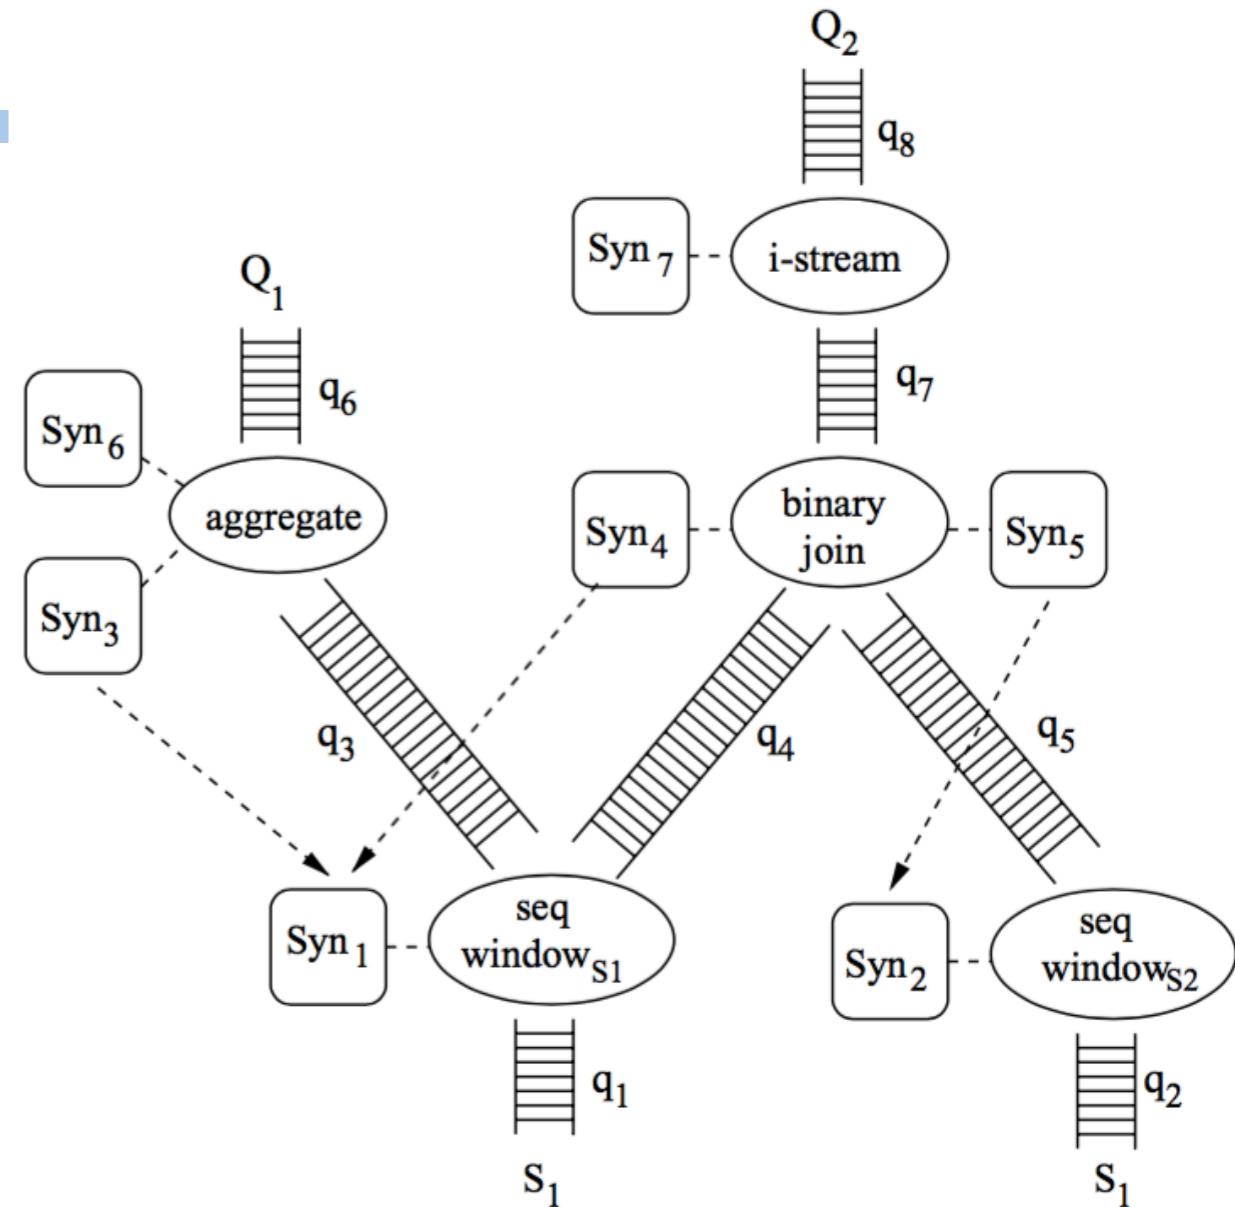
Q1: `select B, max(A)
from S1 [Rows 50000]
group by B;`

Q2: `select Istream(*)
from S1 [Rows 40000]
S2 [Range 600s]
where S1.A = S2.A;`

Example plan

operator i-stream

- ◆ In a naive implementation, whenever an insertion appears in q_7 , it outputs the inserted value to q_8
- ◆ This seem to be exactly the meaning of Istream



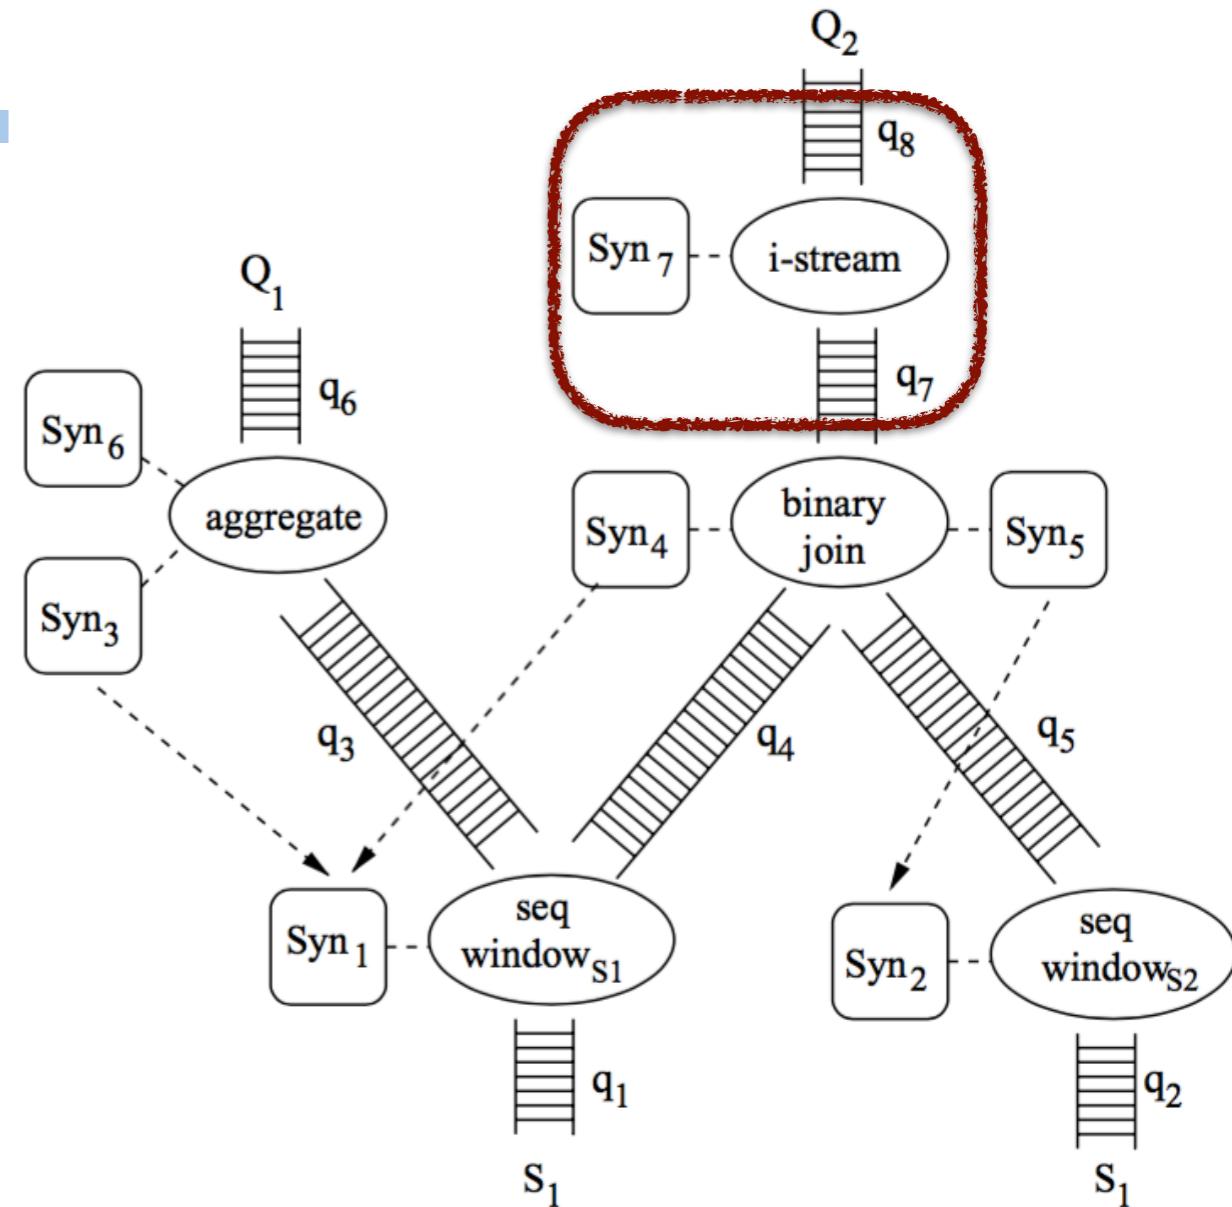
Q1: select B, max(A)
from S1 [Rows 50000]
group by B;

Q2: select Istream(*)
from S1 [Rows 40000]
S2 [Range 600s]
where S1.A = S2.A;

Example plan

operator i-stream

- ◆ In a naive implementation, whenever an insertion appears in q_7 , it outputs the inserted value to q_8
- ◆ This seem to be exactly the meaning of Istream



Q1: select B, max(A)
from S1 [Rows 50000]
group by B;

Q2: select Istream(*)
from S1 [Rows 40000]
S2 [Range 600s]
where S1.A = S2.A;

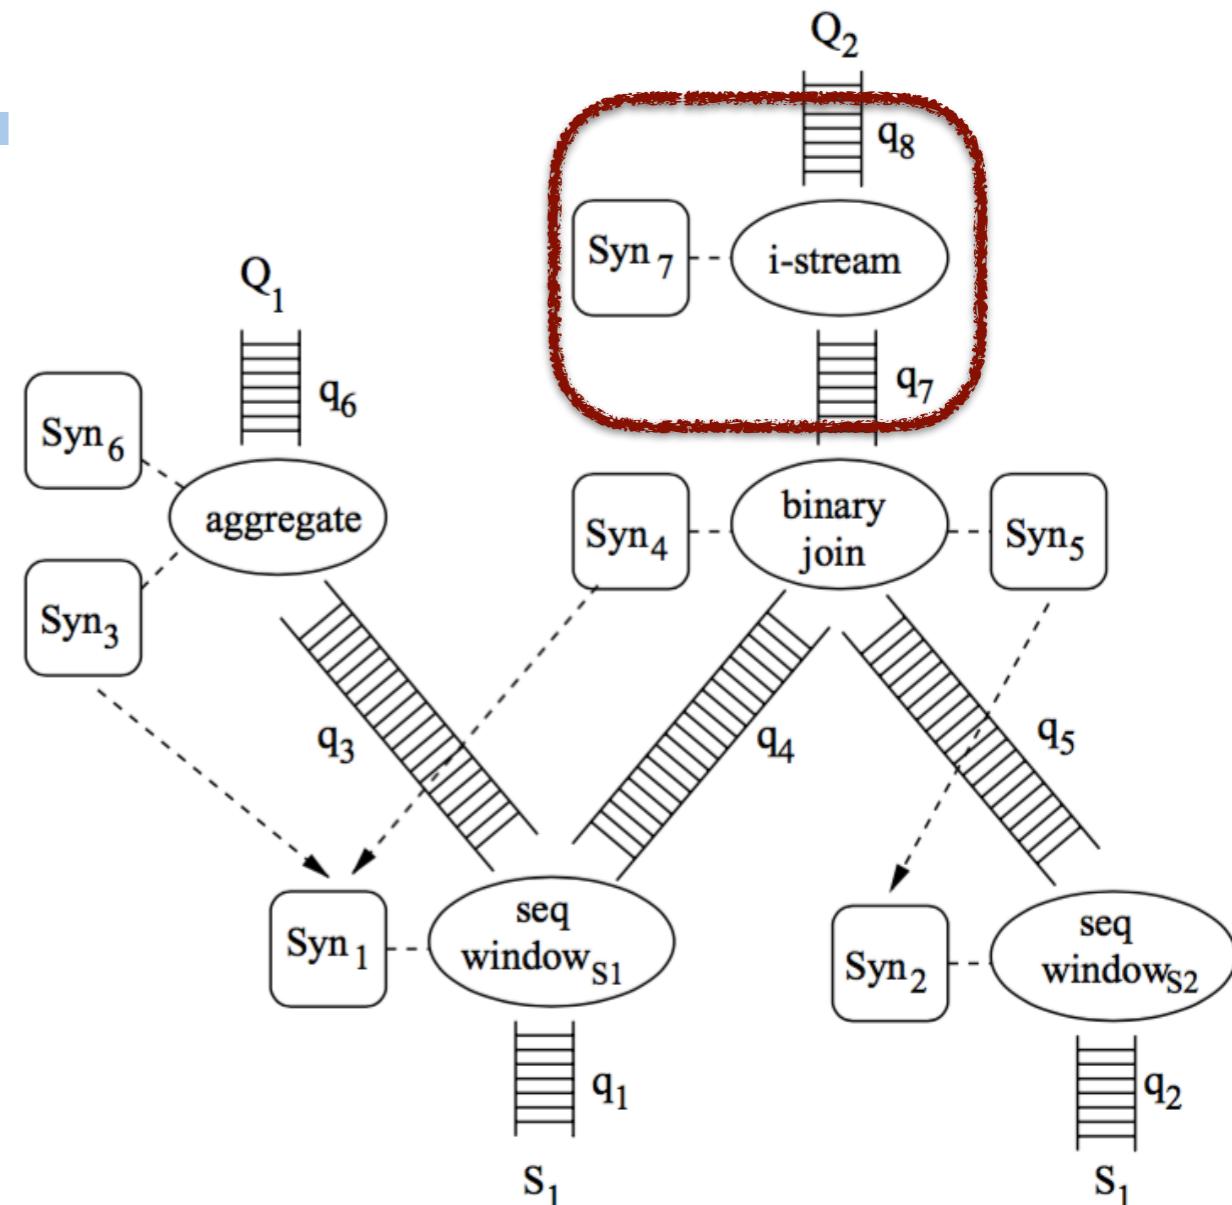
Example plan

operator i-stream

- ◆ In a naive implementation, whenever an insertion appears in q_7 , it outputs the inserted value to q_8
 - ◆ This seem to be exactly the meaning of `Istream`

■ But care must be taken!

- ◆ What if there is both an insertion and a deletion of a same tuple at the same time?



```
Q1: select B, max(A)
      from S1 [Rows 50000]
      group by B;
```

```
Q2: select Istream(*)
      from S1 [Rows 40000]
              S2 [Range 600s]
      where S1.A = S2.A;
```

Example plan

operator i-stream

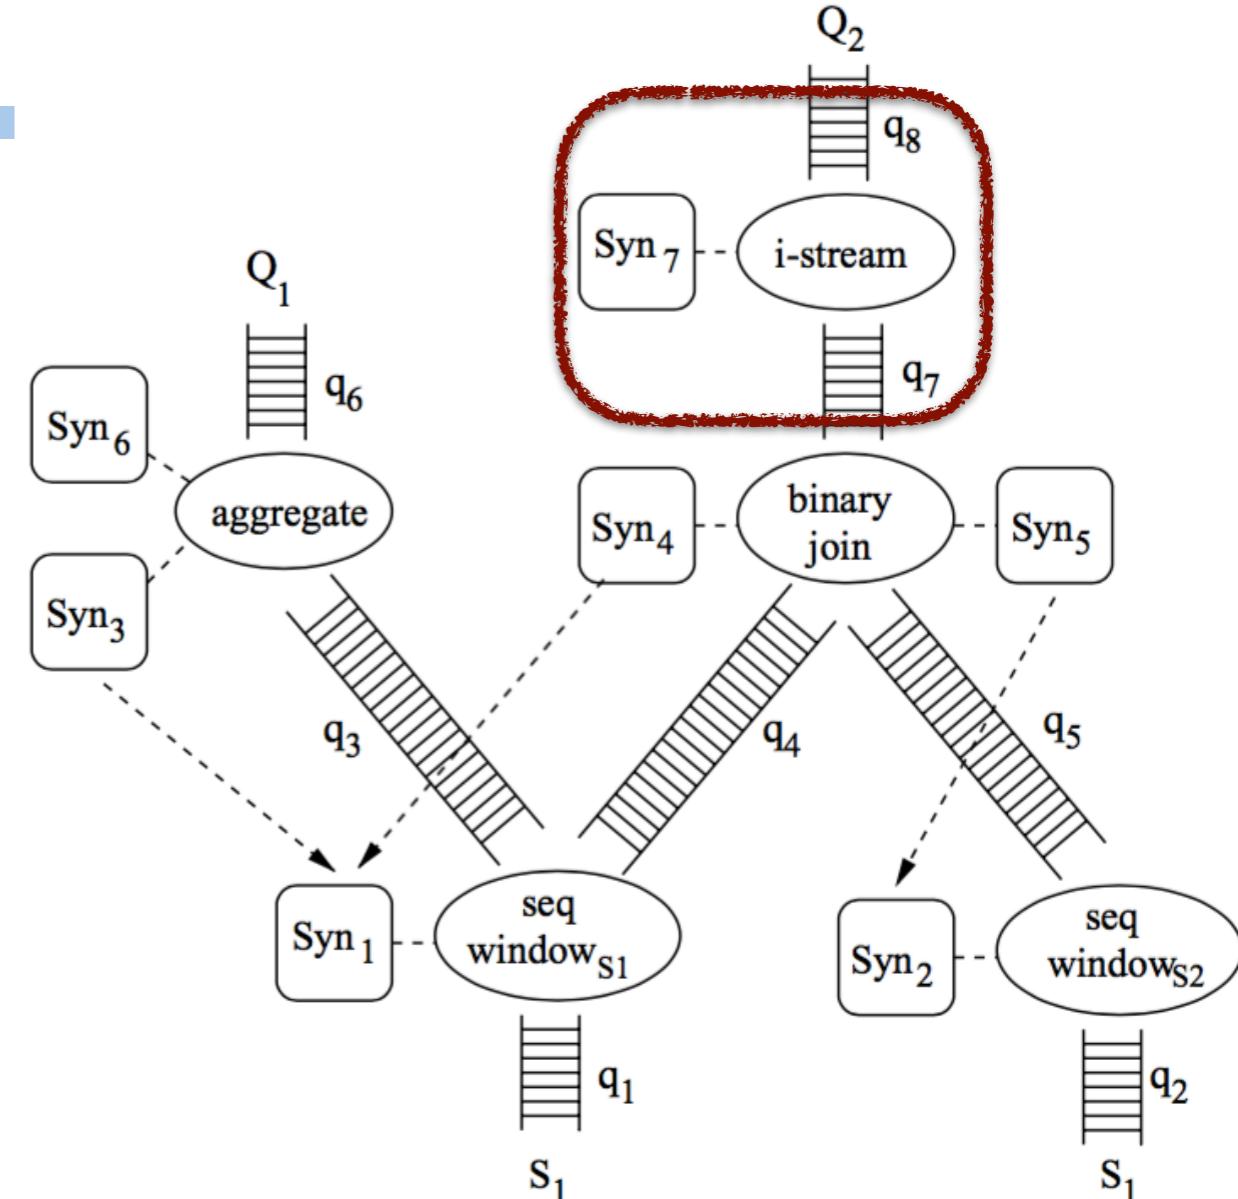
- ◆ In a naive implementation, whenever an insertion appears in q_7 , it outputs the inserted value to q_8
- ◆ This seem to be exactly the meaning of Istream

But care must be taken!

- ◆ What if there is both an insertion and a deletion of a same tuple at the same time?

Syn₇ buffers the insertion and deletion at a given time t

- ◆ The operator outputs insertions without deletions at t to q_8



Q1: select B, max(A)
from S1 [Rows 50000]
group by B;

Q2: select Istream(*)
from S1 [Rows 40000]
S2 [Range 600s]
where S1.A = S2.A;

On generated timestamps

- Special care must be taken with the timestamps that are attached to generated tuples in queues

On generated timestamps

- **Special care must be taken with the timestamps that are attached to generated tuples in queues**
 - ◆ When an input tuple at time t causes a tuple based window to slide, the generated insertion and deletion have timestamp t

On generated timestamps

- **Special care must be taken with the timestamps that are attached to generated tuples in queues**
 - ◆ When an input tuple at time t causes a tuple based window to slide, the generated insertion and deletion have timestamp t
 - ◆ When two tuples are joined, the timestamp is the higher of the two

On generated timestamps

- **Special care must be taken with the timestamps that are attached to generated tuples in queues**
 - ◆ When an input tuple at time t causes a tuple based window to slide, the generated insertion and deletion have timestamp t
 - ◆ When **two tuples are joined, the timestamp is the higher of the two**
 - ◆ When a new aggregation result is generated due to an insertion or deletion at t , the generated insertion and deletion has timestamp t

On generated timestamps

■ Special care must be taken with the timestamps that are attached to generated tuples in queues

- ◆ When an input tuple at time t causes a tuple based window to slide, the generated insertion and deletion have timestamp t
- ◆ When two tuples are joined, the timestamp is the higher of the two
- ◆ When a new aggregation result is generated due to an insertion or deletion at t , the generated insertion and deletion has timestamp t
- ◆ A tuple arriving at t in a stream with a time based window with size T , generates an insertion at t , and a deletion at $t+T+1$

On generated timestamps

- **Special care must be taken with the timestamps that are attached to generated tuples in queues**
 - ◆ When an input tuple at time t causes a tuple based window to slide, the generated insertion and deletion have timestamp t
 - ◆ When **two tuples are joined, the timestamp is the higher of the two**
 - ◆ When a new aggregation result is generated due to an insertion or deletion at t , the generated insertion and deletion has timestamp t
 - ◆ A tuple arriving at t in a stream with a time based window with size T , generates an insertion at t , and a deletion at $t+T+1$
 - The $+1$ depends on the granularity of the system

System operators

- **Real systems have some more operators, to deal with several low-level details. E.g.**

- ◆ To asynchronously receive the tuples from an input stream, and transforming them to the internal representation
- ◆ To manage load balancing
- ◆ To merge plans from views
- ◆ To send stream results to remote clients
- ◆ ...

STREAM operators

Name	Operator Type	Description
seq-window	Stream-to-relation	Implements time-based, tuple-based, and partitioned windows
select	Relation-to-relation	Filters tuples based on predicate(s)
project	Relation-to-relation	Duplicate-preserving projection
binary-join	Relation-to-relation	Joins two input relations
mjoin	Relation-to-relation	Multiway join from [43]
union	Relation-to-relation	Bag union
except	Relation-to-relation	Bag difference
intersect	Relation-to-relation	Bag intersection
antisemijoin	Relation-to-relation	Antisemijoin of two input relations
aggregate	Relation-to-relation	Performs grouping and aggregation
duplicate-eliminate	Relation-to-relation	Performs duplicate elimination
i-stream	Relation-to-stream	Implements Istream semantics
d-stream	Relation-to-stream	Implements Dstream semantics
r-stream	Relation-to-stream	Implements Rstream semantics
stream-shepherd	System operator	Handles input streams arriving over the network
stream-sample	System operator	Samples specified fraction of tuples
stream-glue	System operator	Adapter for merging a stream-producing view into a plan
rel-glue	System operator	Adapter for merging a relation-producing view into a plan
shared-rel-op	System operator	Materializes a relation for sharing
output	System operator	Sends results to remote clients

Query plan generation

■ Apply heuristics to generate a first plan

- ◆ Execute selections before joins (as in DBMSs)
- ◆ Use indexed nested loop joins for joining with synopses
- ◆ Apply filters before ranges
- ◆ Share synopses and operators whenever possible

Query plan generation

■ Apply heuristics to generate a first plan

- ◆ Execute selections before joins (as in DBMSs)
- ◆ Use indexed nested loop joins for joining with synopses
- ◆ Apply filters before ranges
- ◆ Share synopses and operators whenever possible

■ Upon execution, statistical data on what has been executed can be used to rearrange initial query plans

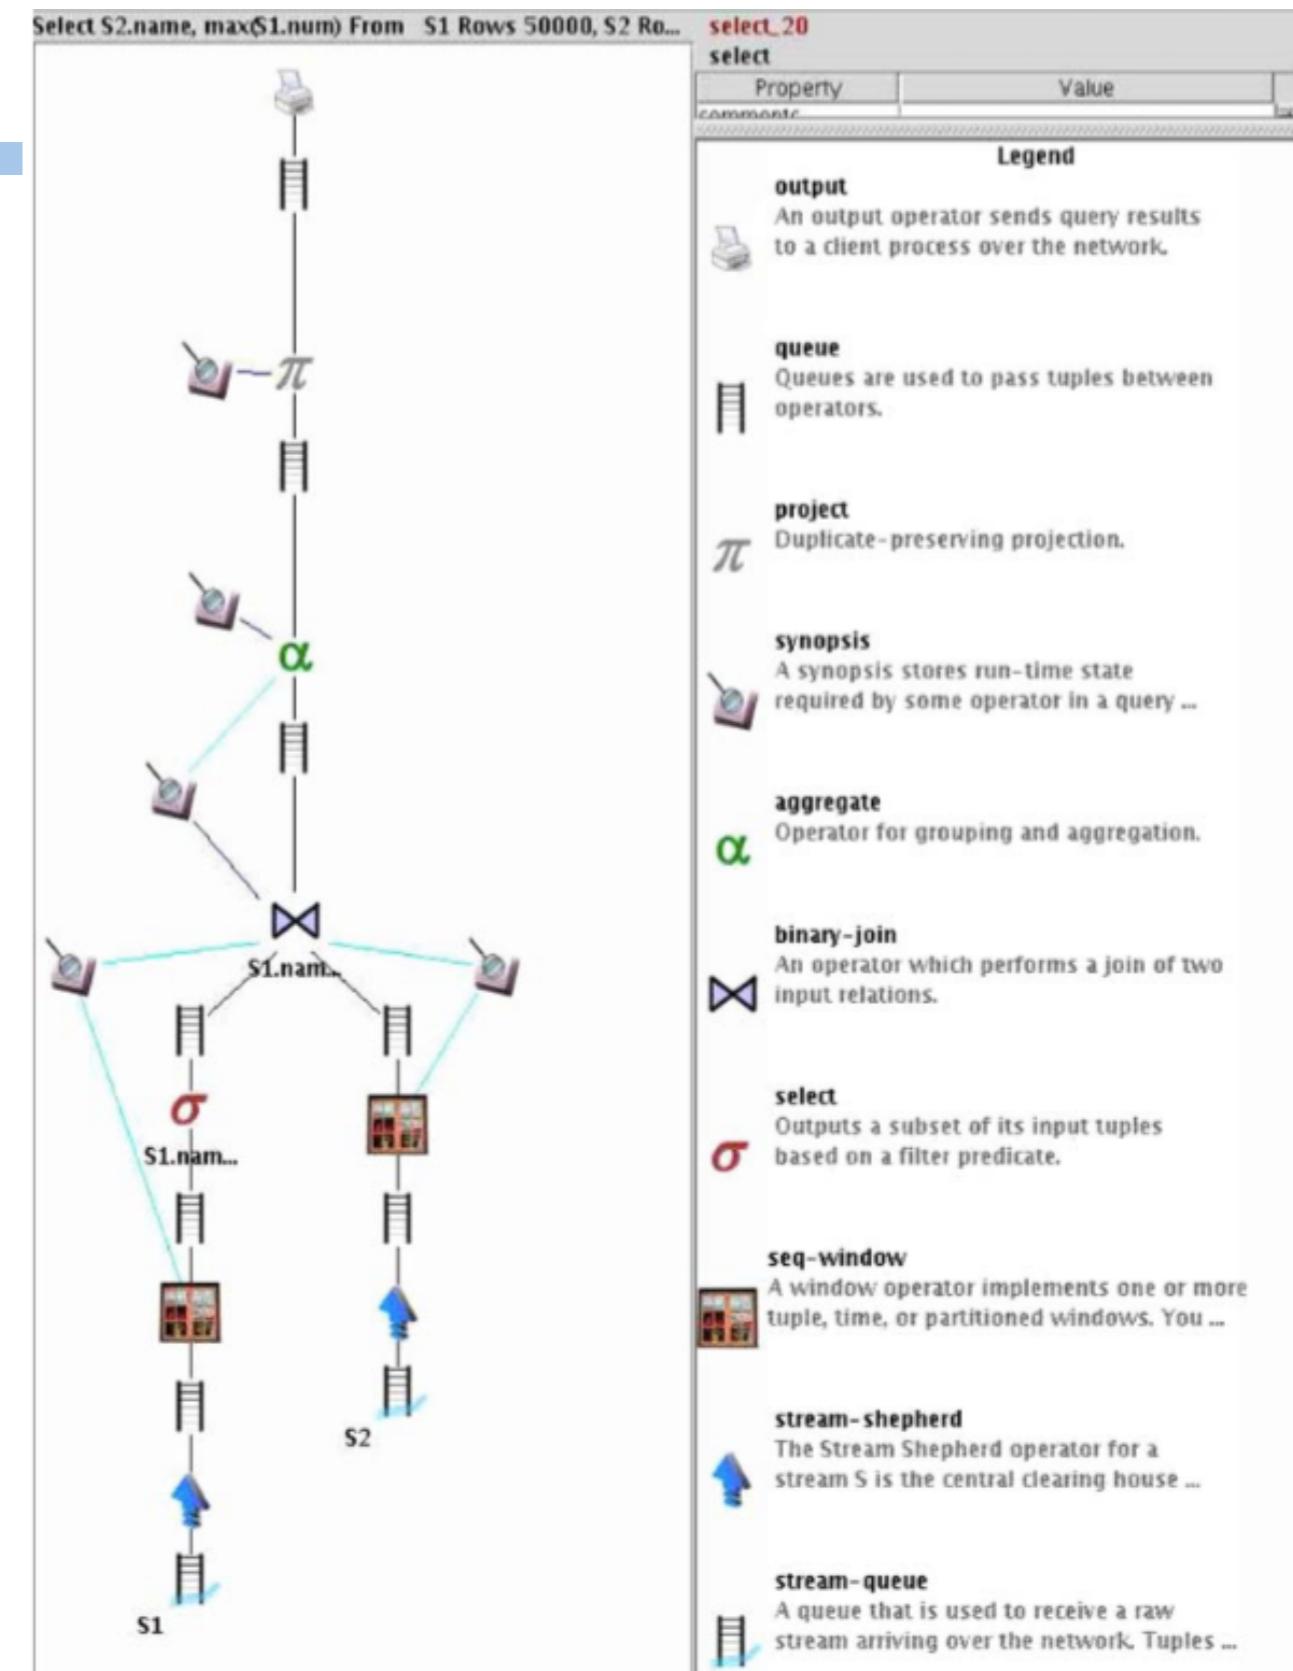
- ◆ Requires a *monitor* for monitoring the plans that are being executed
- ◆ It can use optimisation techniques similar to those of DBMSs, based on what has been monitored

Query execution

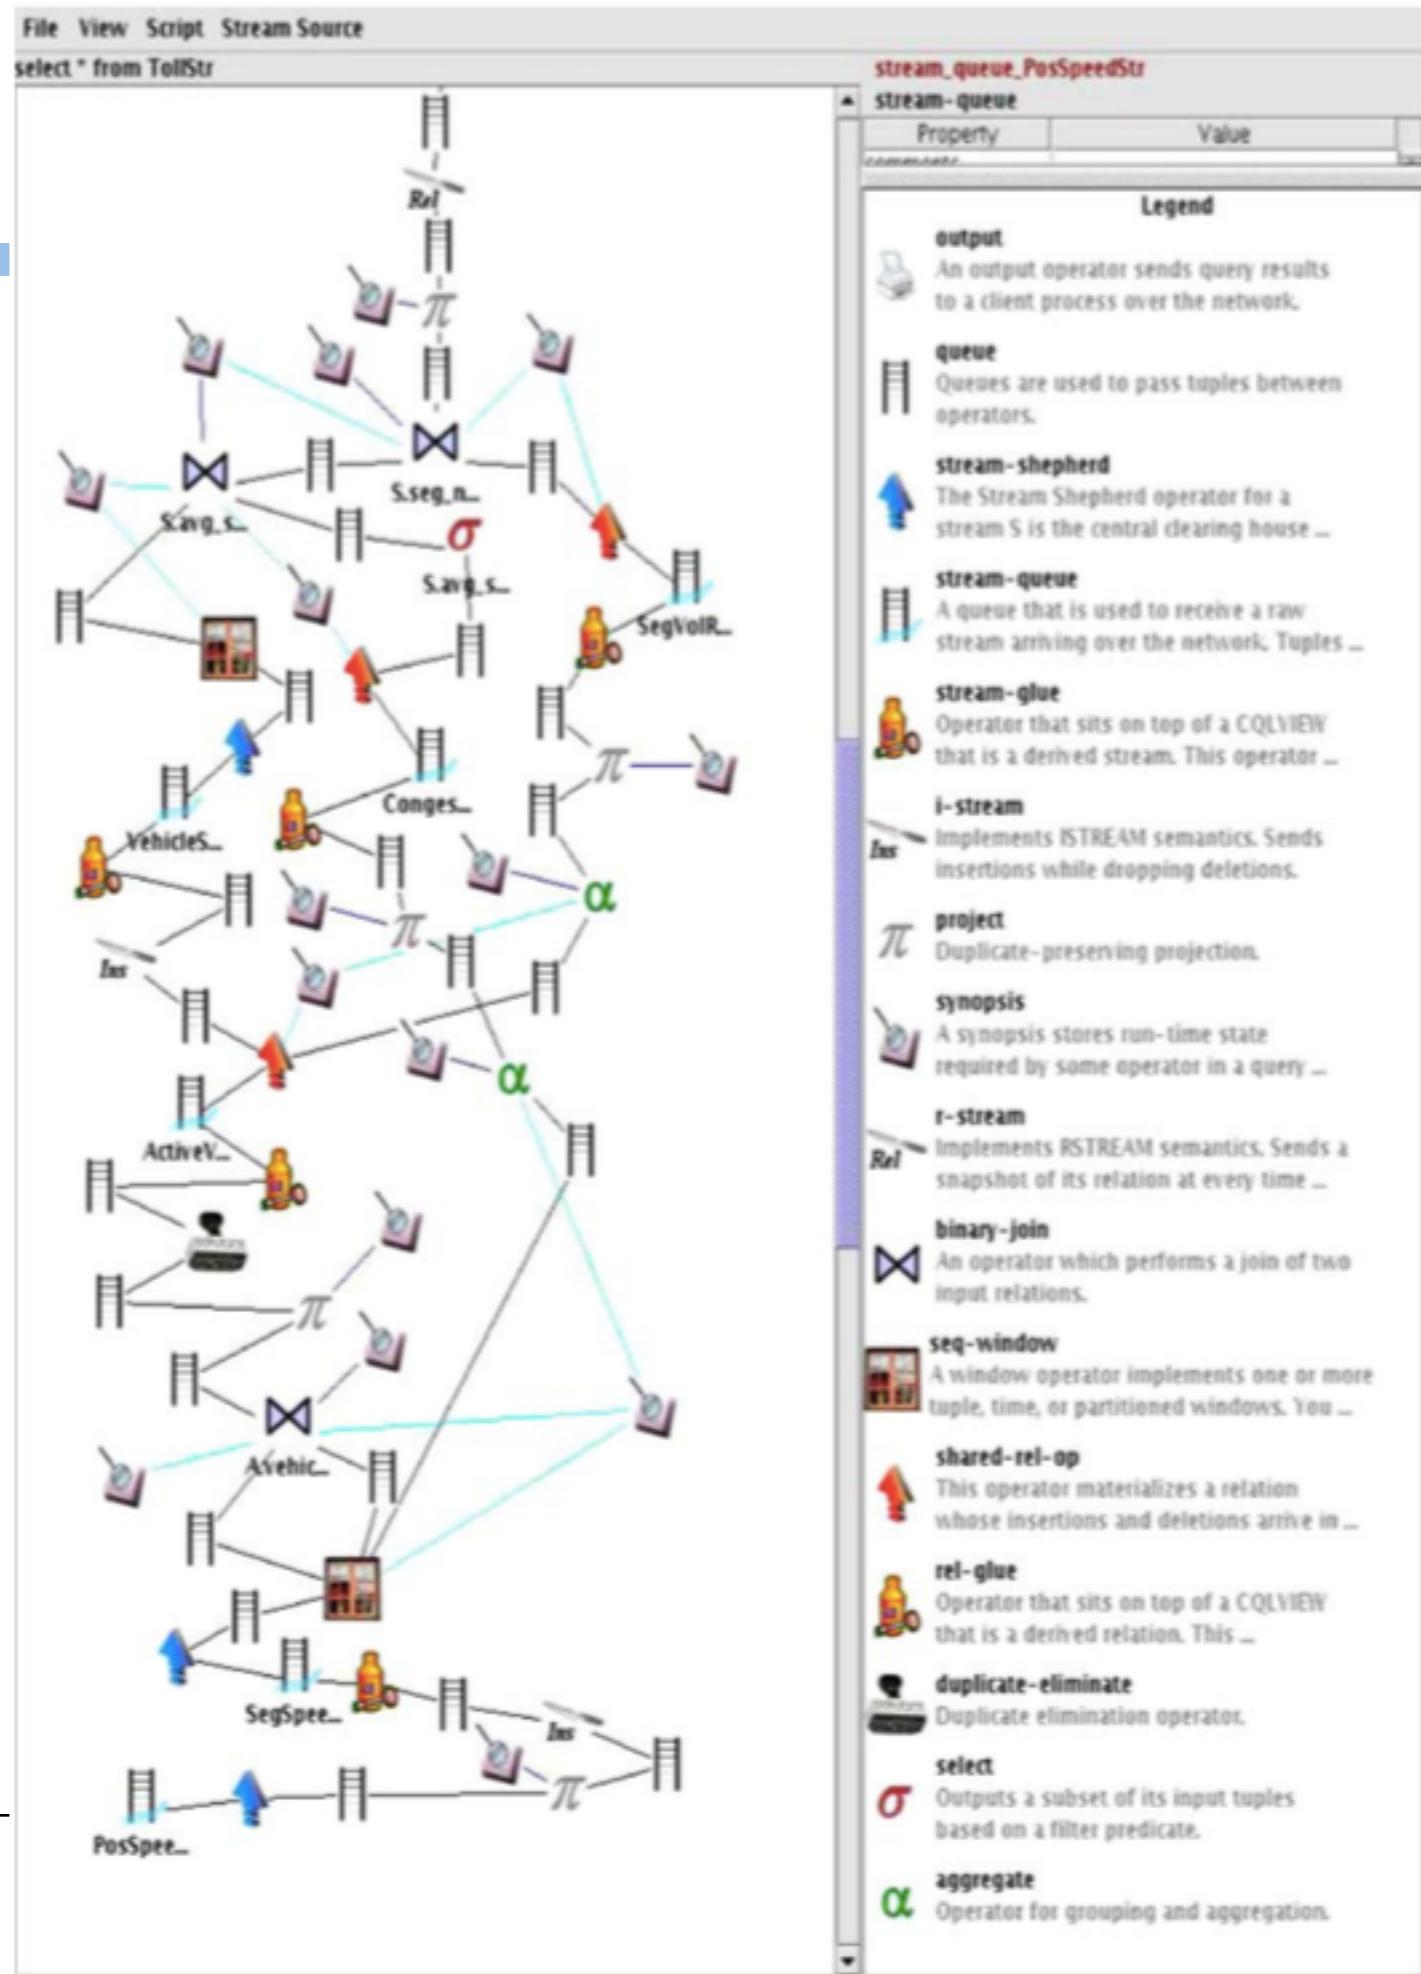
- The execution of a plan is controlled by a global scheduler
- The scheduler is invoked periodically
 - ◆ It selects an operator to execute
 - ◆ Calls a specific procedure for the operator, passing it the required parameters (with data from queues and synopses)
- The scheduler can be a simple round-robin over operators, but may include parallel execution, and more sophisticated scheduling algorithms
 - This is outside the scope of this course

An example plan in STREAM

```
select S2.name, max(S1.n)
from S1 [Rows 50000]
      S2 [Rows 50000]
where S1.name <= 'i' and
      S1.n = S2.n
group by S2.name;
```



Linear Road plan



CQL in Oracle CEP

Oracle Event Processor

■ Java Server for event driven applications

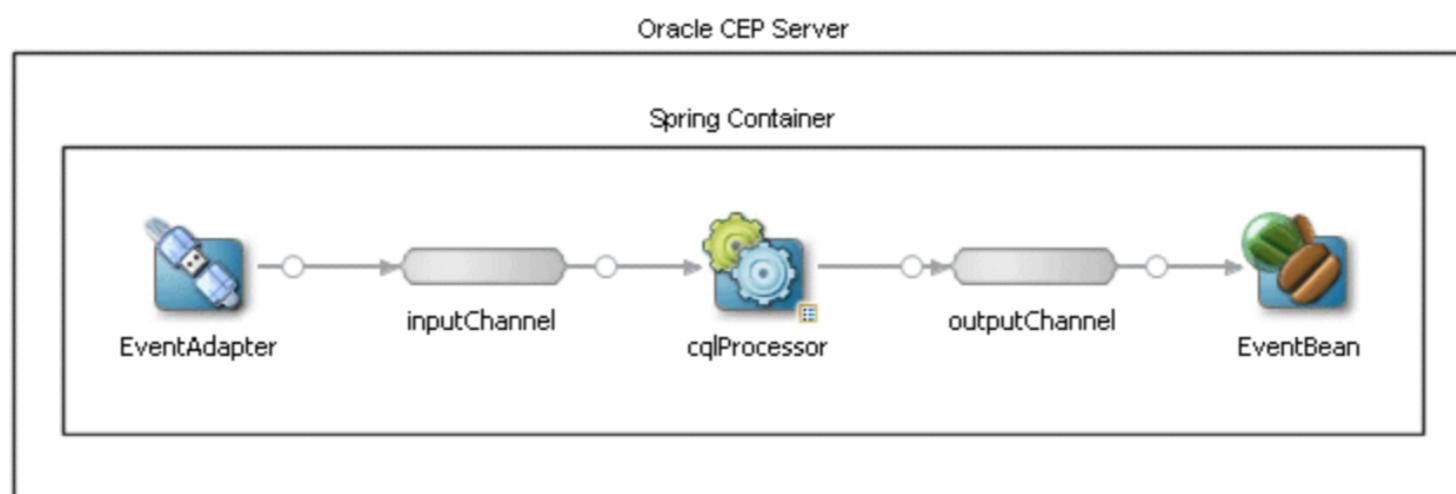
- ◆ A lightweight Java application container
- ◆ Includes a service engine that allows one to write CQL queries

■ Provides tools for developing applications

- ◆ Oracle CEP IDE for Eclipse
- ◆ Oracle CEP Visualizer

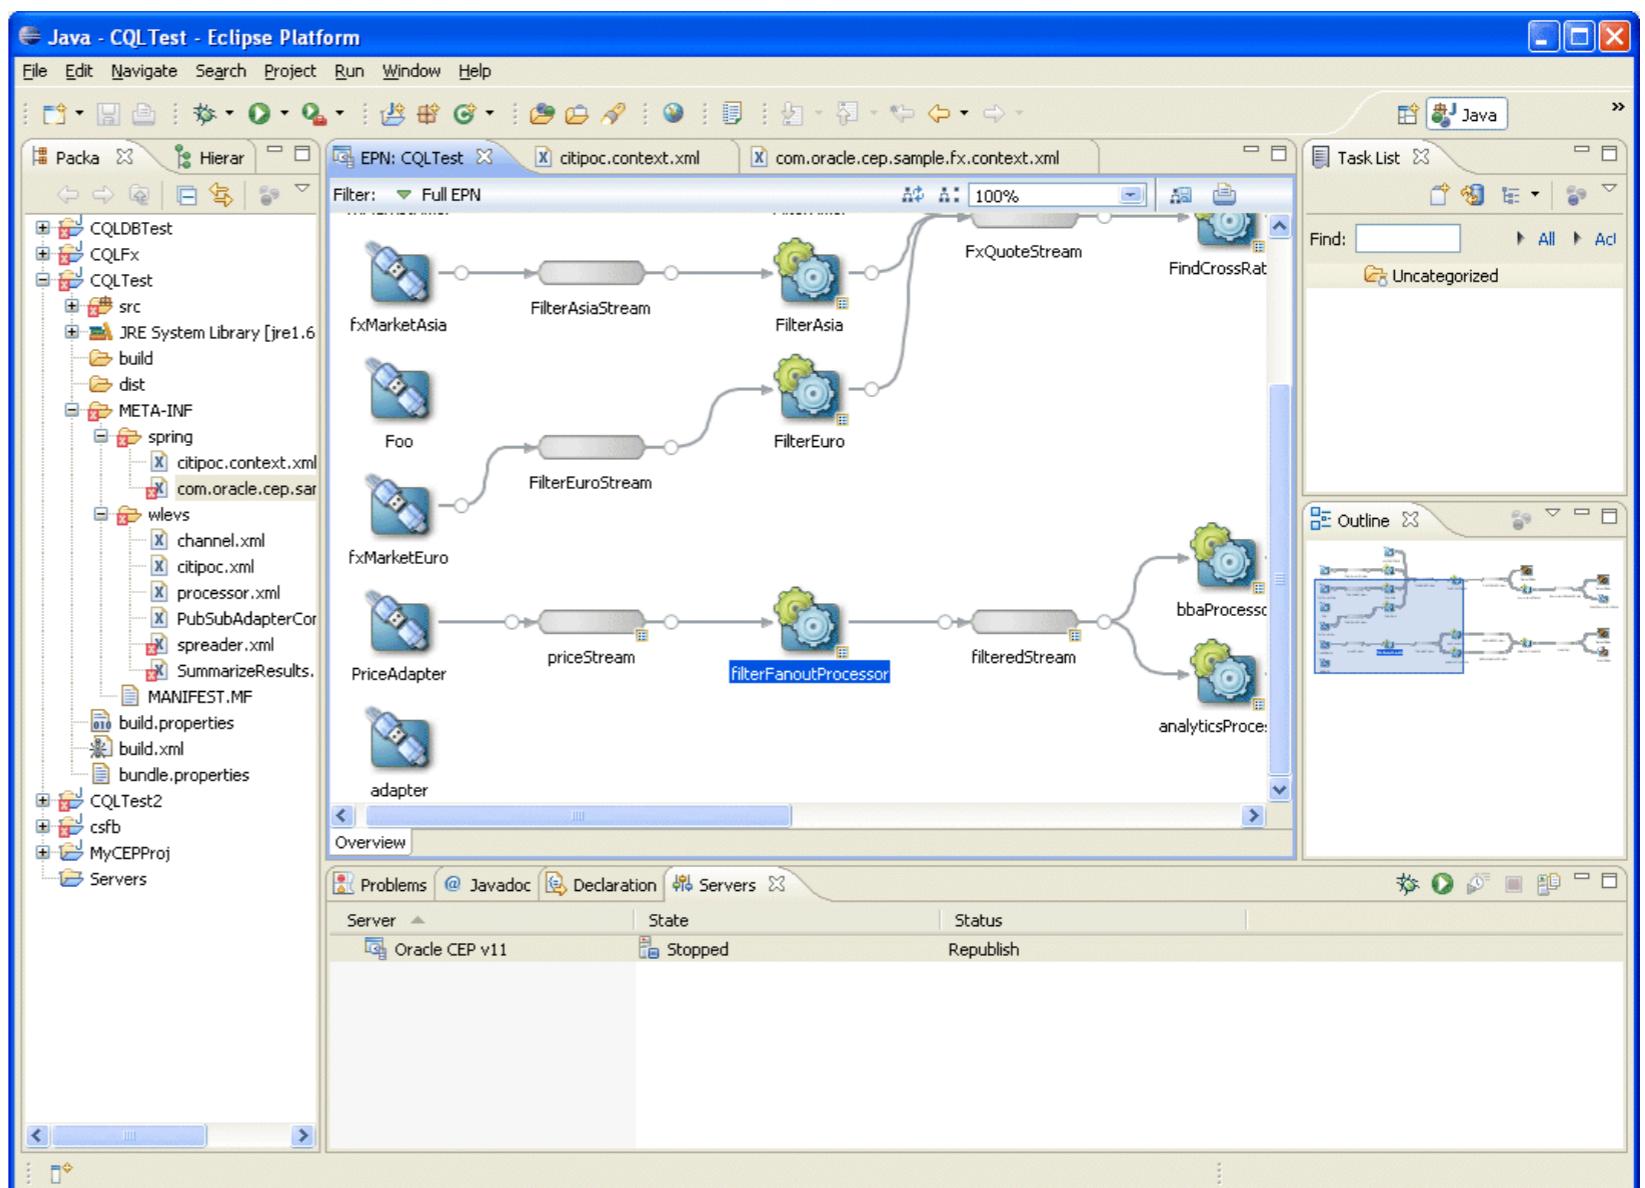
Using Oracle CEP

- With the IDE for Eclipse and the Visualizer, one can:
 - ◆ Create an Event Processing Network
 - ◆ Associate CQL queries to Oracle CQL processors
 - ◆ Package applications and deploy them to Oracle server
- Event adapters support several datasources: JMS, http publish-subscribe, spring beans, relational tables, files, etc



Event Processing Network (EPN)

- One can connect the output of an event processor to the input of another event processor
 - ◆ This way we form an EPN
 - ◆ Each processor contains a CQL query (possibly with views)



Oracle's support of CQL

Oracle's support of CQL

■ Pretty much everything that we covered in the course!

- ◆ With some extra features (details, functions etc)
- ◆ And some (small) differences in the syntax

Putting it to work

- A lot of “bureaucracy” involved :-(

Putting it to work

- A lot of “bureaucracy” involved :-(
 - ◆ Configuring input adapters

Putting it to work

■ A lot of “bureaucracy” involved :-(

- ◆ Configuring input adapters
- ◆ Configuring beans

Putting it to work

■ A lot of “bureaucracy” involved :-(

- ◆ Configuring input adapters
- ◆ Configuring beans
- ◆ Configuring channels

Putting it to work

■ A lot of “bureaucracy” involved :-(

- ◆ Configuring input adapters
- ◆ Configuring beans
- ◆ Configuring channels
- ◆ Configuring CQL server, and Event Processing Language processors

Putting it to work

■ A lot of “bureaucracy” involved :-(

- ◆ Configuring input adapters
- ◆ Configuring beans
- ◆ Configuring channels
- ◆ Configuring CQL server, and Event Processing Language processors
- ◆ Configuring caching

Putting it to work

■ A lot of “bureaucracy” involved :-(

- ◆ Configuring input adapters
- ◆ Configuring beans
- ◆ Configuring channels
- ◆ Configuring CQL server, and Event Processing Language processors
- ◆ Configuring caching
- ◆ ...

Putting it to work

■ A lot of “bureaucracy” involved :-)

- ◆ Configuring input adapters
- ◆ Configuring beans
- ◆ Configuring channels
- ◆ Configuring CQL server, and Event Processing Language processors
- ◆ Configuring caching
- ◆ ...
- ◆ Assembling and deploying applications

Putting it to work

■ A lot of “bureaucracy” involved :-)

- ◆ Configuring input adapters
- ◆ Configuring beans
- ◆ Configuring channels
- ◆ Configuring CQL server, and Event Processing Language processors
- ◆ Configuring caching
- ◆ ...
- ◆ Assembling and deploying applications

■ All of this can be done on Eclipse, which helps...

Event Processor Visualizer

- A browser based tool where one can develop, configure and view aspects of event driven applications

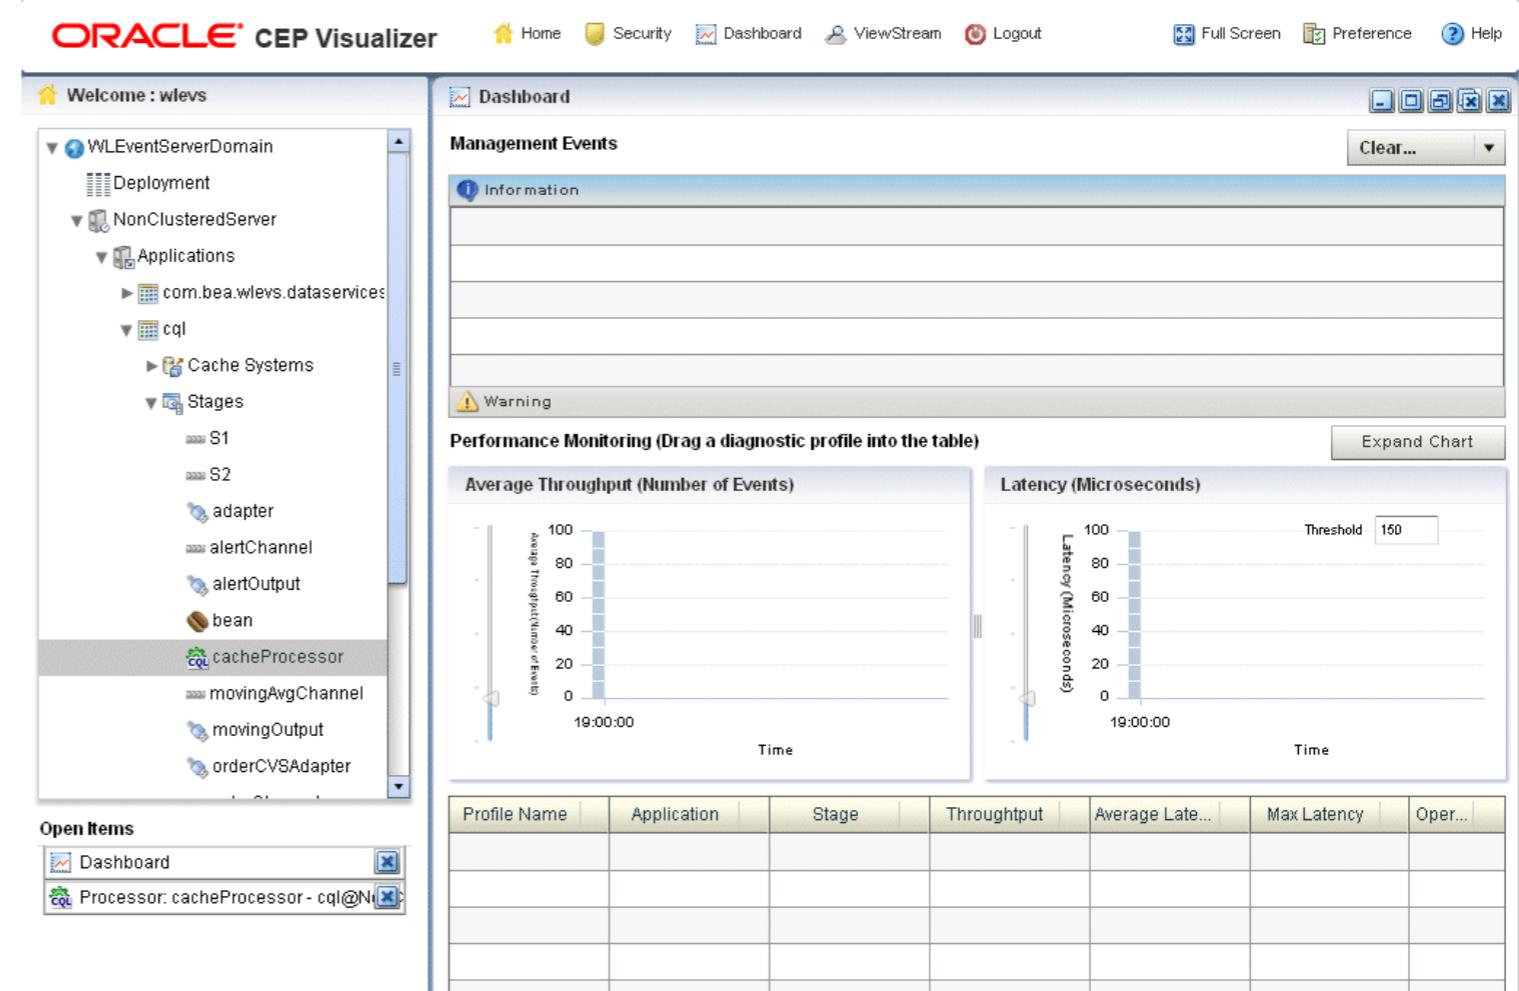
The screenshot shows the Oracle CEP Visualizer interface. The top navigation bar includes links for Home, Security, Dashboard, ViewStream, Logout, Full Screen, Preference, and Help. The left sidebar displays a hierarchical tree structure under 'WLEventServerDomain' for 'NonClusteredServer' and its components: Applications (com.bea.wlevs.dataservices, cql), Cache Systems, Stages (S1, S2), and various adapters and beans. The main dashboard area features a 'Management Events' section with tabs for Information and Warning. Below this is a 'Performance Monitoring' section with two charts: 'Average Throughput (Number of Events)' and 'Latency (Microseconds)'. Both charts show data for the 'cacheProcessor' stage at 19:00:00. A threshold of 150 is indicated for latency. At the bottom, there is a table titled 'Open Items' with two entries: 'Dashboard' and 'Processor: cacheProcessor - cql@N'.

Event Processor Visualizer

- A browser based tool where one can develop, configure and view aspects of event driven applications

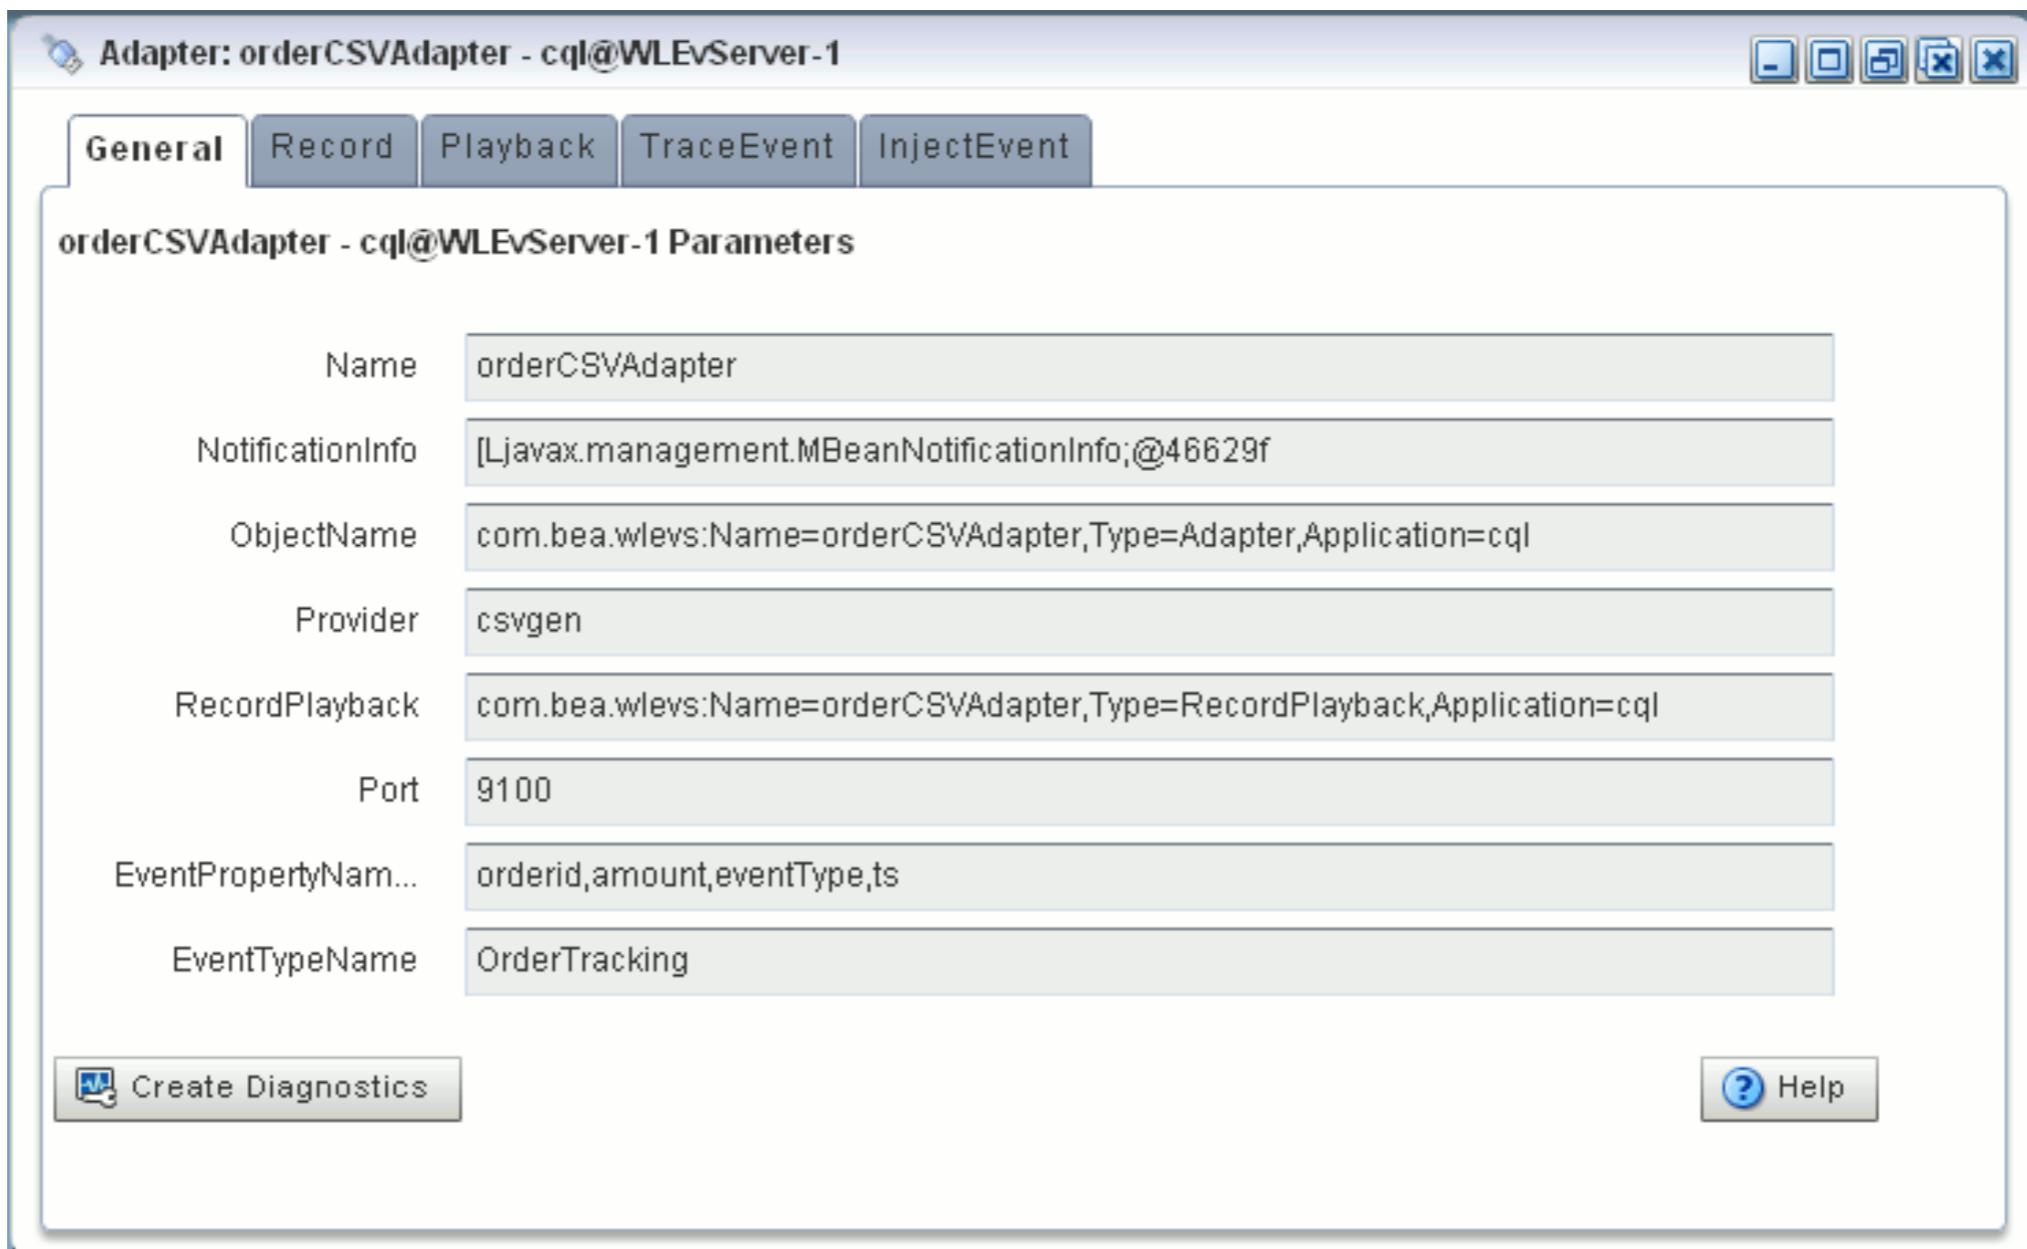
- More specifically, one can:

- ◆ Define EPNs
- ◆ Register CQL queries
 - Including a CQL query wizard
 - (for non-CQL-experts)
 - A query plan viewer
- ◆ Configure event adapters
- ◆ View events in dashboards
- ◆ View event statistics



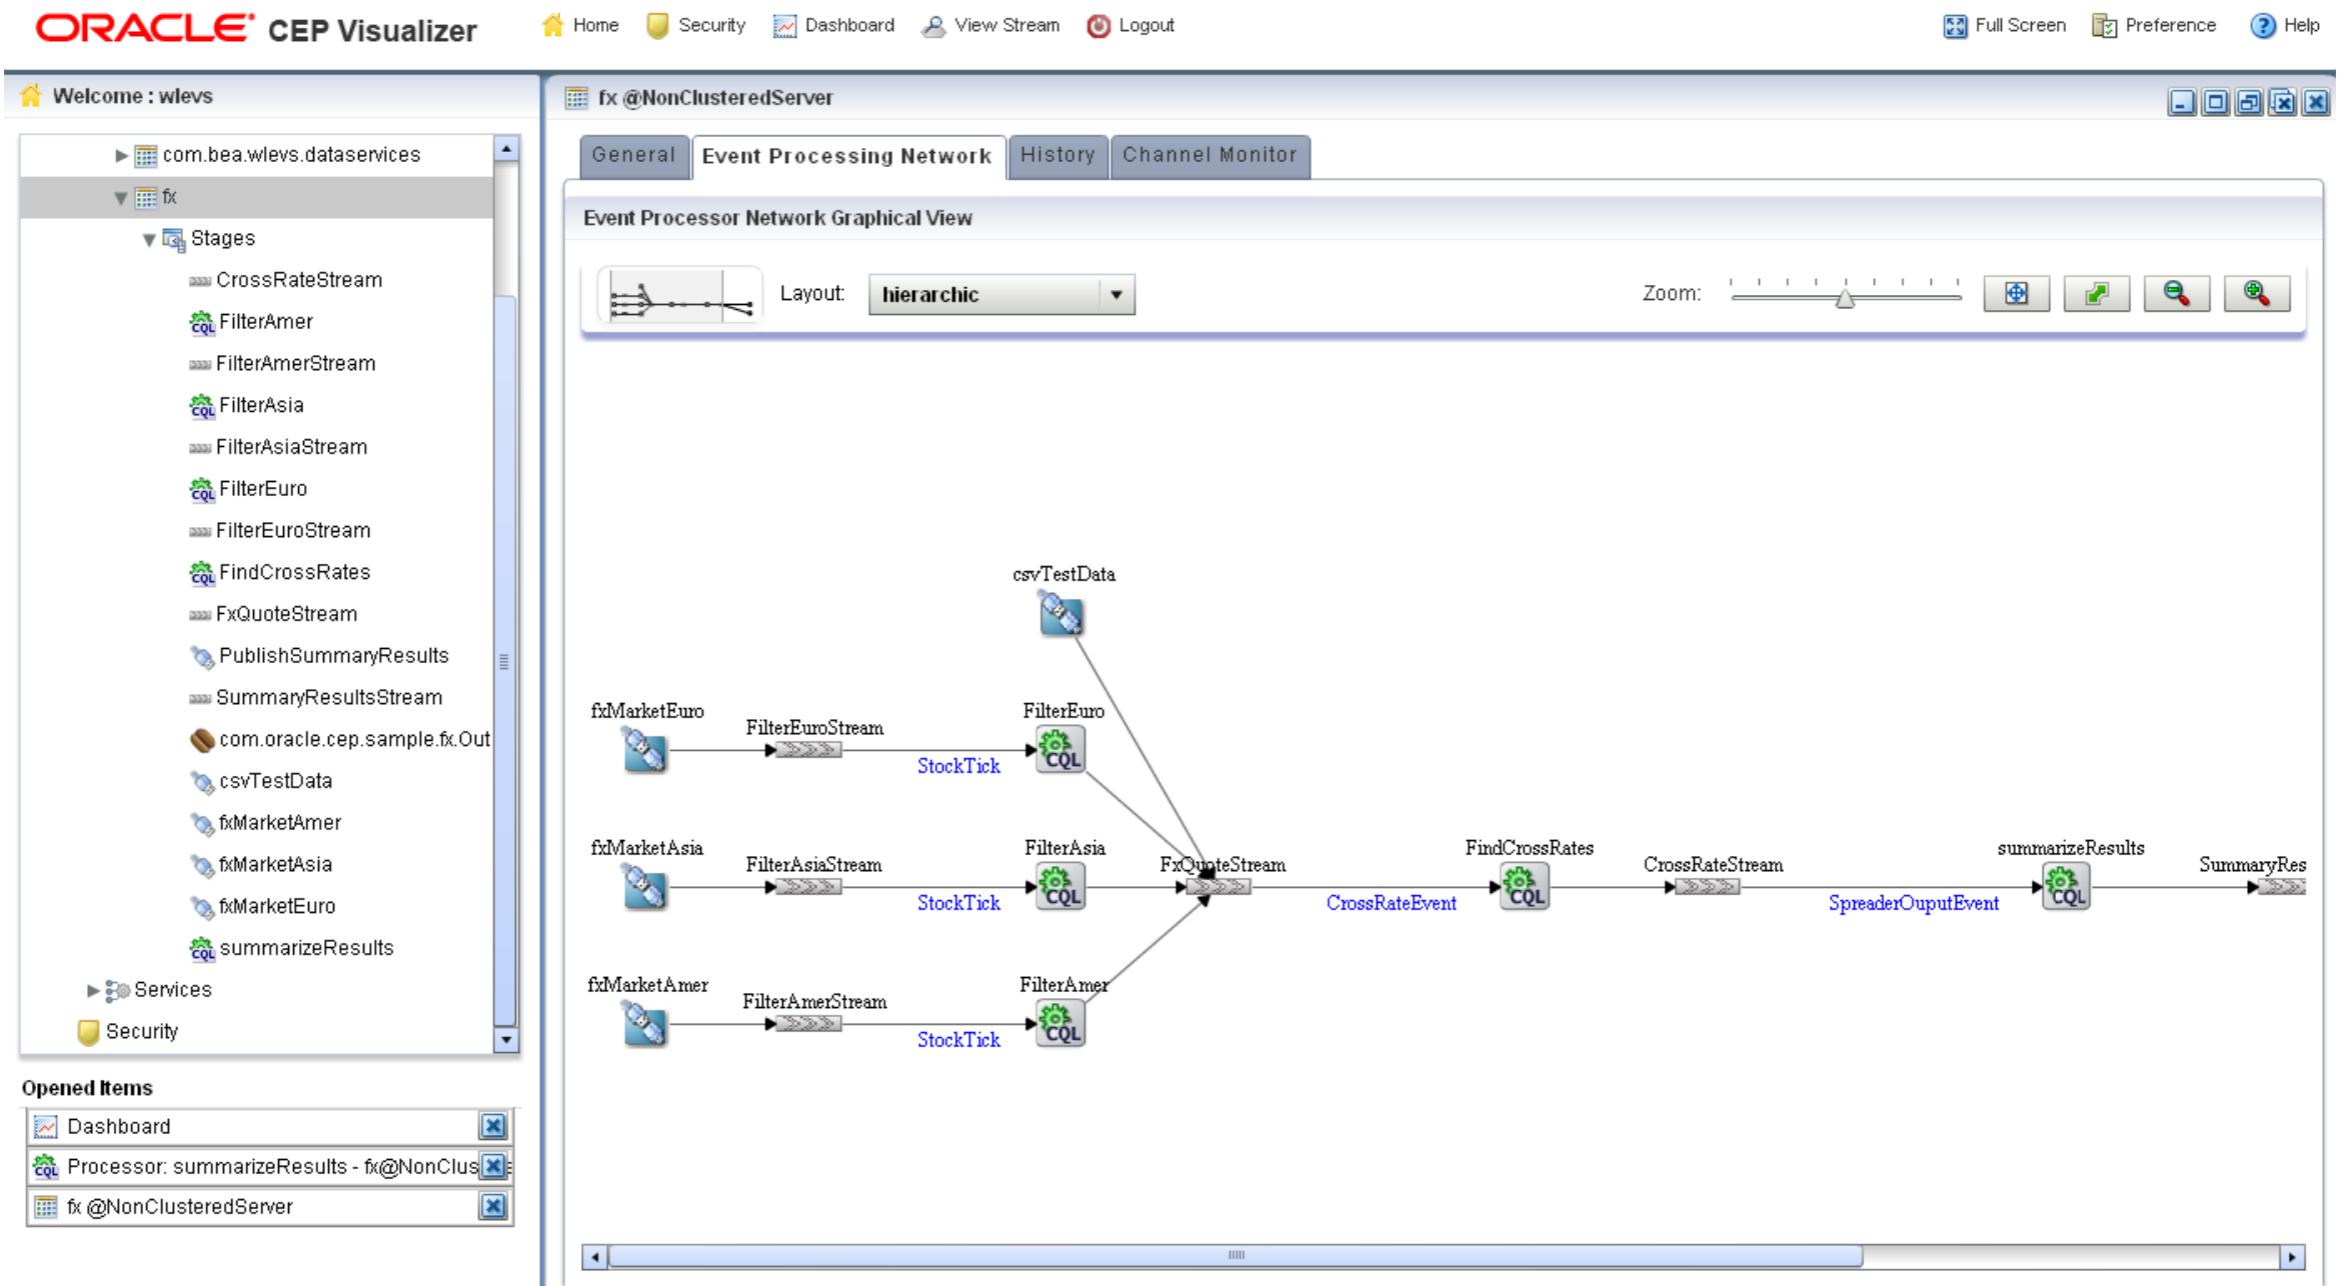
Some Visualizer snapshots

Defining an adapter



Some Visualizer snapshots

Defining an EPN



Some Visualizer snapshots

Defining a CQL query

Processor: summarizeResults - fx@NonClusteredServer

General Record Playback Query Wizard CQL Rules Query Plan Trace Event Inject Event

View Query All Rules

Rule ID	Rule	Type	Running
Rule	<code>select crossRate1 crossRate2 as crossRatePair, count(*) as totalCount, avg(internalPrice) as averageInternalPrice from CrossRateStream group by crossRate1,crossRate2 having count(*) > 0</code>	QUERY	true
SummarizeResultsRule	<code>select crossRate1 crossRate2 as crossRatePair, count(*) as totalCount, :1 as averageInternalPrice from CrossRateStream group by crossRate1,crossRate2 having :2</code>	QUERY	true

Working Area - for Modify and Delete Operation, select a rule from the table

Query ID: SummarizeResultsRule

Query:

```
select crossRate1 || crossRate2 as crossRatePair, count(*) as totalCount, :1 as averageInternalPrice from CrossRateStream group by crossRate1,crossRate2 having :2
```

Enable: true false

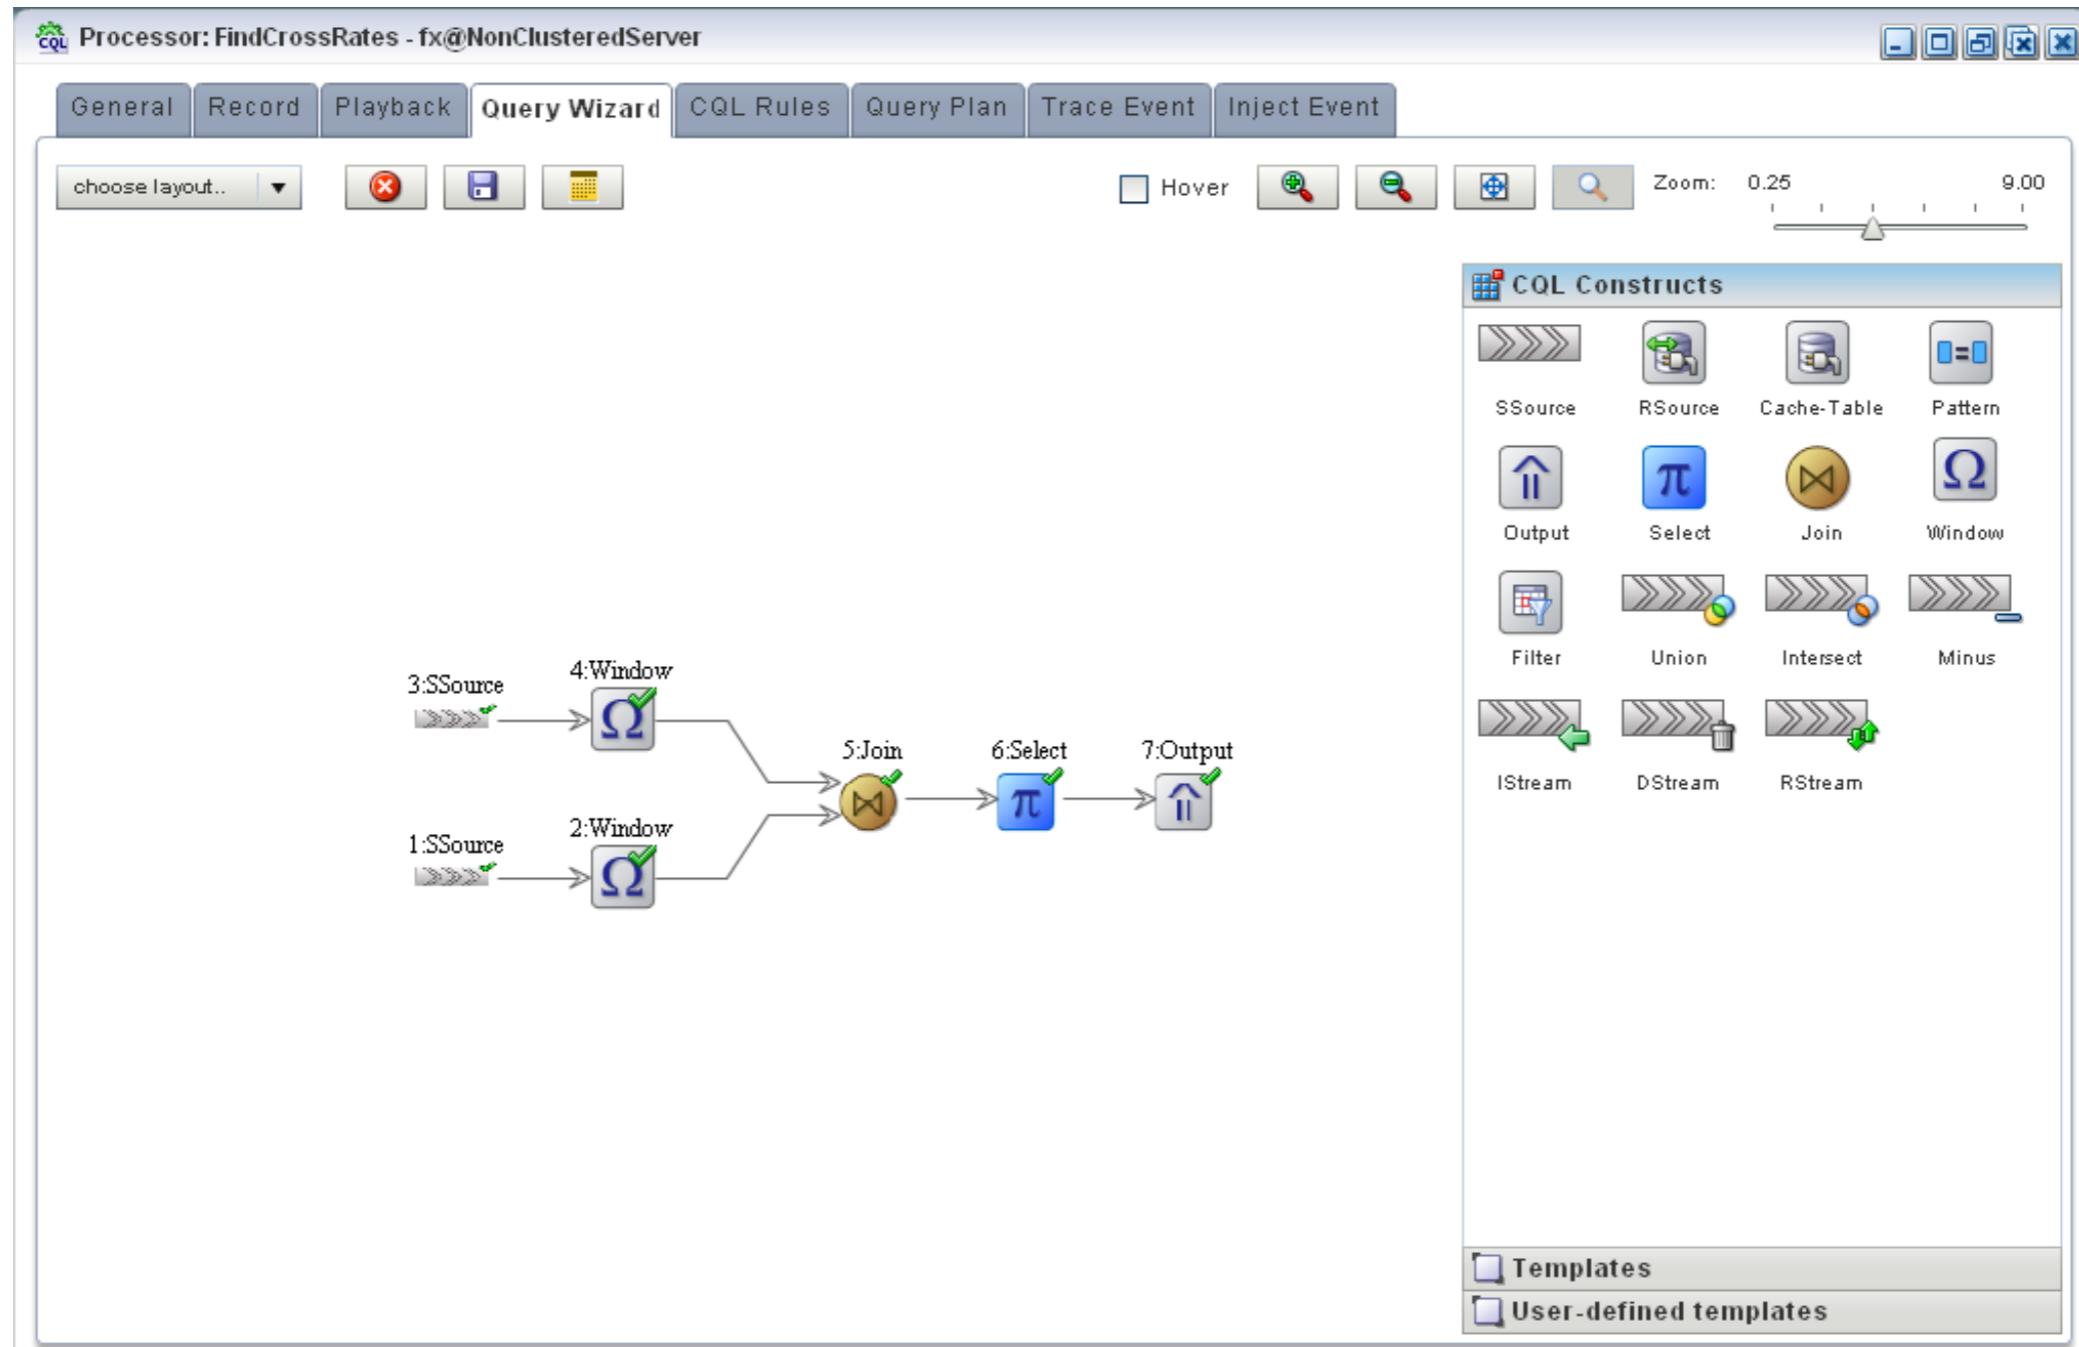
Parameters

Binding Name	1	2
avacount	avg(internalPrice)	count(*) > 0

Delete Query Edit Query Add Parameter Replace Parameter Delete Parameter Stop Save Cancel Help

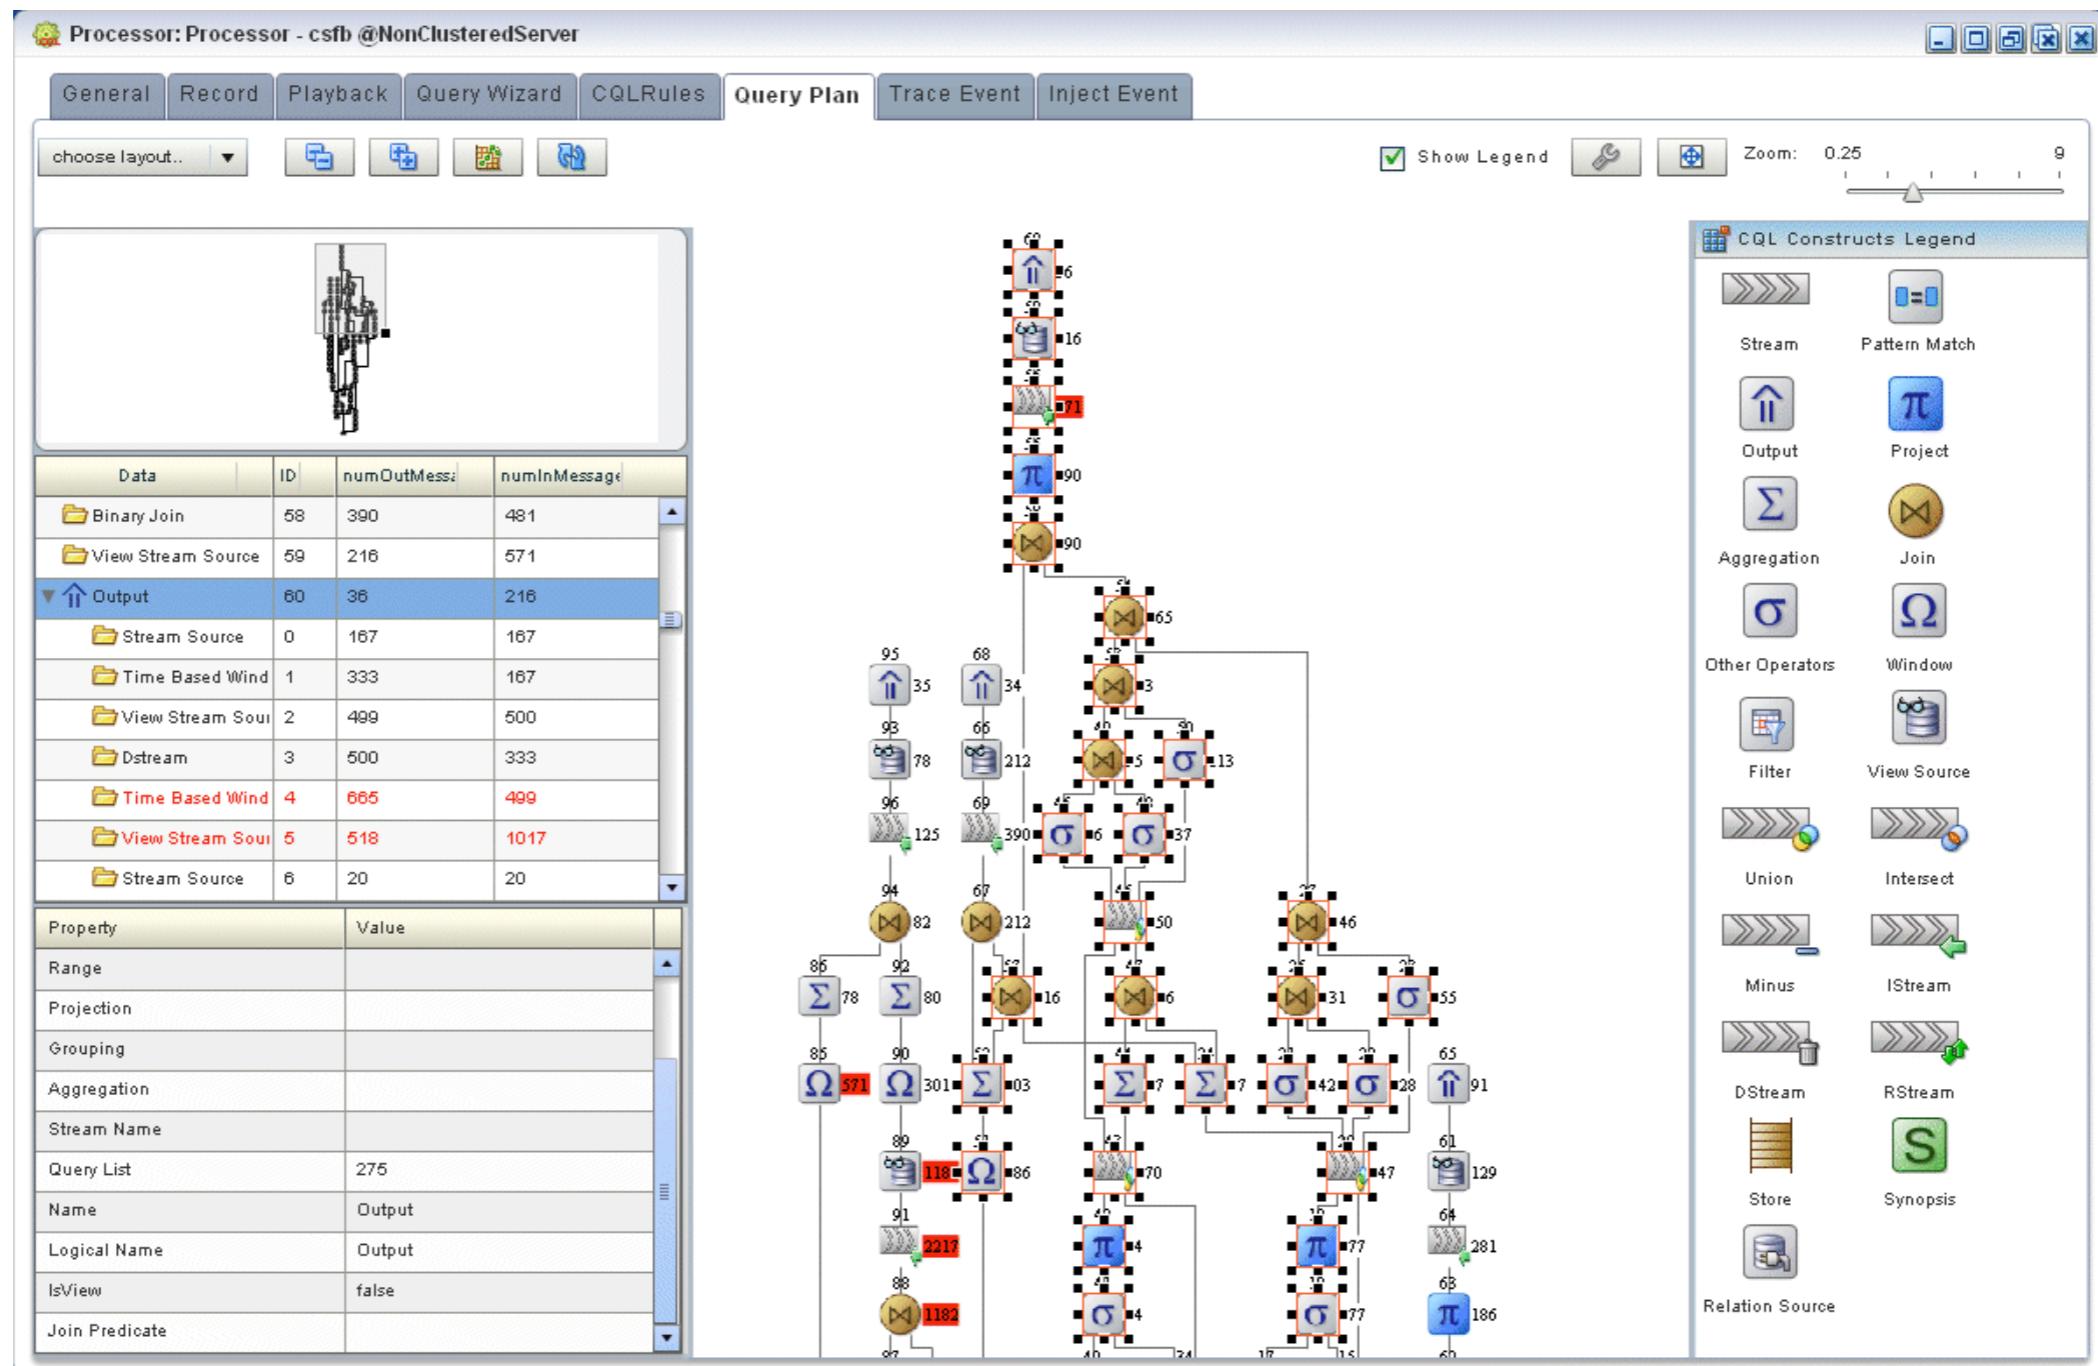
Some Visualizer snapshots

CQL “for dummies” (not for you!)



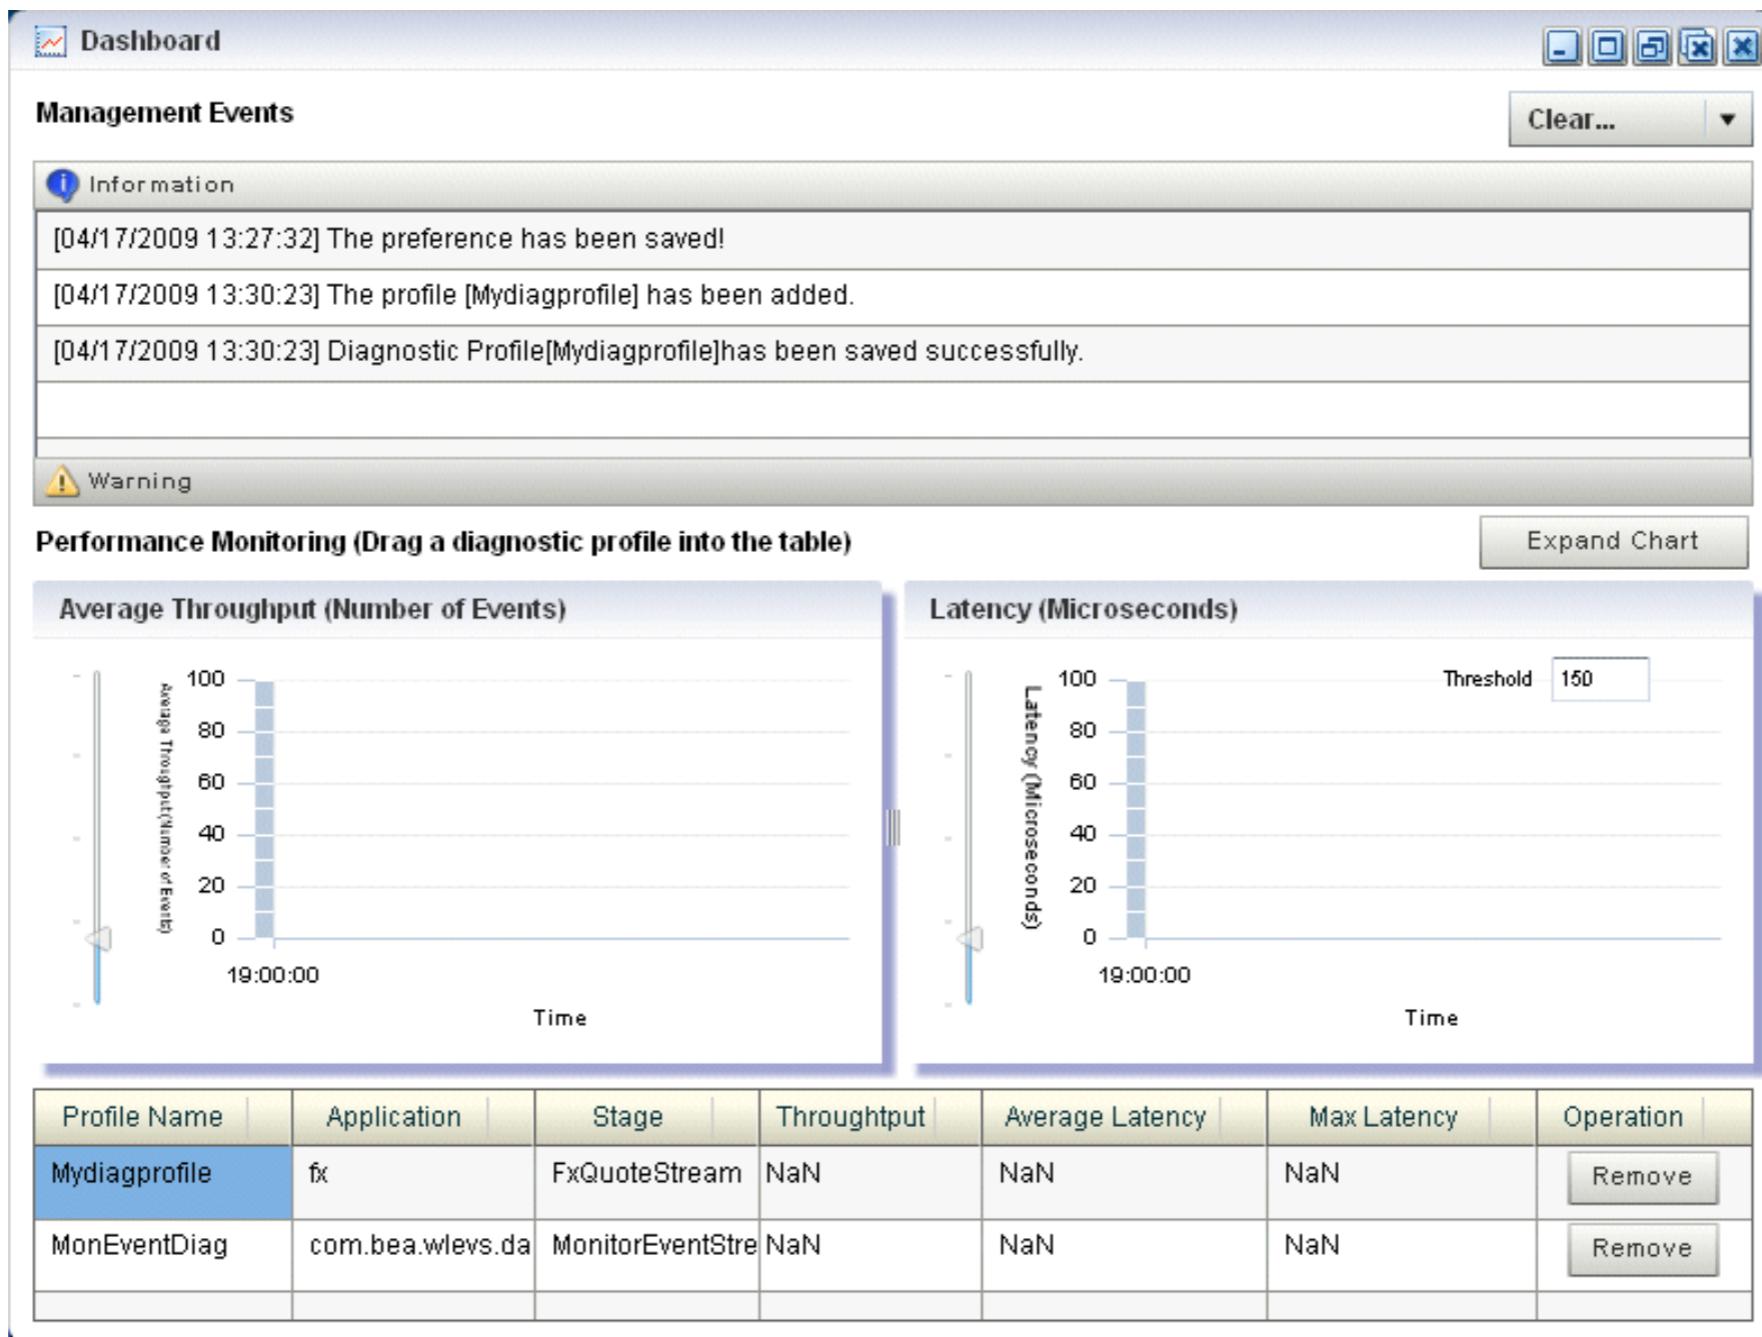
Some Visualizer snapshots

Query plans



Some Visualizer snapshots

Event dashboard and statistics



Readings

■ Oracle CEP CQL Language Reference

- ◆ https://docs.oracle.com/cd/E16764_01/doc.1111/e12048/toc.htm

■ Oracle CEP IDE Developer's Guide for Eclipse

- ◆ https://docs.oracle.com/cd/E16764_01/doc.1111/e14301/toc.htm

■ Oracle Event Processing Visualizer User Guide

- ◆ http://docs.oracle.com/cd/E28280_01/doc.1111/e14302/toc.htm

■ CURRENT DOC

<https://docs.oracle.com/middleware/1213/eventprocessing/index.html>

Further Reading and Summary



Q&A

Further Reading

■ Recommend Readings

- ◆ Arvind Arasu, Shivnath Babu, Jennifer Widom: The CQL continuous query language: semantic foundations and query execution. VLDB J. 15(2): 121-142 (2006)

■ Supplemental readings:

- ◆ Arvind Arasu et al.: STREAM: The Stanford Data Stream Management System. Stanford InfoLab Tech. Rep. (2004)
- ◆ Arvind Arasu et al.: Linear Road: A Stream Data Management Benchmark. VLDB 2004: 480-491
- ◆ Arvind Arasu, Jennifer Widom: A Denotational Semantics for Continuous Queries over Streams and Relations. SIGMOD Record 33(3): 6-12 (2004)

Further Reading and Summary



Q&A