

# Stream Processing

Lecture 8

2018/2019

# Exercises

- Combine two streams
- Combine a stream with an RDD
- Use the storedata

# Table of Contents

- Structured Streaming Programming
- Programming Spark streaming

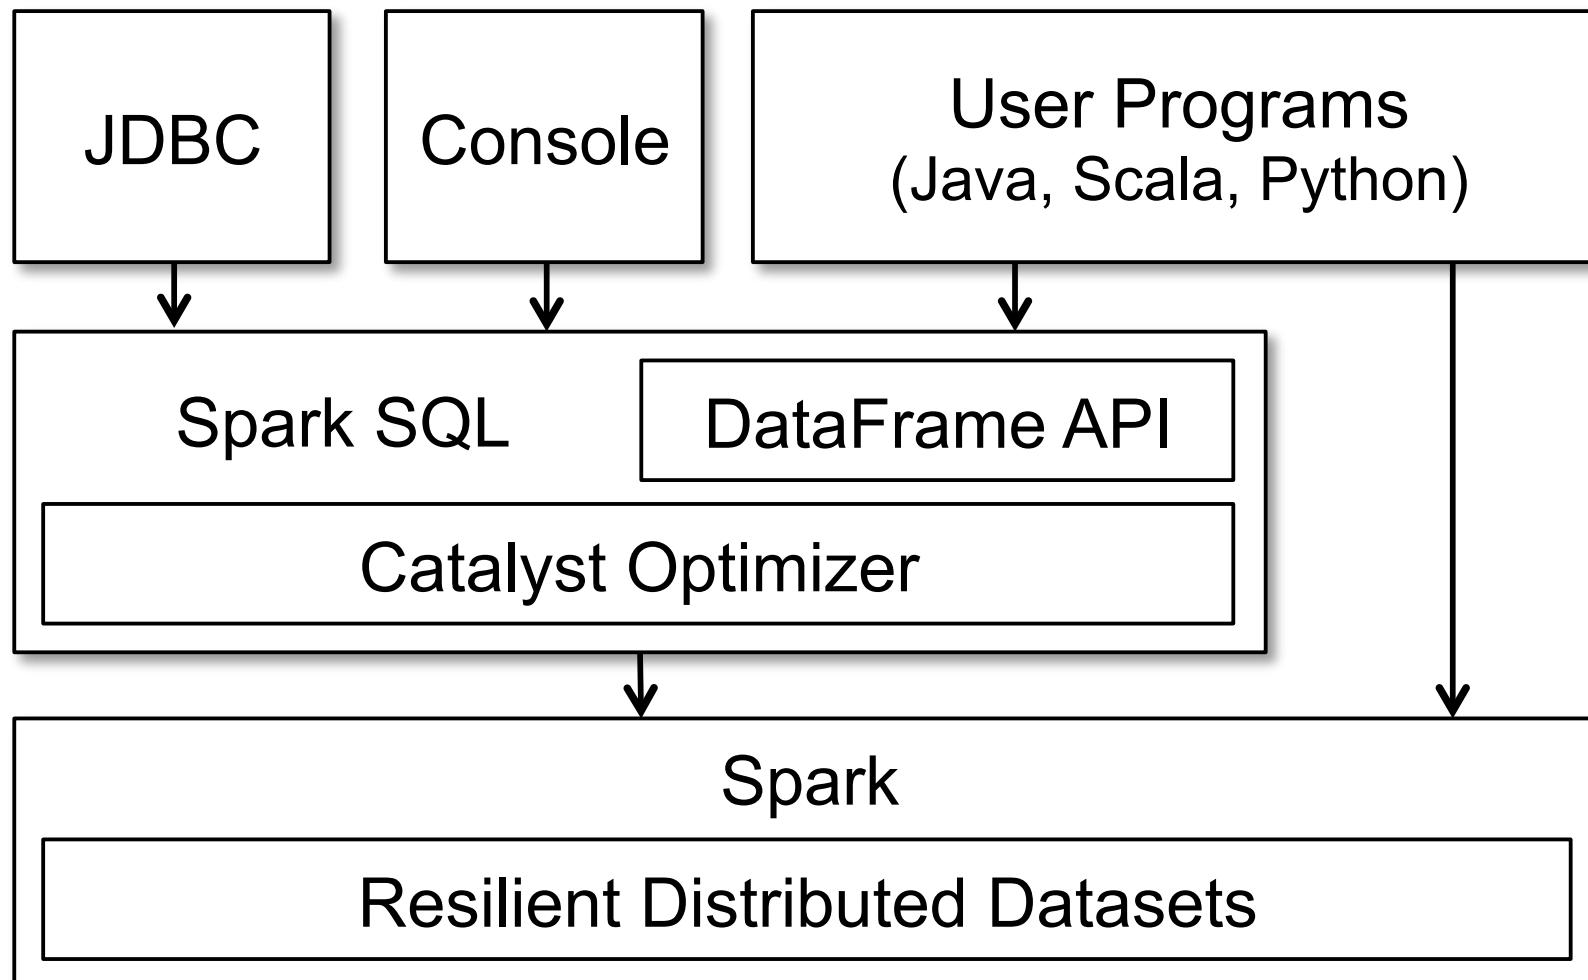
# Spark Streaming

- Interface
  - Discretized stream, with each mini-batch composed of RDDs
    - RDD: distributed collections.
    - RDDs manipulated through transformation operators (e.g., map, filter, reduce, etc.).
- Execution
  - Mini-batch of RDDs evaluated periodically.

# Goals for Spark SQL

1. Support relational processing both within Spark programs and on external data sources using a programmer-friendly API.
2. Provide high performance using established DBMS optimization techniques.
3. Easily support new data sources, including semi-structured data and external databases amenable to query federation.
4. Enable extension with advanced analytics algorithms such as graph processing and machine learning.

# SparkSQL Architecture



# Spark DataFrames

- DataFrames are distributed collections of data that is grouped into named columns.
- DataFrames can be seen as RDDs with a schema that names the fields of the underlying tuples.
- How to create a DataFrame:
  - Import data from a file: JSON, CSV, parquet, etc.;
  - Import data from other systems: SQL DBs, Hive;
  - Convert a RDD into a DataFrame by supplying a suitable schema.

# DataFrame Operations

- DataFrames provide a DSL for executing relational operations, as available in frameworks like Python Pandas.
- Some operations:
  - `select( cols )`
  - `filter(condition)`
  - `join( RDD, on, how )`
  - `groupBy( cols )`
  - `sort( cols, )`

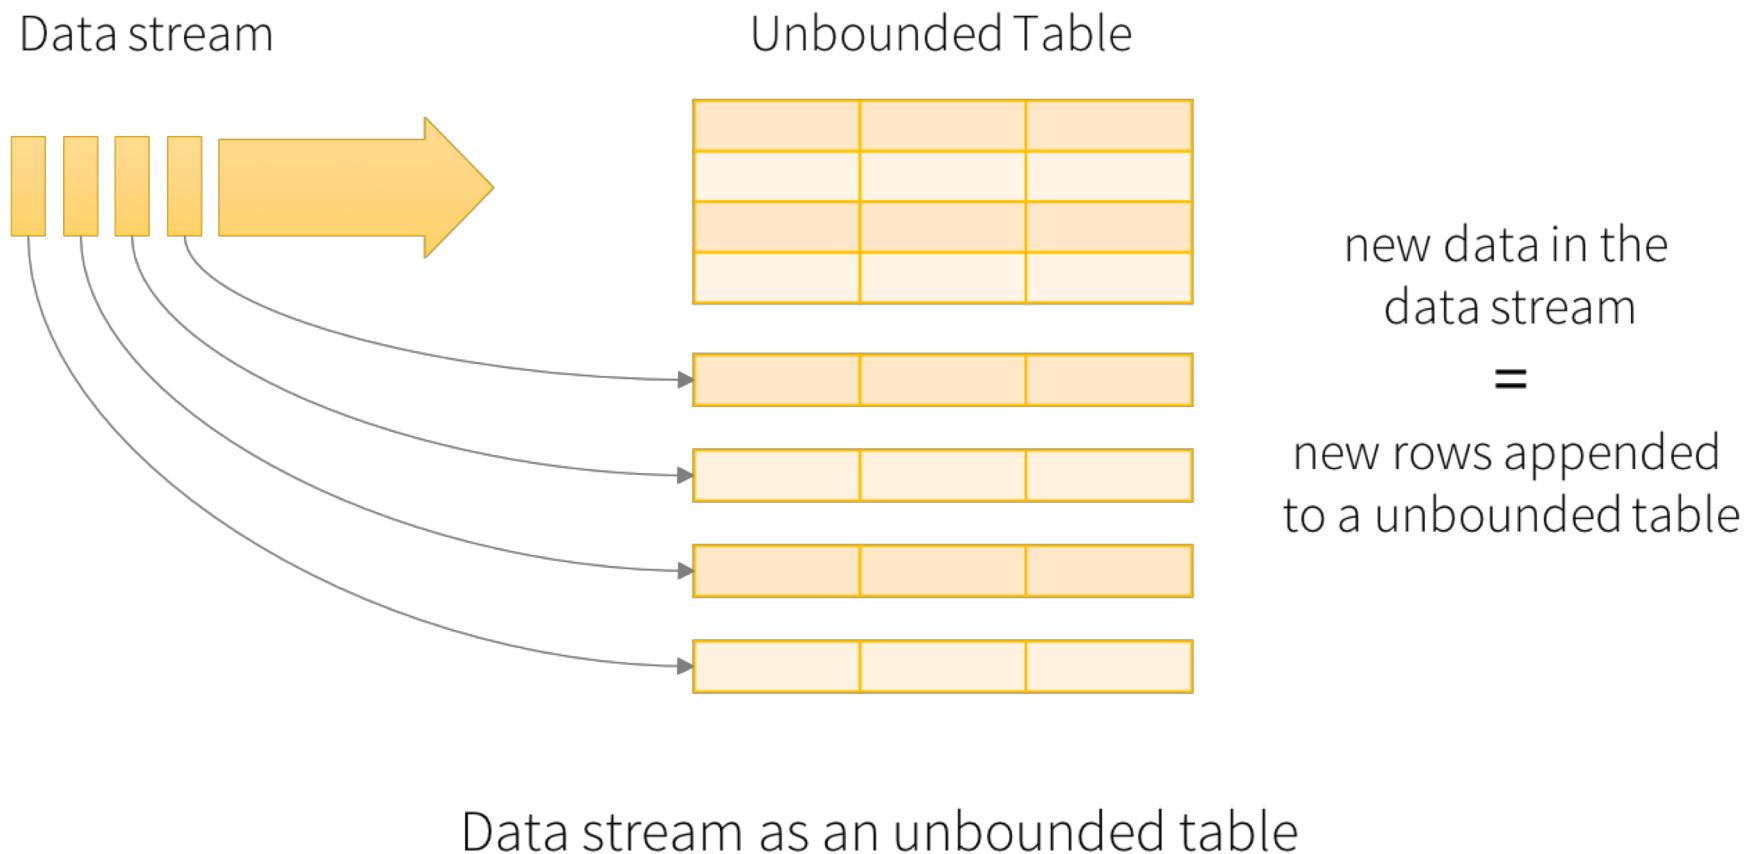
# Spark : DataFrame advantages

- Spark programs based on DataFrames are more readable due to its higher-level API.
- API close to relational operator of SQL.
- Some common programming patterns are exposed as high-level operations on DataFrames, also leading to shorter programs.

# Structured Streaming

- Key idea is to treat a live data stream as a table that is being continuously appended.
  - Similar to the batch processing model.
- Express streaming computation as standard batch-like query as on a static table, and Spark runs it as an *incremental* query on the *unbounded* input table.

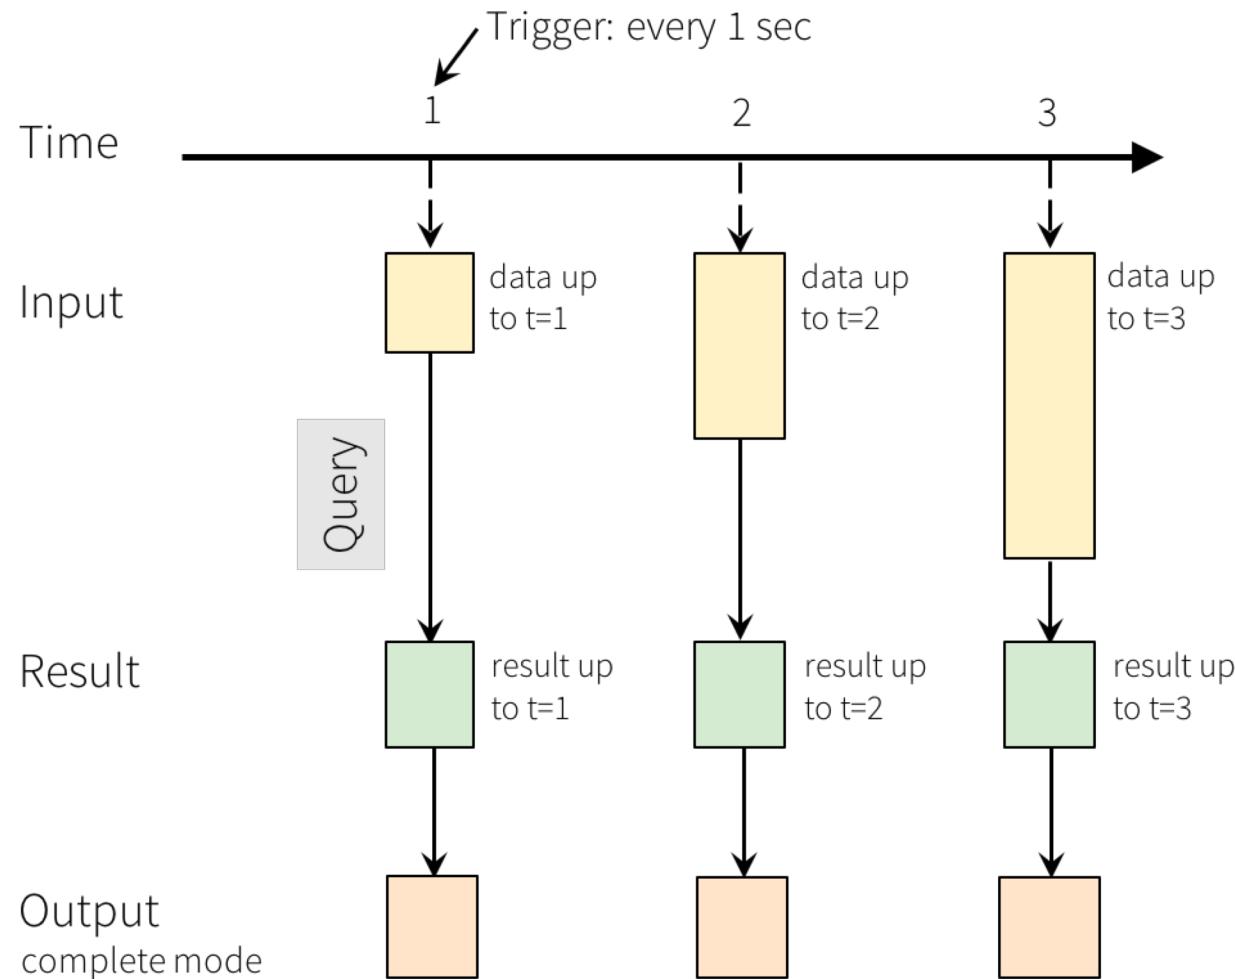
# Data stream model



# Overall execution model

- Source provides rows that are appended to the Input Table every trigger interval.
- A query on the input will generate the “Result Table”.
- Whenever the result table gets updated, the changes can be sent to an external sink.

# Execution model (cont.)



Programming Model for Structured Streaming

# Output modes

- ***Complete Mode*** - The entire updated Result Table will be written to the external storage.
- ***Append Mode*** - Only the new rows appended in the Result Table since the last trigger will be written to the external storage.
- ***Update Mode*** - Only the rows that were updated in the Result Table since the last trigger will be written to the external storage.

# Example

```
from pyspark.sql import SparkSession
from pyspark.sql.functions import explode
from pyspark.sql.functions import split

spark = SparkSession \
    .builder \
    .appName("StructuredNetworkWordCount") \
    .getOrCreate()
```

# Example (cont)

```
# Create DataFrame representing the stream of input
# lines from connection to localhost:9999
lines = spark.readStream.format("socket") \
    .option("host", "localhost").option("port", 9999) \
    .load()

# Split the lines into words
words = lines.select(explode( split(lines.value, " "))) \
    .alias("word") )

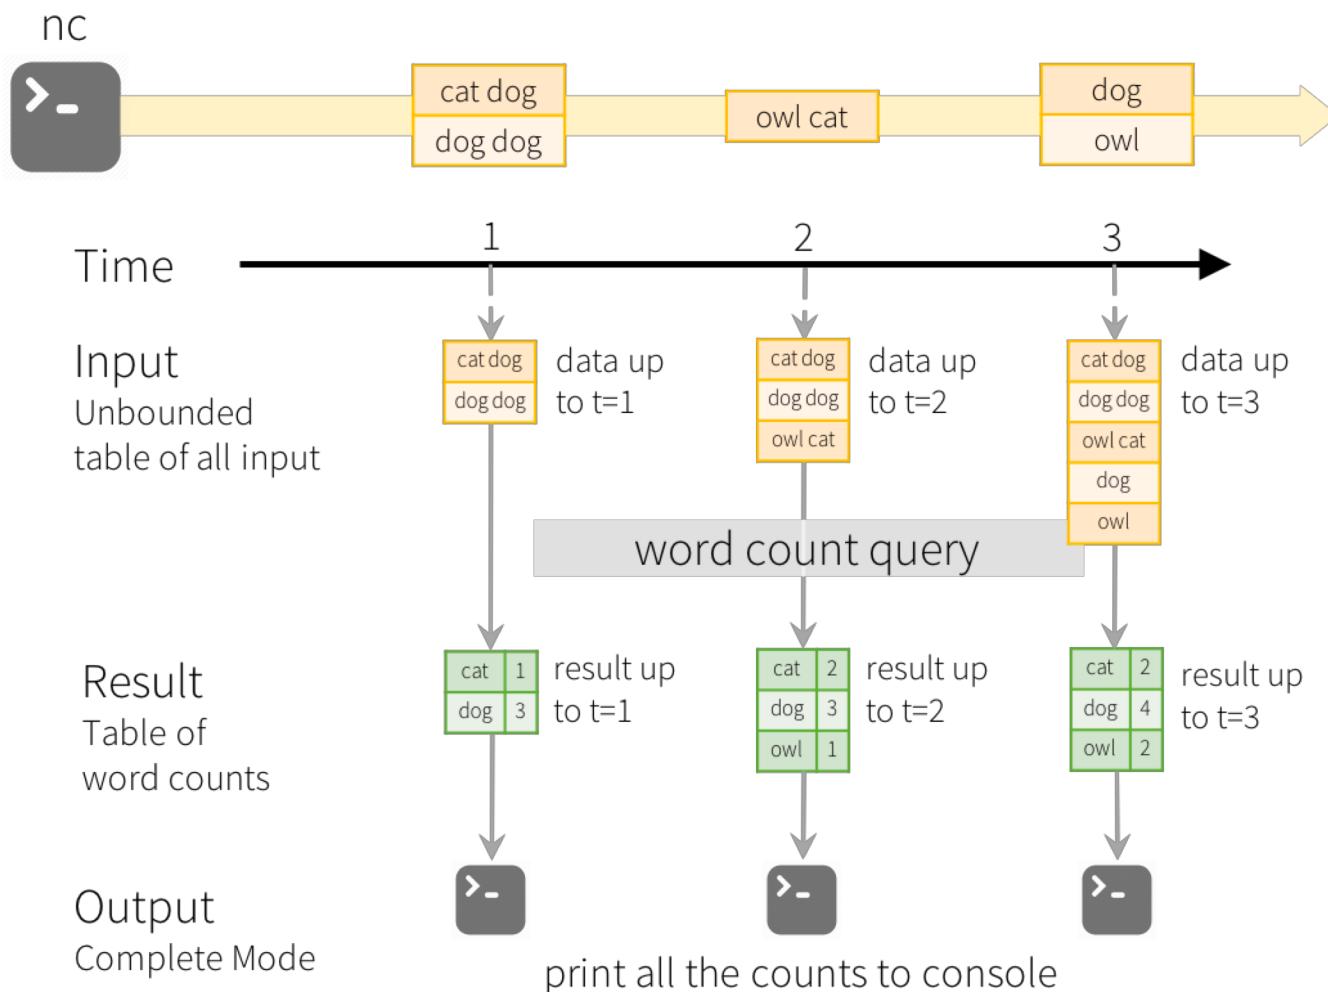
# Generate running word count
wordCounts = words.groupBy("word").count()
```

# Example (cont)

```
# Start running the query that prints the
# running counts to the console
query = wordCounts \
    .writeStream \
    .outputMode("complete") \
    .format("console") \
    .start()

query.awaitTermination()
```

# Example (cont)



Model of the Quick Example

# Incremental execution

- Spark Streaming processing:
  - reads the latest available data from the input;
  - process the data incrementally to update the result;
  - Discards the input data, keeping only minimal data to update the result.
- No need to maintain running aggregation or reason about fault-tolerance and data consistency.

# Exactly-once semantics

- Exactly-once semantics by combining:
  - Replayable sources
    - Spark uses checkpointing and write-ahead logs to record the offset range.
  - Idempotent sinks

# Input sources

- **File source** - Reads files written in a directory as a stream of data.
- **Kafka source** - Reads data from Kafka.
- **Socket source (for testing)** - Reads UTF8 text data from a socket connection. No fault-tolerance guarantees.

# Example

- In the example, we will consider a set of log files containing information for web accesses.

date IP\_source return\_value operation URL time

2016-12-06T08:58:35.318+0000 37.139.9.11 404 GET /codemove/TTCENCUFMH3C 0.026

# Operation: select

```
df.select(what).where(condition)
```

- Select columns for rows matching a given condition

```
log = ...  
# streaming DataFrame of schema  
# { timestamp: Timestamp, IP: string, ... , time: float}
```

```
log.select("IP").where("time > 1.0")
```

```
log.groupBy("IP").count()
```

# Operation: aggregation

`df.groupBy(what).aggregation()`

- Group rows by a given column and execute some aggregation function.

`log = ...`

```
# streaming DataFrame of schema  
# { timestamp: Timestamp, IP: string, ... , time: float}
```

`log.groupBy("IP").count()`

# Using SQL

- It is possible to use SQL by registering data frames as tables

```
df.createOrReplaceTempView(name)
```

- Create a view from a data frame

```
spark.sql(sql_statement)
```

# Event time vs. received time

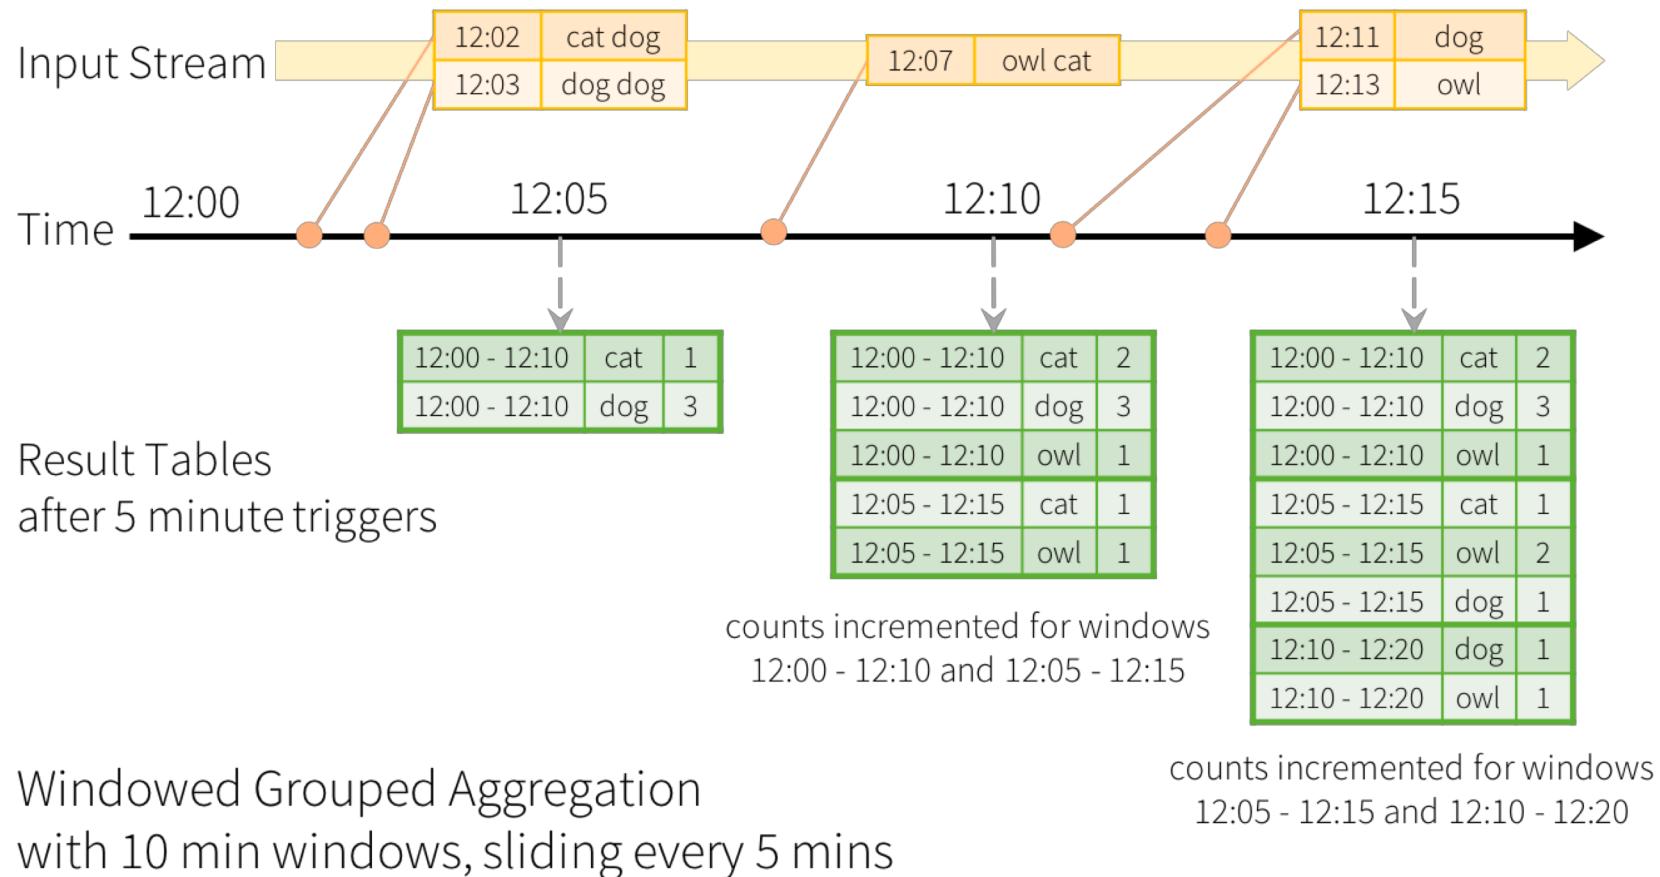
- Window operation allow to group elements based on event time

```
stream.groupBy( window(timestamp,  
    windowDuration, slideDuration), what)
```

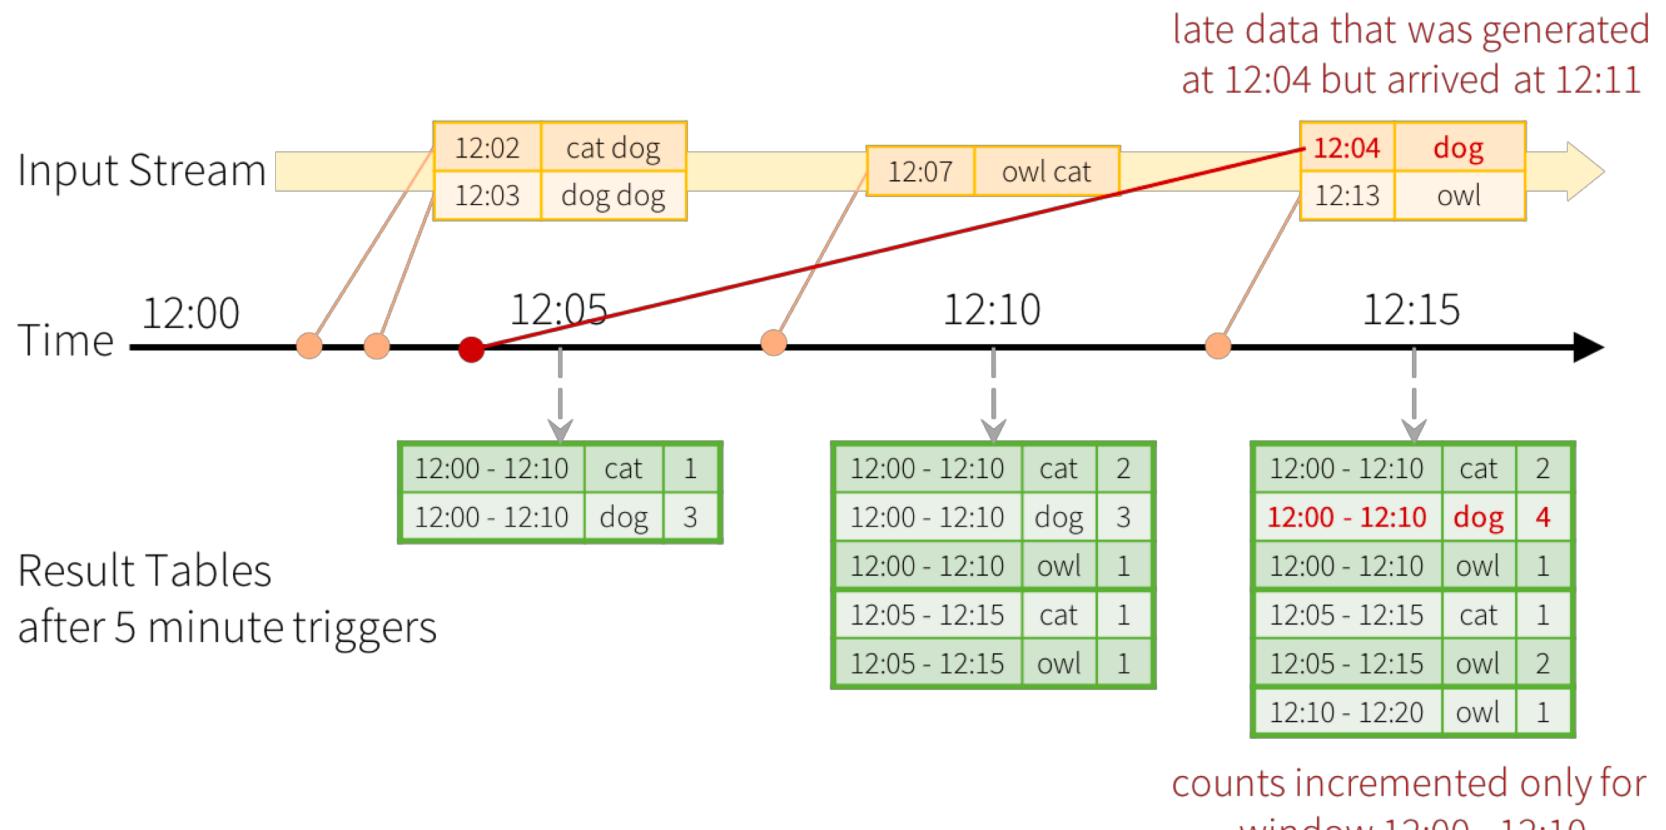
# Event time (example)

```
words = ...  
# streaming DataFrame of schema  
# { timestamp: Timestamp, word: String }  
  
# Group the data by window and word  
# and compute the count of each group  
windowedCounts = words.groupBy( \  
    window(words.timestamp, "10 minutes", "5 minutes"), \  
    words.word  
).count()
```

# Event time (example, cont.)



# Handling late data



Late data handling in  
Windowed Grouped Aggregation

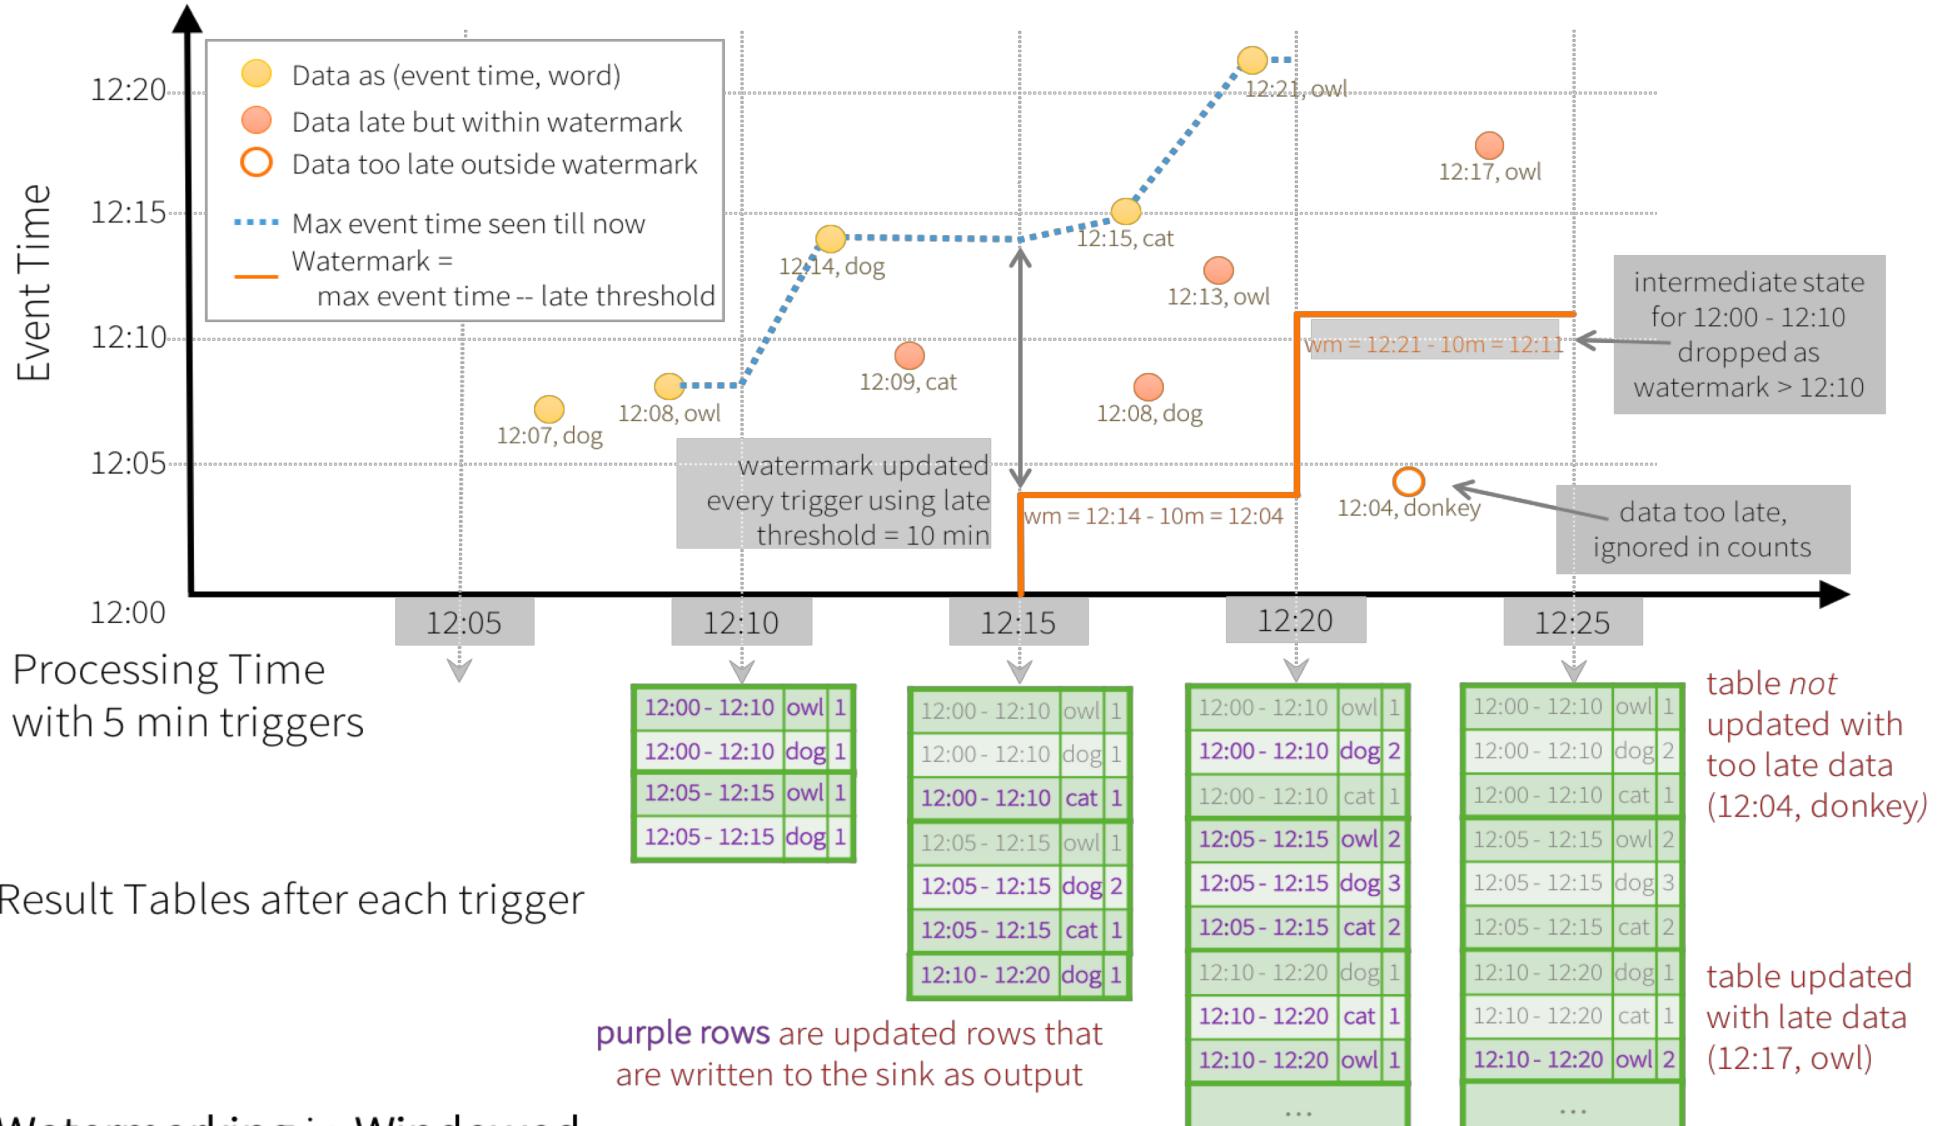
# Handling late data (cont.)

- Problem?
  - Handling late data requires keeping data for
- Watermarking: define the threshold on how late the data is expected to be in terms of event time
  - Late data within the threshold will be aggregated, but data later than the threshold will start getting dropped

# Watermarking (example)

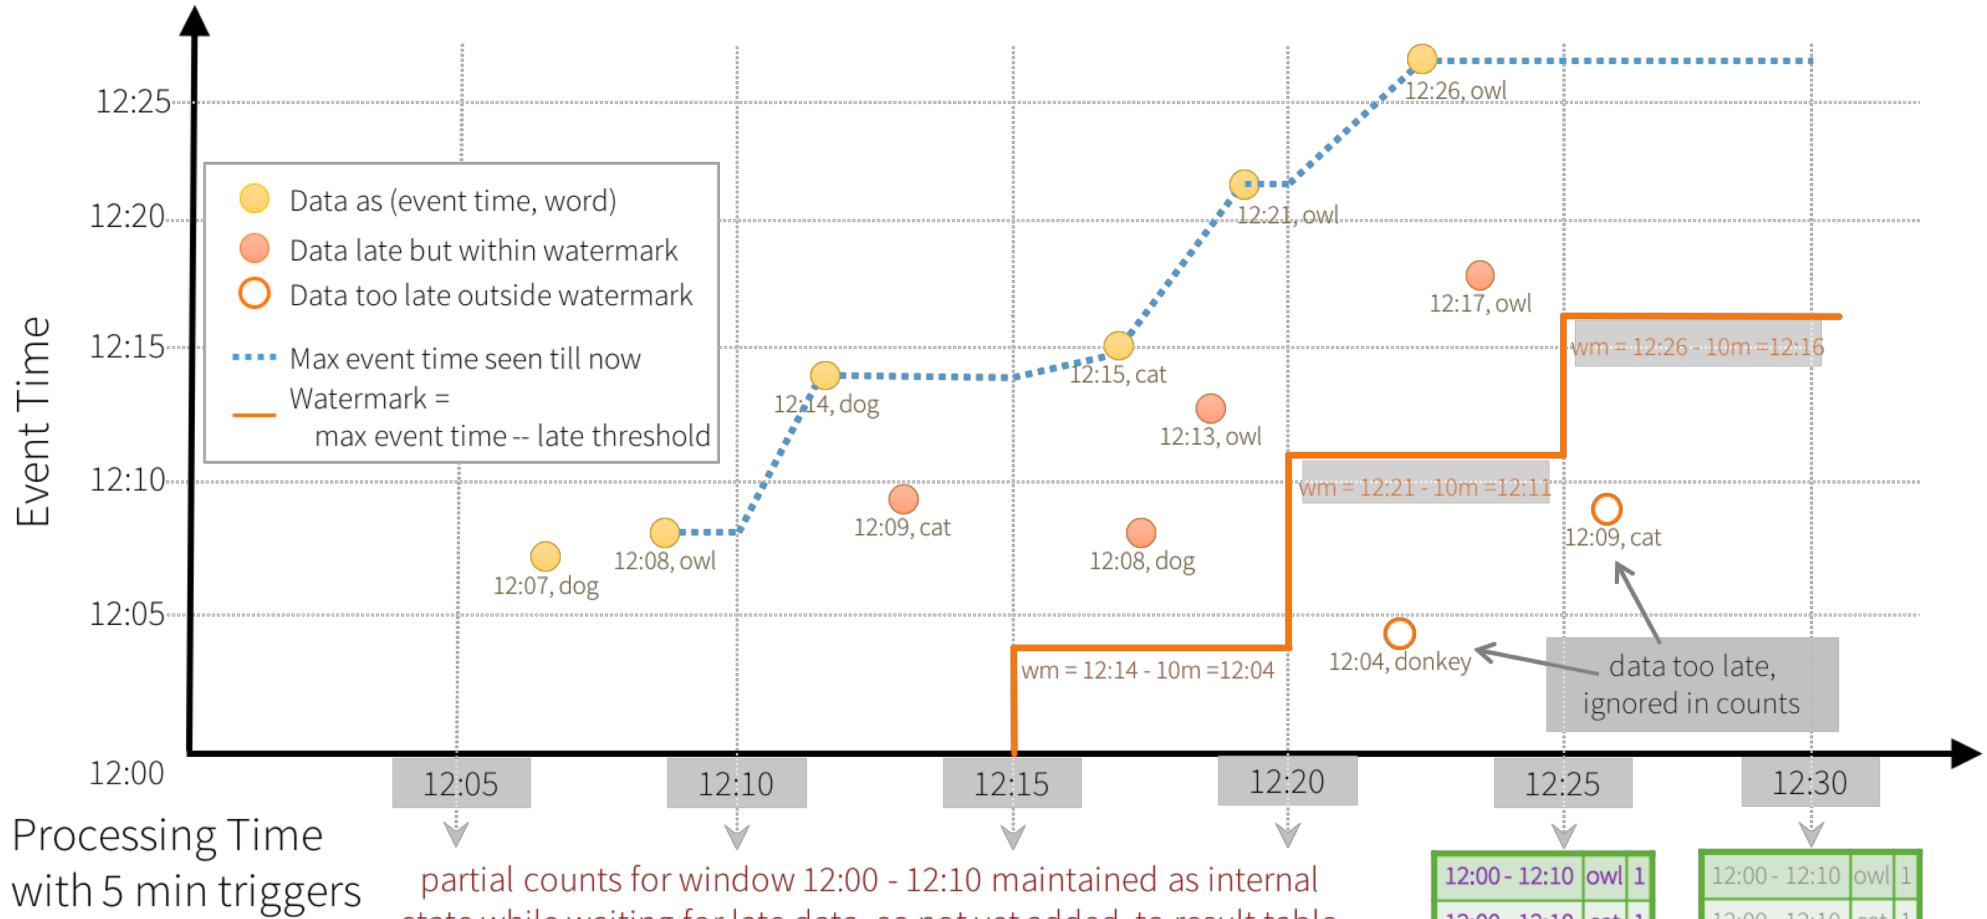
```
words = ...  
# streaming DataFrame of schema  
# { timestamp: Timestamp, word: String }  
  
# Group the data by window and word  
# and compute the count of each group  
windowedCounts = words \  
    .withWatermark("timestamp", "10 minutes") \  
    .groupByKey( \  
        window(words.timestamp, "10 minutes", "5 minutes"), \  
        words.word) \  
    .count()
```

# Watermarking (example)



Watermarking in Windowed  
Grouped Aggregation with Update Mode

# Watermarking (example)



Watermarking in Windowed  
Grouped Aggregation with Append Mode

*final counts* for 12:00 - 12:10 added to table  
when watermark > 12:10, late data counted,  
and intermediate state for window dropped

12:00 - 12:10	owl	1
12:00 - 12:10	cat	1
12:00 - 12:10	dog	2

12:00 - 12:10	owl	2
12:05 - 12:15	cat	2
12:05 - 12:15	dog	3

Result Tables after  
each trigger

# Join operations

- Supports joining:
  - a streaming Dataset/DataFrame with a static Dataset/DataFrame
  - a streaming Dataset/DataFrame with another streaming Dataset/DataFrame
- The result of the streaming join is generated incrementally.

# Stream-static joins

```
staticDf = spark.read. . .
```

```
streamingDf = spark.readStream. . .
```

```
streamingDf.join(staticDf, "type")
```

*# inner equi-join with a static DF*

```
streamingDf.join(staticDf, "type", "right_join")
```

*# right outer join with a static DF*

# Stream-stream join

- Challenge: at any given point, the view of the data is incomplete. Need to buffer past input for matching with new input.
- Solution: use watermarking to guarantee that data is not buffered forever.

# Stream-stream joins

```
from pyspark.sql.functions import expr

impressions = spark.readStream. ...
clicks = spark.readStream. ...

# Apply watermarks on event-time columns
impressionsWithWatermark = impressions.withWatermark("impressionTime",
"2 hours")
clicksWithWatermark = clicks.withWatermark("clickTime", "3 hours")

# Join with event-time constraints
impressionsWithWatermark.join( \
    clicksWithWatermark, \
    expr("""
        clickAdId = impressionAdId AND
        clickTime >= impressionTime AND
        clickTime <= impressionTime + interval 1 hour
    """) \
)
```

# Output sink

- **File sink** - Stores the output to a directory.
- **Kafka sink** - Stores the output to one or more topics in Kafka.
- **Foreach sink** - Runs arbitrary computation on the records in the output.
- **Console sink (for debugging)** - Prints the output to the console/stdout every time there is a trigger.
- **Memory sink (for debugging)** - The output is stored in memory as an in-memory table.

# Table of Contents

- Structured Streaming Programming
- Programming Spark streaming

# Example

```
import socket
from pyspark import SparkContext
from pyspark.streaming import StreamingContext

ssc = StreamingContext(sc, 5)
lines = ssc.socketTextStream("localhost", 7777)

lines = lines.window(15,5)
counts = lines.filter( lambda line : len(line) > 0 ) \
    .map(lambda line: (line.split(' ')[1],1)) \
    .reduceByKey(lambda a, b: a+b) \
    .transform( \
        lambda rdd: rdd.sortBy(lambda x: x[1], \
                               ascending=False))
counts.pprint()

ssc.start()
ssc.awaitTermination(120)
ssc.stop()
```

# Exercise 1

- Dump the number of request per 5 seconds if there is an IP address with more than 100 requests

# Exercise 2

- Assume you have a CSV file with the mapping of IP to country
- Return the number of access per country

# Bibliography

- <https://spark.apache.org/docs/latest/structured-streaming-programming-guide.html>